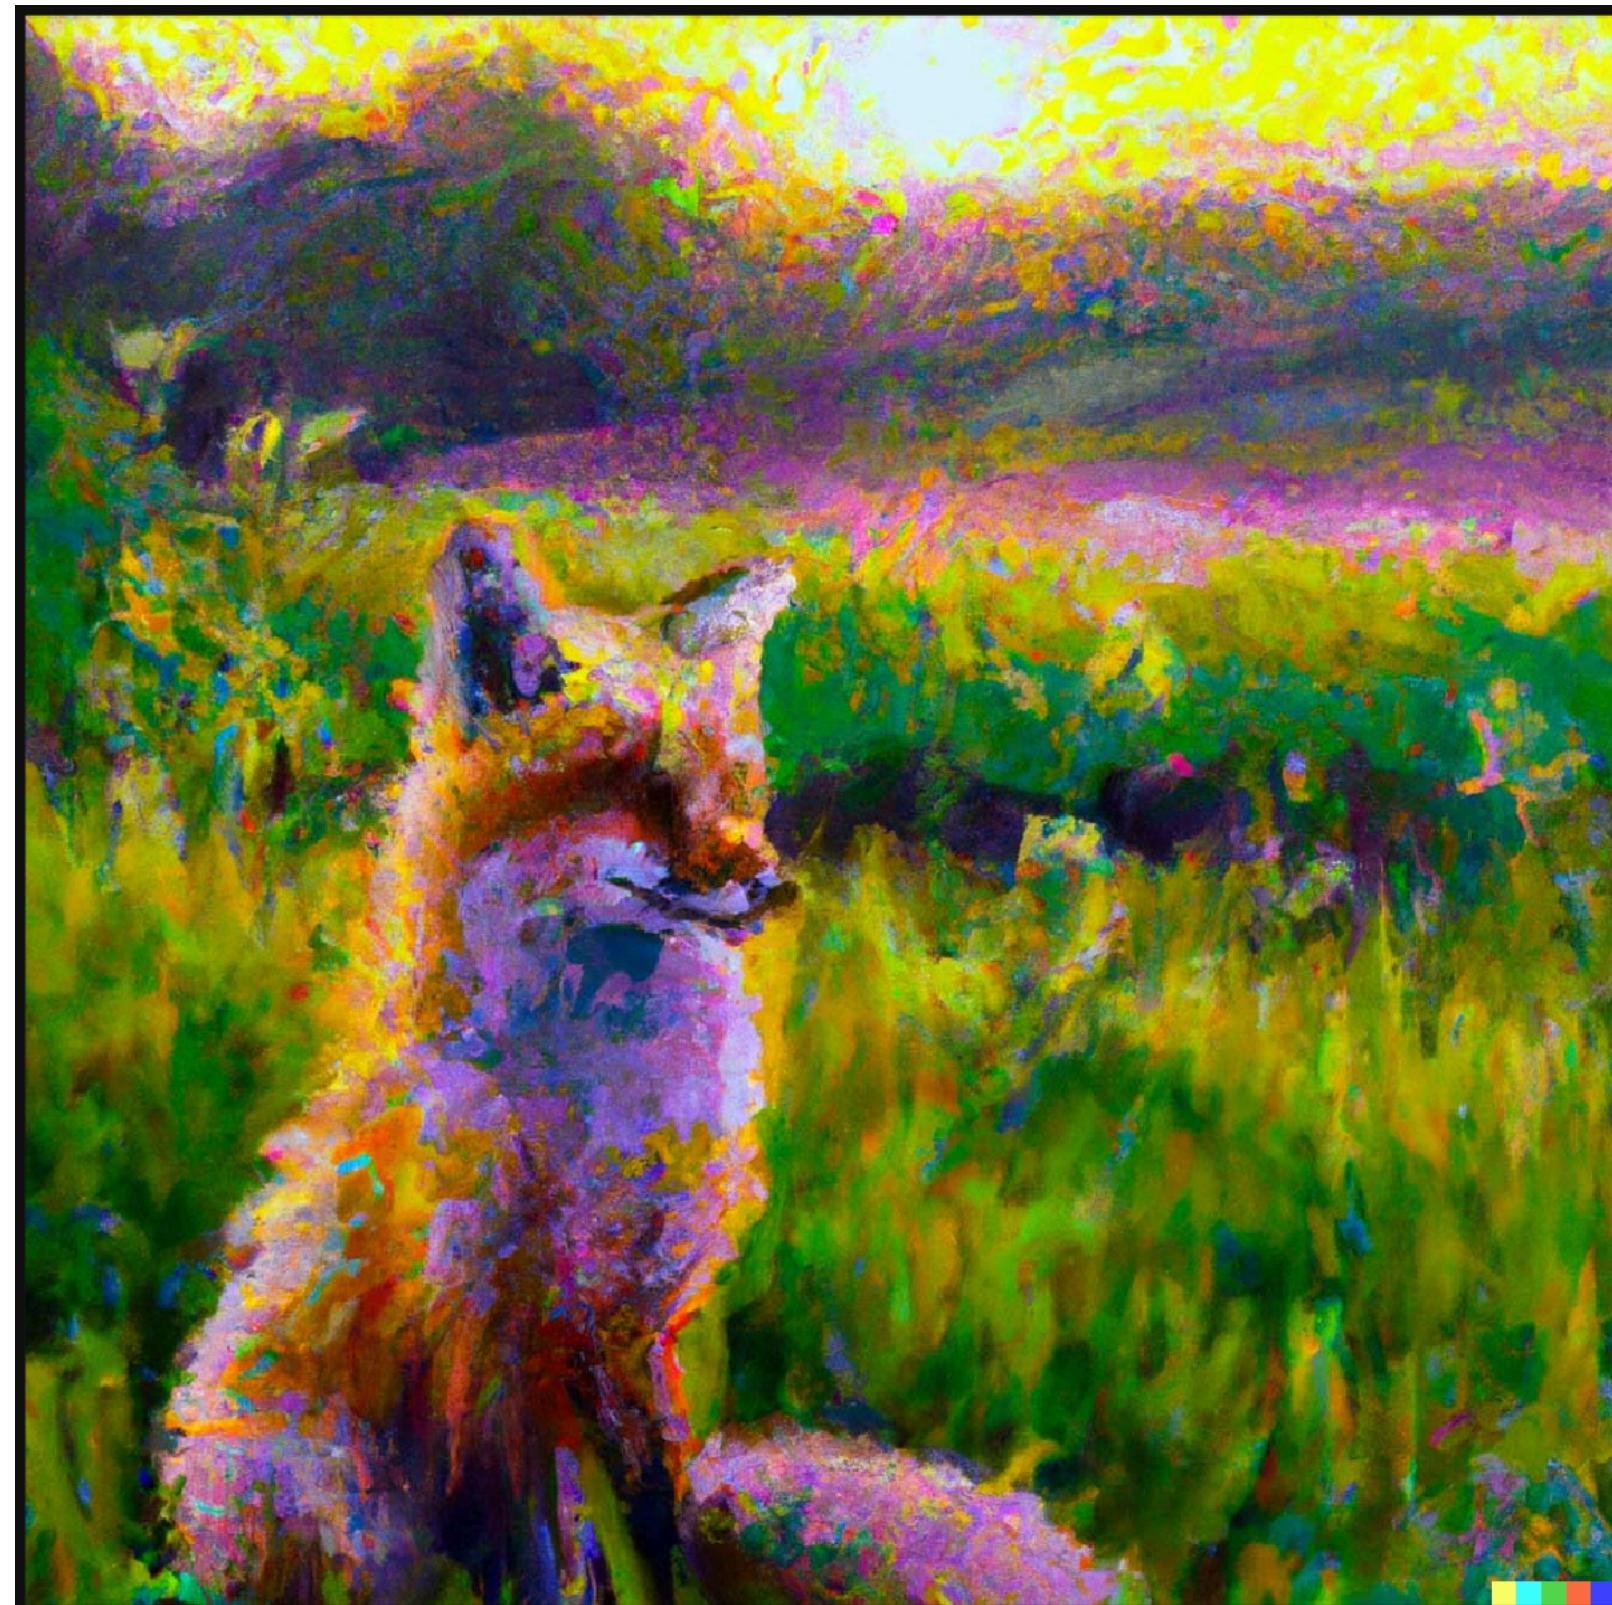


A central Pomeranian dog wears a golden crown, set against a red background. It is flanked by two other dogs, one on each side, also wearing ornate blue and gold crowns. The dogs have white and tan fur. The overall aesthetic is regal and whimsical.

Generative Artificial Intelligence

JAKUB M. TOMCZAK
VRIJE UNIVERSITEIT AMSTERDAM

Recent breakthroughs



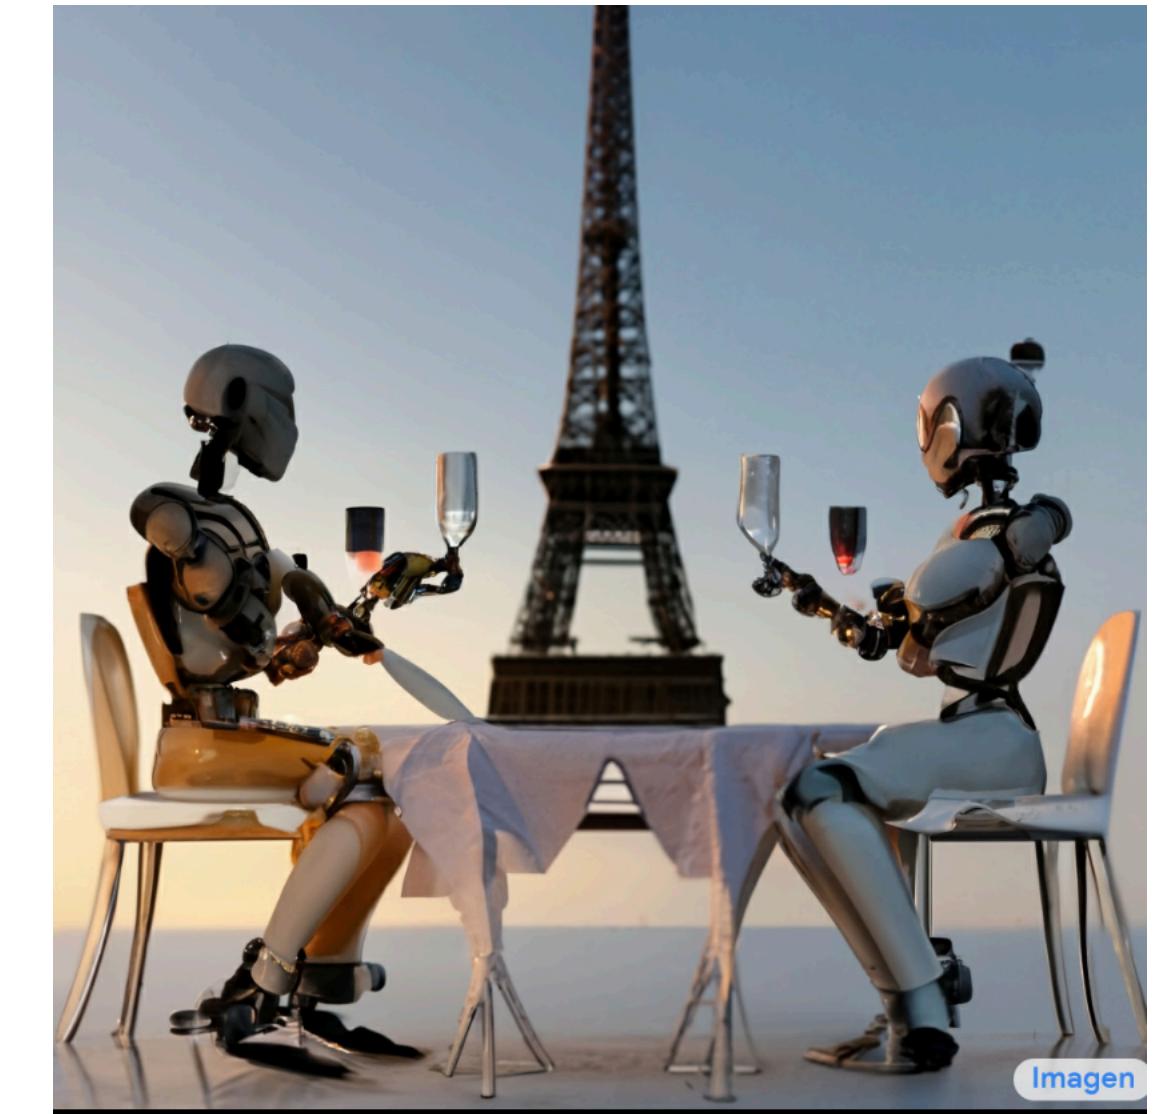
"a painting of a fox sitting in a field at sunrise in the style of Claude Monet"

Dalle-2



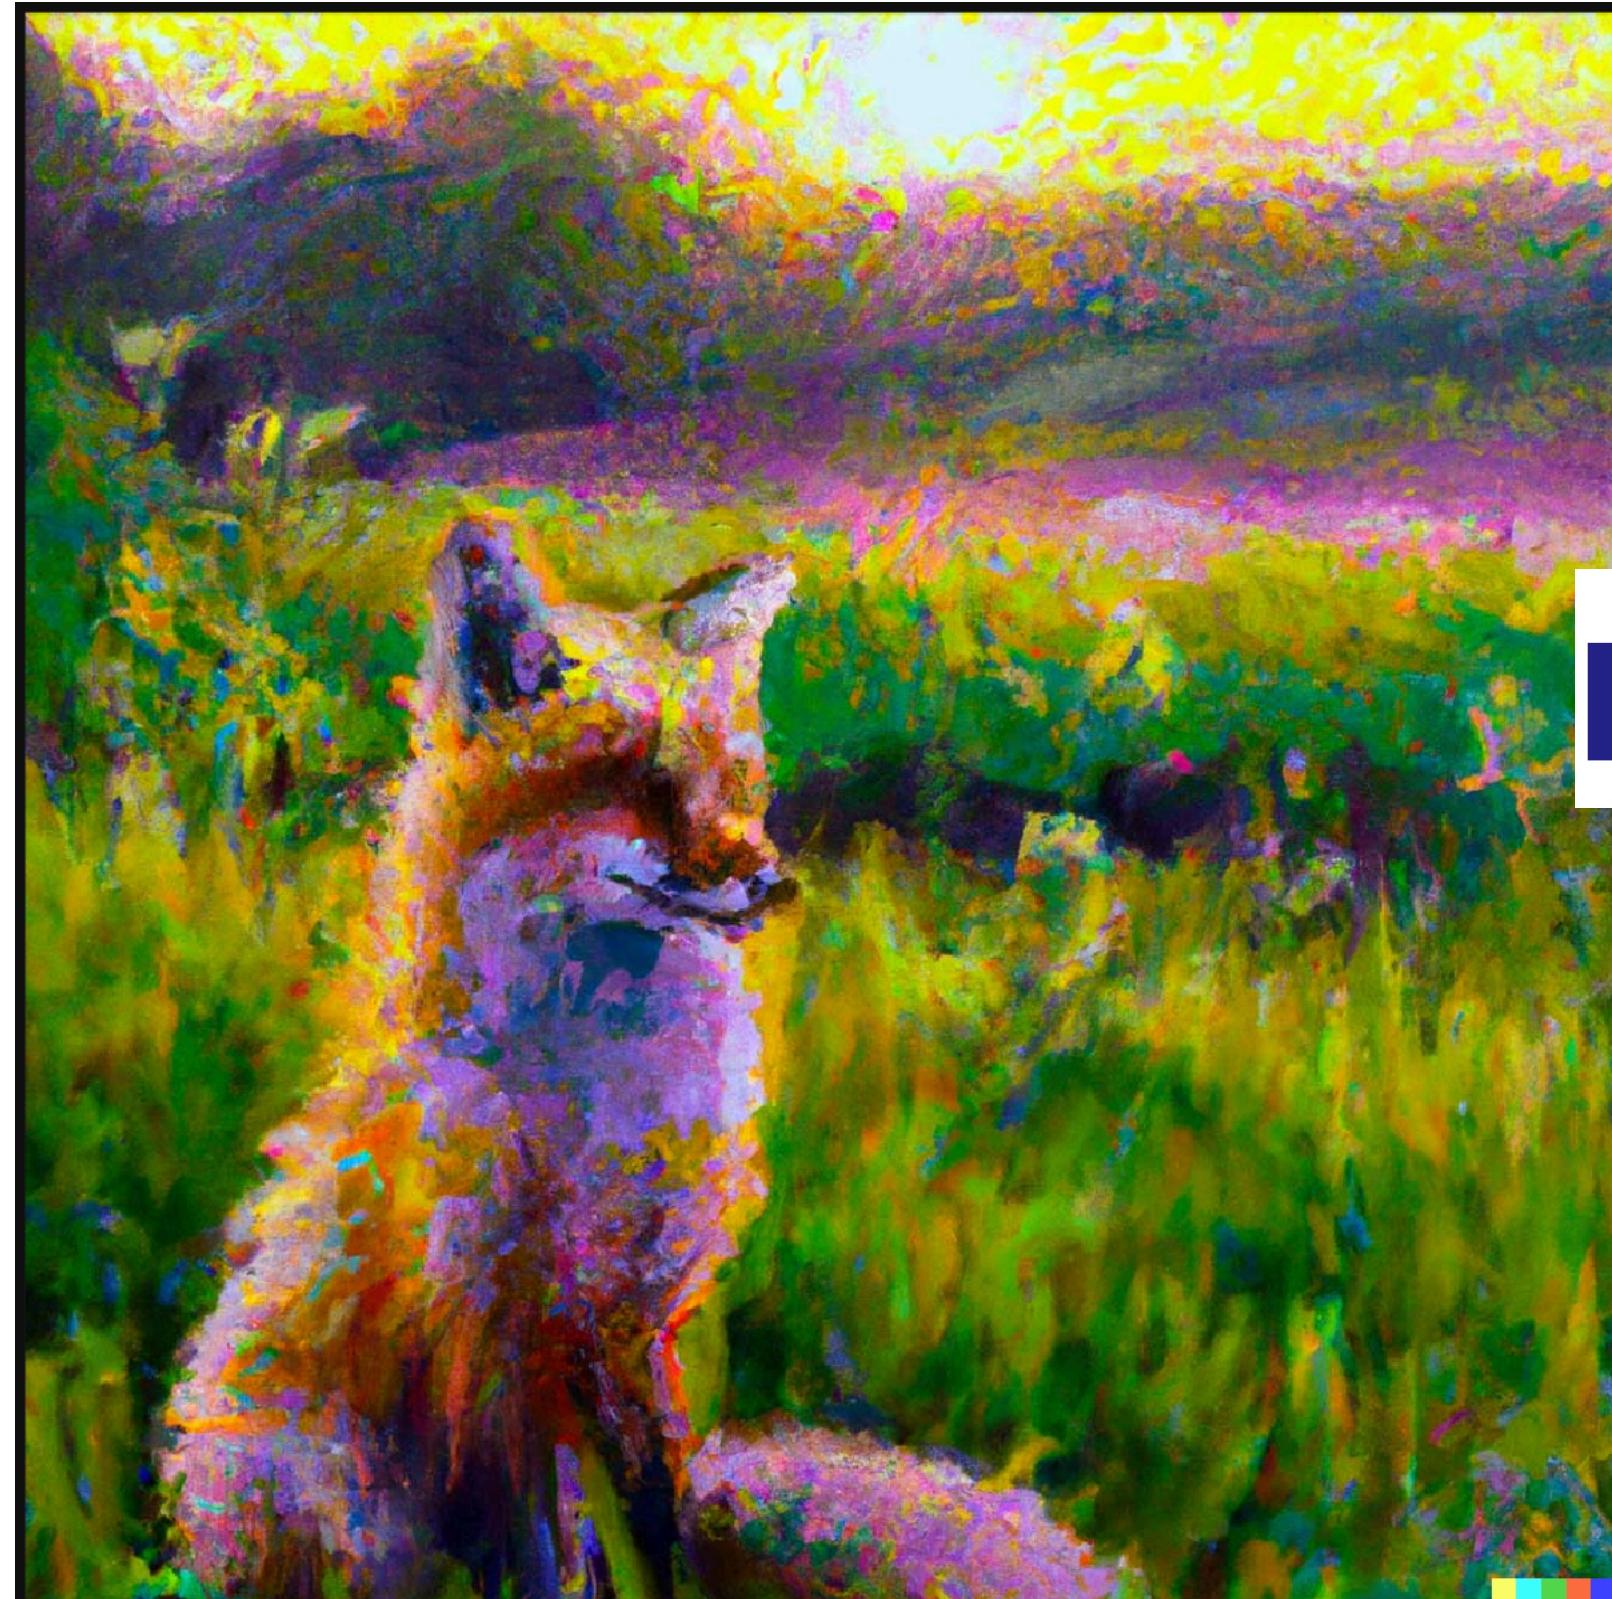
A Pomeranian is sitting on the Kings throne wearing a crown. Two tiger soldiers are standing next to the throne.

Imagen



A robot couple fine dining with Eiffel Tower in the background.

Recent breakthroughs



"a painting of a fox sitting in a field at sunrise in the style of Claude Monet"

Fantastic results!



A Pomeranian is sitting on the King's throne wearing a crown. Two tiger soldiers are standing next to the throne.



A robot couple fine dining with Eiffel Tower in the background.

Dalle·2

Imagen

Recent breakthroughs



How big are the models?

Dalle·2

"a painting of a fox sitting in a field at sunrise in the style of Claude Monet"

Fantastic results!



A Pomeranian is sitting on the King's throne wearing a crown. Two tiger soldiers are standing next to the throne.



Imagen

Recent breakthroughs



"a painting of a fox sitting in a field at sunrise in the style of Claude Monet"

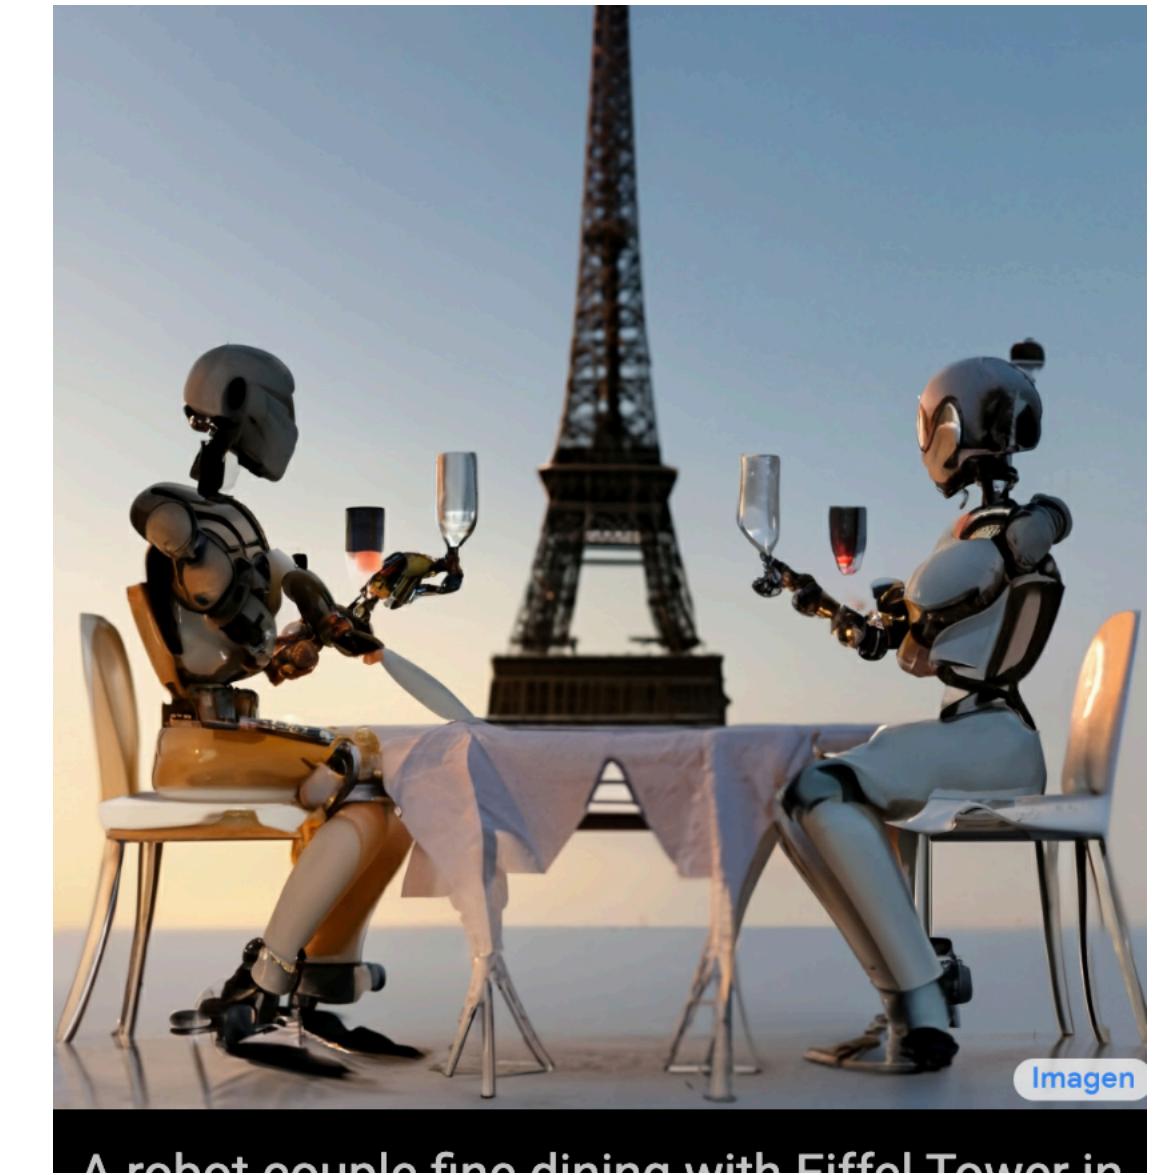
Fantastic results!

How big are the models?

Are the samples cherry-picked?

Can we use these models beyond data synthesis?

How long did it take to train these models?



Modeling: Discriminative vs. Generative

Modeling: Discriminative vs. Generative

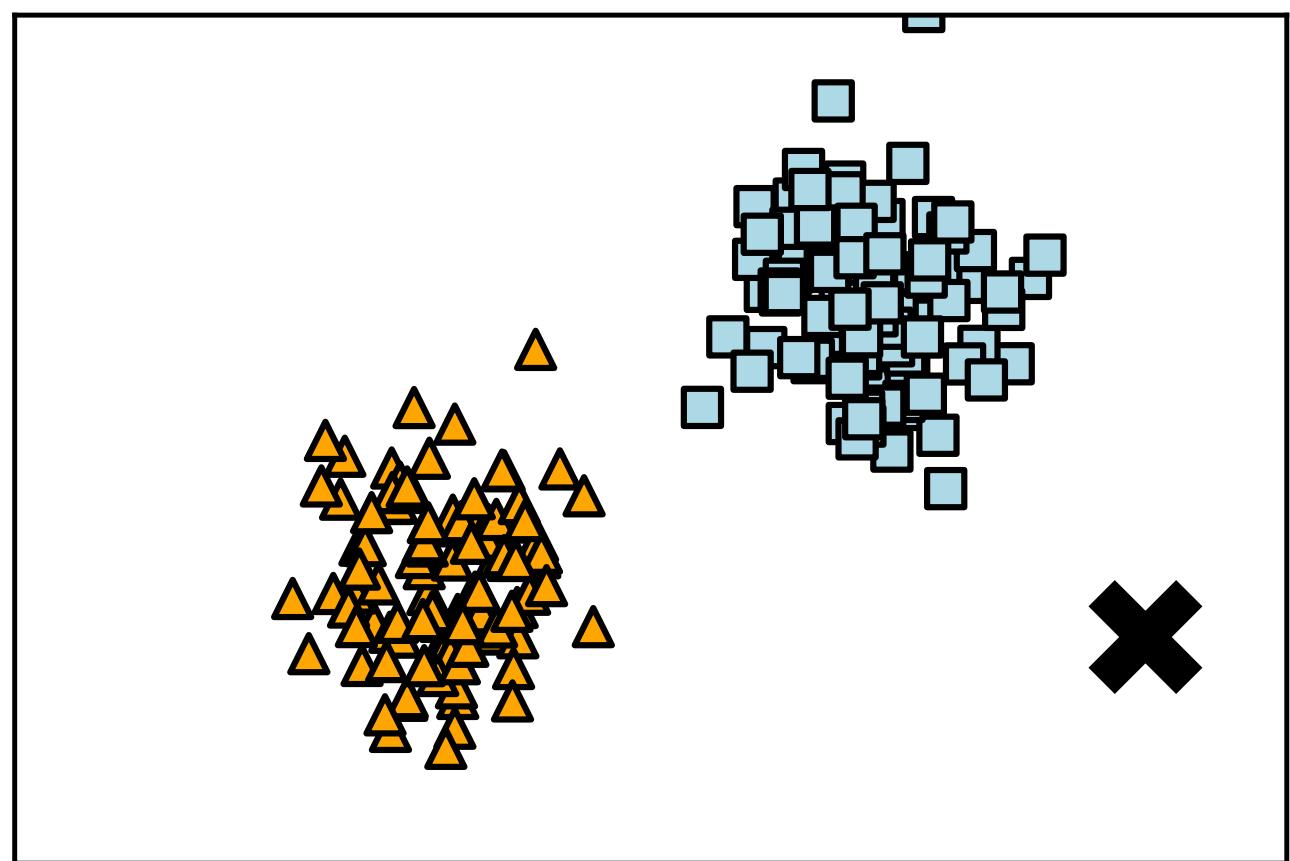
- **Discriminative models:** drawing boundaries in the data space, $p(y | \mathbf{x})$.
- **Generative models:** explaining how the data was generated, $p(\mathbf{x}, y)$.

Modeling: Discriminative vs. Generative

- **Discriminative models:** drawing boundaries in the data space, $p(y | \mathbf{x})$.
- **Generative models:** explaining how the data was generated, $p(\mathbf{x}, y)$.
- In ML, many models are **generative**:
 - Naive Bayes, Linear Discriminant Analysis
 - Bayesian networks & Markov random fields
 - (Gaussian) Mixture Models, Latent Dirichlet Allocation, Factor Analysis, PCA
 - Chinese restaurant process, Indian buffet process

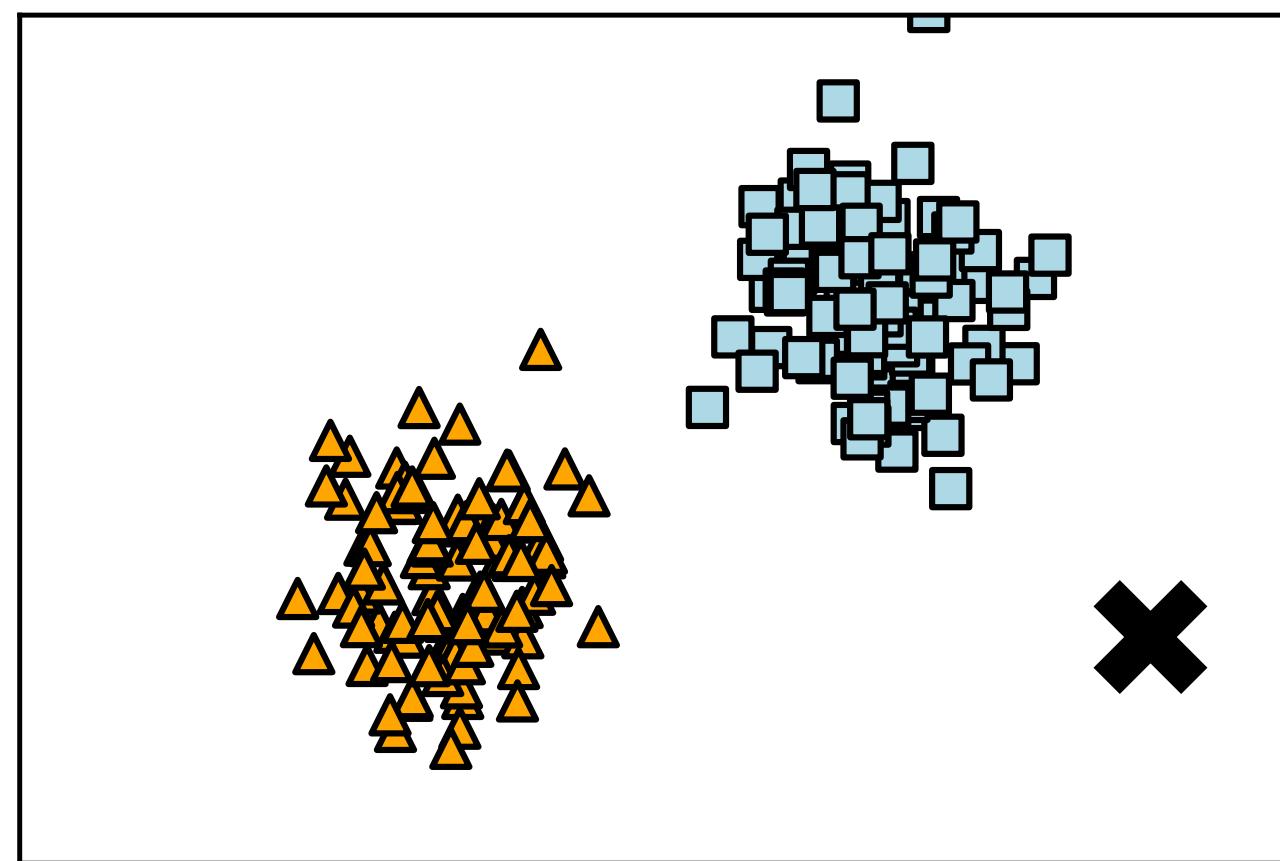
Generative = How data is generated

Generative = How data is generated

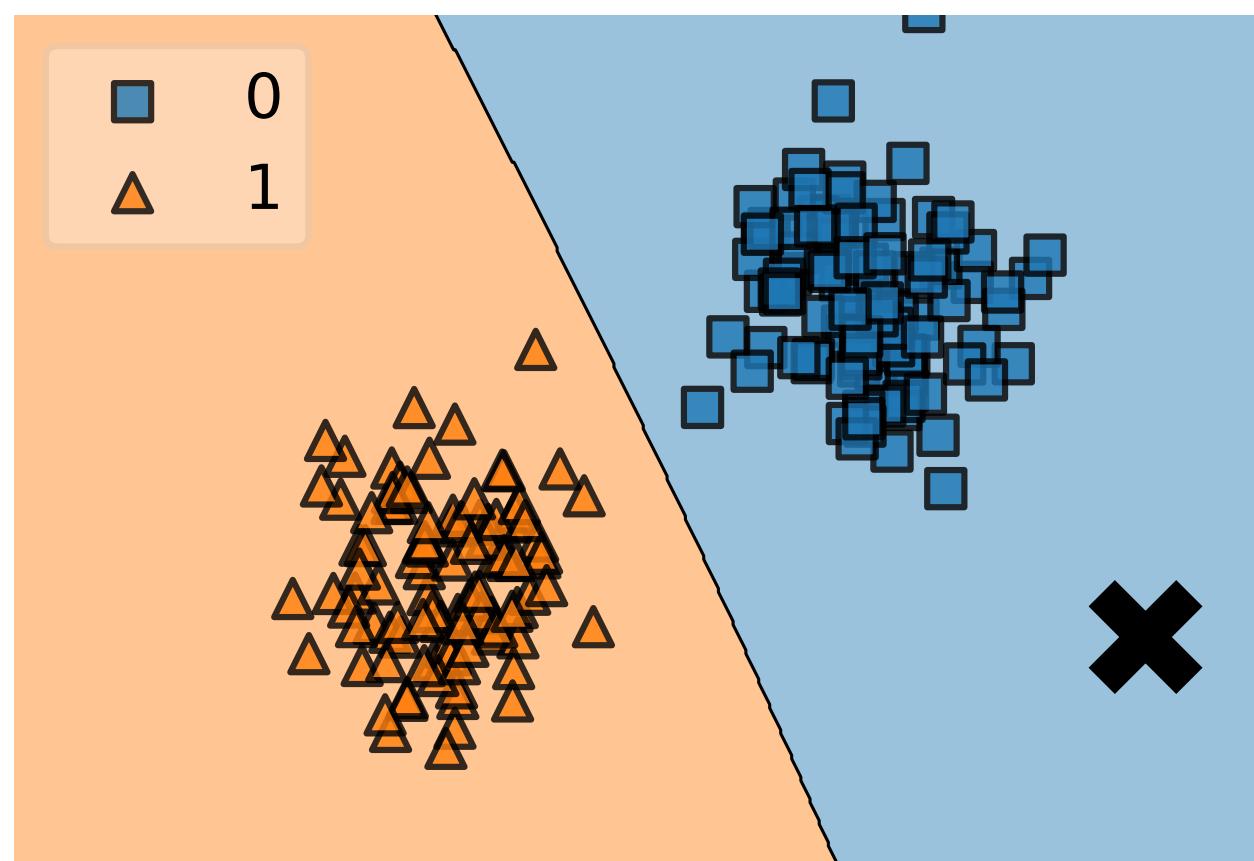


Data

Generative = How data is generated



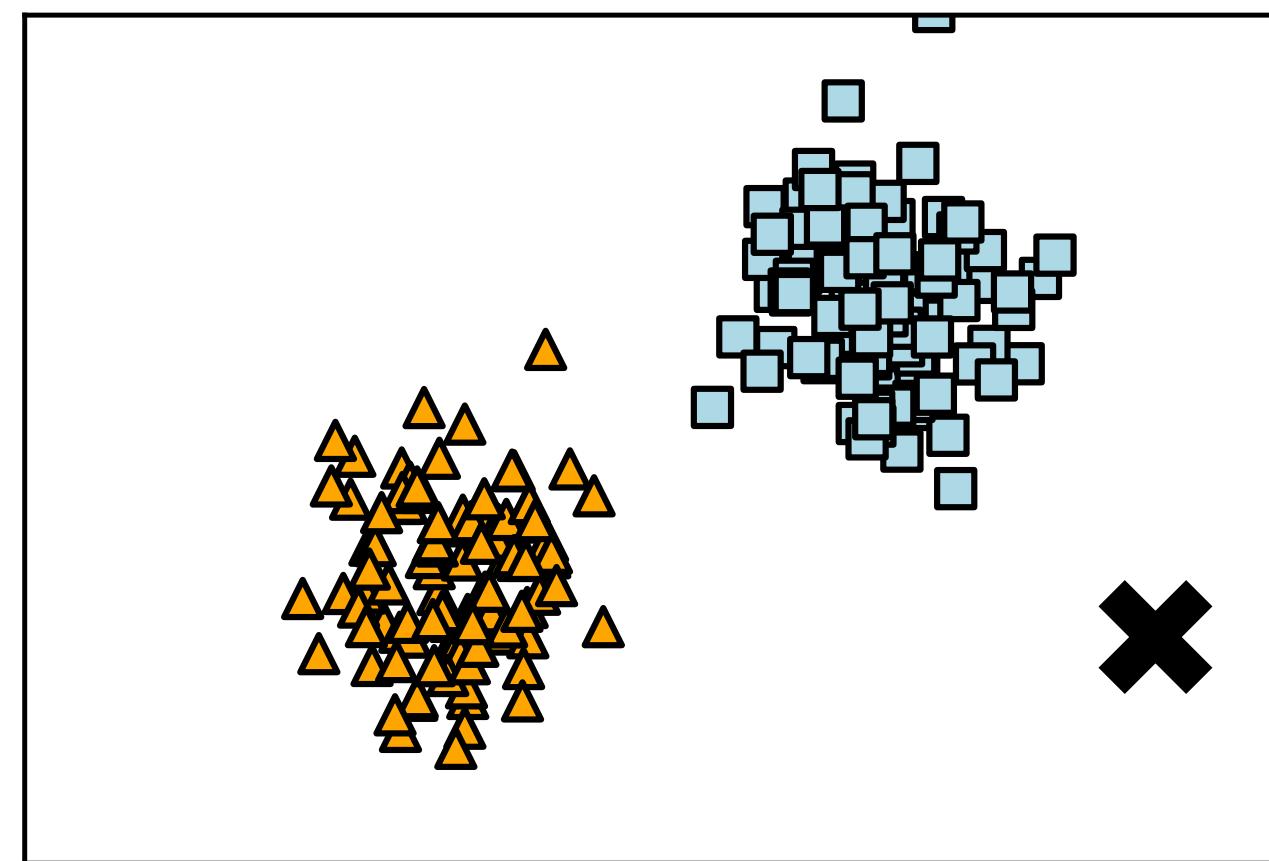
Data



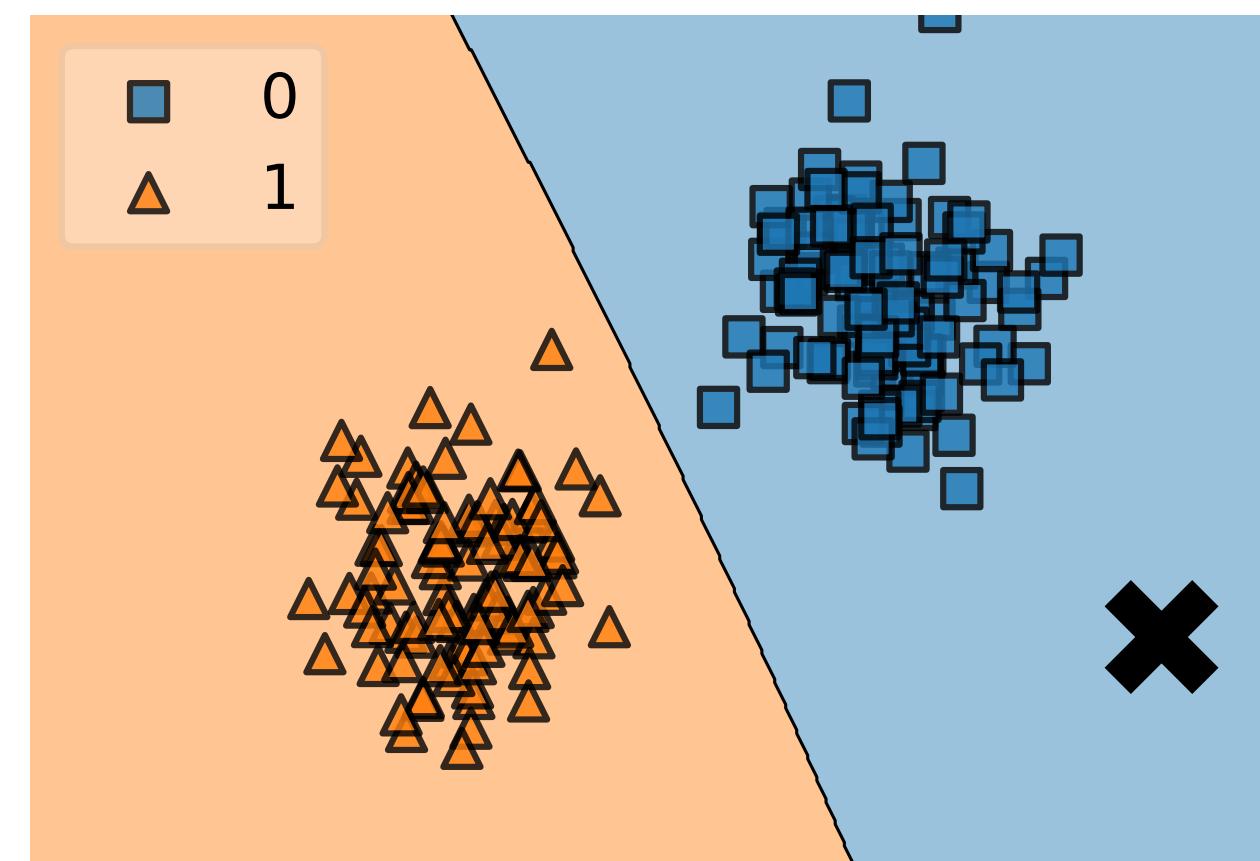
$p(y|\mathbf{x})$

$p(\text{blue}|\mathbf{x})$ is high
= certain decision!

Generative = How data is generated

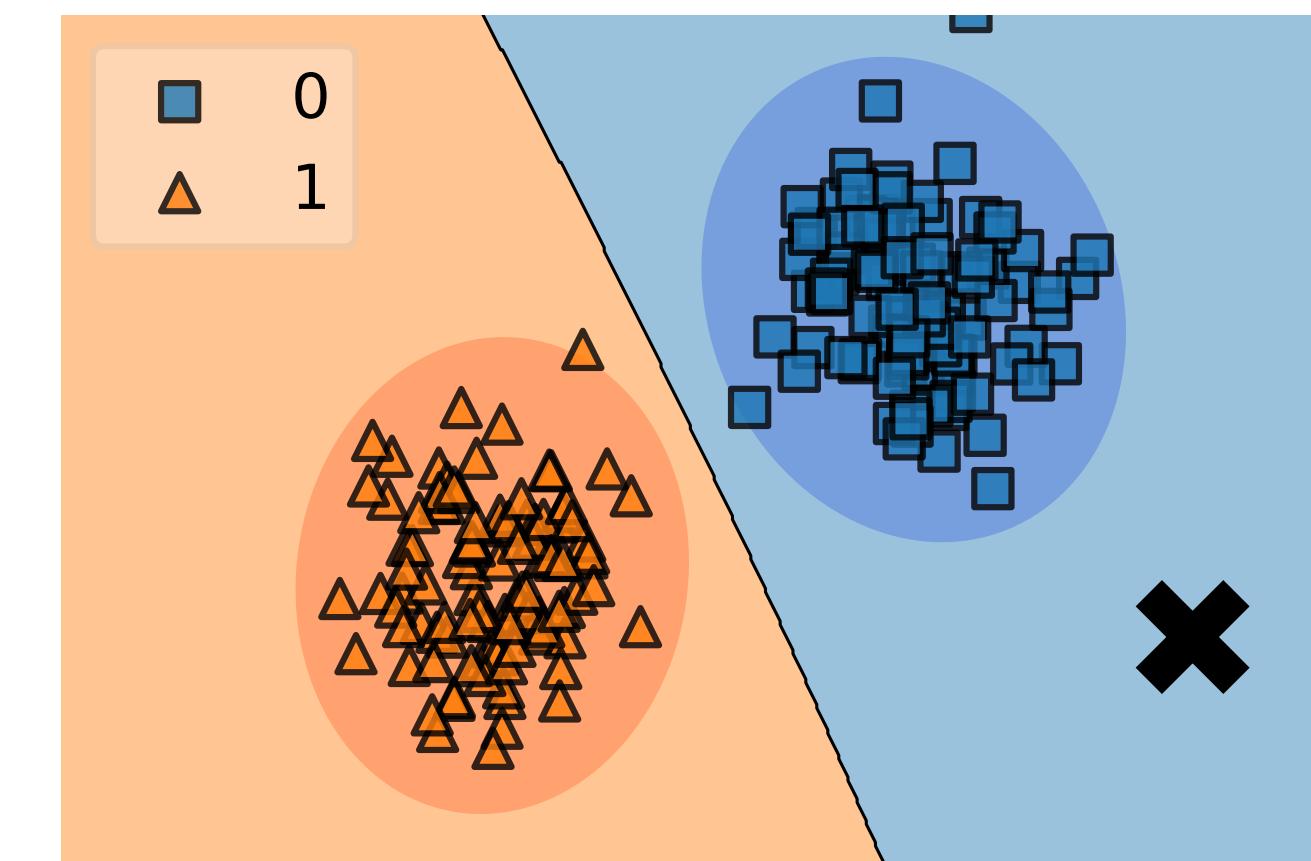


Data



$p(y|\mathbf{x})$

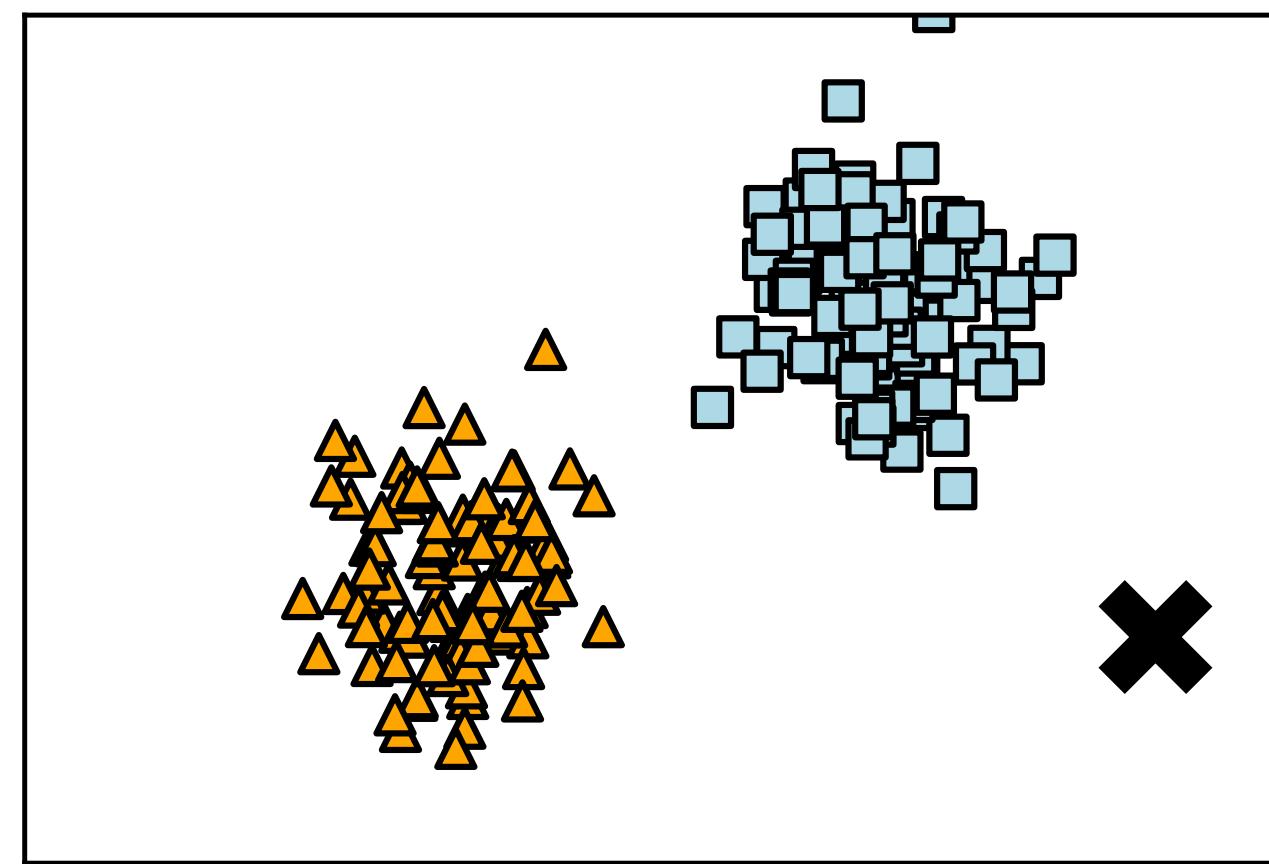
$p(\text{blue}|\mathbf{x})$ is high
= certain decision!



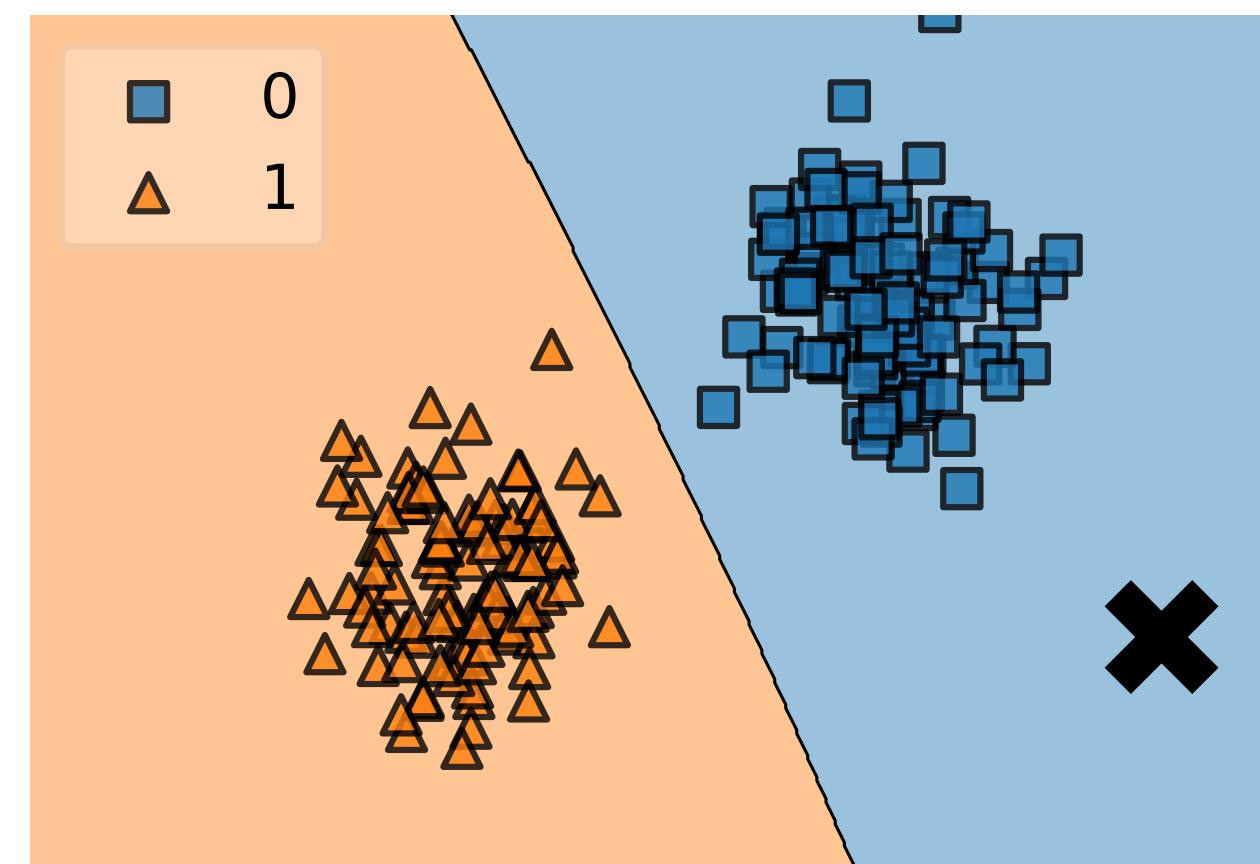
$p(\mathbf{x}, y) = p(y|\mathbf{x}) p(\mathbf{x})$

$p(\text{blue}|\mathbf{x})$ is high
and $p(\mathbf{x})$ is low
= uncertain decision!

Generative = How data is generated

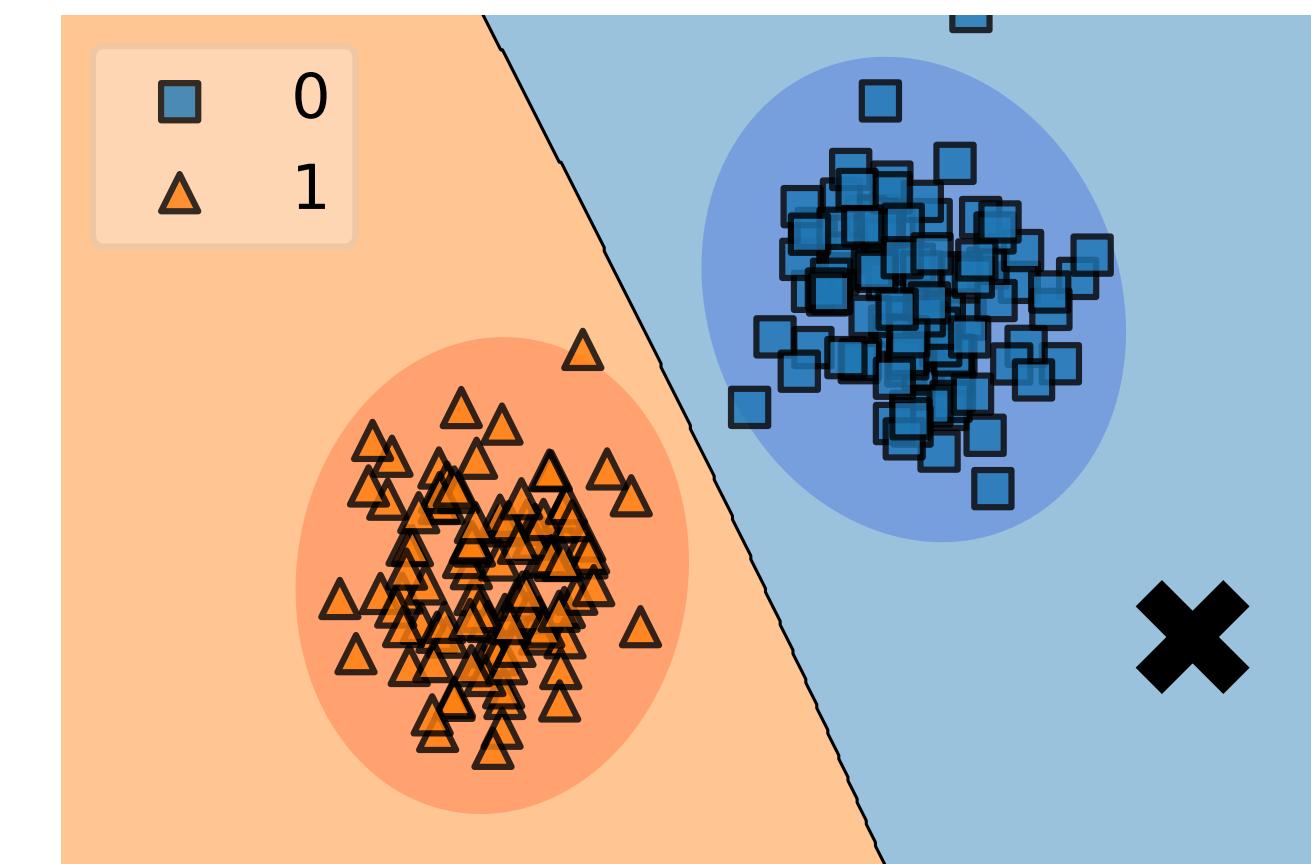


Data



$p(y|\mathbf{x})$

$p(\text{blue}|\mathbf{x})$ is high
= certain decision!



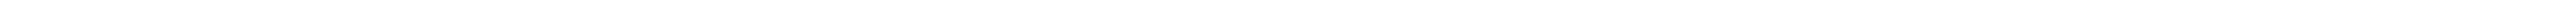
$p(\mathbf{x}, y) = p(y|\mathbf{x}) p(\mathbf{x})$

$p(\text{blue}|\mathbf{x})$ is high
and $p(\mathbf{x})$ is low
= uncertain decision!

Knowing the joint distribution tells us a lot about the phenomenon!

Why Generative AI?

Human intelligence is intrinsically generative



Why Generative AI?

Human intelligence is intrinsically generative

GAI = AGI?

or

AGI through GAI

Generative AI

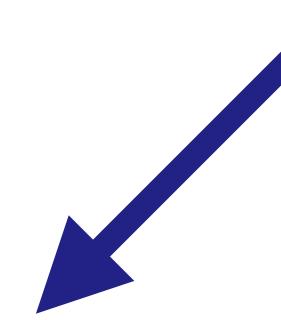
$$p(\mathbf{x}, y) = p(y | \mathbf{x}) \, p(\mathbf{x})$$

Generative AI

$$p(\mathbf{x}, y) = p(y | \mathbf{x}) p(\mathbf{x})$$

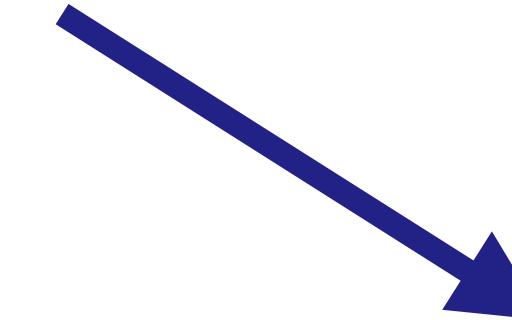
Any deep learning predictor

Relatively easy



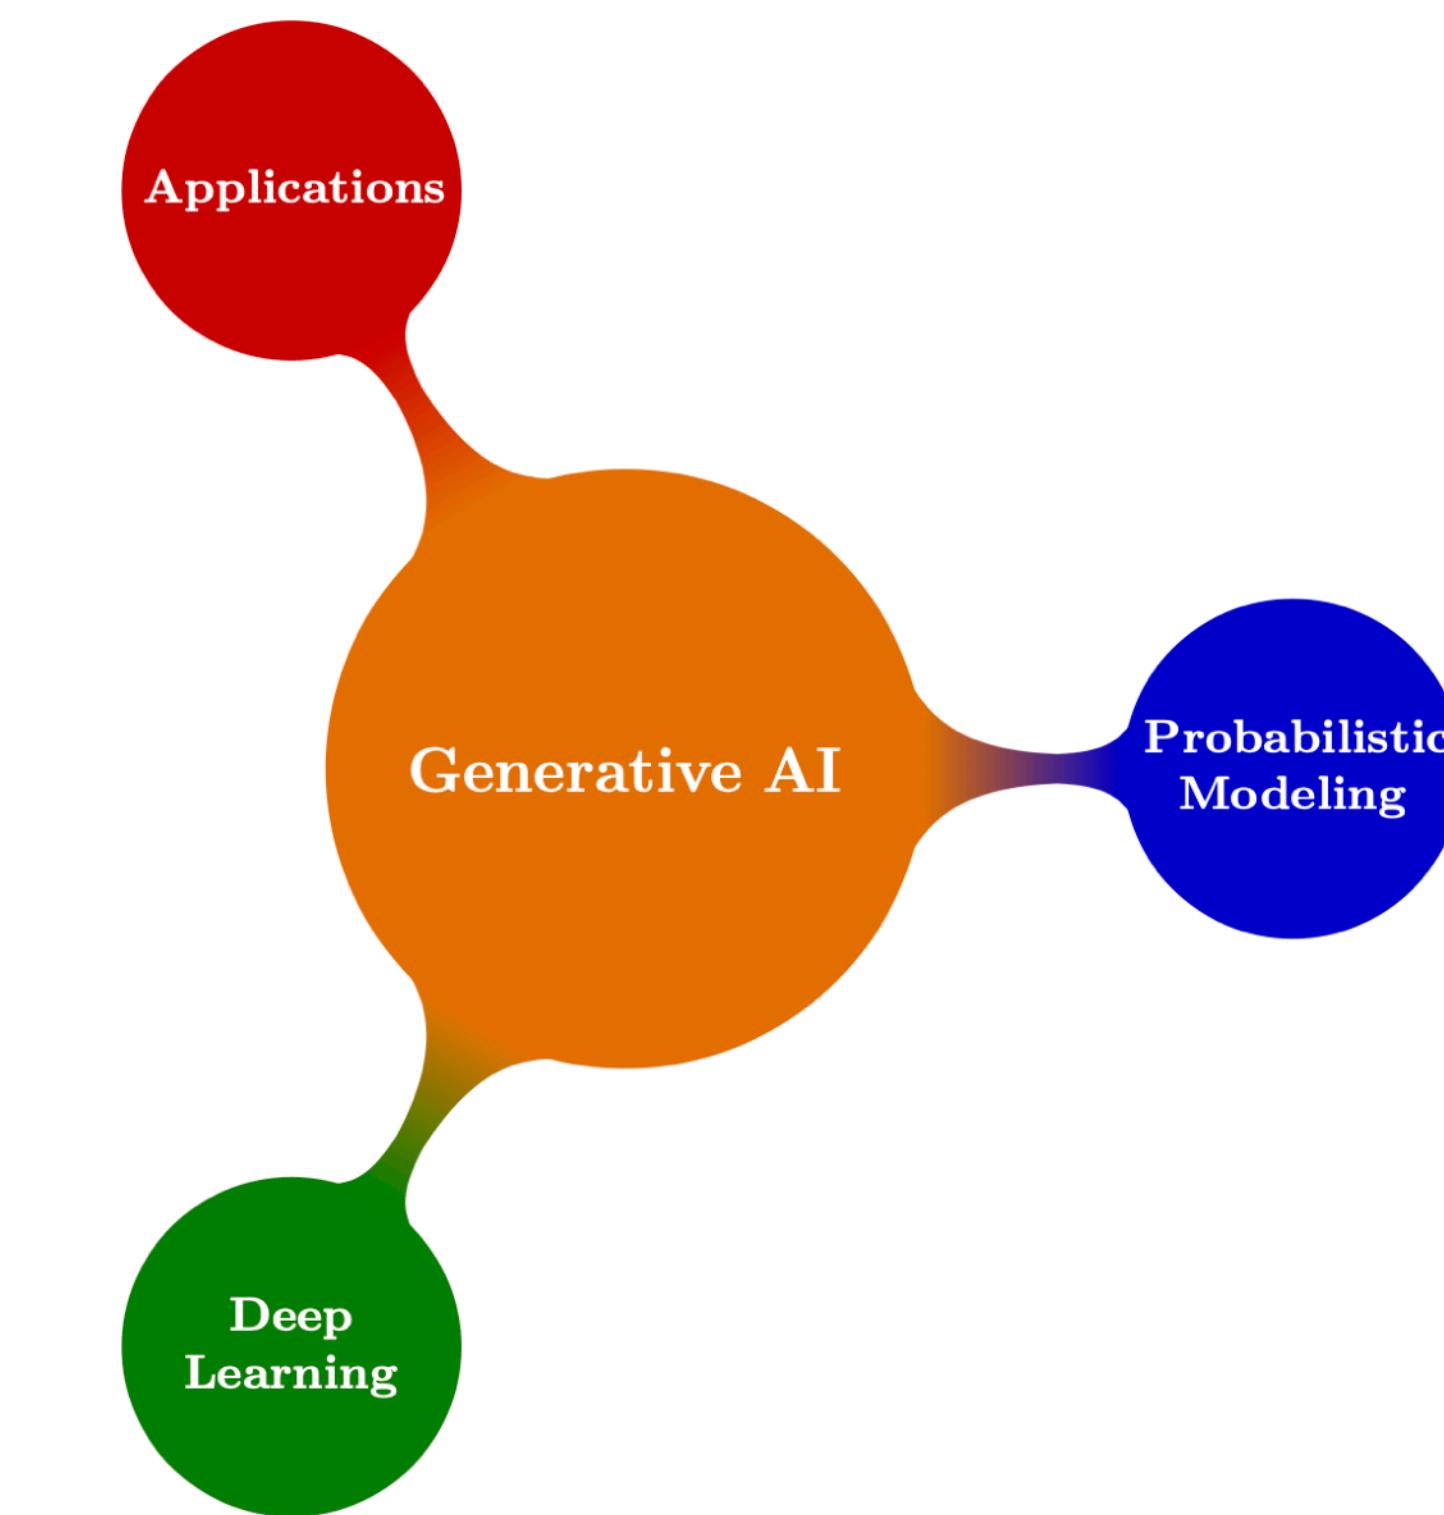
A density estimator

A probabilistic model

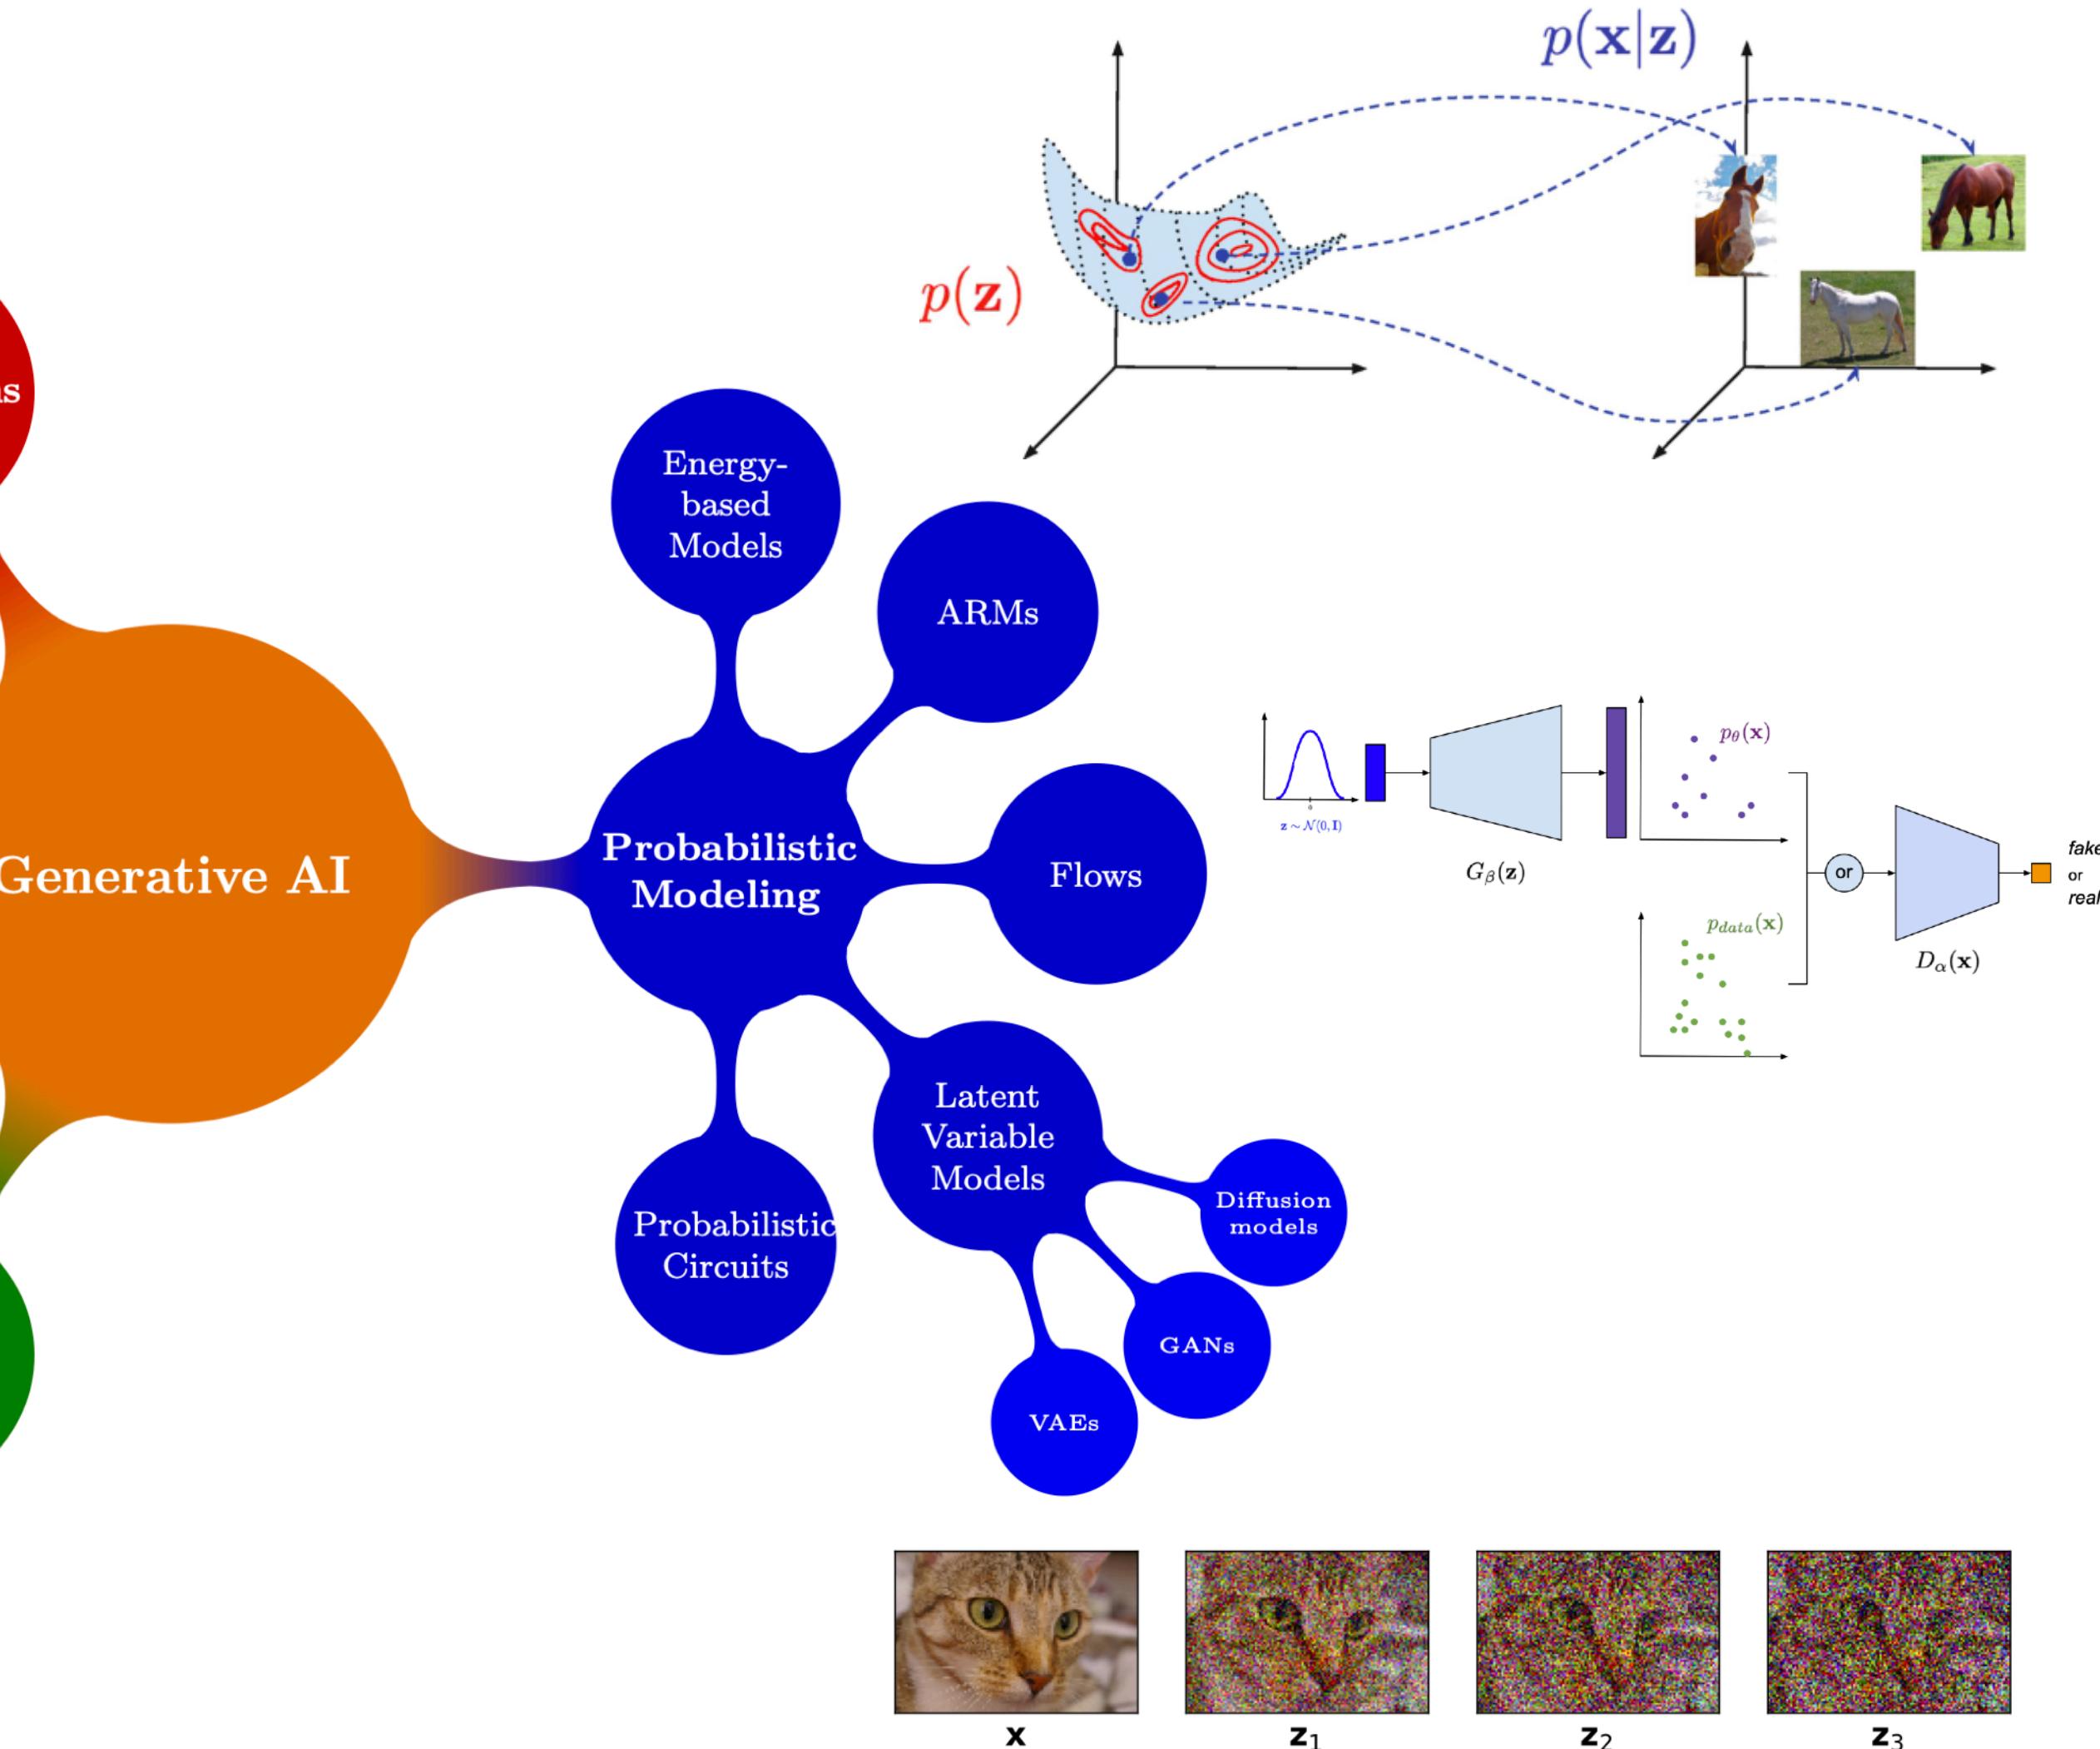


Challenging!

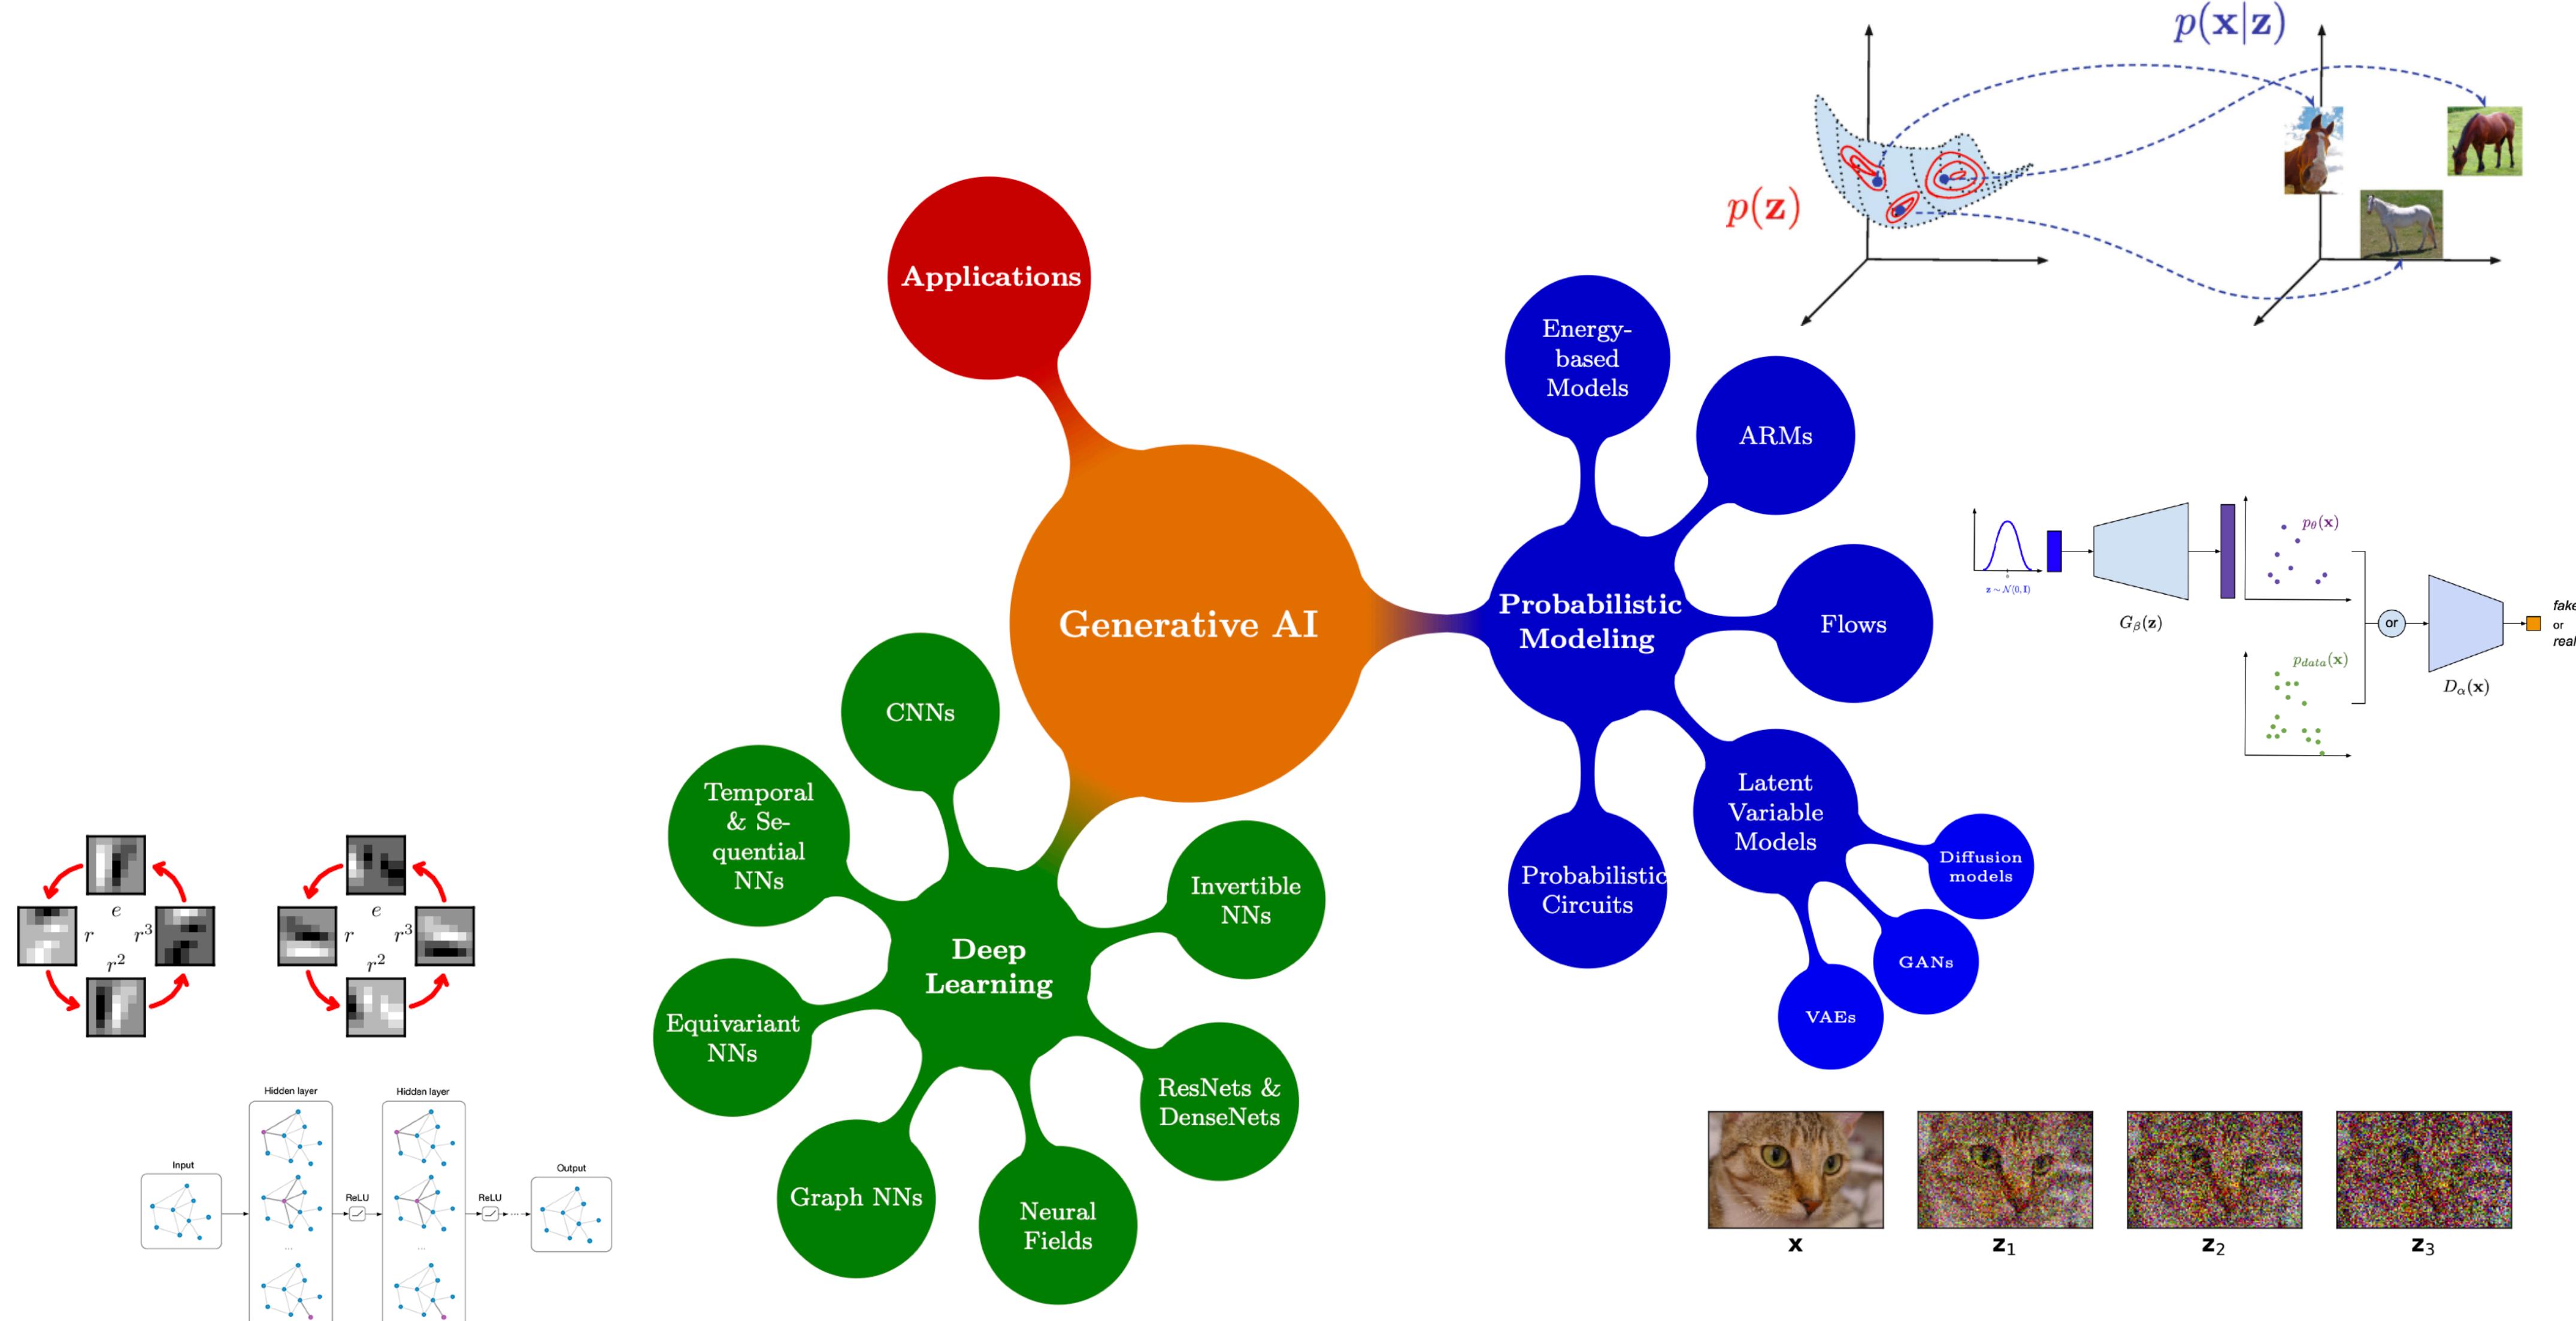
Overview of Generative AI



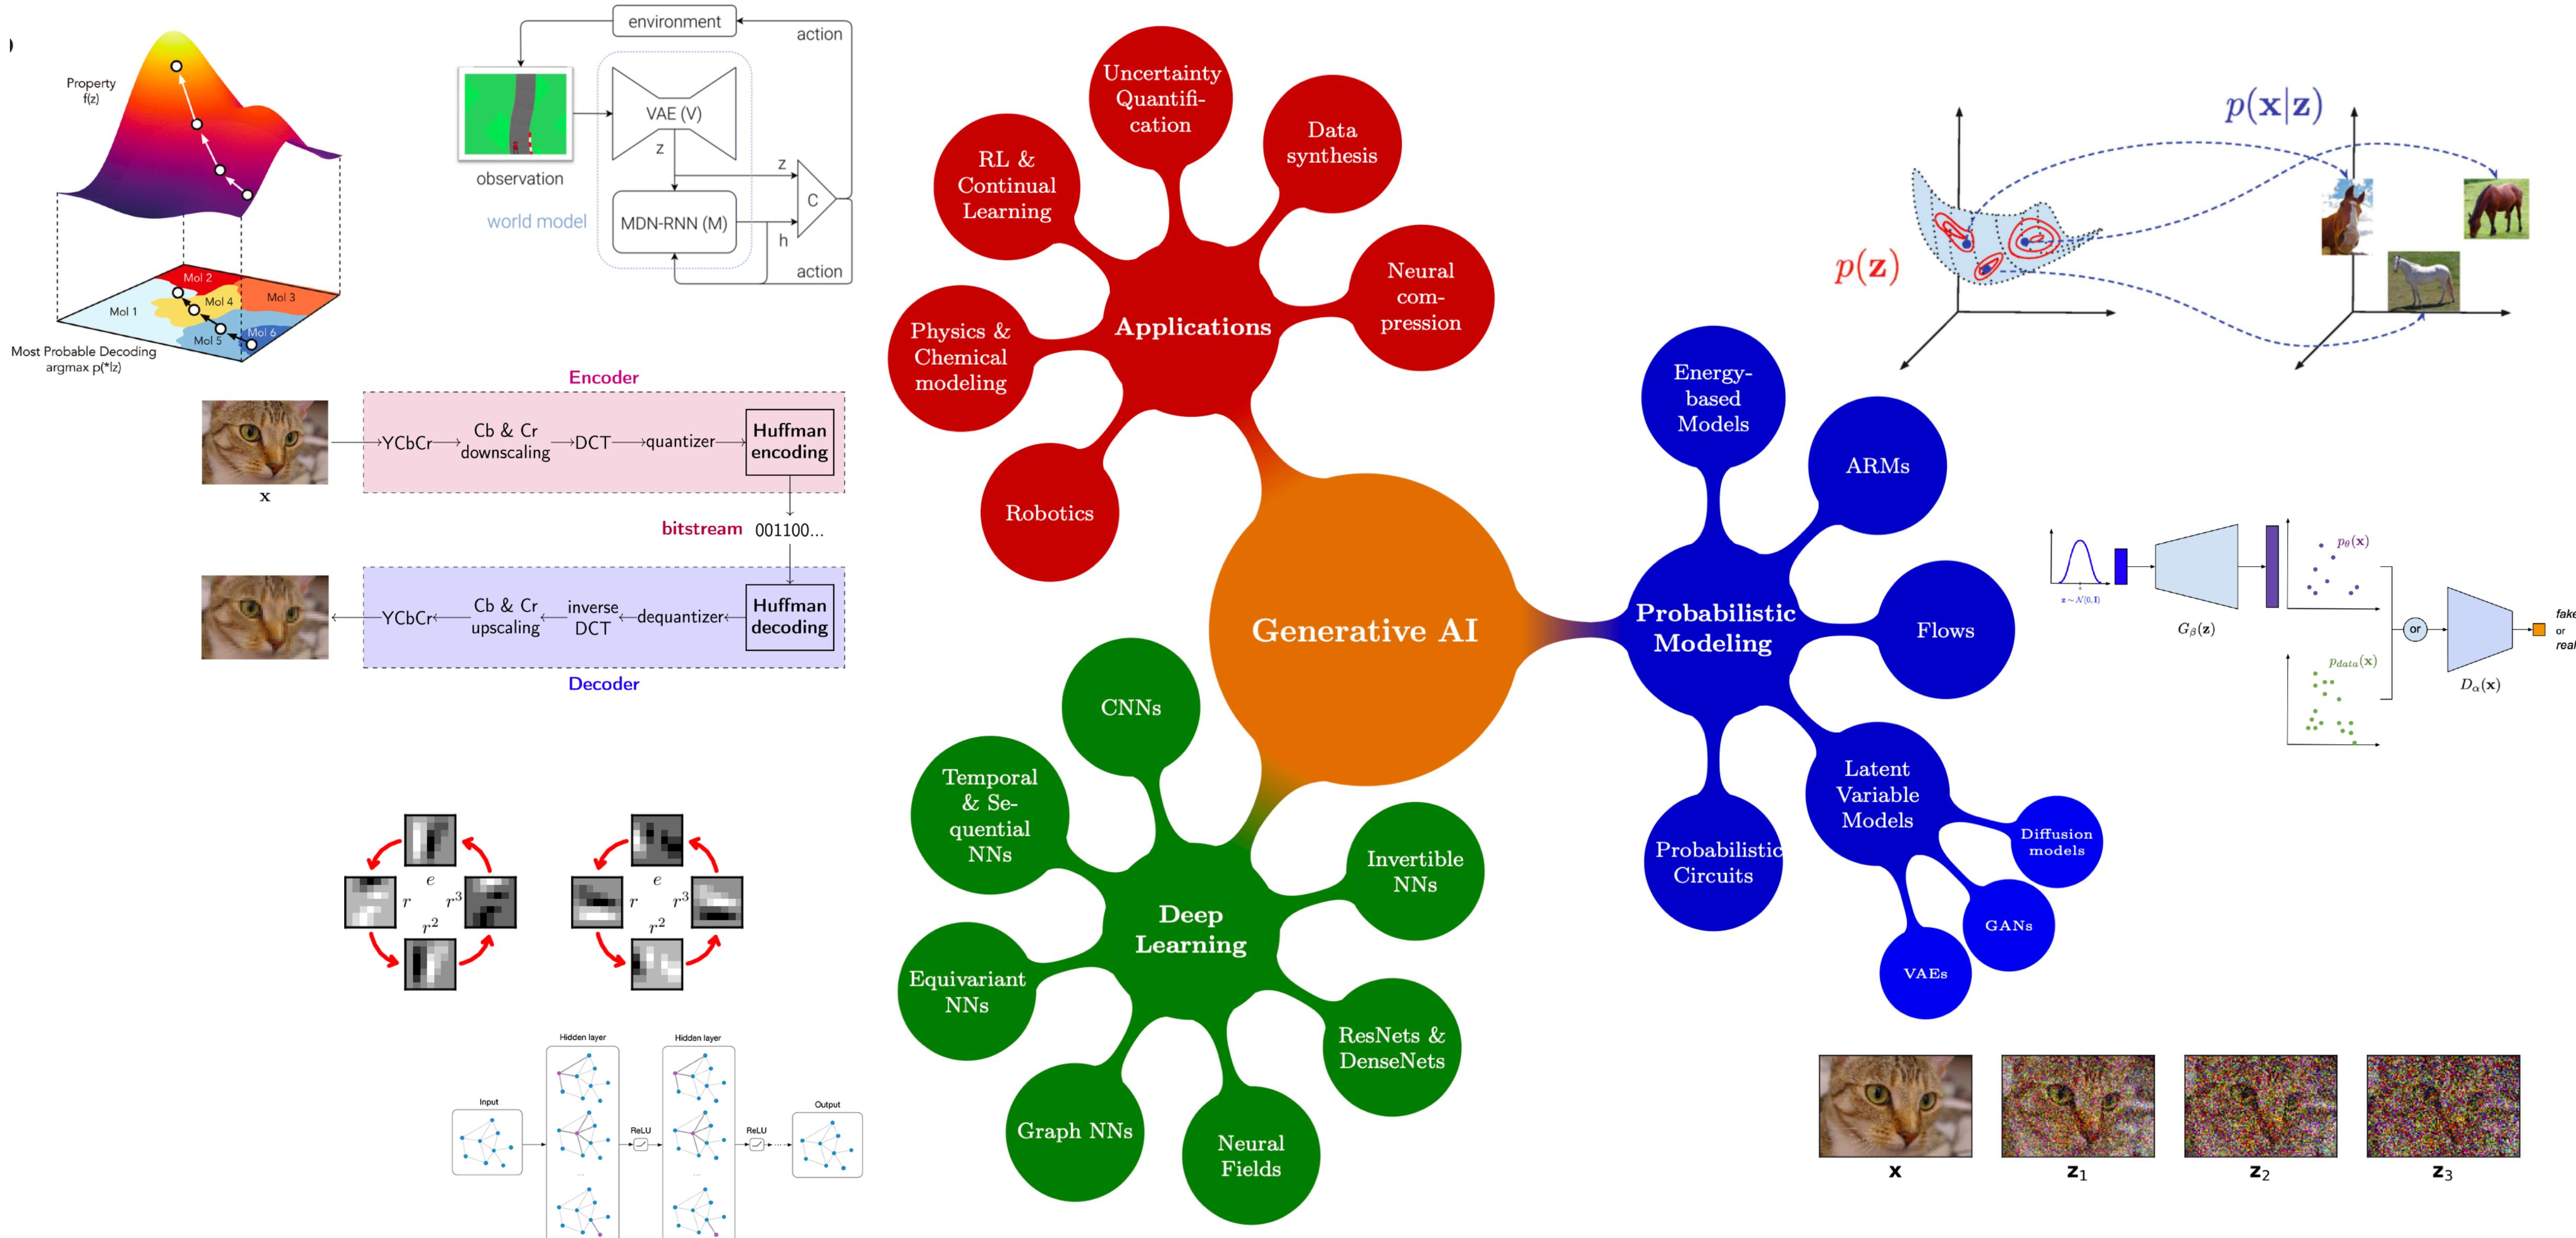
Overview of Generative AI



Overview of Generative AI



Overview of Generative AI



Generative AI and (spherical) cows



A high-dim object

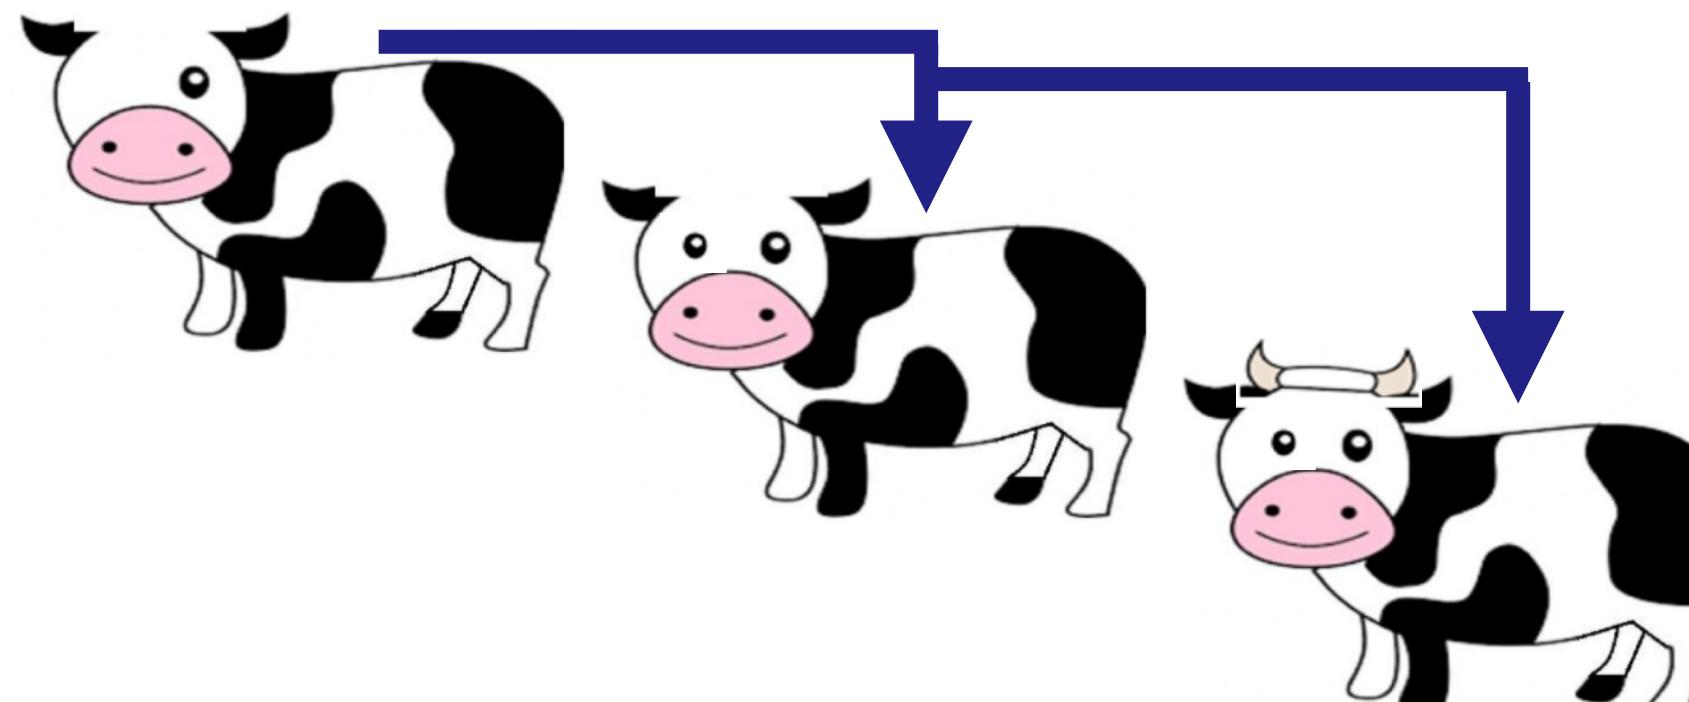


Latent Variable Models

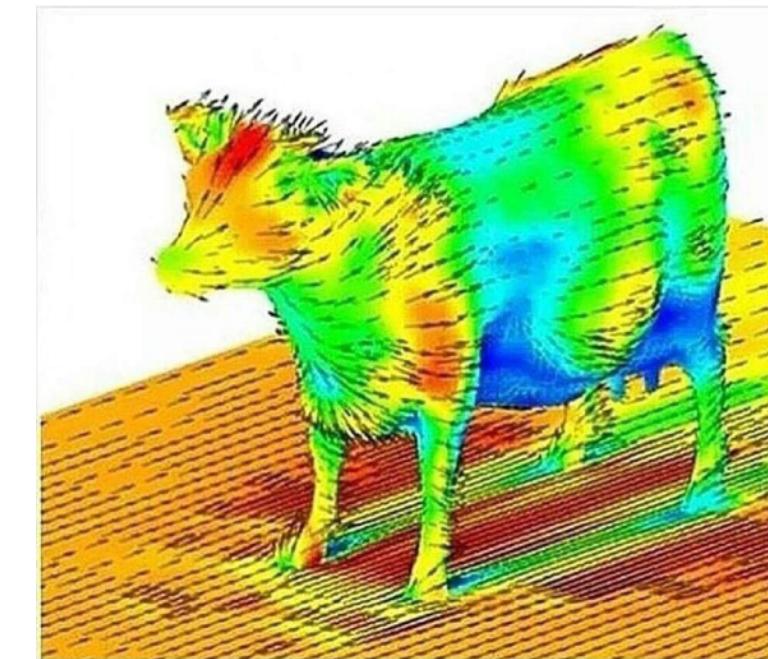


Flow-based models

Goal: $p(\mathbf{x})$



Autoregressive Models

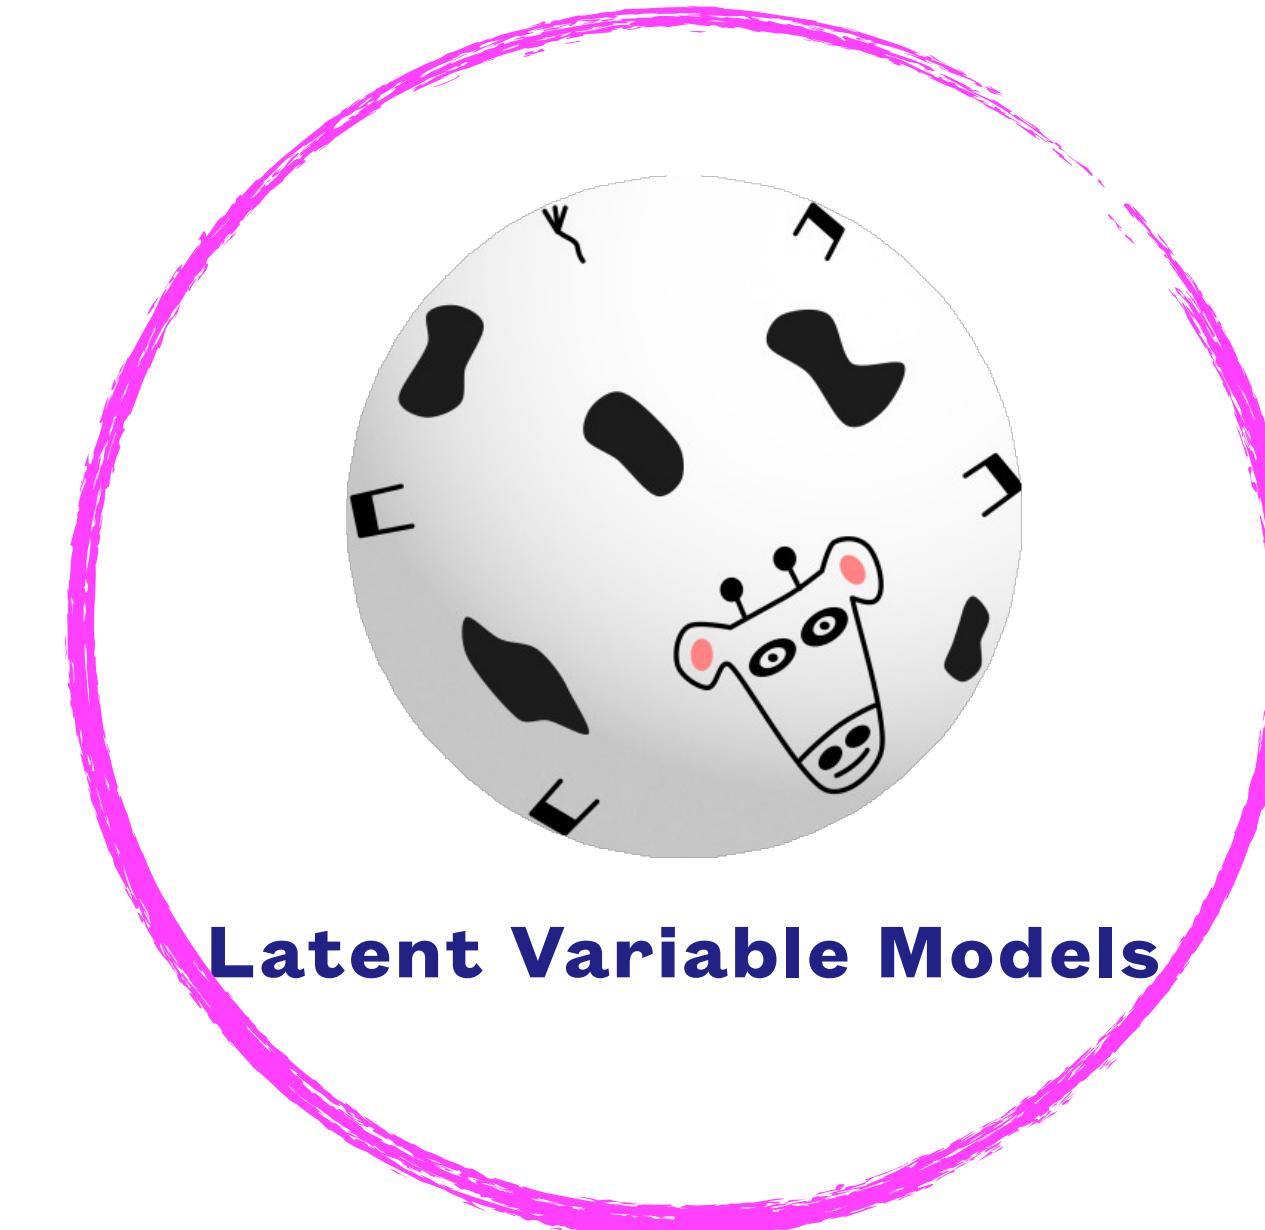


Diffusion models
Energy-based models

Generative AI and (spherical) cows



A high-dim object

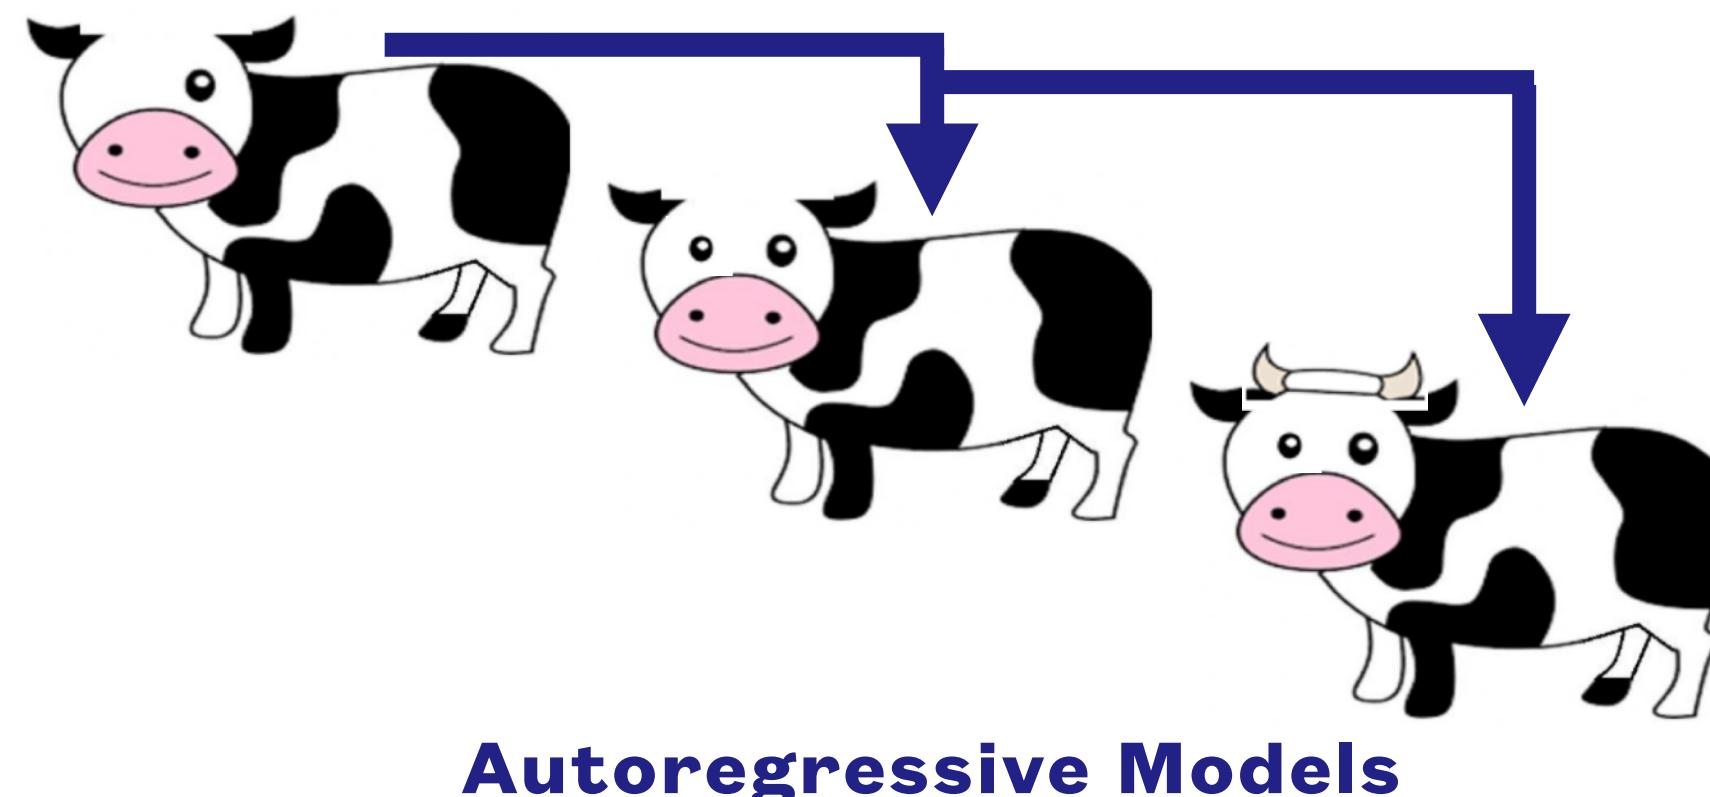


Latent Variable Models

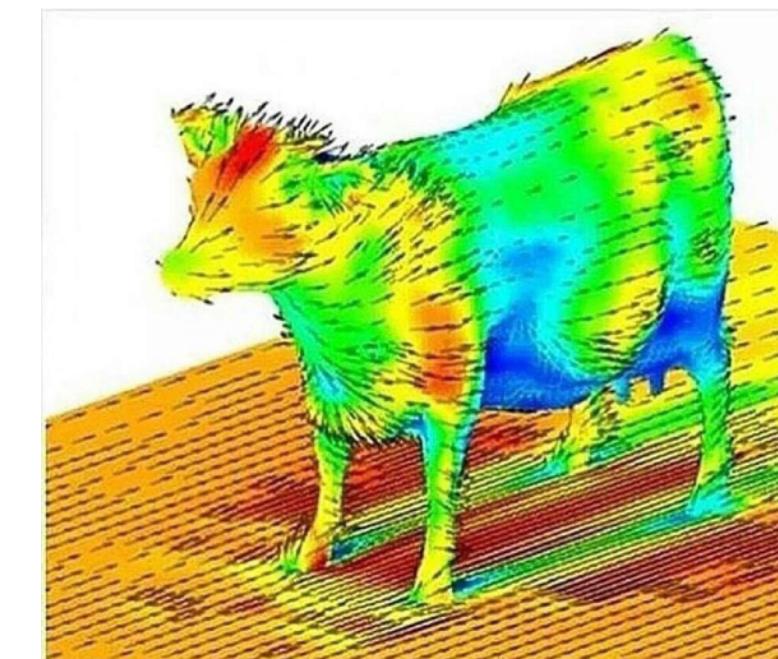


Flow-based models

Goal: $p(\mathbf{x})$



Autoregressive Models



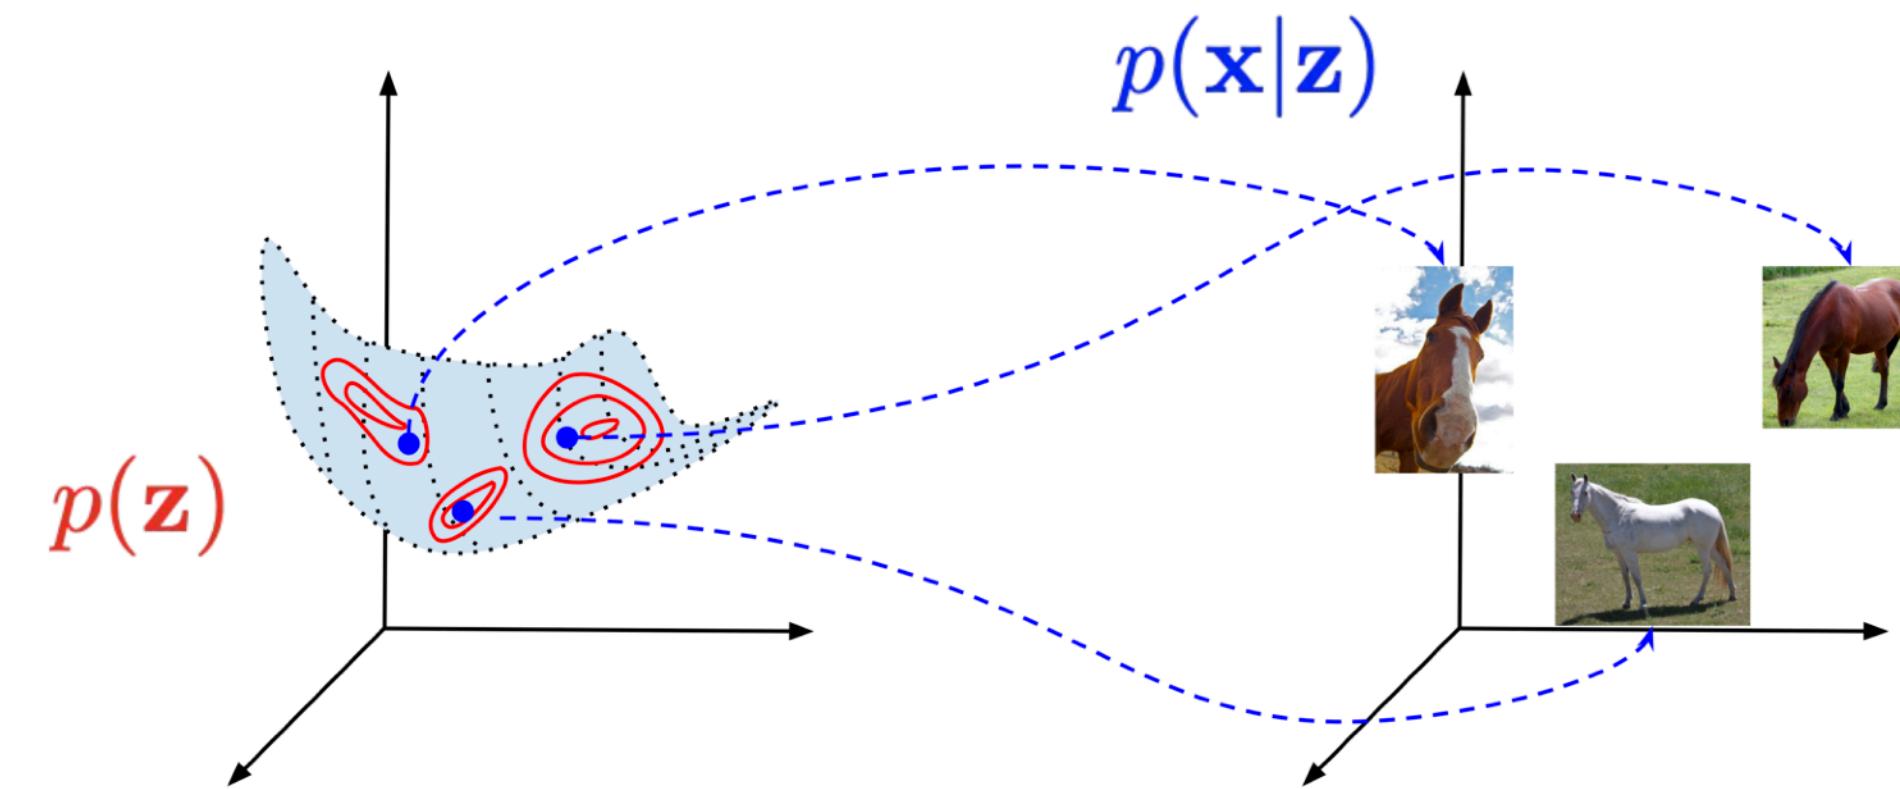
Diffusion models
Energy-based models

Variational Auto-encoders

Let's consider a latent variable model where we distinguish:

- latent variables $\mathbf{z} \in \mathcal{Z}^M$
- observable variables $\mathbf{x} \in \mathcal{X}^D$

Latent variables lie on a low-dimensional manifold.



Generative process:

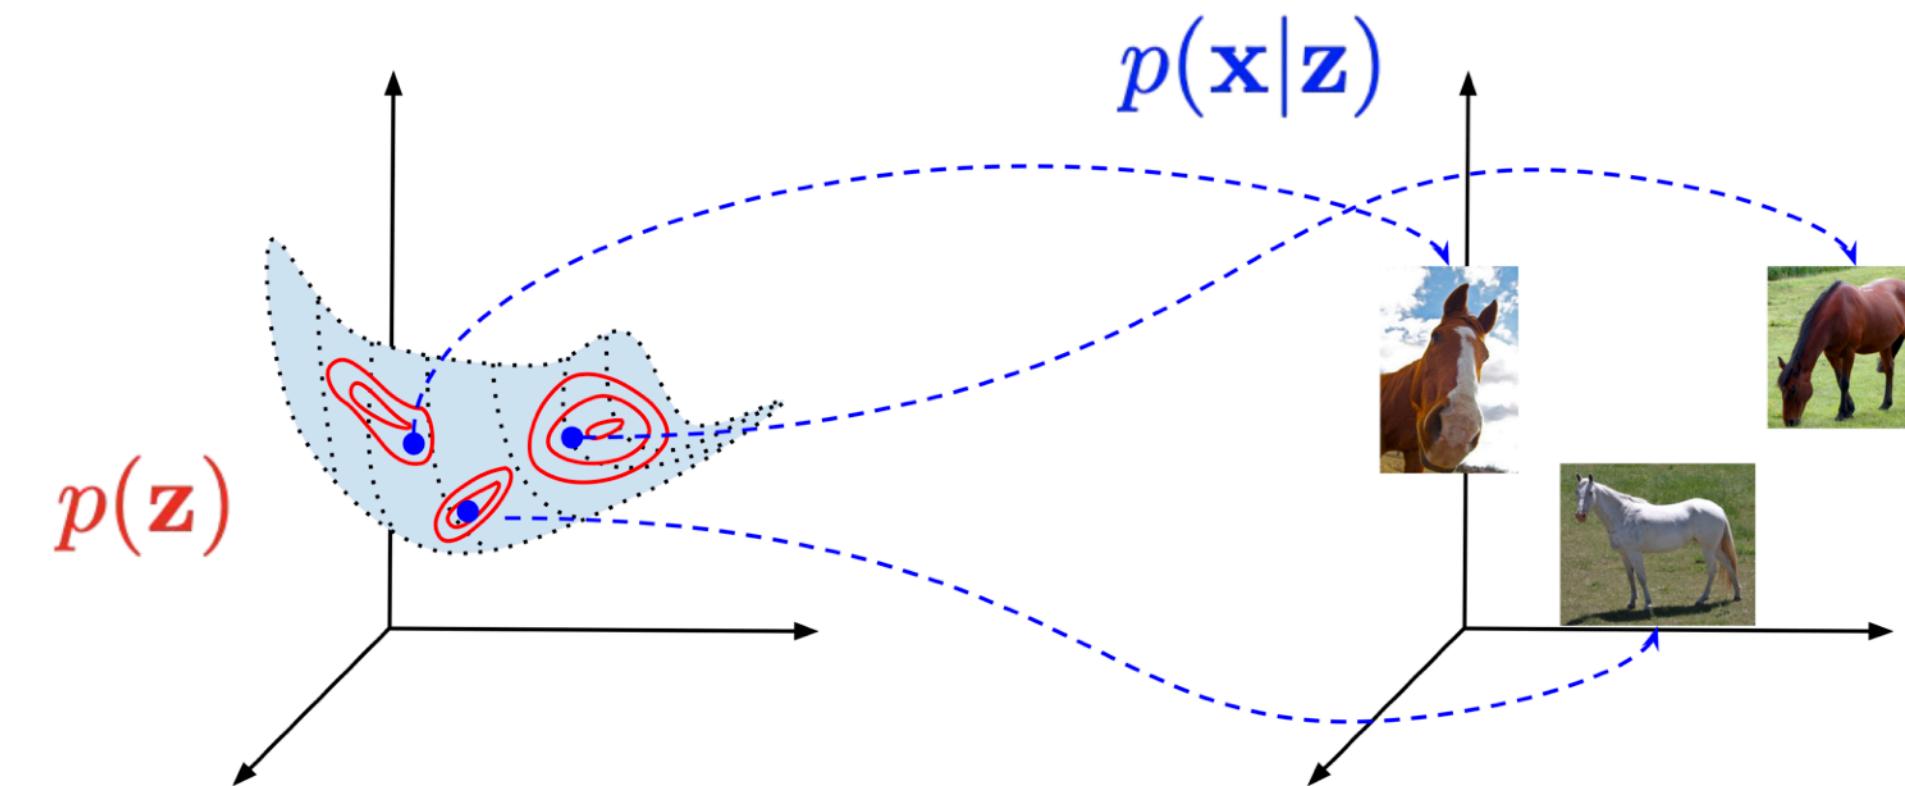
1. $\mathbf{z} \sim p(\mathbf{z})$
2. $\mathbf{x} \sim p(\mathbf{x} | \mathbf{z})$

Variational Auto-encoders

Let's consider a latent variable model where we distinguish:

- latent variables $\mathbf{z} \in \mathcal{Z}^M$
- observable variables $\mathbf{x} \in \mathcal{X}^D$

Latent variables lie on a low-dimensional manifold.



Generative process:

1. $\mathbf{z} \sim p(\mathbf{z})$
2. $\mathbf{x} \sim p(\mathbf{x} | \mathbf{z})$

The objective function:

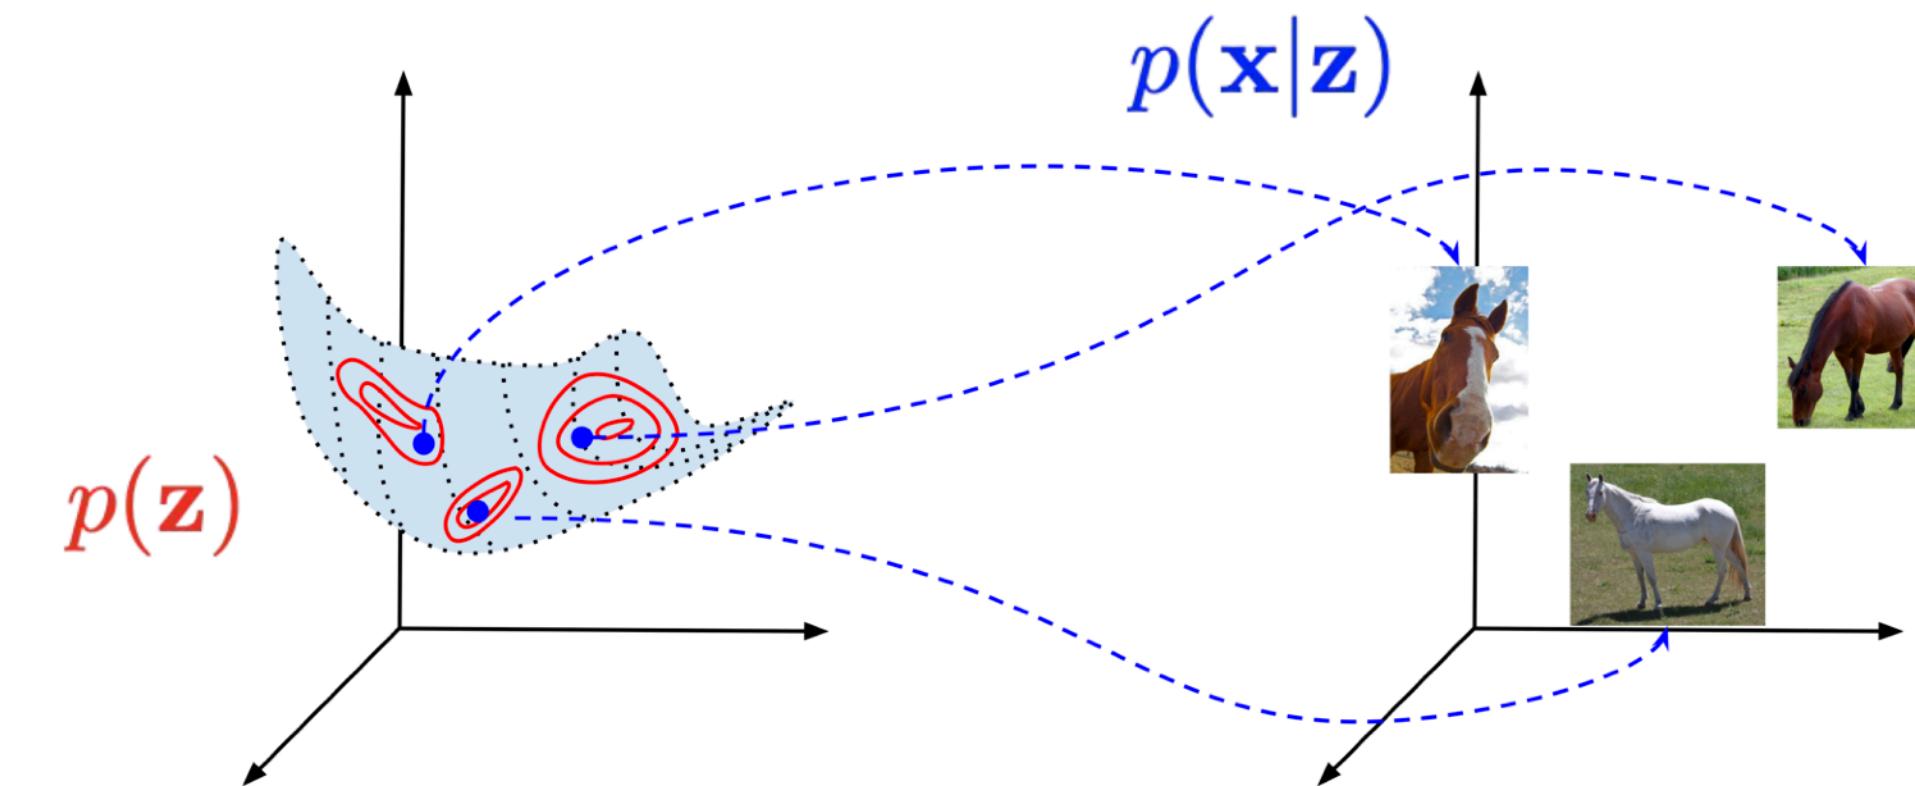
$$\ln p(\mathbf{x}) = \ln \int p(\mathbf{x} | \mathbf{z}) p(\mathbf{z}) \, d\mathbf{z}$$

Variational Auto-encoders

Let's consider a latent variable model where we distinguish:

- latent variables $\mathbf{z} \in \mathcal{Z}^M$
- observable variables $\mathbf{x} \in \mathcal{X}^D$

Latent variables lie on a low-dimensional manifold.



Generative process:

1. $\mathbf{z} \sim p(\mathbf{z})$
2. $\mathbf{x} \sim p(\mathbf{x} | \mathbf{z})$

The objective function:

$$\ln p(\mathbf{x}) = \ln \int p(\mathbf{x} | \mathbf{z}) p(\mathbf{z}) \, d\mathbf{z}$$

The integral is intractable...

Variational Auto-encoders

$$\ln p(\mathbf{x}) \geq \mathbb{E}_{\mathbf{z} \sim q_\phi(\mathbf{z}|\mathbf{x})} [\ln p(\mathbf{x}|\mathbf{z})] - \mathbb{E}_{\mathbf{z} \sim q_\phi(\mathbf{z}|\mathbf{x})} [\ln q_\phi(\mathbf{z}|\mathbf{x}) - \ln p(\mathbf{z})]$$

ELBO: Evidence Lower Bound

Variational Auto-encoders

$$\ln p(\mathbf{x}) \geq \mathbb{E}_{\mathbf{z} \sim q_\phi(\mathbf{z} | \mathbf{x})} [\ln p(\mathbf{x} | \mathbf{z})] - \mathbb{E}_{\mathbf{z} \sim q_\phi(\mathbf{z} | \mathbf{x})} [\ln q_\phi(\mathbf{z} | \mathbf{x}) - \ln p(\mathbf{z})]$$

We consider amortized inference: $q_\phi(\mathbf{z} | \mathbf{x})$

In other words, a single parameterization for each new input \mathbf{x} .

Variational Auto-encoders

$$\ln p(\mathbf{x}) \geq \mathbb{E}_{\mathbf{z} \sim q_\phi(\mathbf{z} | \mathbf{x})} [\ln p(\mathbf{x} | \mathbf{z})] - \mathbb{E}_{\mathbf{z} \sim q_\phi(\mathbf{z} | \mathbf{x})} [\ln q_\phi(\mathbf{z} | \mathbf{x}) - \ln p(\mathbf{z})]$$

We consider amortized inference: $q_\phi(\mathbf{z} | \mathbf{x})$

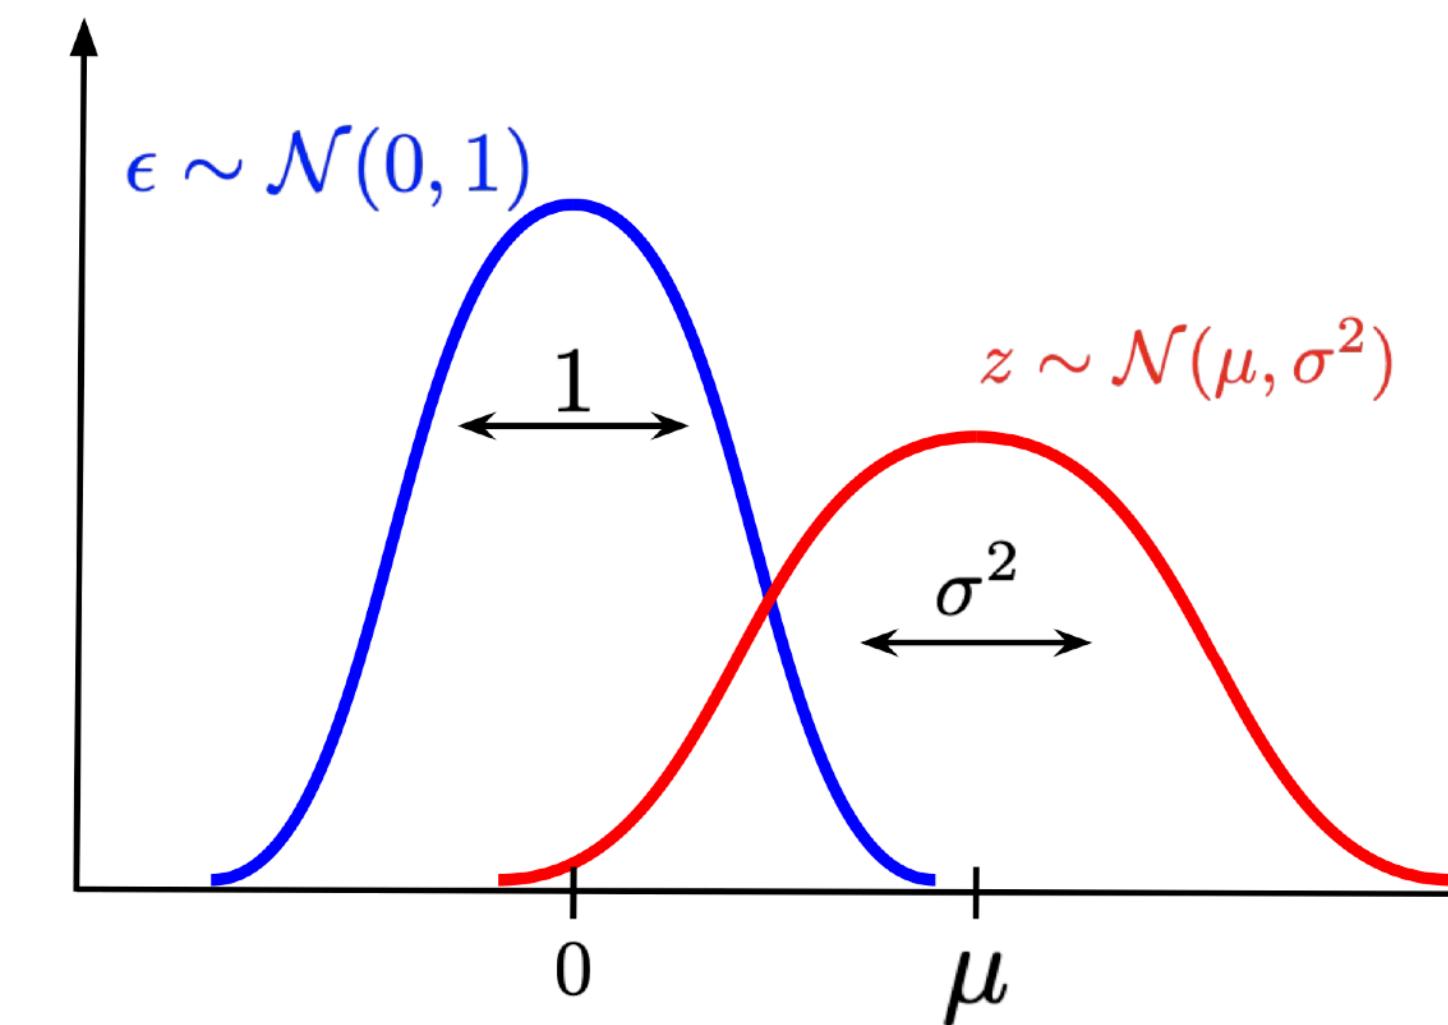
In other words, a single parameterization for each new input \mathbf{x} .

Moreover, we use the **reparameterization trick**.

Every Gaussian variable could be defined as:

$$z = \mu + \sigma \cdot \epsilon$$

where $\epsilon \sim \mathcal{N}(0, 1)$



Variational Auto-encoders

$$\ln p(\mathbf{x}) \geq \mathbb{E}_{\mathbf{z} \sim q_\phi(\mathbf{z} | \mathbf{x})} [\ln p(\mathbf{x} | \mathbf{z})] - \mathbb{E}_{\mathbf{z} \sim q_\phi(\mathbf{z} | \mathbf{x})} [\ln q_\phi(\mathbf{z} | \mathbf{x}) - \ln p(\mathbf{z})]$$

We consider amortized inference: $q_\phi(\mathbf{z} | \mathbf{x})$

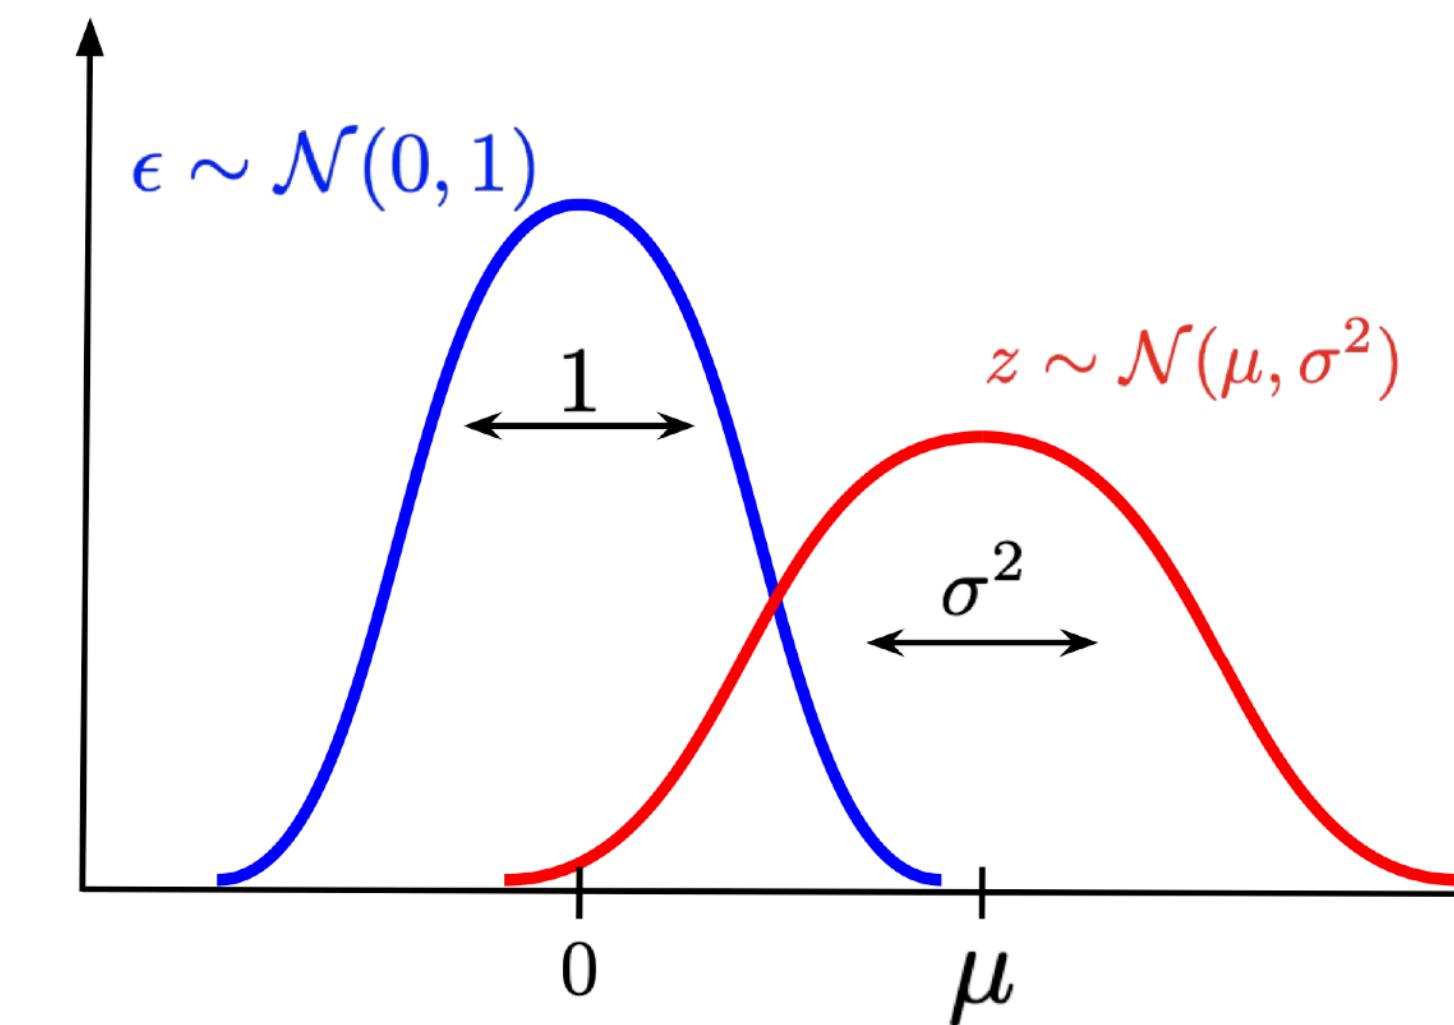
In other words, a single parameterization for each new input \mathbf{x} .

Moreover, we use the **reparameterization trick**.

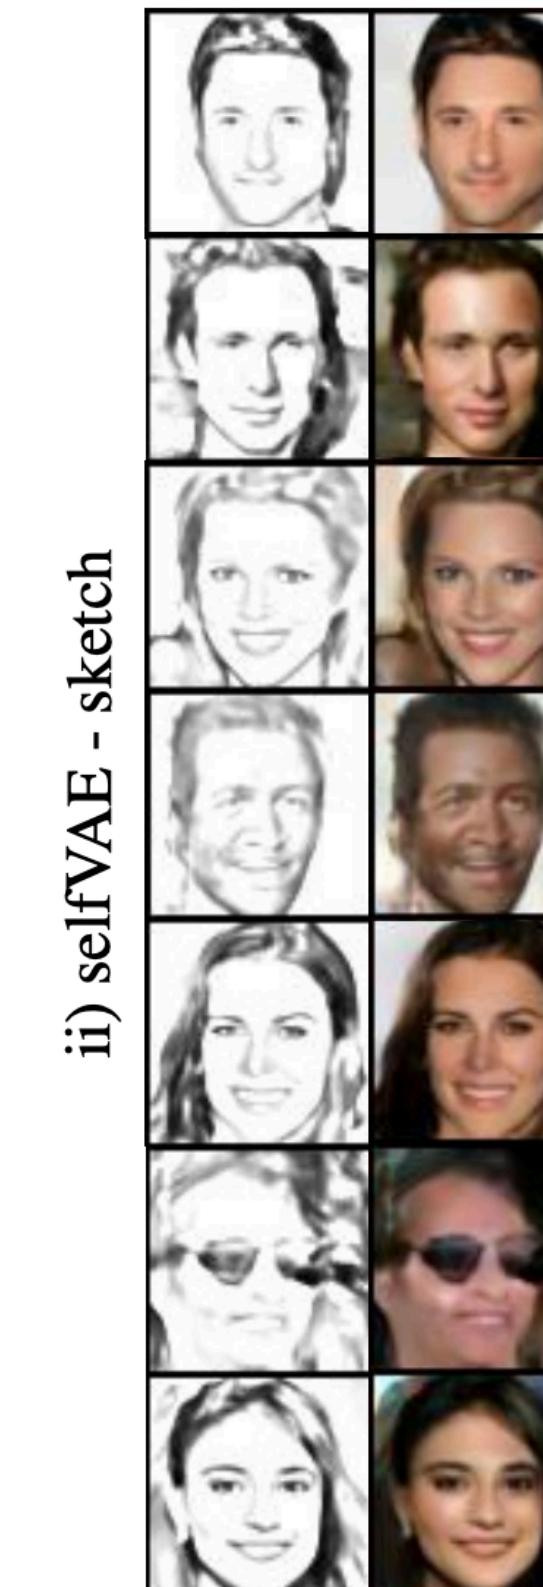
It reduces the variance of the gradients.

It allows to get randomness outside \mathbf{z} .

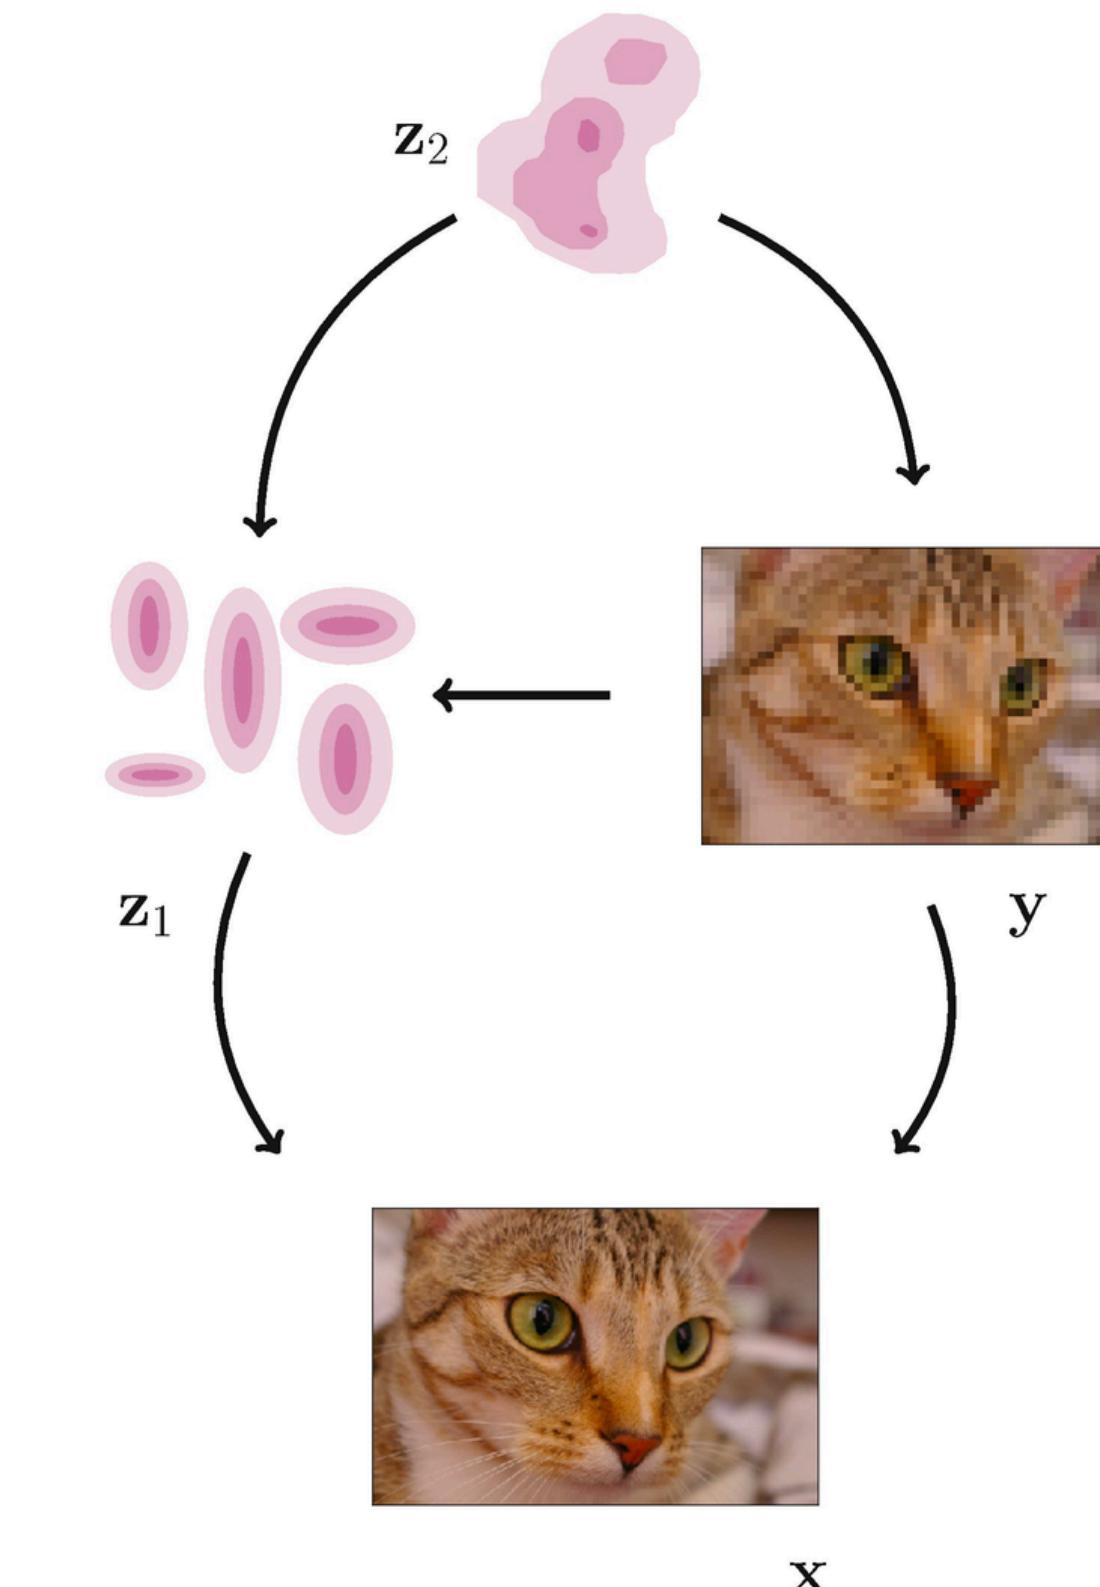
$$z = \mu + \sigma \cdot \epsilon$$



Hierarchical Variational Auto-encoders



Generations

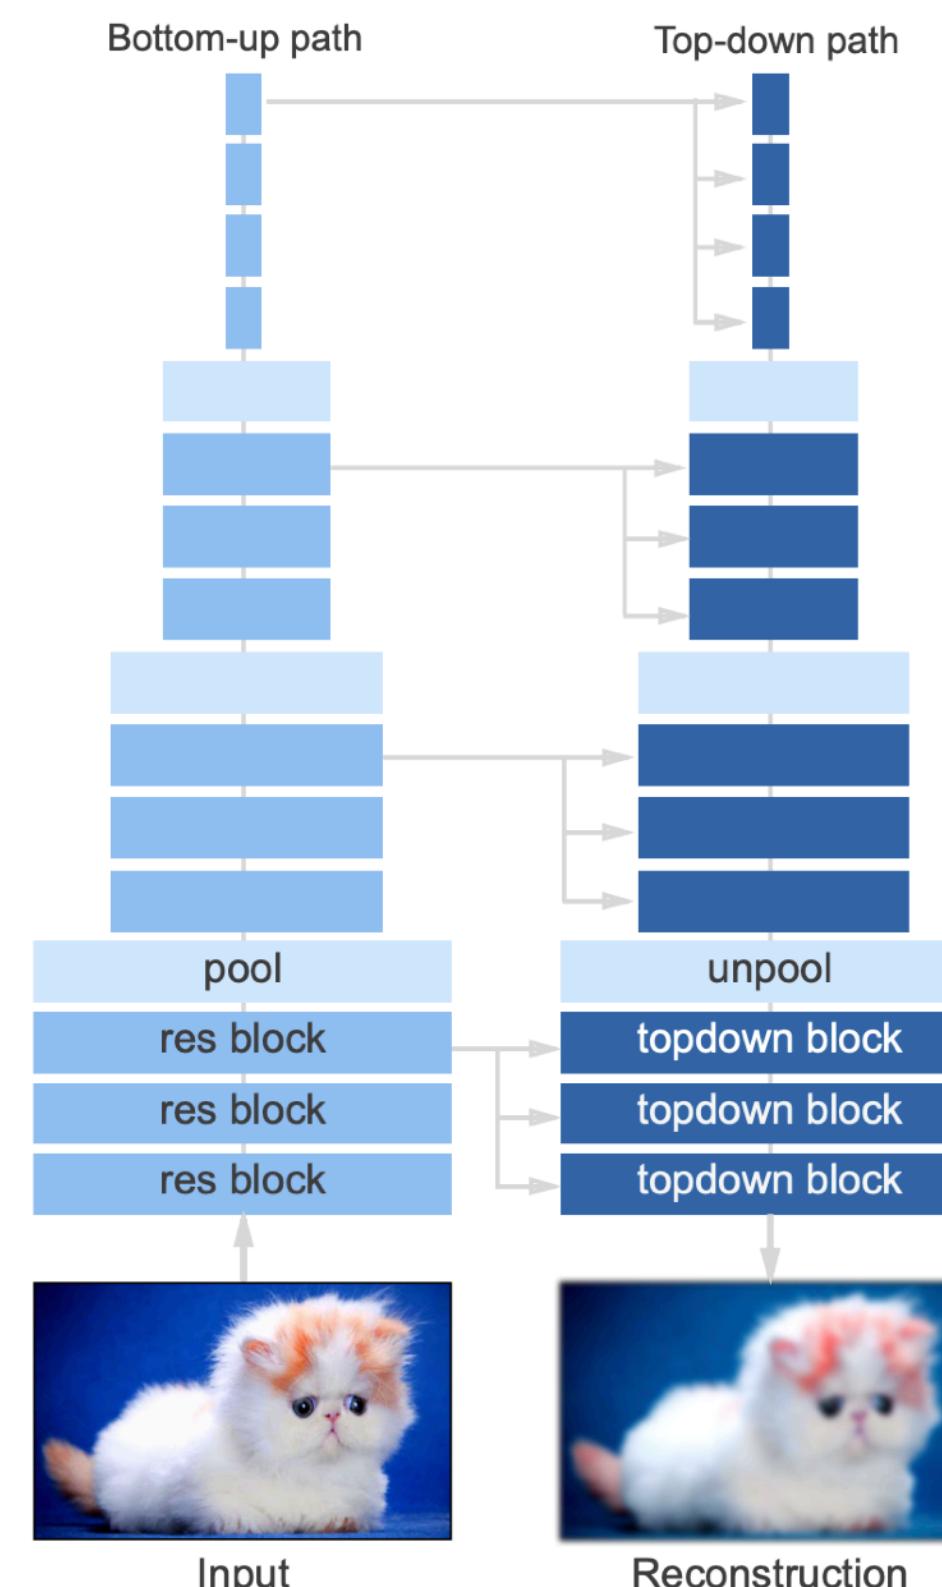


Hierarchical VAE

Top-down Variational Auto-encoders



Generations

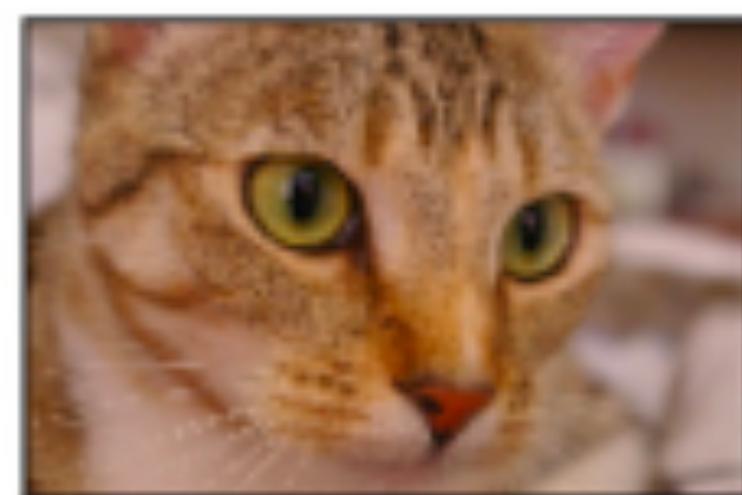


Hierarchical VAE

Diffusion-based deep generative models

$$q_{\phi}(\mathbf{z}_i \mid \mathbf{z}_{i-1}) = \mathcal{N}\left(\mathbf{z}_i \mid \sqrt{1 - \beta_i} \mathbf{z}_{i-1}, \beta_i \mathbf{I}\right)$$

Forward diffusion (FIXED!)



x



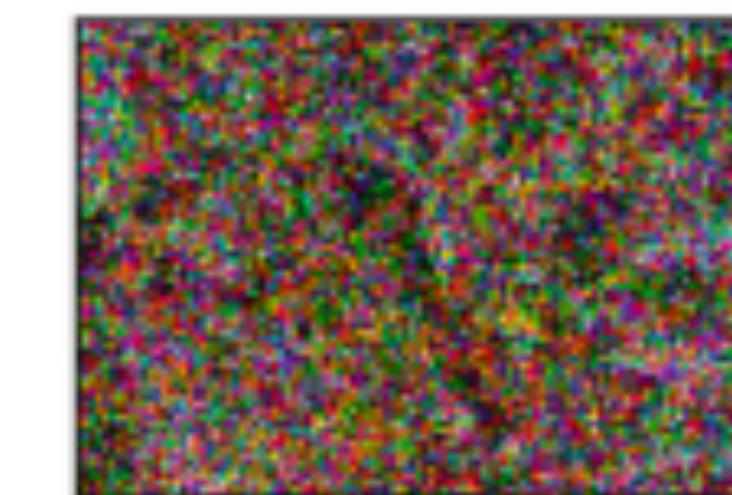
z₁



z₂



z₃

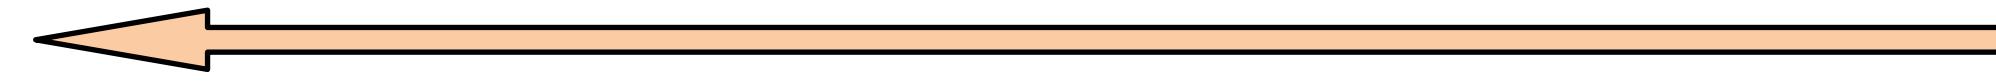


z₄



z₅

Backward diffusion (learnable)



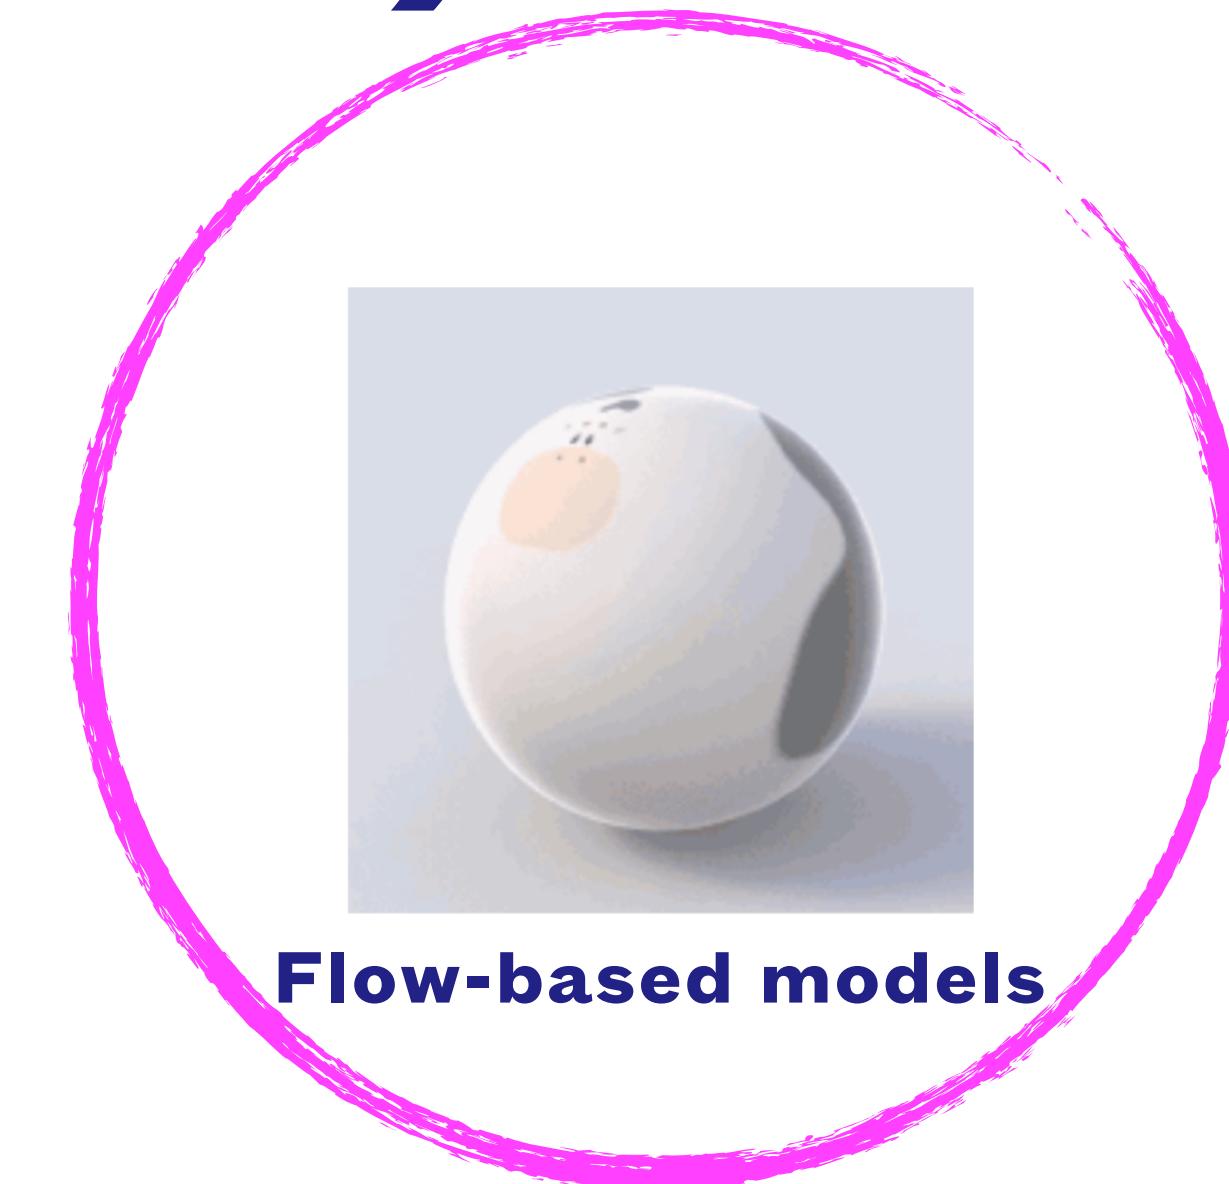
Generative AI and (spherical) cows



A high-dim object

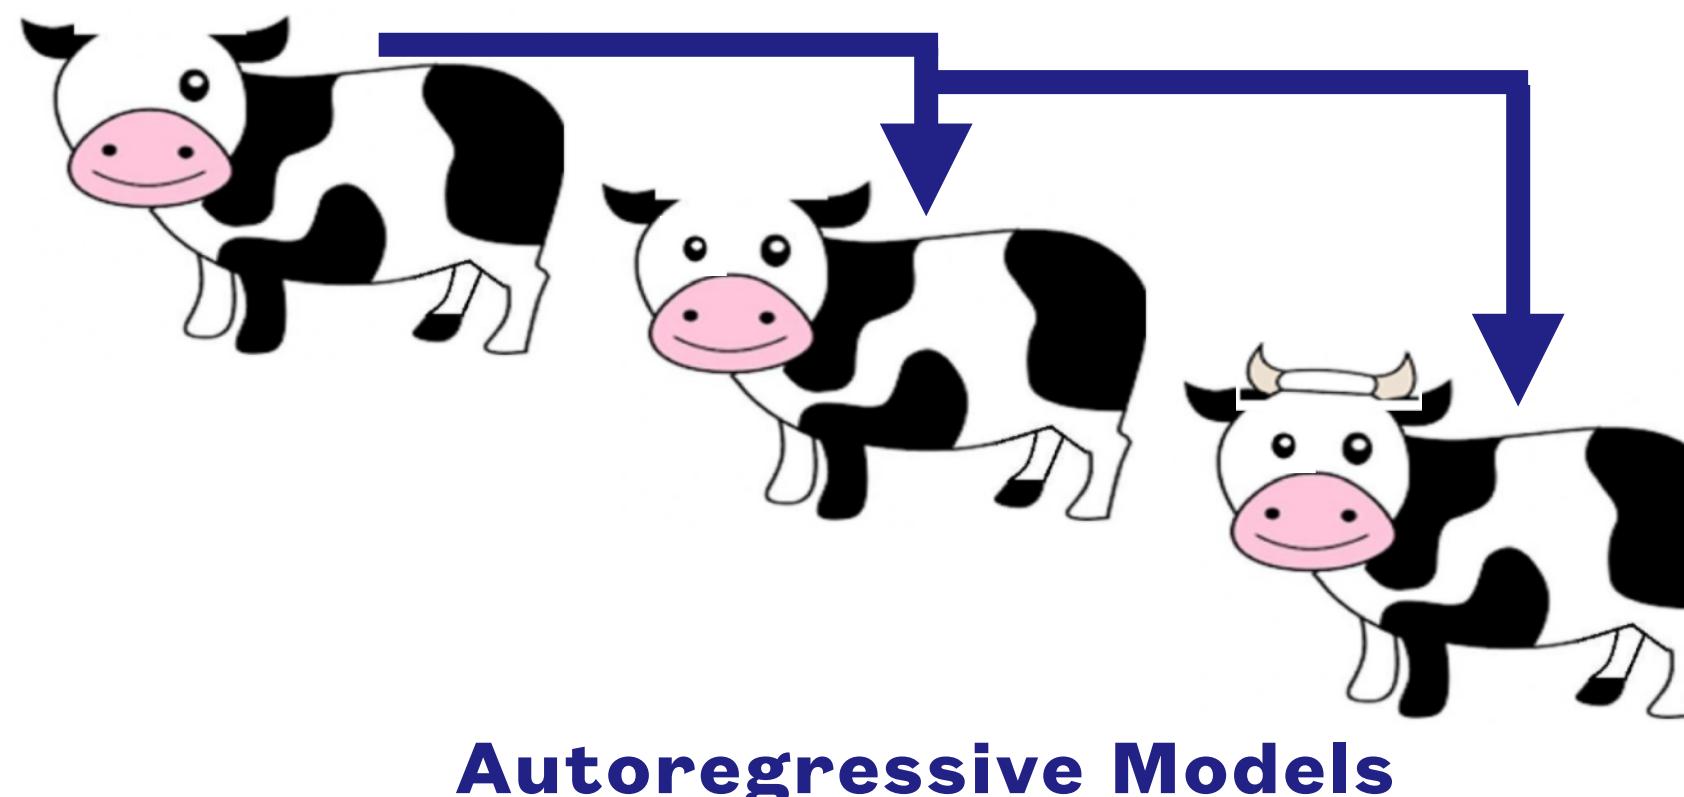


Latent Variable Models

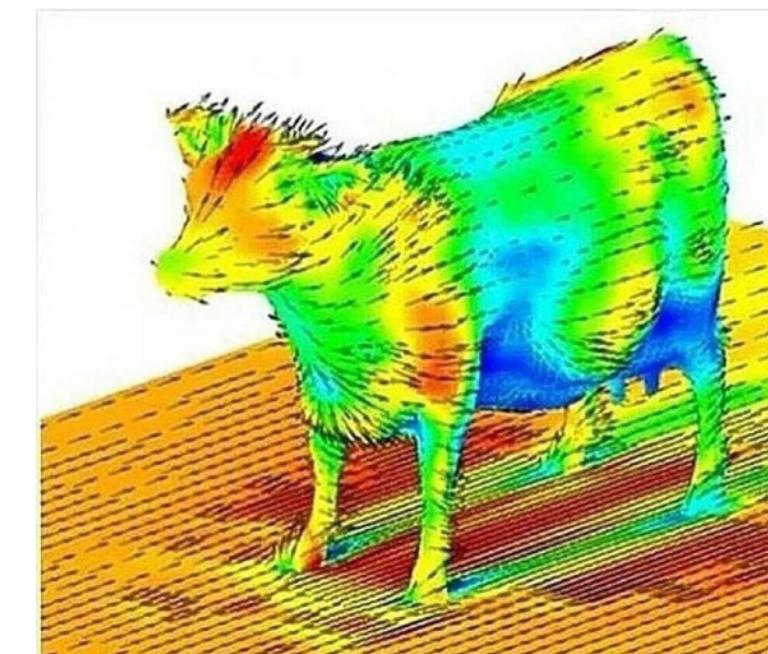


Flow-based models

Goal: $p(\mathbf{x})$



Autoregressive Models



Diffusion models
Energy-based models

Flows (Flow-based models)

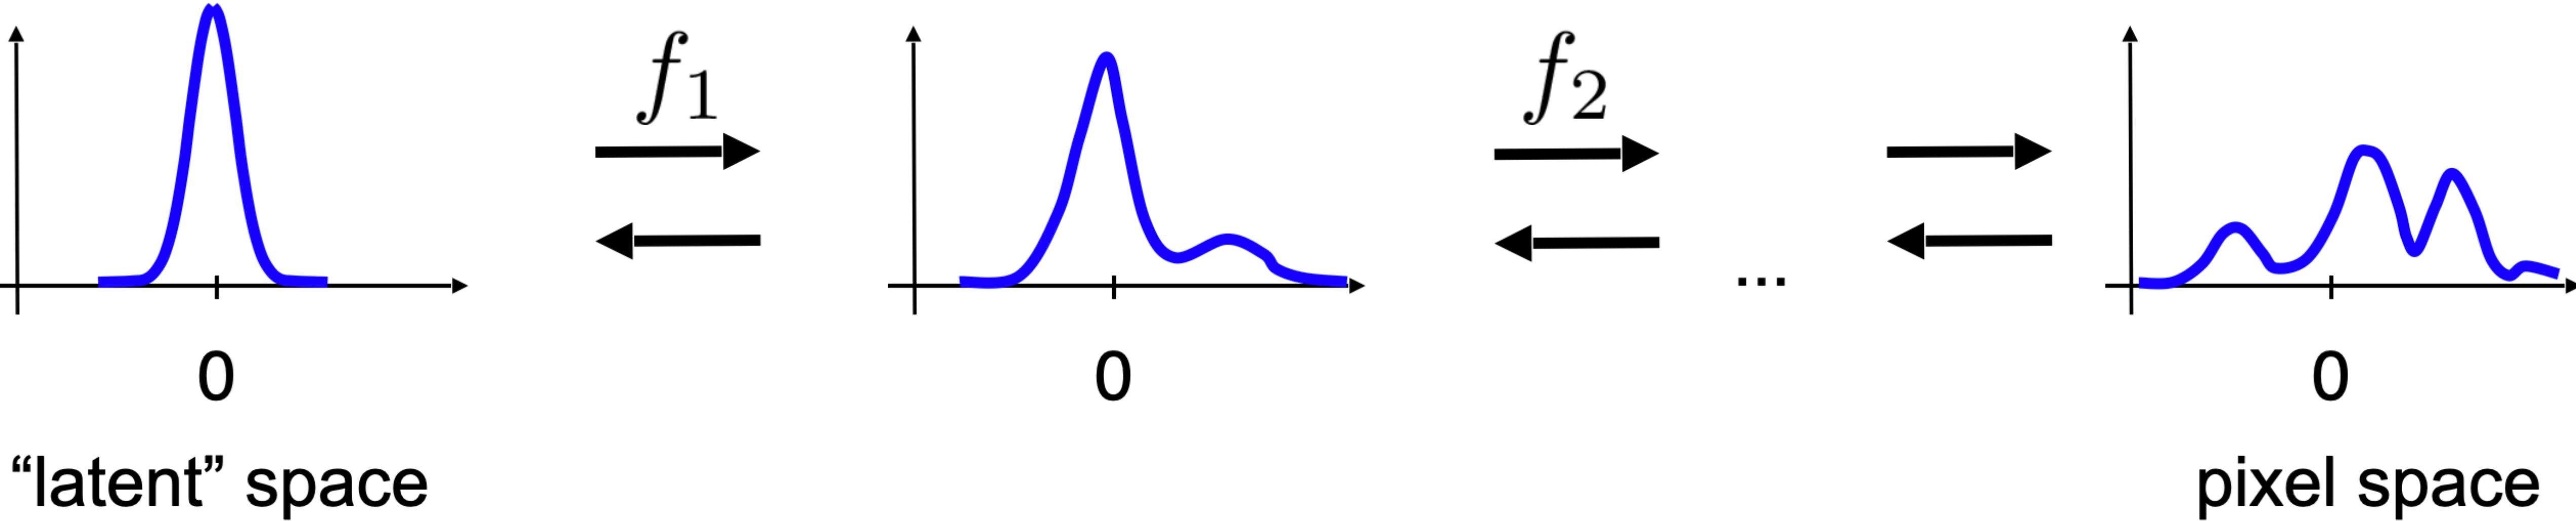
- We change a random variable \mathbf{x} to another random variable \mathbf{z} using invertible transformations, $\mathbf{x}, \mathbf{z} \in \mathbb{R}^D$:

$$p(\mathbf{x}) = \pi(\mathbf{z}_0 = f^{-1}(\mathbf{x})) \prod_{i=1}^K \left| \mathbf{J}_{f_i}(z_{i-1}) \right|^{-1}$$

Flows (Flow-based models)

- We change a random variable \mathbf{x} to another random variable \mathbf{z} using invertible transformations, $\mathbf{x}, \mathbf{z} \in \mathbb{R}^D$:

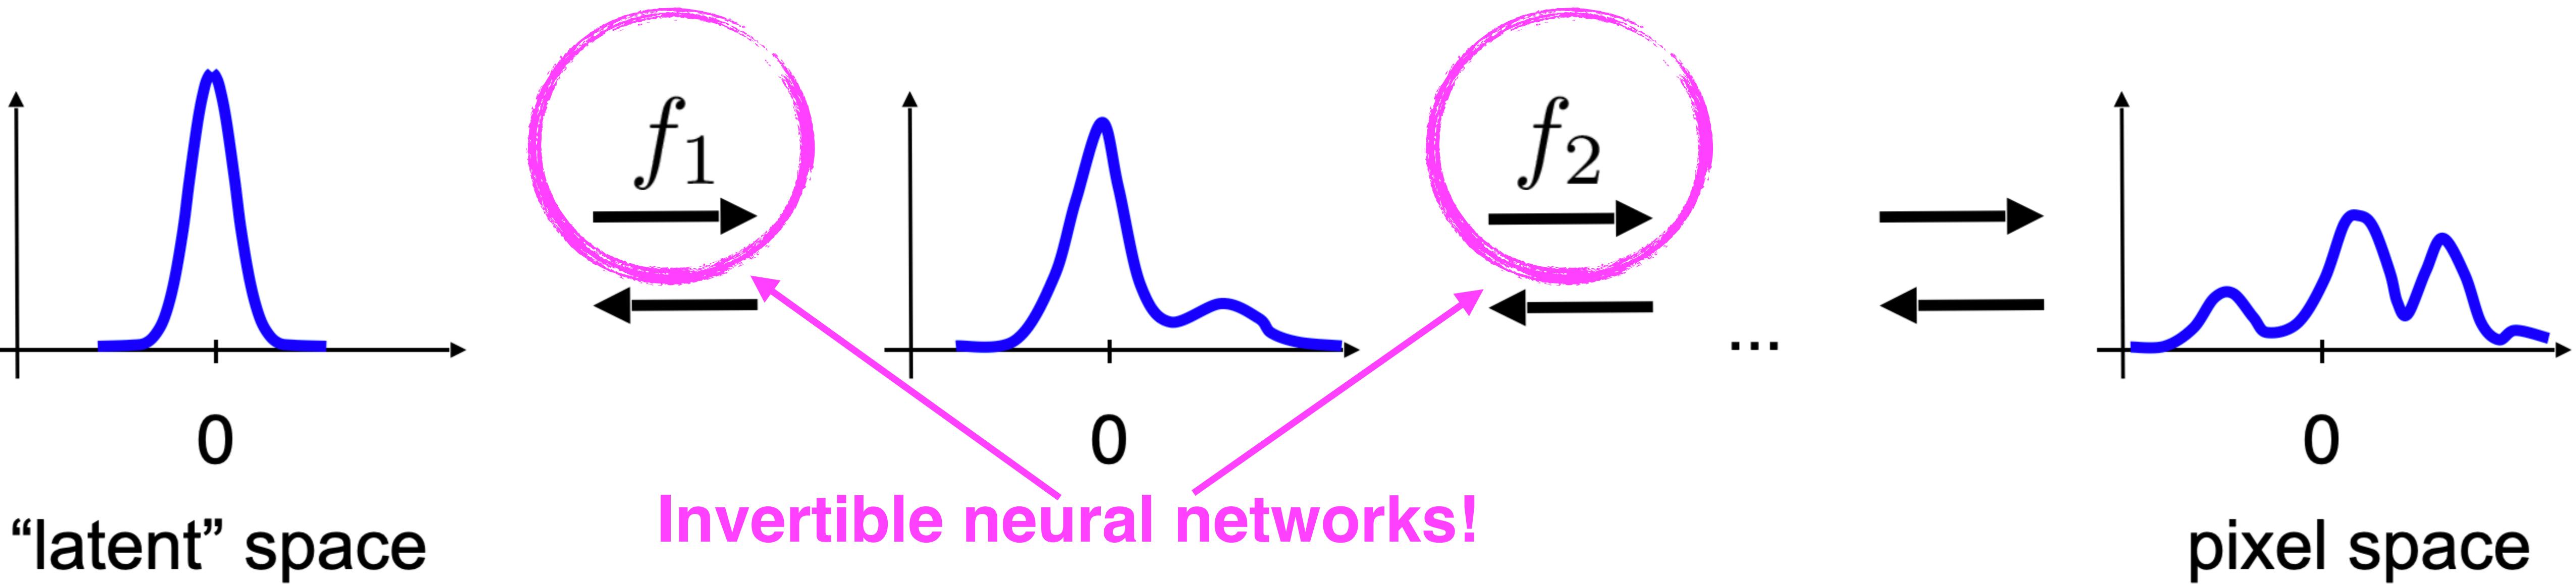
$$p(\mathbf{x}) = \pi(\mathbf{z}_0 = f^{-1}(\mathbf{x})) \prod_{i=1}^K \left| \mathbf{J}_{f_i}(z_{i-1}) \right|^{-1}$$



Flows (Flow-based models)

- We change a random variable \mathbf{x} to another random variable \mathbf{z} using invertible transformations, $\mathbf{x}, \mathbf{z} \in \mathbb{R}^D$:

$$p(\mathbf{x}) = \pi(\mathbf{z}_0 = f^{-1}(\mathbf{x})) \prod_{i=1}^K \left| \mathbf{J}_{f_i}(z_{i-1}) \right|^{-1}$$

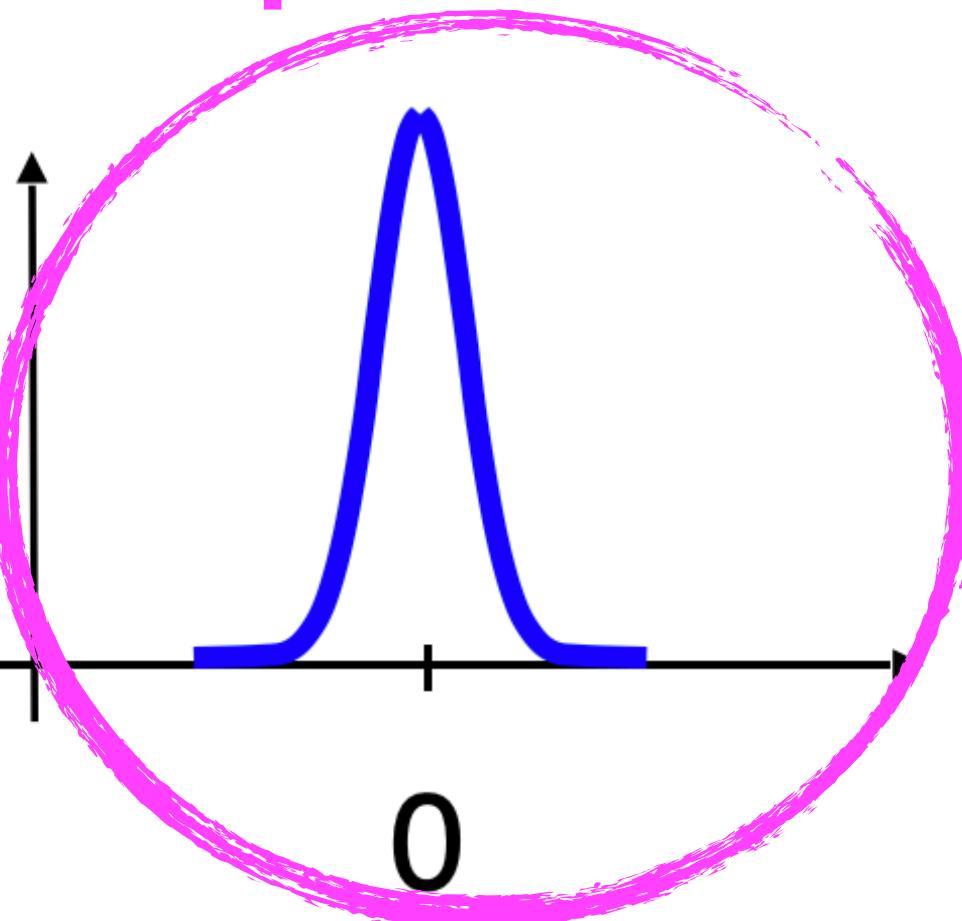


Flows (Flow-based models)

- We change a random variable \mathbf{x} to another random variable \mathbf{z} using invertible transformations, $\mathbf{x}, \mathbf{z} \in \mathbb{R}^D$:

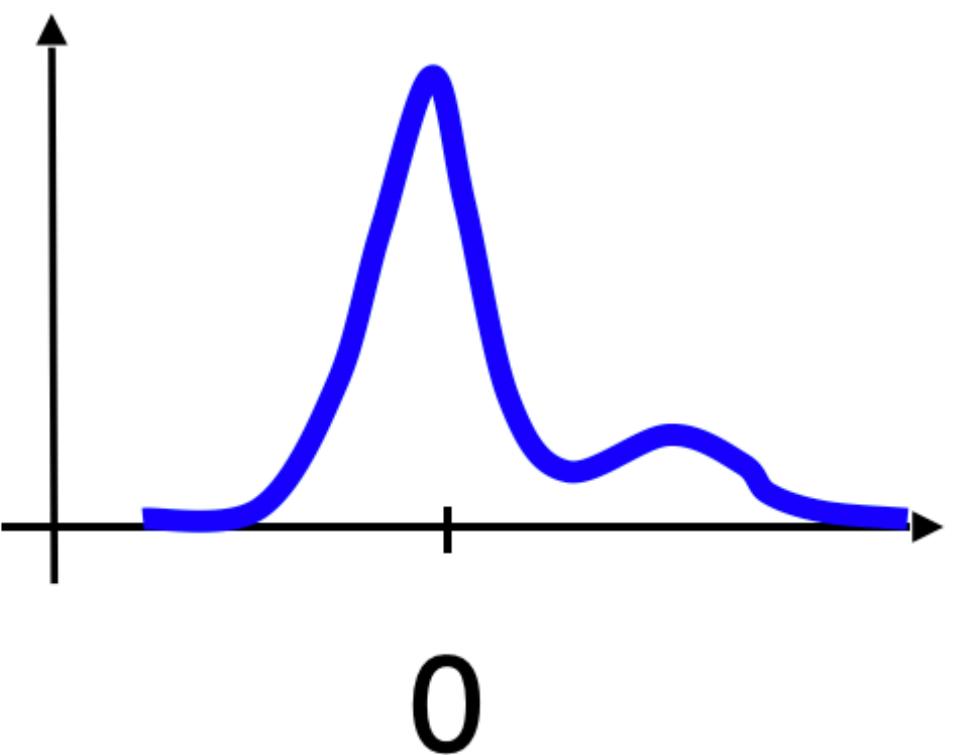
$$p(\mathbf{x}) = \pi(\mathbf{z}_0 = f^{-1}(\mathbf{x})) \prod_{i=1}^K \left| \mathbf{J}_{f_i}(z_{i-1}) \right|^{-1}$$

Simple distribution



"latent" space

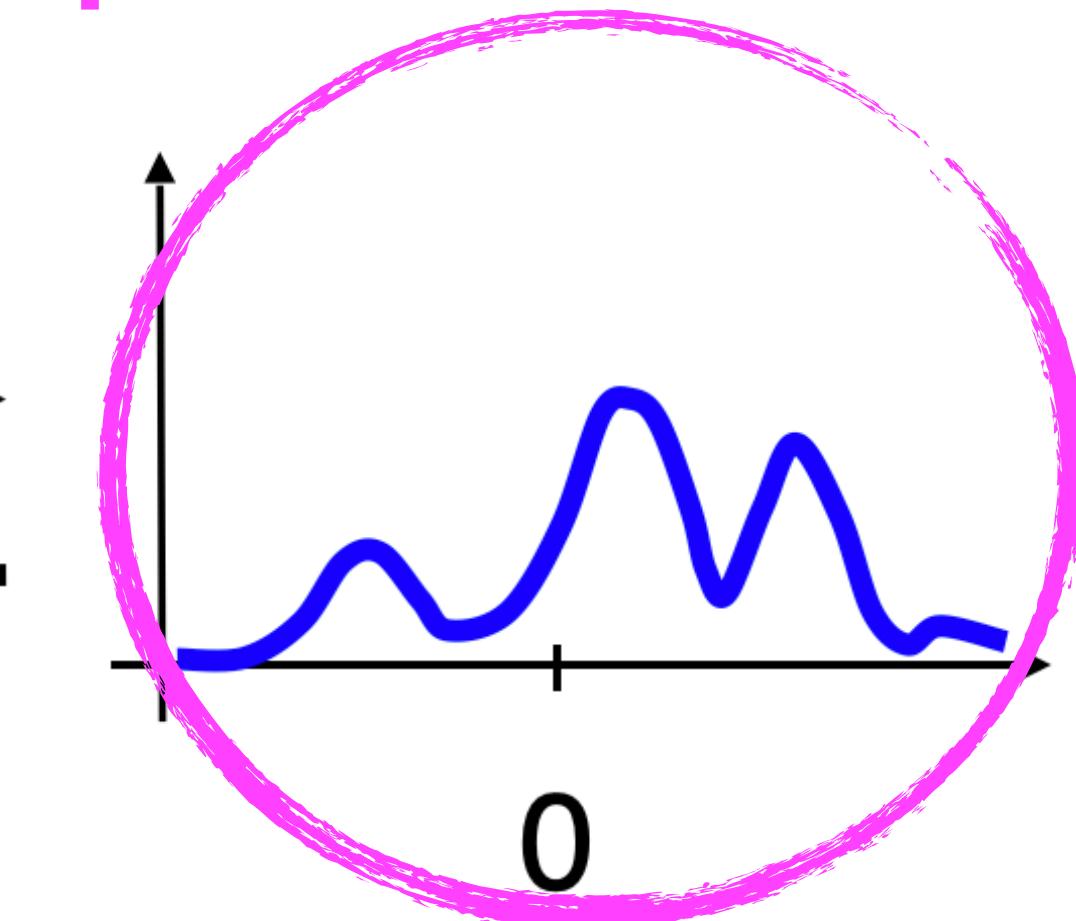
$$f_1$$



$$f_2$$

...

Complex distribution



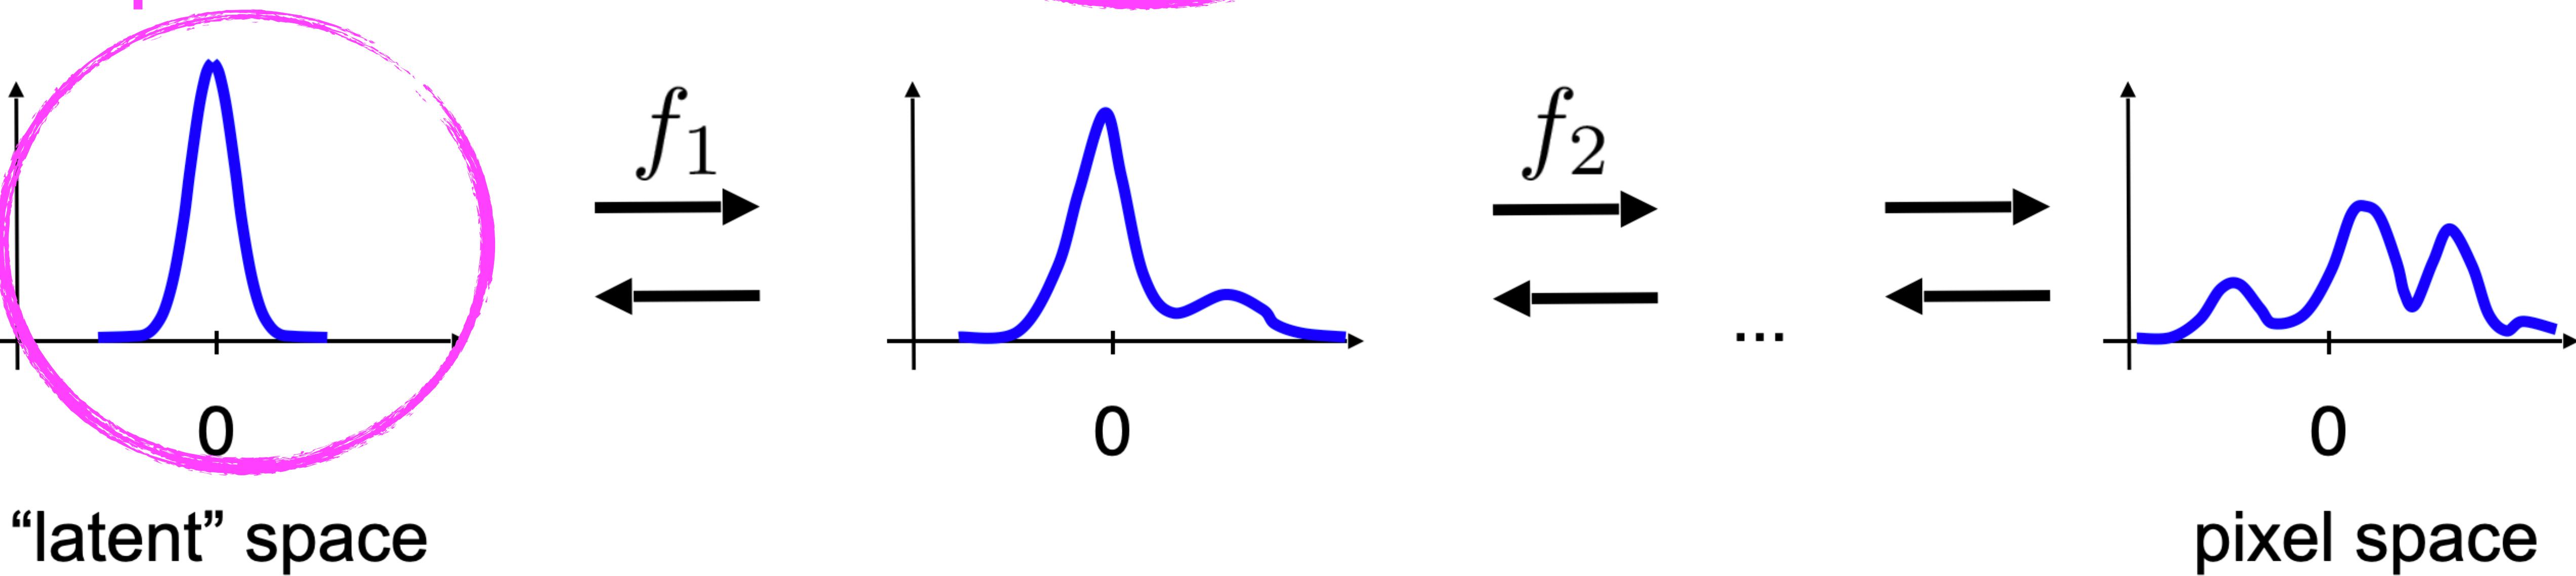
pixel space

Flows (Flow-based models)

- We change a random variable \mathbf{x} to another random variable \mathbf{z} using invertible transformations, $\mathbf{x}, \mathbf{z} \in \mathbb{R}^D$: **Known (Gaussian)**

$$p(\mathbf{x}) = \pi(\mathbf{z}_0 = f^{-1}(\mathbf{x})) \prod_{i=1}^K \left| \mathbf{J}_{f_i}(z_{i-1}) \right|^{-1}$$

Simple distribution

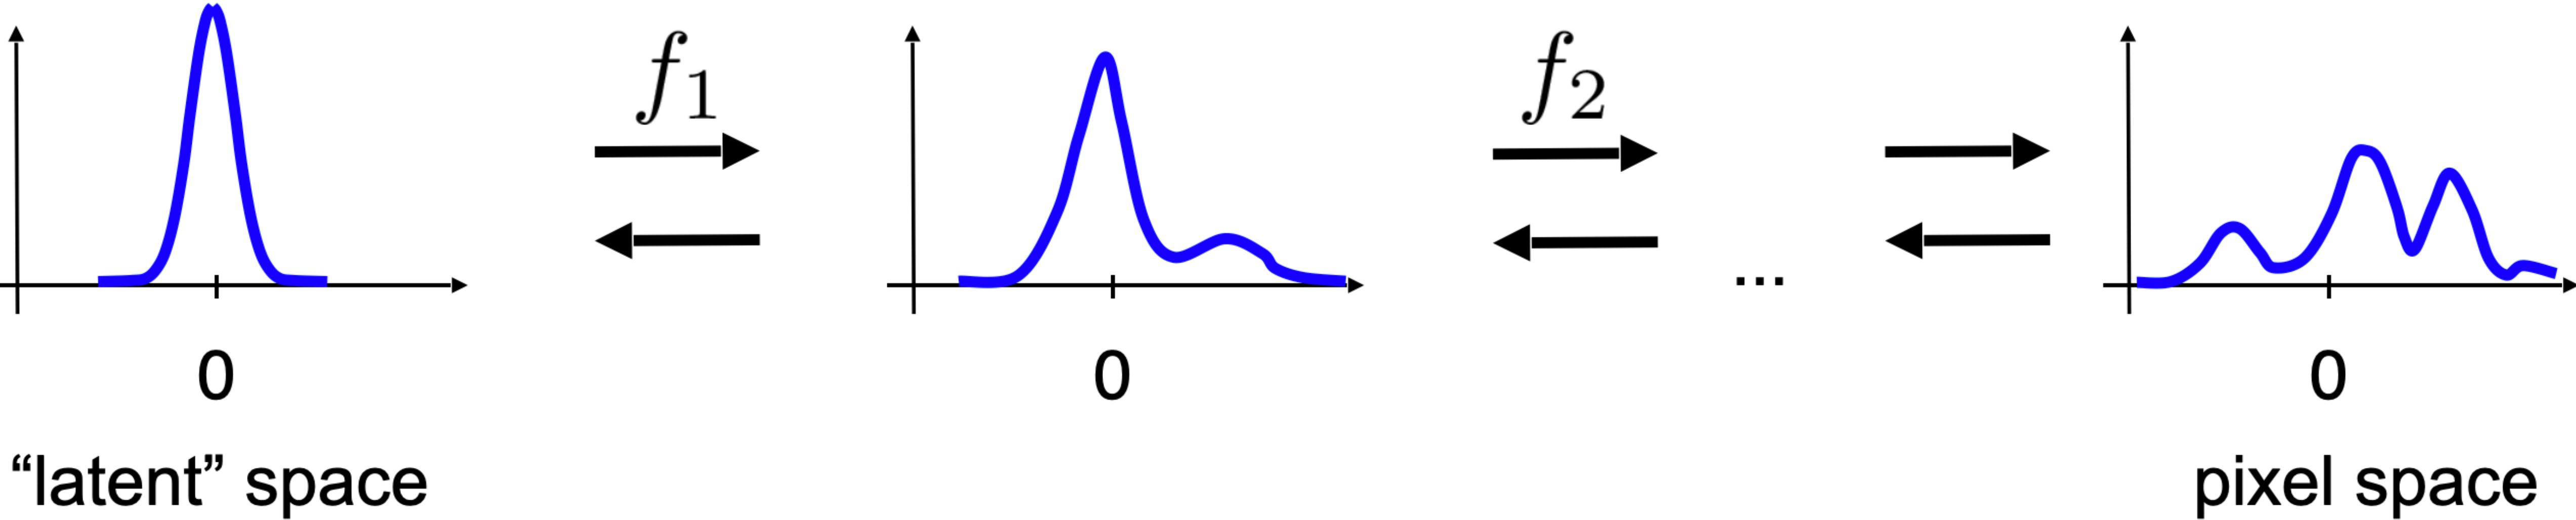


Flows (Flow-based models)

- We change a random variable \mathbf{x} to another random variable \mathbf{z} using invertible transformations, $\mathbf{x}, \mathbf{z} \in \mathbb{R}^D$:

$$p(\mathbf{x}) = \pi(\mathbf{z}_0 = f^{-1}(\mathbf{x})) \prod_{i=1}^K \left| \mathbf{J}_{f_i}(z_{i-1}) \right|^{-1}$$

Jacobian must be tractable



Flows (Flow-based models)

- We change a random variable \mathbf{x} to another random variable \mathbf{z} using invertible transformations, $\mathbf{x}, \mathbf{z} \in \mathbb{R}^D$:

$$p(\mathbf{x}) = \pi(\mathbf{z}_0 = f^{-1}(\mathbf{x})) \prod_{i=1}^K \left| \mathbf{J}_{f_i}(z_{i-1}) \right|^{-1}$$

- Training objective:

$$\ln p(\mathbf{x}) = \ln \pi(\mathbf{z}_0 = f^{-1}(\mathbf{x})) - \sum_{i=1}^K \ln \left| \mathbf{J}_{f_i}(z_{i-1}) \right|$$

Flows (Flow-based models): Invertible layers

Two main components

1) Coupling layer:

$$\mathbf{y}_a = \mathbf{x}_a$$

$$\mathbf{y}_b = \exp(s(\mathbf{x}_a)) \odot \mathbf{x}_b + t(\mathbf{x}_a)$$

is invertible by design:

$$\mathbf{x}_b = (\mathbf{y}_b - t(\mathbf{y}_a)) \odot \exp(-s(\mathbf{y}_a))$$

$$\mathbf{x}_a = \mathbf{y}_a$$

2) Permutation layer

Flows (Flow-based models): Invertible layers

Two main components

1) Coupling layer:

$$\mathbf{y}_a = \mathbf{x}_a$$

$$\mathbf{y}_b = \exp(s(\mathbf{x}_a)) \odot \mathbf{x}_b + t(\mathbf{x}_a)$$

is invertible by design:

$$\mathbf{x}_b = (\mathbf{y}_b - t(\mathbf{y}_a)) \odot \exp(-s(\mathbf{y}_a))$$

$$\mathbf{x}_a = \mathbf{y}_a$$

2) Permutation layer

Jacobian is tractable!

$$\det(\mathbf{J}) = \prod_{j=1}^{D-d} \exp\left(s(\mathbf{x}_a)\right)_j = \exp\left(\sum_{j=1}^{D-d} s(\mathbf{x}_a)_j\right)$$

$$\det(\mathbf{J}) = 1$$

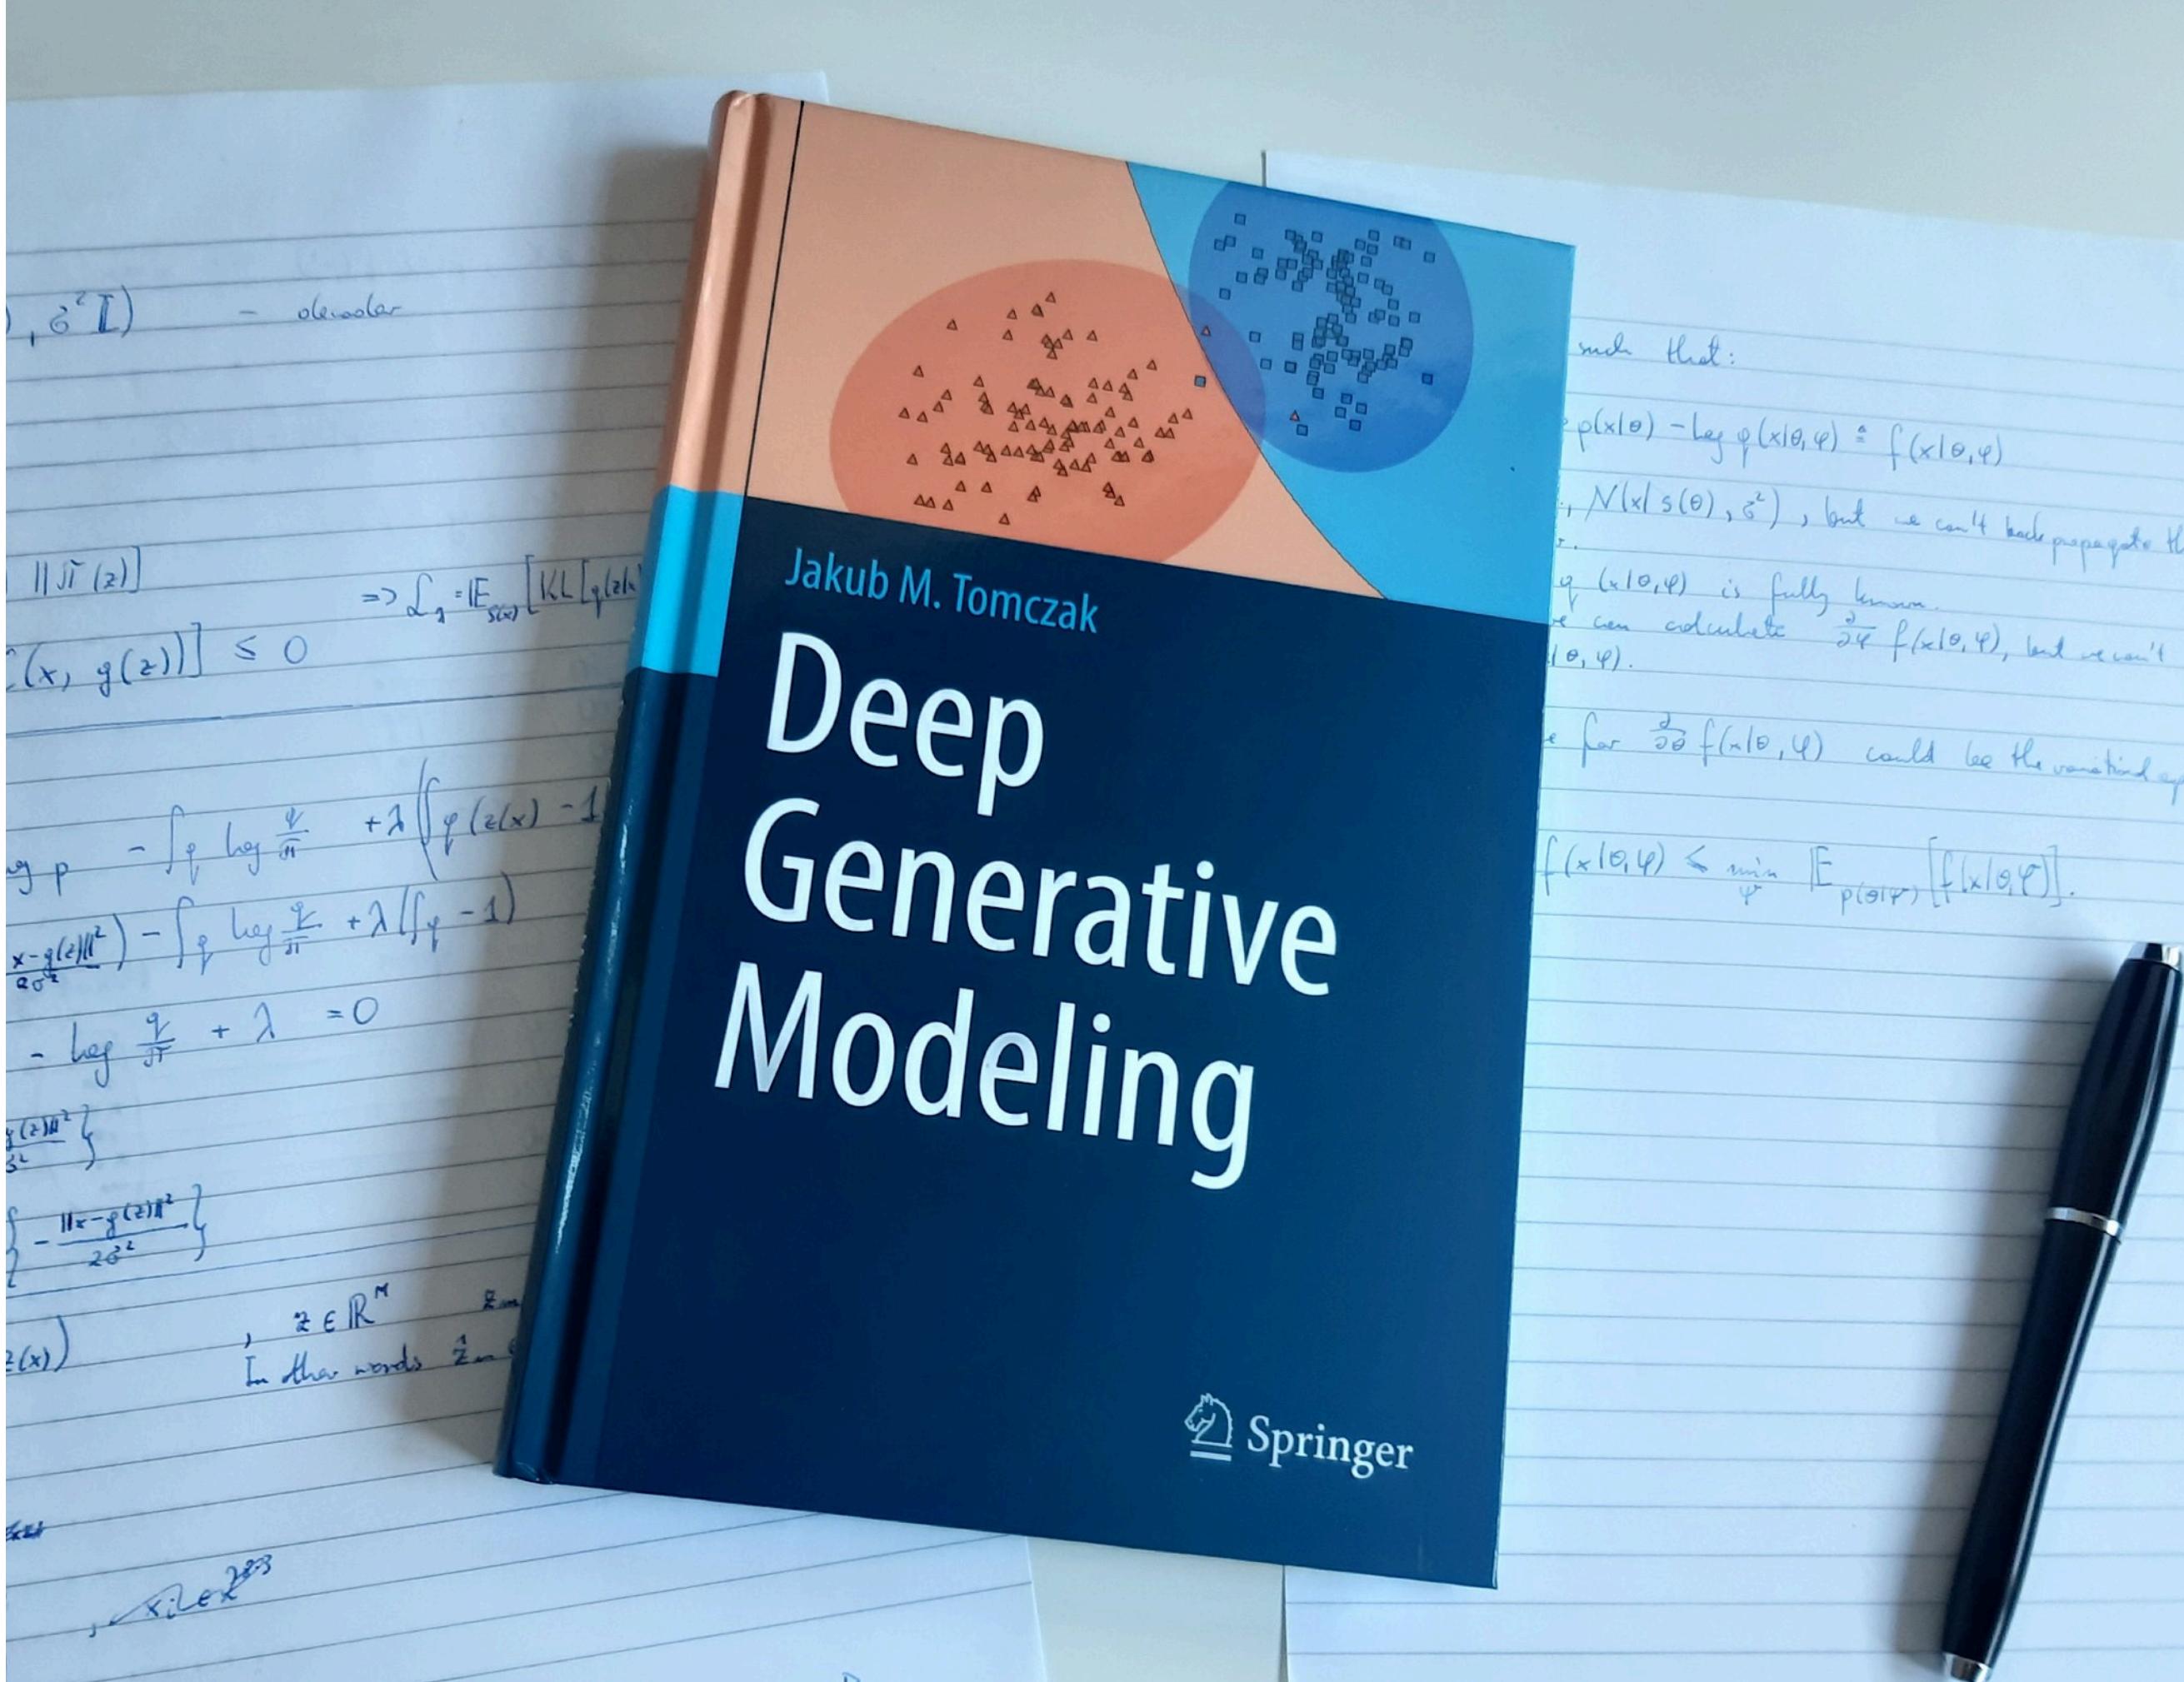
Flows (Flow-based models)



Deep Generative Modeling

The first comprehensive book on Generative AI.

Tomczak, J.M., (2022), "Deep Generative Modeling", Springer Cham

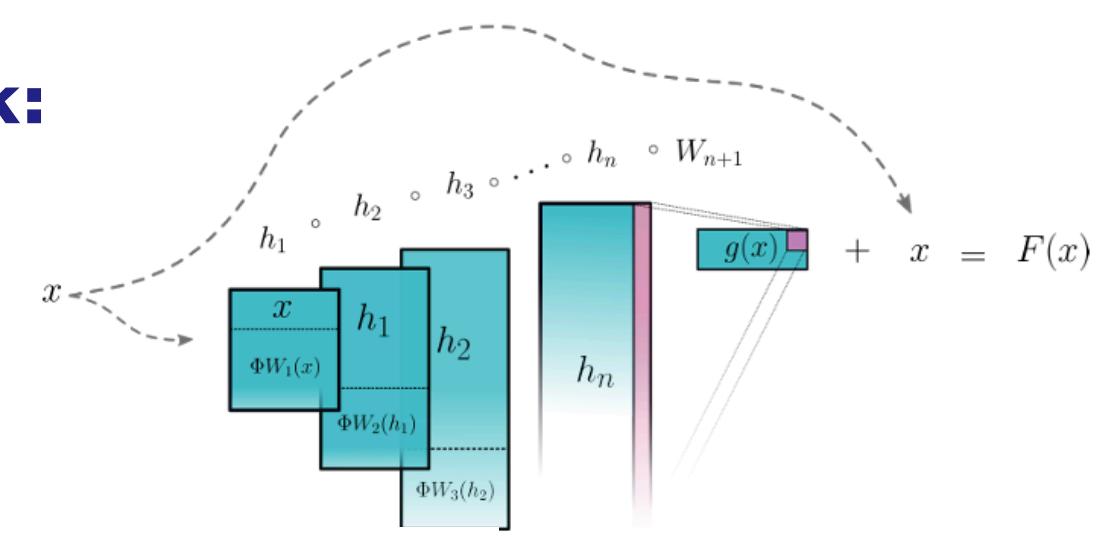


Invertible DenseNets with Concatenated LipSwish

The change-of-variables formula:

$$\ln p_X(x) = \ln p_Z(z) + \ln |\det J_F(x)|$$

DenseNet block:



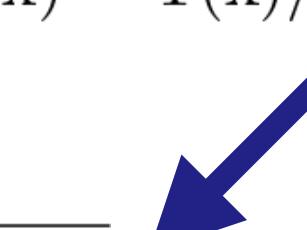
$$g(x) = W_{n+1} \circ h_n \circ \dots \circ h_1(x)$$

$$h_1(x) = \begin{bmatrix} x \\ \phi(W_1x) \end{bmatrix}, \quad h_2(h_1(x)) = \begin{bmatrix} h_1(x) \\ \phi(W_2h_1(x)) \end{bmatrix}$$

Concatenated LipSwish:

$$\Phi(x) = \begin{bmatrix} \phi_1(x) \\ \phi_2(x) \end{bmatrix} = \begin{bmatrix} \text{LipSwish}(x) \\ \text{LipSwish}(-x) \end{bmatrix}, \quad \text{CLipSwish}(x) = \Phi(x)/\text{Lip}(\Phi)$$

$$\sup_x \sqrt{\left(\frac{\partial \phi_1(x)}{\partial x} \right)^2 + \left(\frac{\partial \phi_2(x)}{\partial x} \right)^2}$$



Example generations:



(a) Real CIFAR10 images.



(b) Samples of i-DenseNets trained on CIFAR10.



(c) Real ImageNet32 images.



(d) Samples of i-DenseNets trained on ImageNet32.

Classification:

Model \ Evaluation	$\lambda = 0$	$\lambda = \frac{1}{D}$	$\lambda = 1$		
	Acc \uparrow	Acc \uparrow	bpd \downarrow	Acc \uparrow	bpd \downarrow
Coupling	89.77%	87.58%	4.30	67.62%	3.54
+ 1 \times 1 conv	90.82%	87.96%	4.09	67.38%	3.47
Residual Blocks (full)	91.78%	90.47%	3.62	70.32%	3.39
Dense Blocks (full)	92.40%	90.79%	3.49	75.67%	3.31



Yura
Perugachi-Diaz



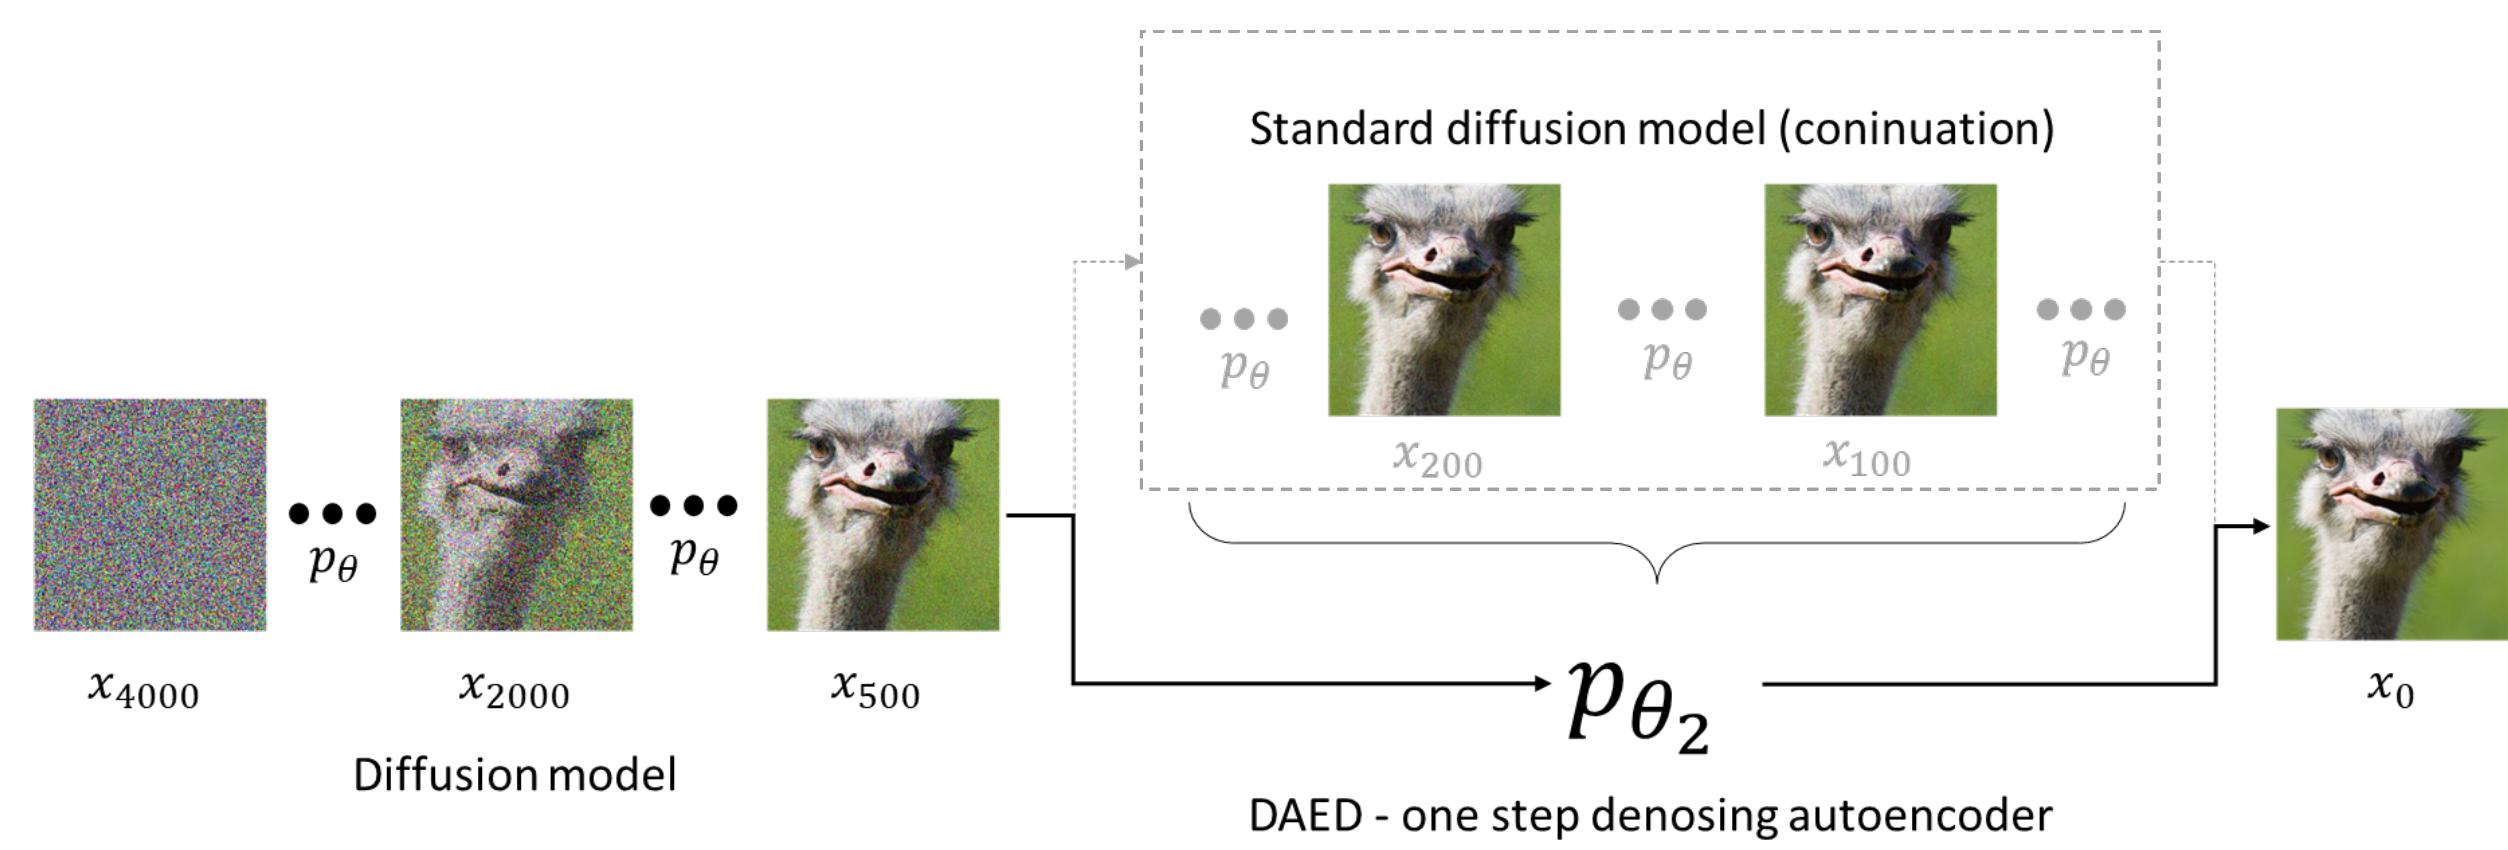
Sandjai
Bhulai



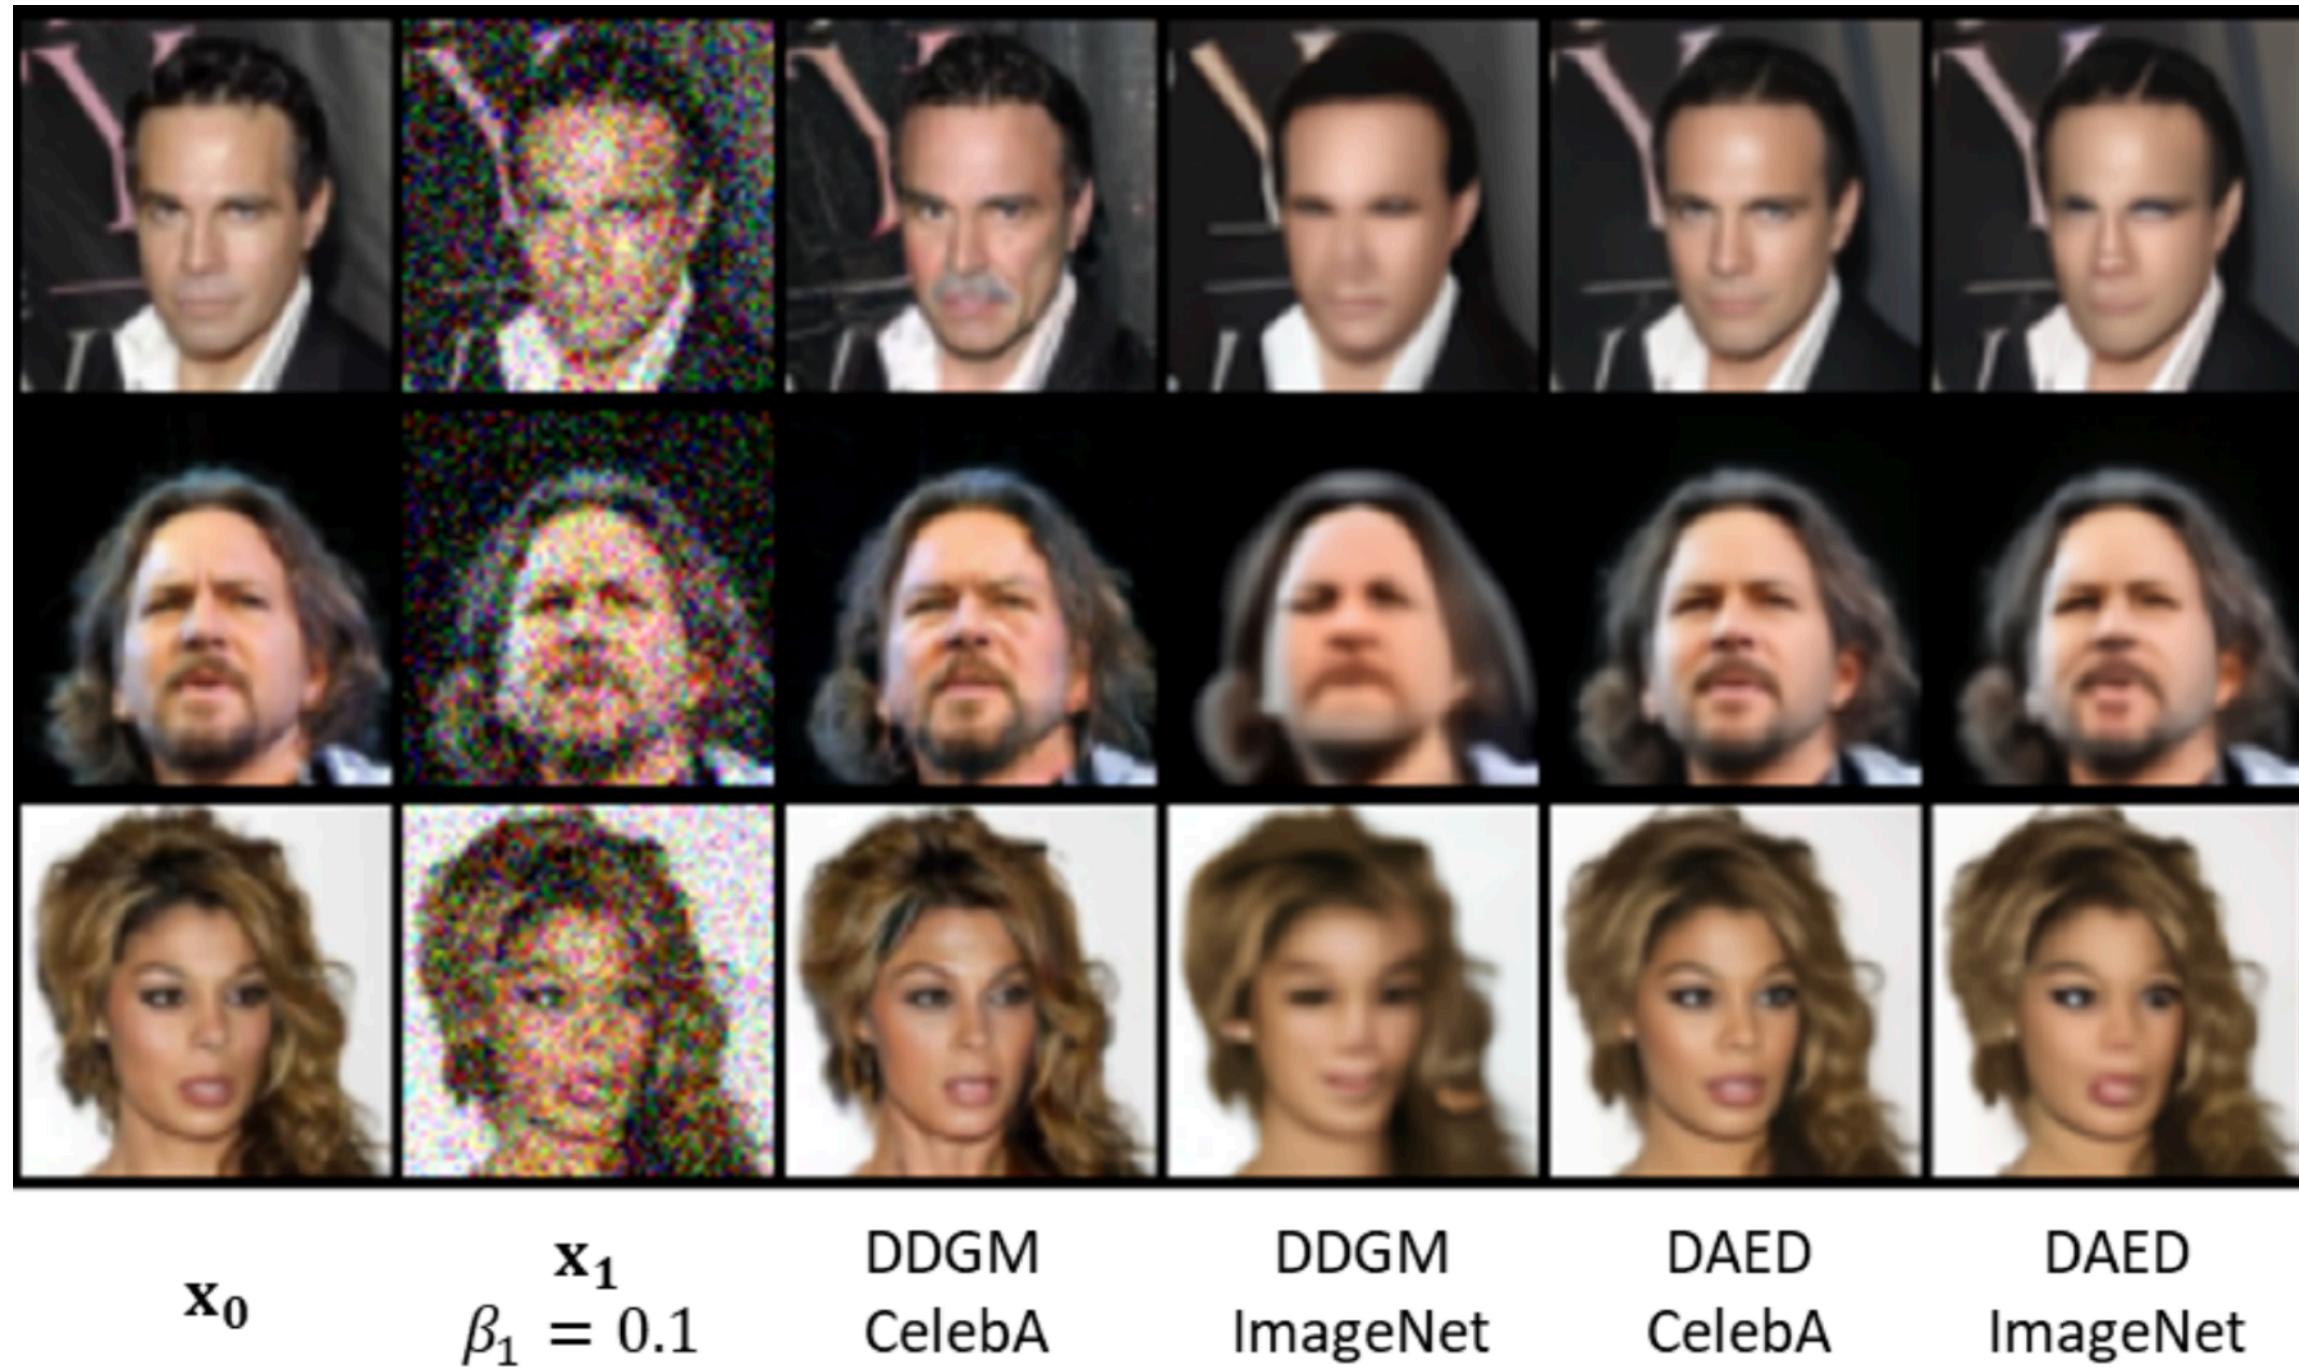
Jakub M.
Tomczak

On Analyzing Generative and Denoising Capabilities of Diffusion-based Deep Generative Models

Dividing a diffusion model into a denoiser and a generator:



Example results:



$$\begin{aligned} \bar{\ell}(\mathbf{x}_0; \varphi, \theta) &= \mathbb{E}_{\mathbf{x}_1 \sim q(\mathbf{x}_1 | \mathbf{x}_0)} [\ln p(\mathbf{x}_0 | f_\varphi(\mathbf{x}_1)) + \ln p_\theta(\mathbf{x}_1)] \\ &\geq \underbrace{\mathbb{E}_{\mathbf{x}_1 \sim q(\mathbf{x}_1 | \mathbf{x}_0)} [\ln p(\mathbf{x}_0 | f_\varphi(\mathbf{x}_1))]}_{\ell_{\text{DAE}}(\mathbf{x}_0; \varphi)} + \underbrace{\mathbb{E}_{q(\mathbf{x}_2, \dots, \mathbf{x}_T | \mathbf{x}_1)} \left[\frac{\ln p_\theta(\mathbf{x}_1, \dots, \mathbf{x}_T)}{q(\mathbf{x}_1, \dots, \mathbf{x}_T | \mathbf{x}_0)} \right]}_{\ell_{\text{D}}(\mathbf{x}_0; \theta)} \end{aligned}$$

Alleviating Adversarial Attacks on Variational Autoencoders with MCMC

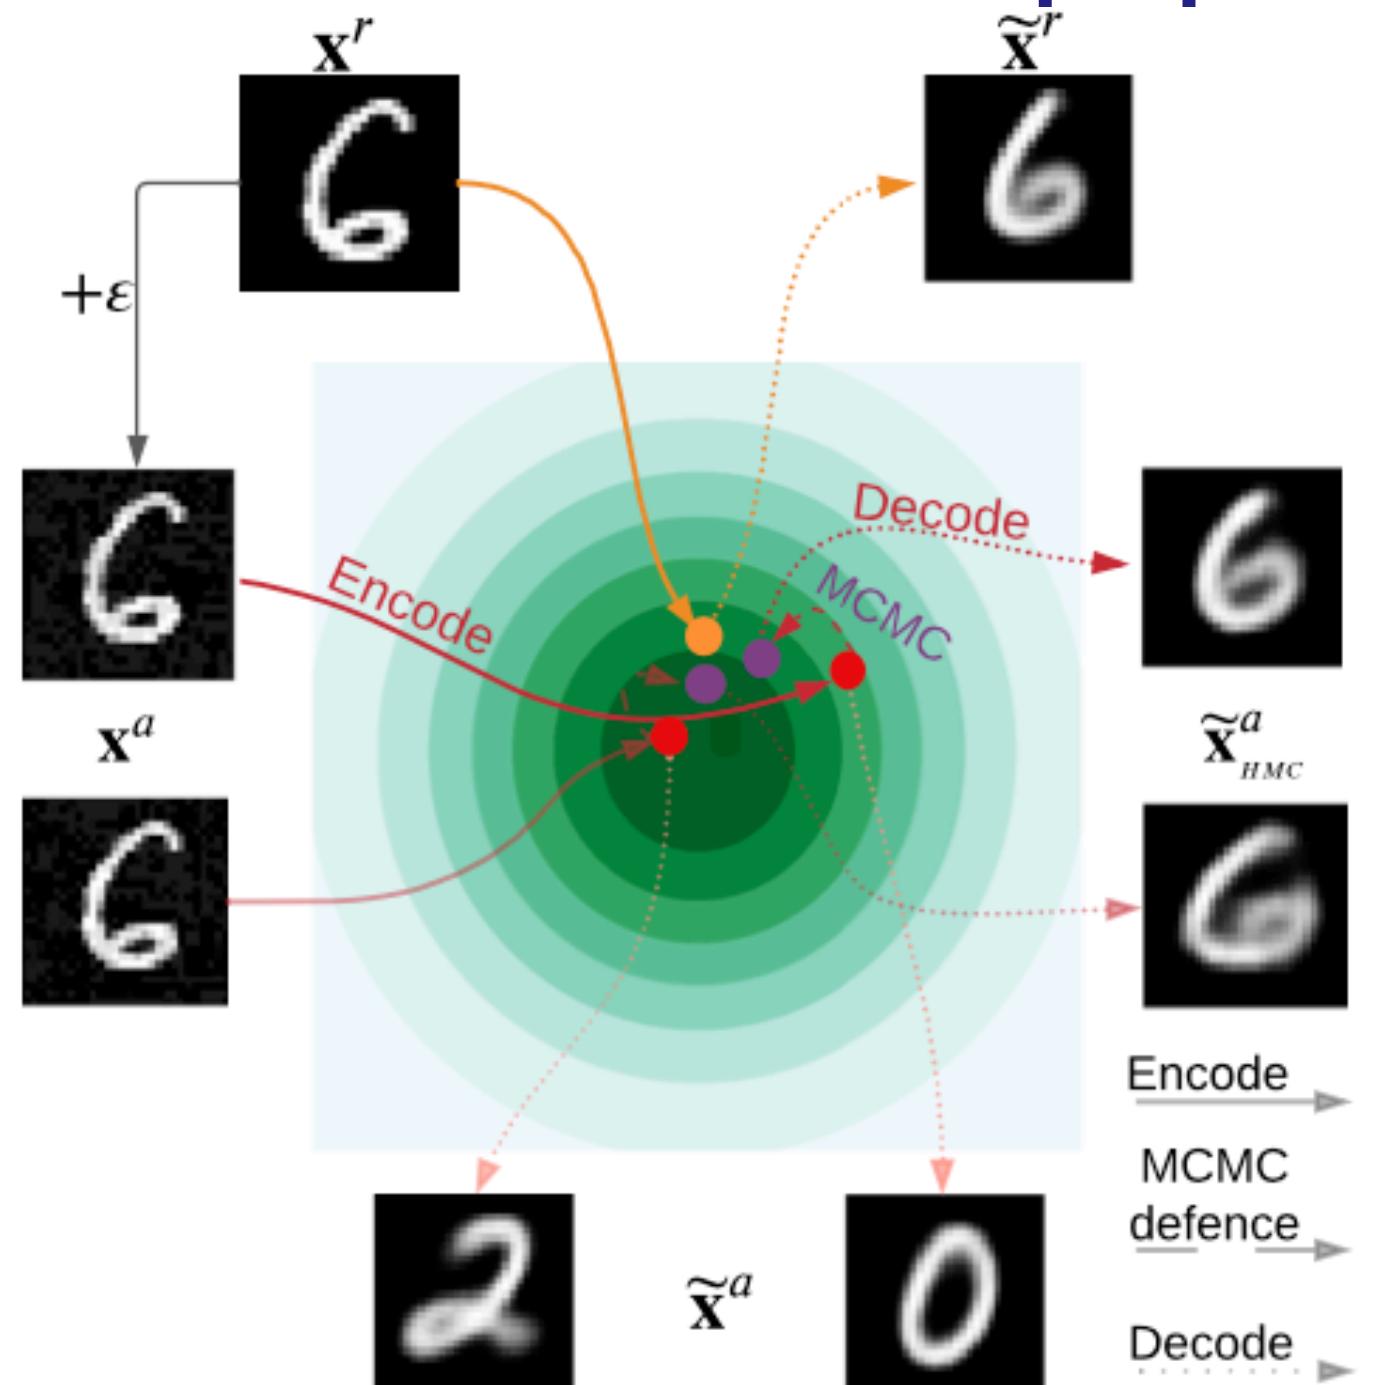


Anna
Kuzina

Max
Welling

Jakub M.
Tomczak

Unsupervised attacks on VAE and the proposed defense:



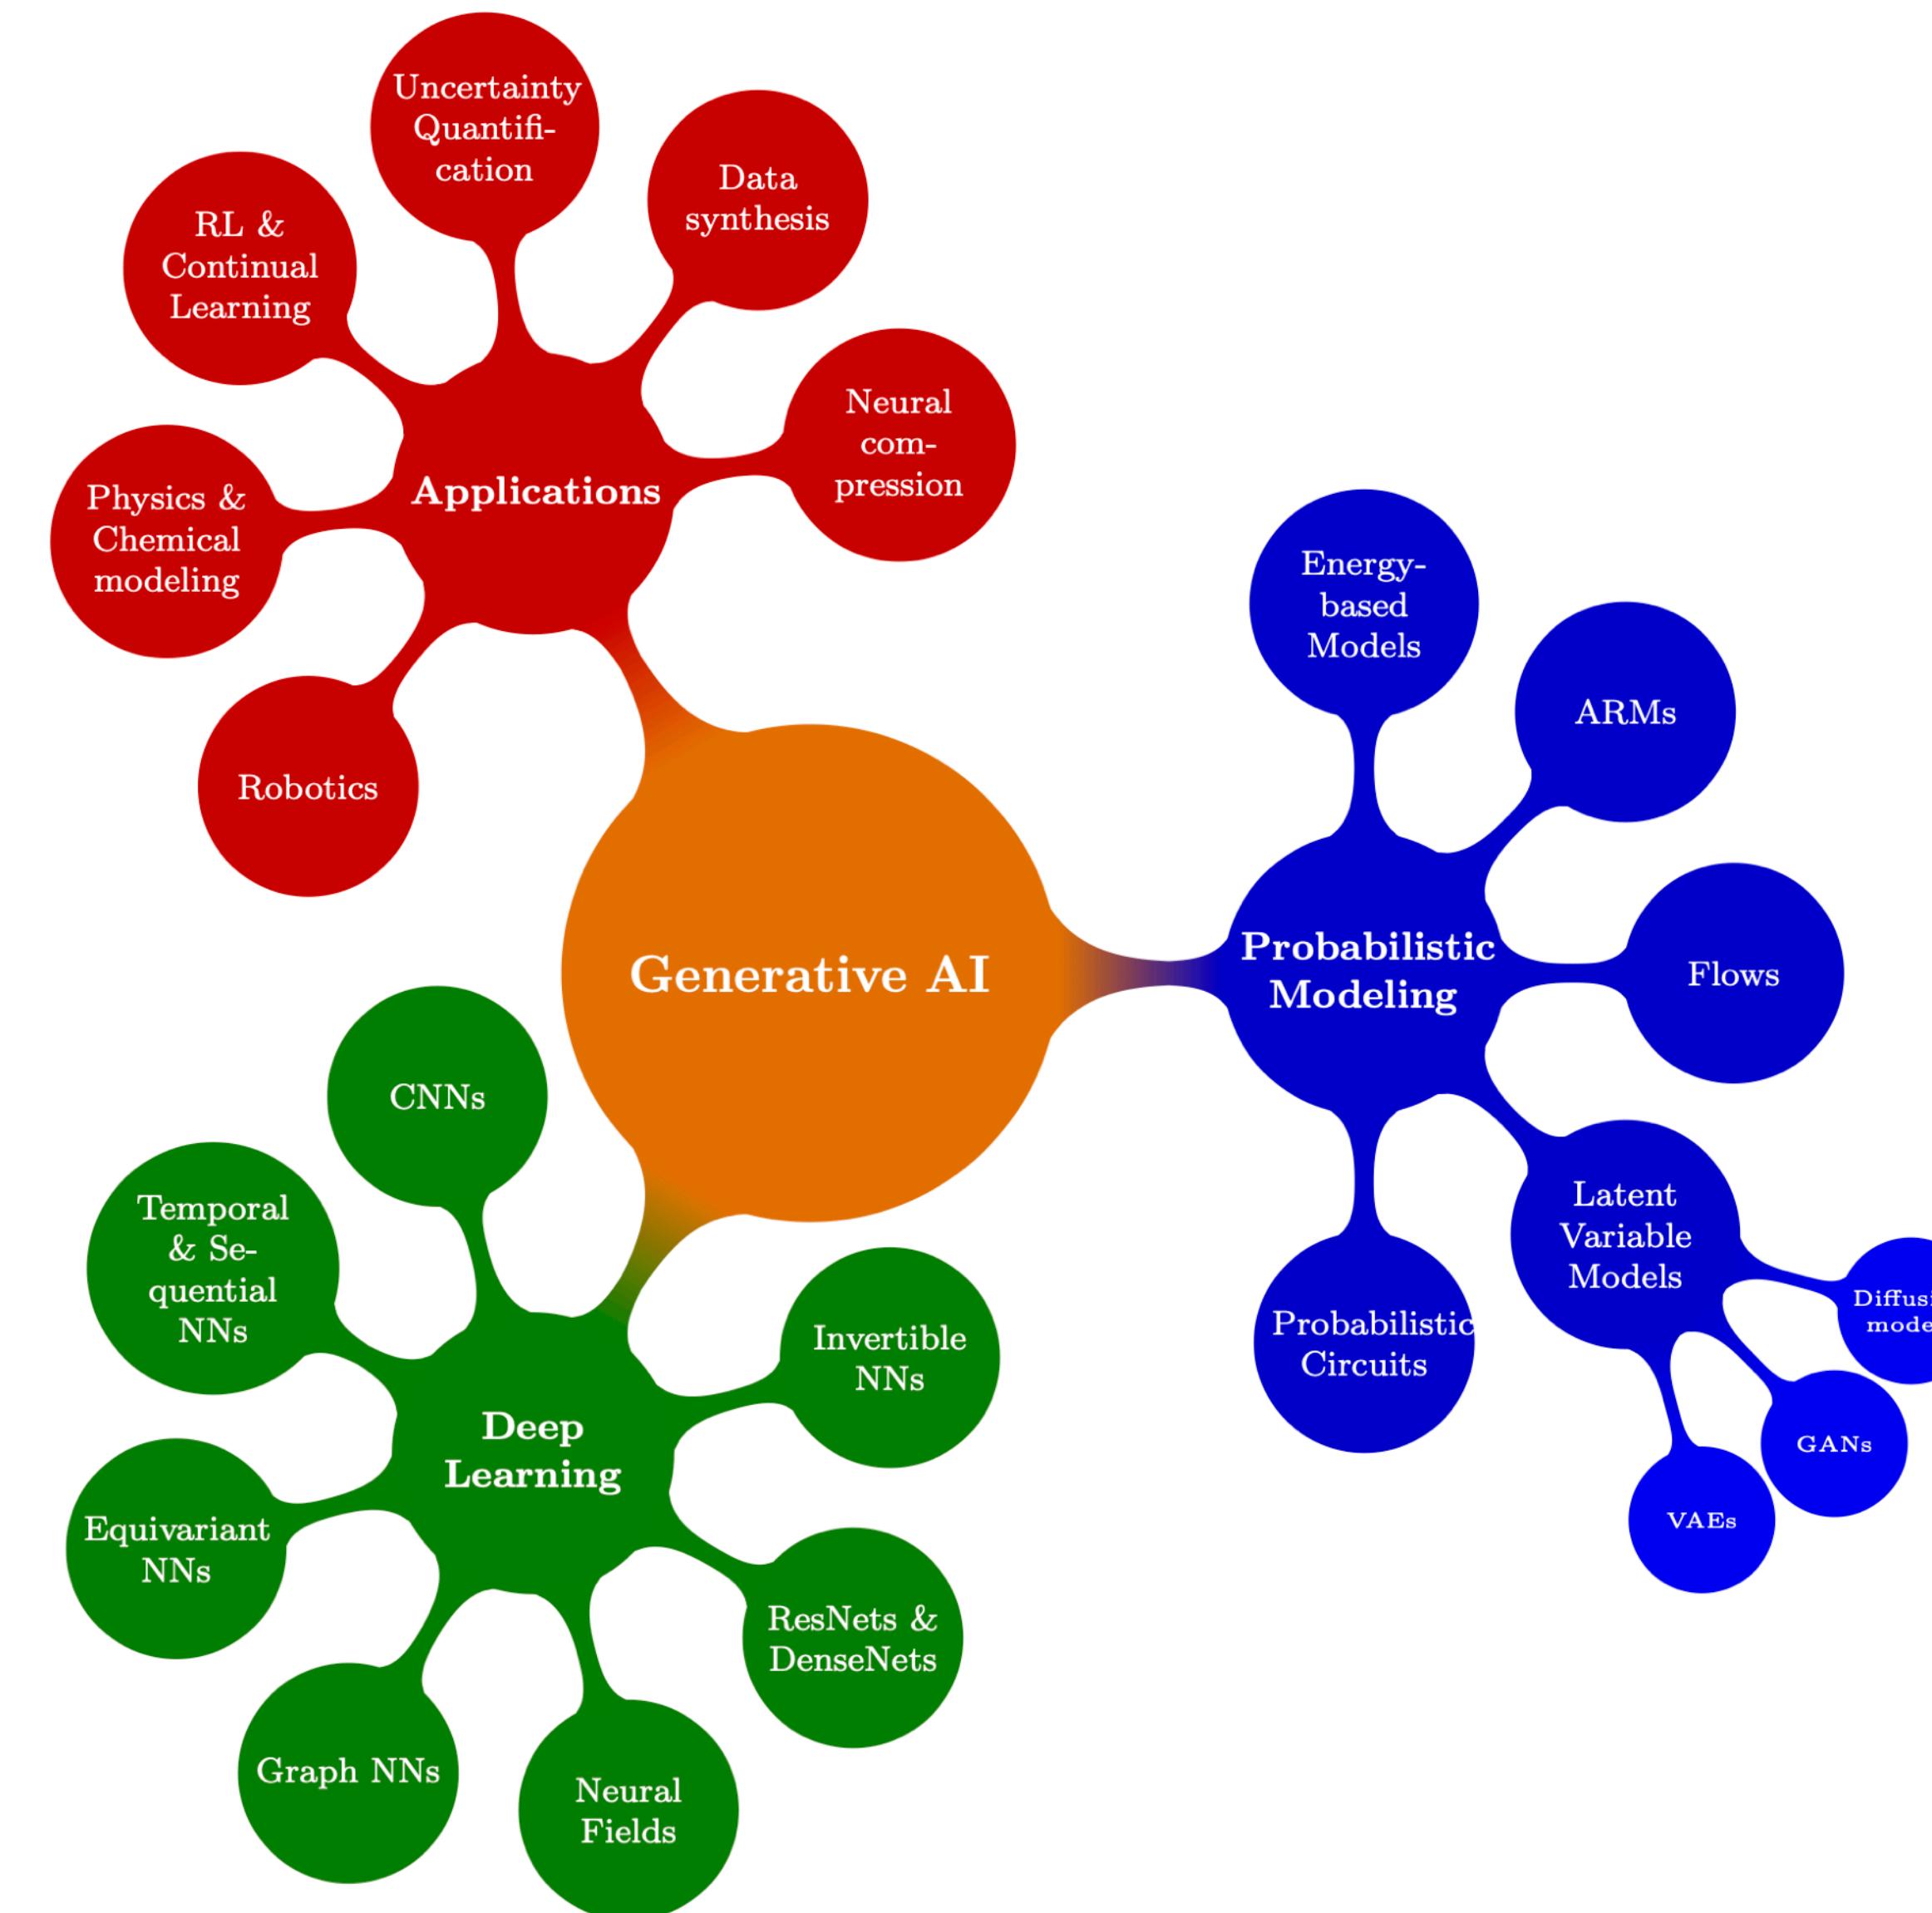
Example results:



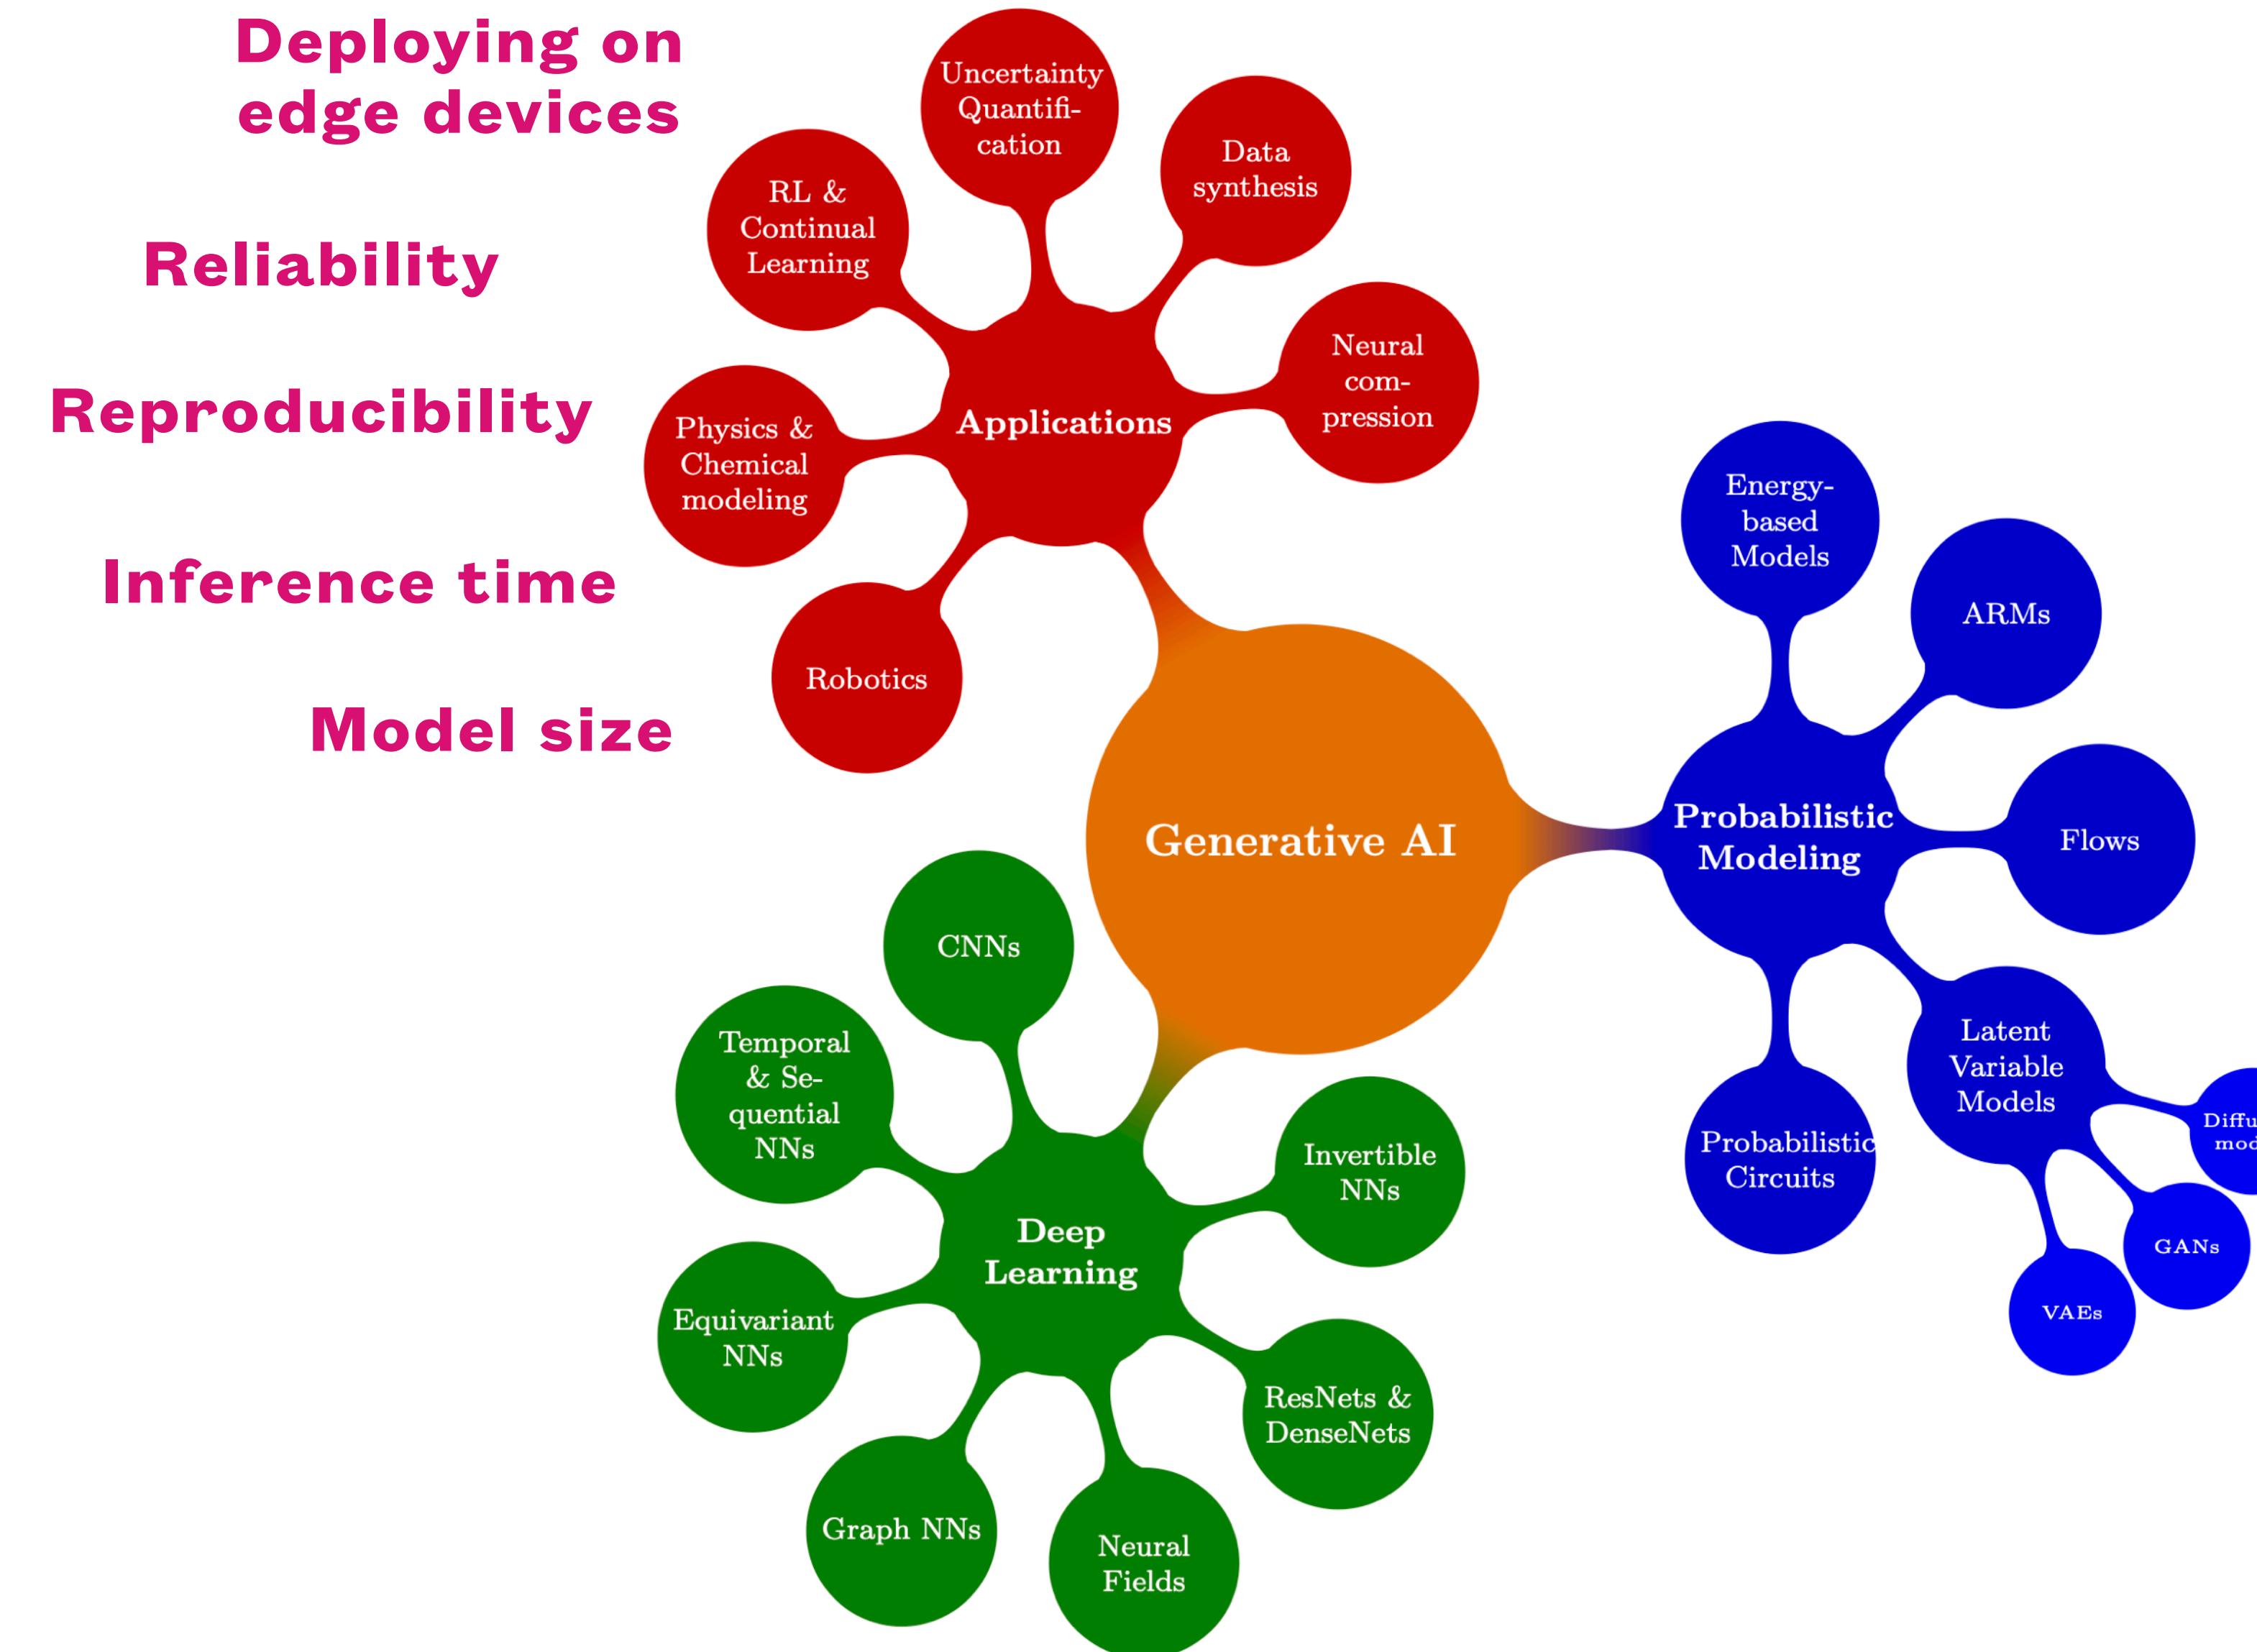
Theorem 1 The upper bound on the total variation distance between samples from MCMC for a given corrupted point \mathbf{x}^a , $q^{(t)}(\mathbf{z}|\mathbf{x}^a)$, and the variational posterior for the given real point \mathbf{x}^r , $q_\phi(\mathbf{z}|\mathbf{x}^r)$, is the following:

$$\begin{aligned} \text{TV} \left[q^{(t)}(\mathbf{z}|\mathbf{x}^a), q_\phi(\mathbf{z}|\mathbf{x}^r) \right] &\leq \text{TV} \left[q^{(t)}(\mathbf{z}|\mathbf{x}^a), p_\theta(\mathbf{z}|\mathbf{x}^a) \right] \\ &+ \sqrt{\frac{1}{2} D_{\text{KL}} [q_\phi(\mathbf{z}|\mathbf{x}^r) \| p_\theta(\mathbf{z}|\mathbf{x}^r)]} \\ &+ o(\sqrt{\|\varepsilon\|}). \end{aligned}$$

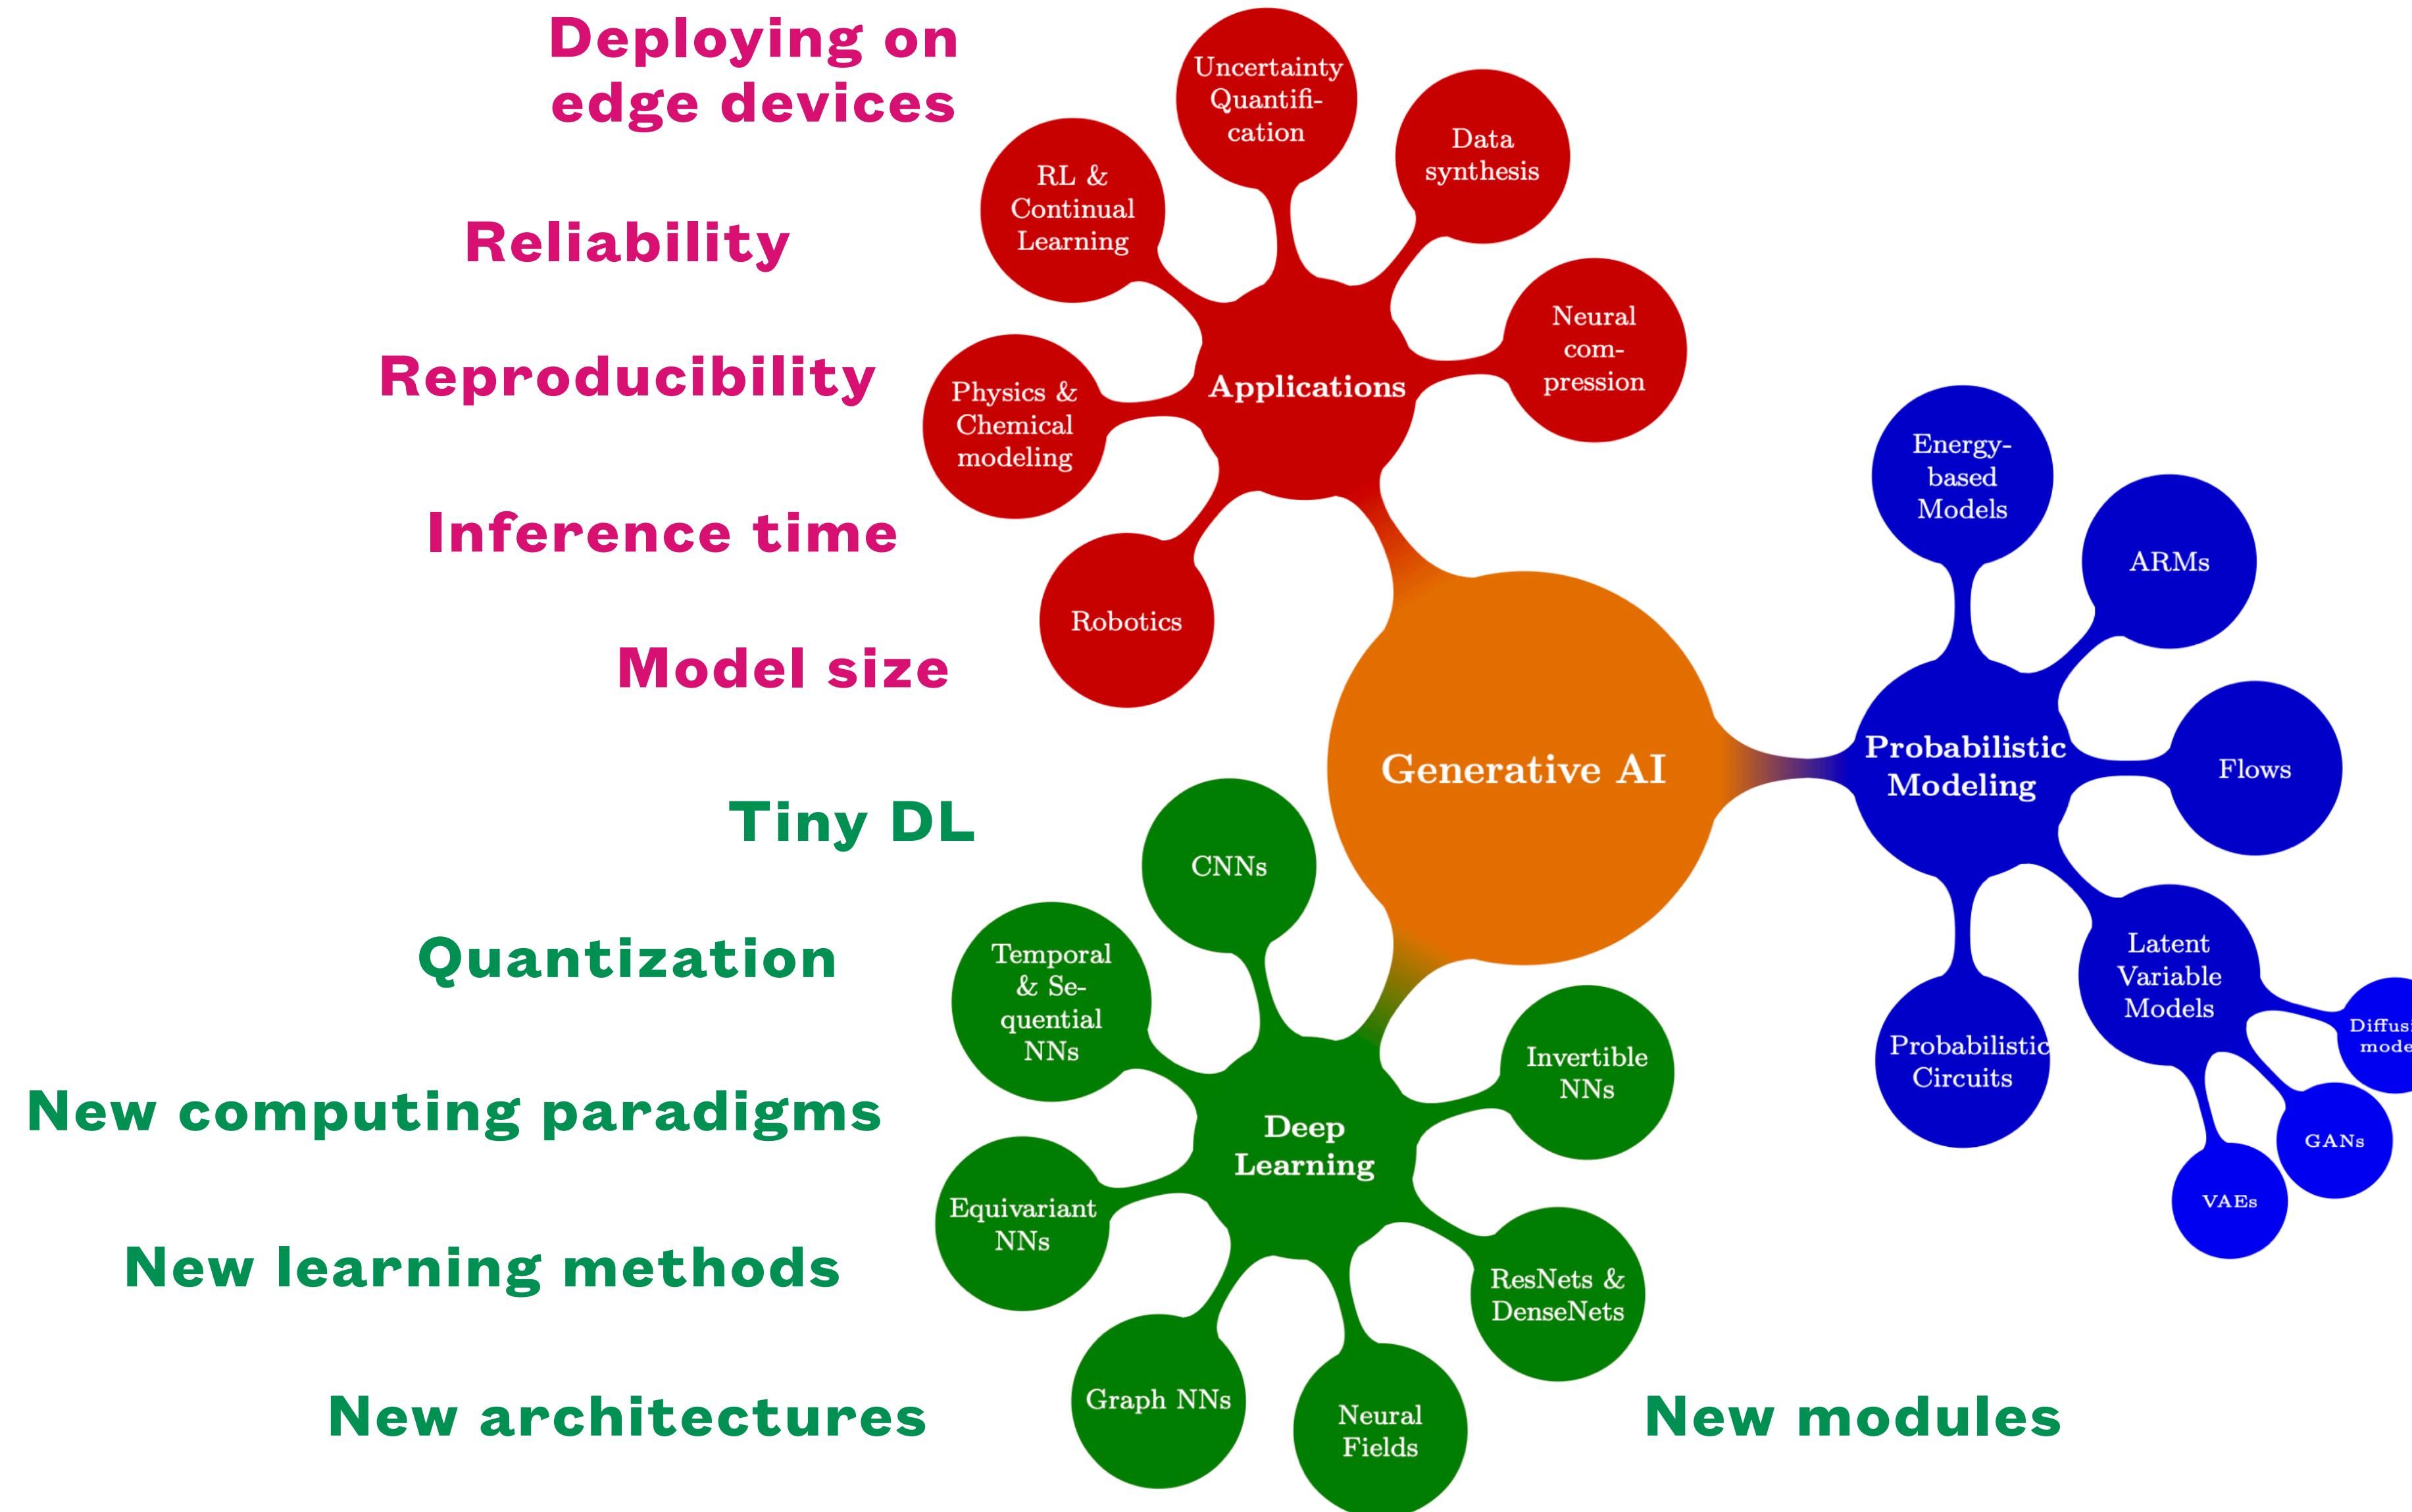
Challenges in Generative AI



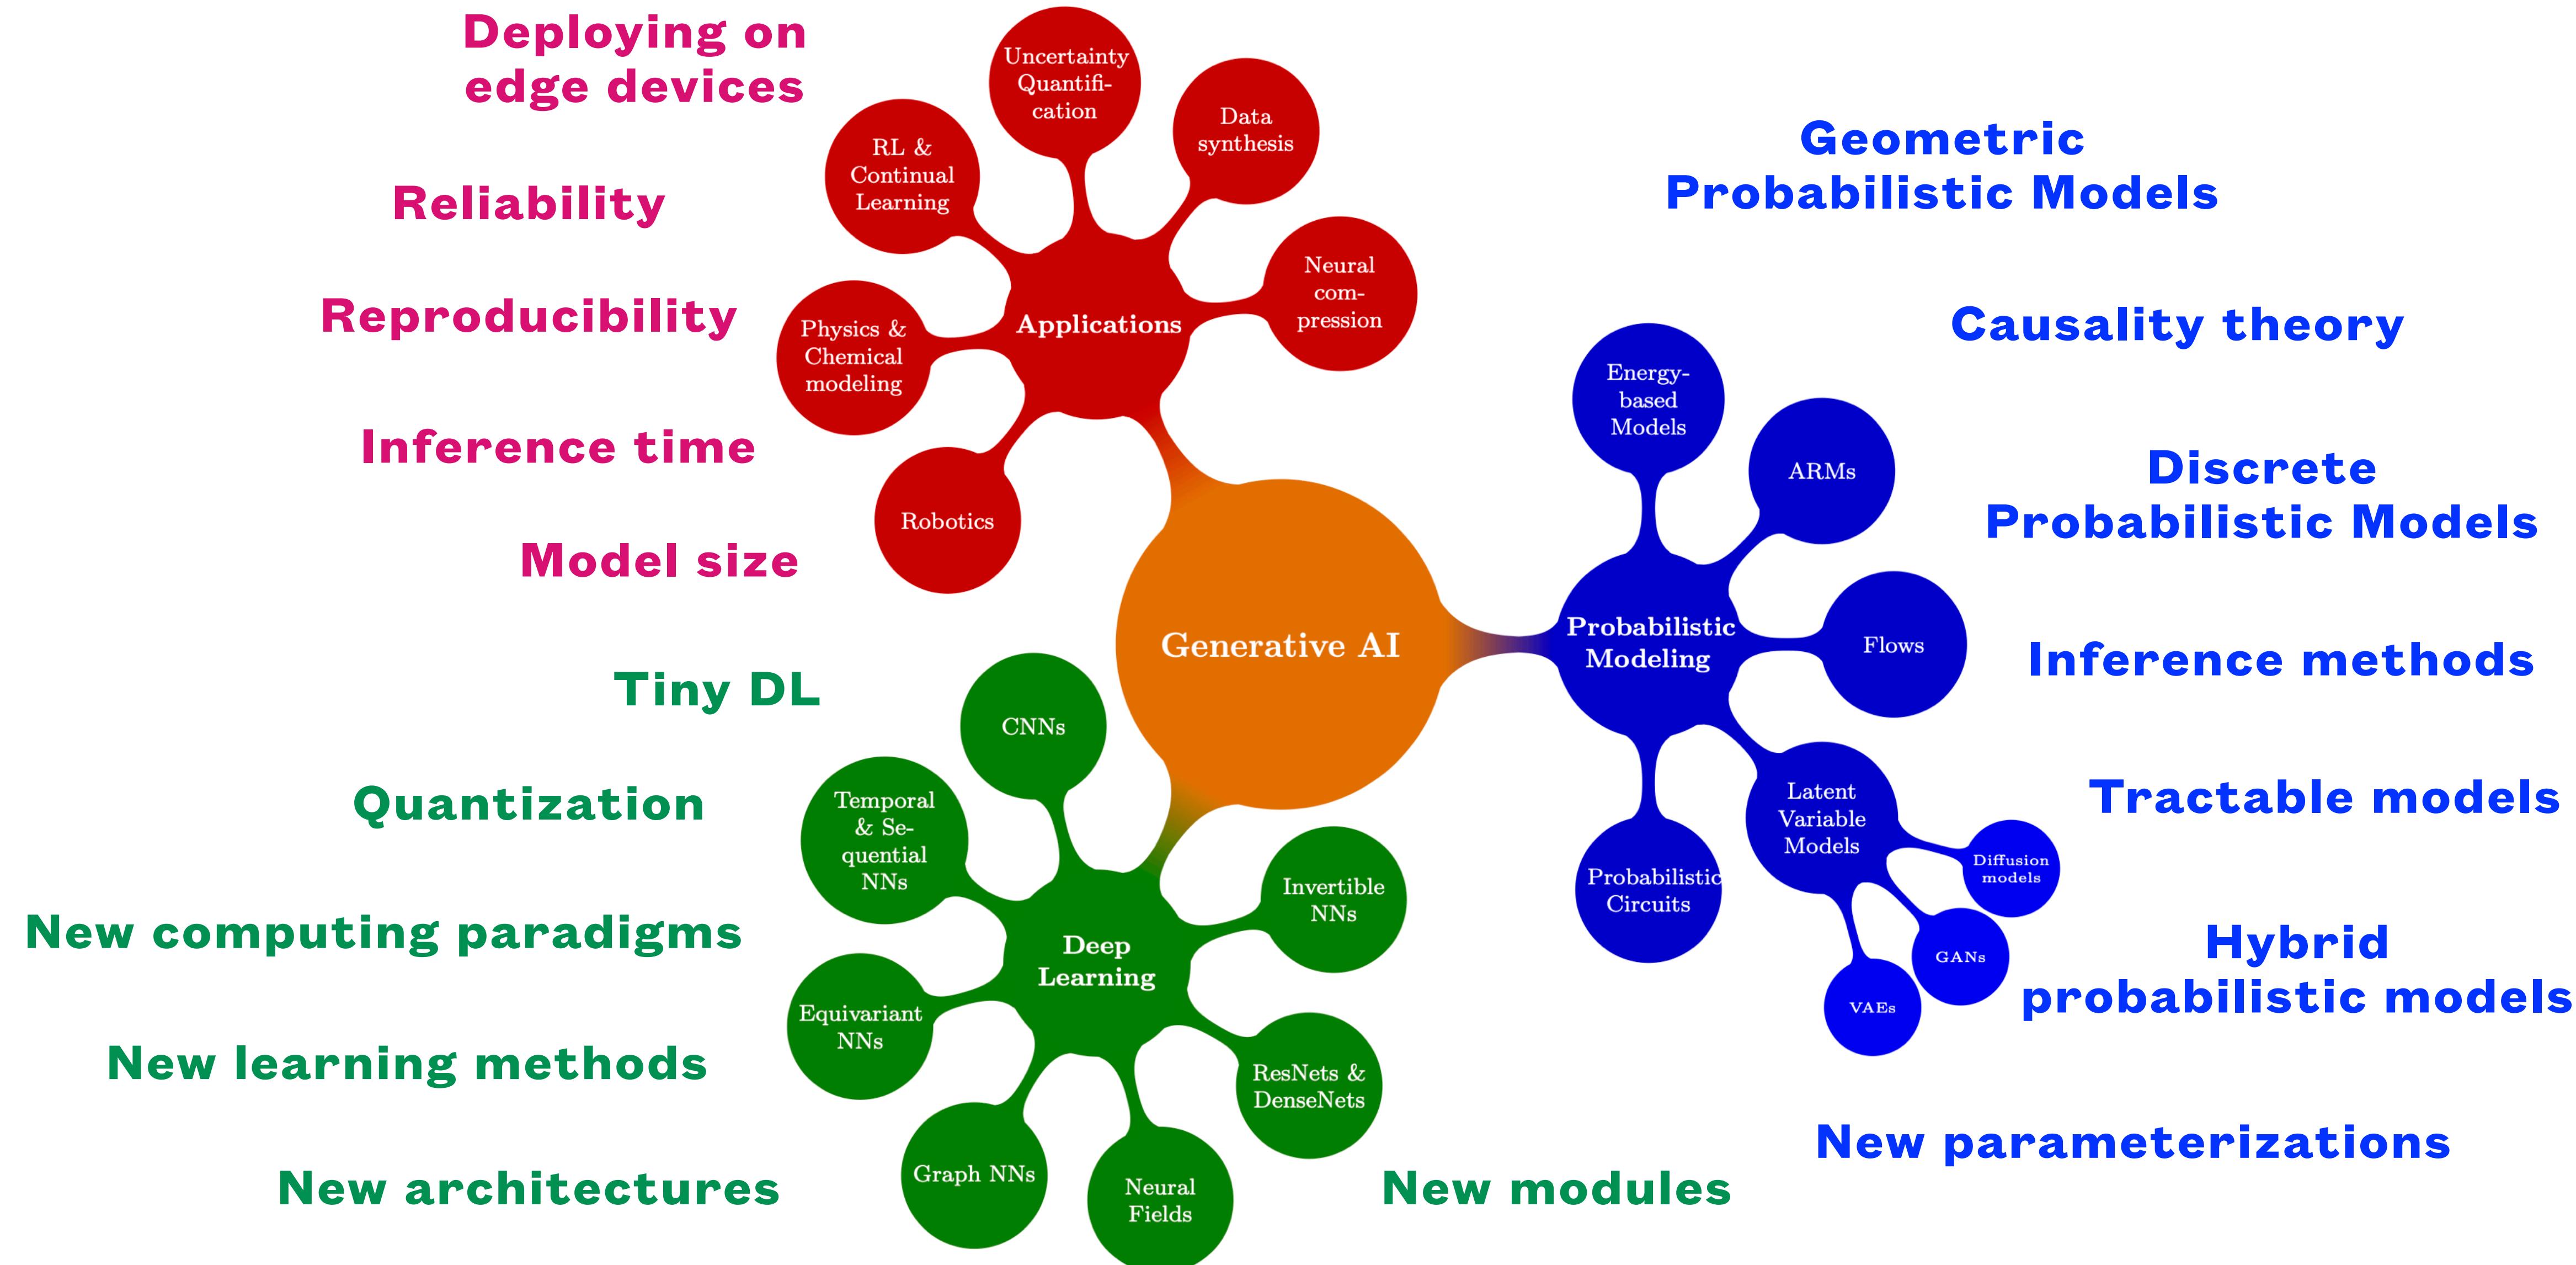
Challenges in Generative AI



Challenges in Generative AI



Challenges in Generative AI



Thank you!

Jakub M. Tomczak
assistant professor of AI
Vrije Universiteit Amsterdam

-  jmk.tomczak@gmail.com
 -  <https://jmtomczak.github.io/>
 -  <https://github.com/jmtomczak>
-