

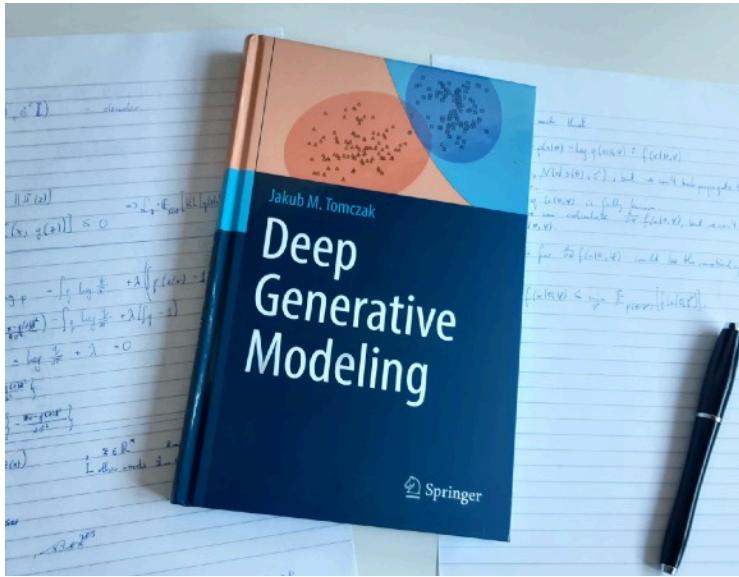
Deep Generative Modeling is a key to unlocking AI potential

Jakub M. Tomczak

MORE ABOUT DEEP GENERATIVE MODELING

Blog: <https://jmtomczak.github.io/blog.html>

Book: <https://link.springer.com/book/10.1007/978-3-030-93158-2>



What is **intelligence?**

What is **intelligence**?

...

INFORMATION, INTELLIGENCE AND ARTIFICIAL INTELLIGENCE

What is **intelligence**?

...

What is **artificial intelligence**?

INFORMATION, INTELLIGENCE AND ARTIFICIAL INTELLIGENCE

What is **intelligence**?

...

What is **artificial intelligence**?



INFORMATION, INTELLIGENCE AND ARTIFICIAL INTELLIGENCE

What is **intelligence**?

...

What is **artificial intelligence**?

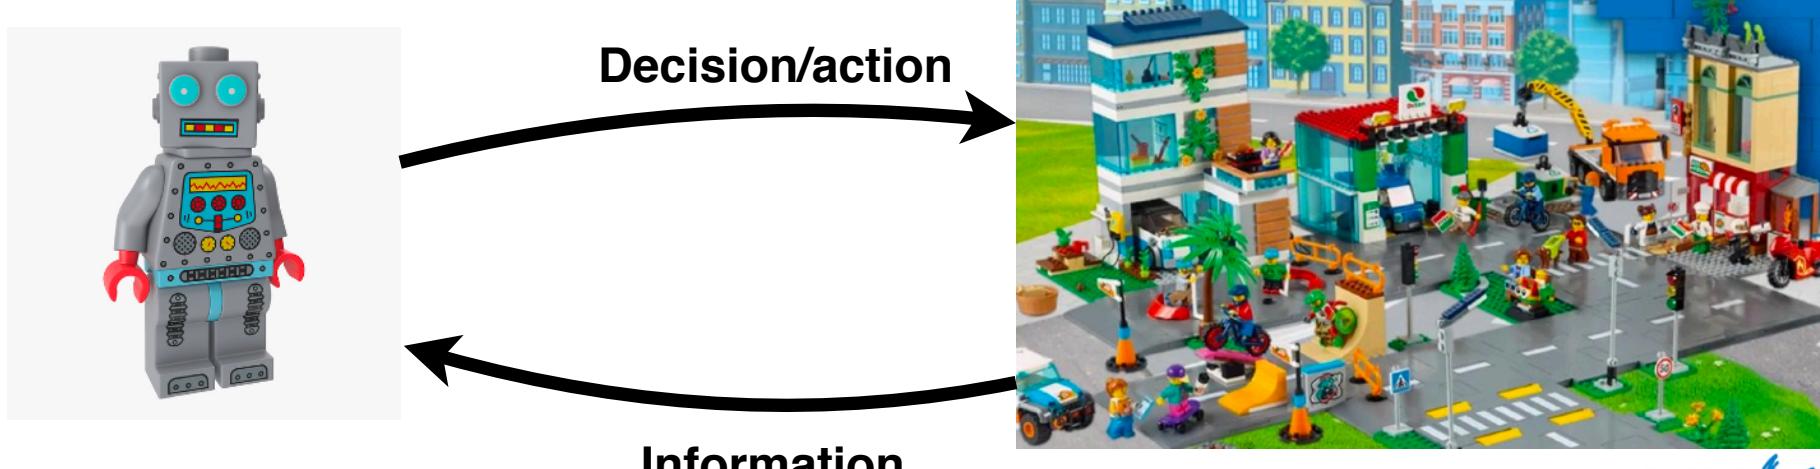


INFORMATION, INTELLIGENCE AND ARTIFICIAL INTELLIGENCE

What is **intelligence**?

...

What is **artificial intelligence**?

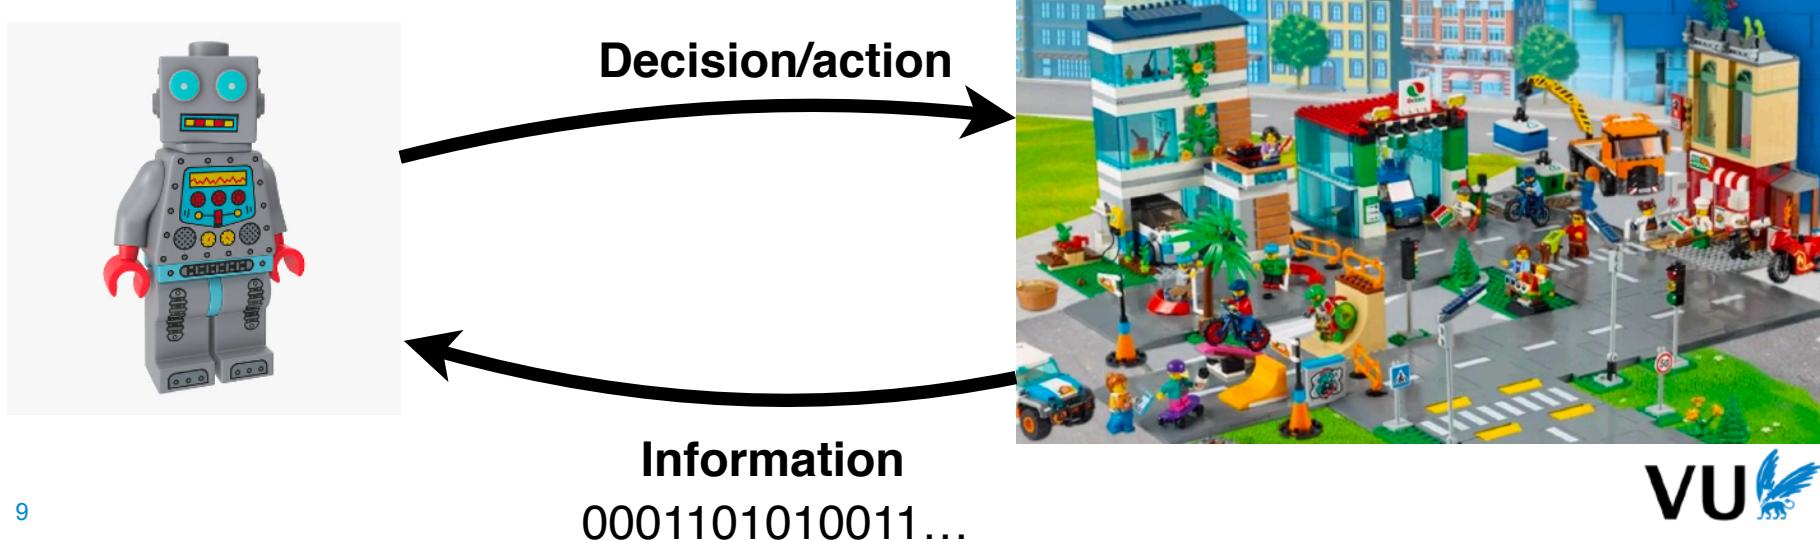


INFORMATION, INTELLIGENCE AND ARTIFICIAL INTELLIGENCE

What is **intelligence**?

...

What is **artificial intelligence**?



What is **artificial intelligence**?

- **Information** processing
- **Information** storing
- **Information** transmission



What is **artificial intelligence**?

- **Information** processing
- **Information** storing
- **Information** transmission
- **Decision** making



INFORMATION, INTELLIGENCE AND ARTIFICIAL INTELLIGENCE

What is **artificial intelligence**?

- **Information** processing
 - **Information** storing
 - **Information** transmission
 - **Decision** making
- } Learning
Knowledge representation
Models...



What is **artificial intelligence**?

- **Information** processing
 - **Information** storing
 - **Information** transmission
 - **Decision** making
- } Learning
Knowledge representation
Models...



The question is how to formalize the problem of AI?

INFORMATION, INTELLIGENCE AND ARTIFICIAL INTELLIGENCE

Information (a quick recap)



Claude Shannon

Information (a quick recap)

We have a random source of data x .



Claude Shannon

Information (a quick recap)

We have a random source of data x .

We can quantify the **uncertainty** of this source by calculating **the entropy**:



Claude Shannon

$$\mathbb{H}[x] = - \sum_x p(x) \log p(x)$$

Information (a quick recap)

We have a random source of data x .

We can quantify the **uncertainty** of this source by calculating **the entropy**:



Claude Shannon

$$\mathbb{H}[x] = - \sum_x p(x) \log p(x)$$

Entropy is max if all x 's are equiprobable.

Entropy is min if the probability of one value is 1.

Information (a quick recap)

We have a random source of data x .

We can quantify the **uncertainty** of this source by calculating **the entropy**:



Claude Shannon

$$\mathbb{H}[x] = - \sum_x p(x) \log p(x)$$

Optimal message length \approx the entropy.

Information (a quick recap)

We have two random sources: x and y .

We can quantify the **uncertainty** of them by calculating
the joint entropy:



Claude Shannon

$$\mathbb{H}[x, y] = - \sum_{x,y} p(x, y) \log p(x, y)$$

or **the conditional entropy**:

$$\mathbb{H}[y | x] = - \sum_{x,y} p(x, y) \log p(y | x)$$

Mutual Information (a quick recap)

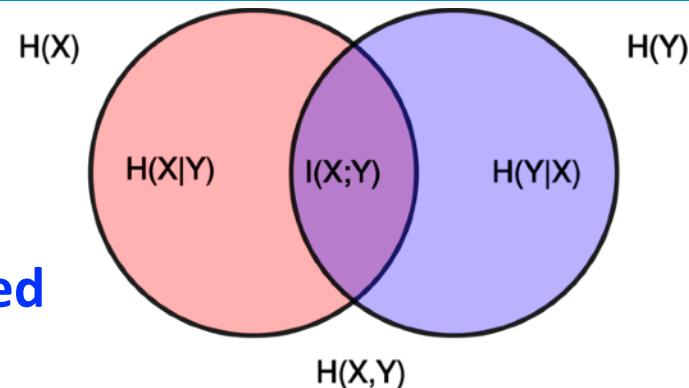
We have two random sources: x and y .

Mutual Information (a quick recap)

We have two random sources: x and y .

We can quantify how much **information is shared
by the two sources:**

$$I[x; y] = H[y] - H[y | x]$$

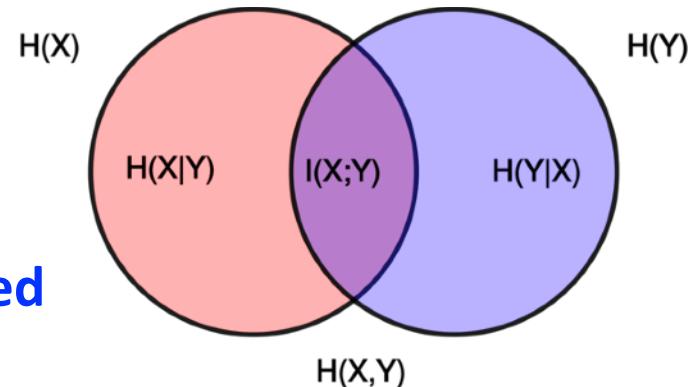


INFORMATION, INTELLIGENCE AND ARTIFICIAL INTELLIGENCE

Mutual Information (a quick recap)

We have two random sources: x and y .

We can quantify how much **information is shared by the two sources:**



$$I[x; y] = H[y] - H[y | x]$$

or **how much knowing one source reduces uncertainty about the other.**

We have two random sources: x (e.g., images) and y (e.g., decisions).

We have two random sources: x (e.g., images) and y (e.g., decisions).

We have also a **model** m (a representation of a world).

We have two random sources: x (e.g., images) and y (e.g., decisions).

We have also a **model** m (a representation of a world).

The **goal** of AI: **maximize** the **mutual information** between (x, y) and m :

$$\mathbb{I}[(x, y); m] = \mathbb{H}[x, y] - \mathbb{H}[x, y | m]$$

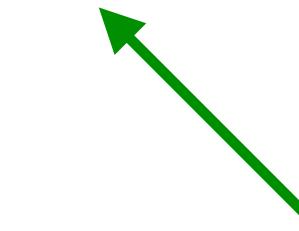
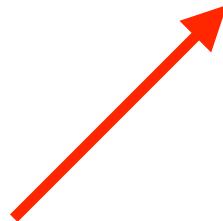
We have two random sources: x (e.g., images) and y (e.g., decisions).

We have also a **model** m (a representation of a world).

The **goal** of AI: **maximize the mutual information** between (x, y) and m :

$$\mathbb{I}[(x, y); m] = \mathbb{H}[x, y] - \mathbb{H}[x, y | m]$$

Entropy of the world
(model has no influence on that)



That's the “real” goal!

The **goal** of AI: **maximize** the **mutual information** between (x, y) and m (or minimize $\mathbb{H}[x, y \mid m]$, i.e., minimize uncertainty of the world):

$$\mathbb{H}[x, y \mid m] = \sum_{x,y,m} p(x, y, m) [\log p(y \mid x, m) + \log p(x \mid m)]$$

The **goal** of AI: **maximize** the **mutual information** between (x, y) and m (or minimize $\mathbb{H}[x, y \mid m]$, i.e., minimize uncertainty of the world):

$$\mathbb{H}[x, y \mid m] = \sum_{x, y, m} p(x, y, m) [\log p(y \mid x, m) + \log p(x \mid m)]$$



A model for
decision making



A model for
understanding
the world.

The **goal** of AI: **maximize the mutual information** between (x, y) and m (or minimize $\mathbb{H}[x, y \mid m]$, i.e., minimize uncertainty of the world).

In order to achieve that, AI should focus on learning **two models**:

- **A model for decision making:** $p(y \mid x, m)$
- **A model for understanding the world:** $p(x \mid m)$

WHAT HAPPENS IF WE LEARN ONLY DECISION MAKING

The bulk of AI is focused **only** on the decision making part!

WHAT HAPPENS IF WE LEARN ONLY DECISION MAKING

The bulk of AI is focused **only** on the decision making part!

Example: Let's say we have a model that is well trained.



$$p(y = \text{cat}|\mathbf{x}) = 0.90$$

$$p(y = \text{dog}|\mathbf{x}) = 0.05$$

$$p(y = \text{horse}|\mathbf{x}) = 0.05$$

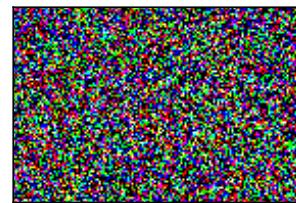
WHAT HAPPENS IF WE LEARN ONLY DECISION MAKING

The bulk of AI is focused **only** on the decision making part!

Example: Let's say we have a model that is well trained.



+



=



$$\begin{aligned} p(y = \text{cat}|\mathbf{x}) &= 0.90 \\ p(y = \text{dog}|\mathbf{x}) &= 0.05 \\ p(y = \text{horse}|\mathbf{x}) &= 0.05 \end{aligned}$$

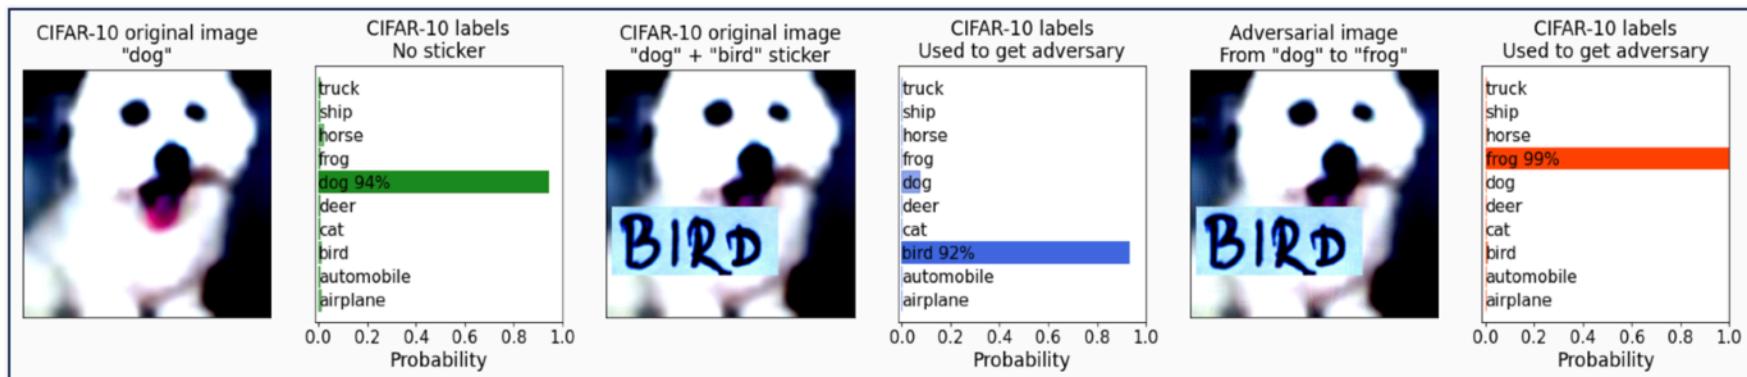
$$\begin{aligned} p(y = \text{cat}|\mathbf{x}) &= 0.05 \\ p(y = \text{dog}|\mathbf{x}) &= 0.05 \\ p(y = \text{horse}|\mathbf{x}) &= 0.90 \end{aligned}$$

But after adding a little noise it could fail completely...

IS LEARNING CLASSIFIERS ENOUGH?

Let's assume we have a perfectly trained neural net.

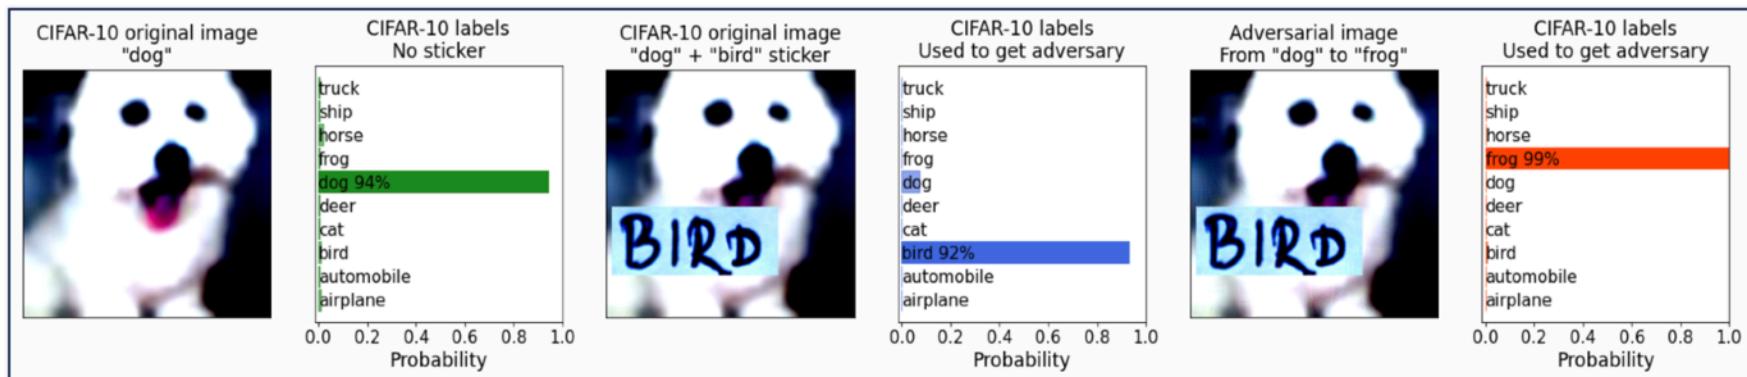
What happens if we add (adversarial) noise to an image?



IS LEARNING CLASSIFIERS ENOUGH?

Let's assume we have a perfectly trained neural net.

What happens if we add (adversarial) noise to an image?

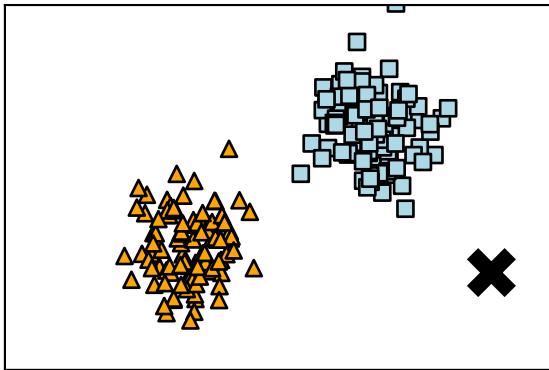


It fails completely...

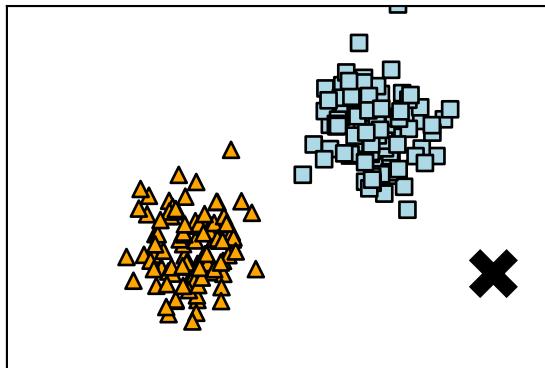
³⁴S. Fort, "Pixels still beat text: Attacking the OpenAI CLIP model with text patches and adversarial pixel perturbations", [\[Link\]](#)

DEEP GENERATIVE MODELING: WHY DO WE NEED IT?

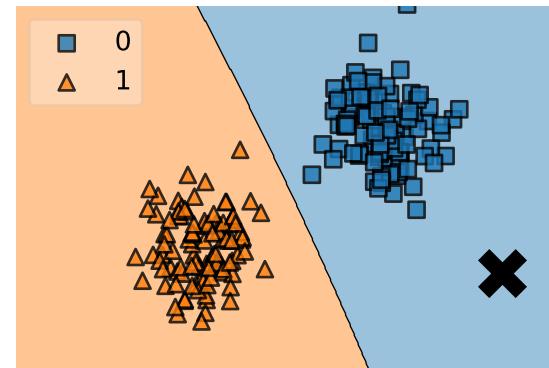
DEEP GENERATIVE MODELING: WHY DO WE NEED IT?



DEEP GENERATIVE MODELING: WHY DO WE NEED IT?



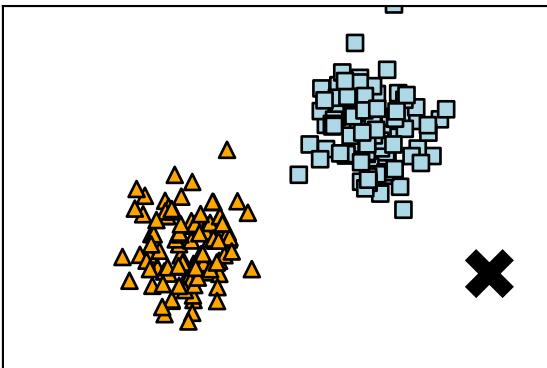
Data



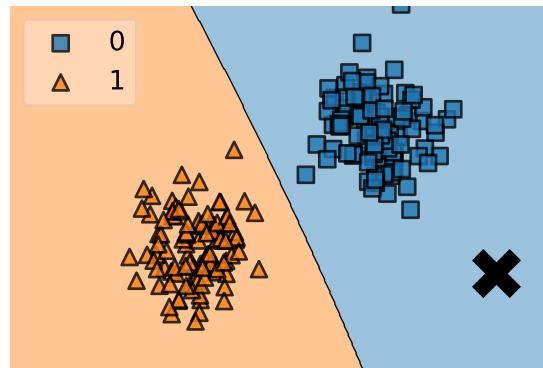
$p(y|\mathbf{x})$

$p(\text{blue}|\mathbf{x})$ is high
= certain decision!

DEEP GENERATIVE MODELING: WHY DO WE NEED IT?

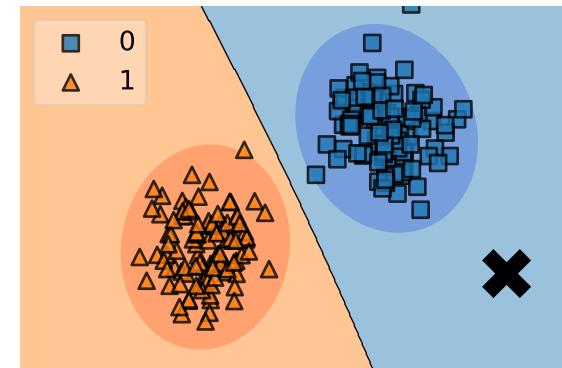


Data



$p(y|\mathbf{x})$

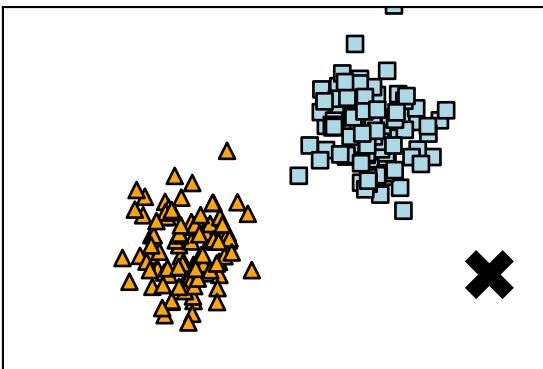
$p(\text{blue}|\mathbf{x})$ is high
= certain decision!



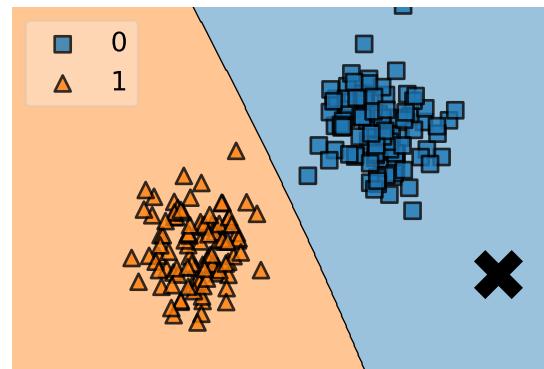
$p(\mathbf{x}, y) = p(y|\mathbf{x}) p(\mathbf{x})$

$p(\text{blue}|\mathbf{x})$ is high
and $p(\mathbf{x})$ is low
= uncertain decision!

DEEP GENERATIVE MODELING: WHY DO WE NEED IT?

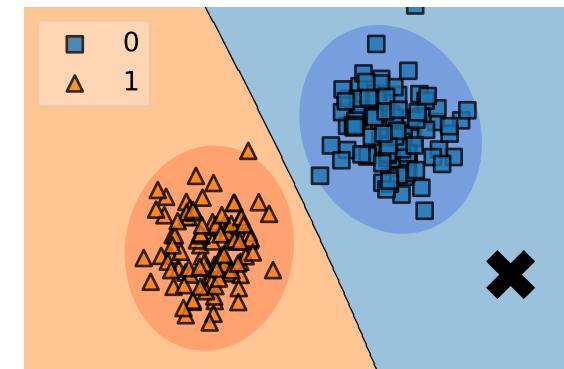


Data



$p(y|\mathbf{x})$

$p(\text{blue}|\mathbf{x})$ is high
= certain decision!

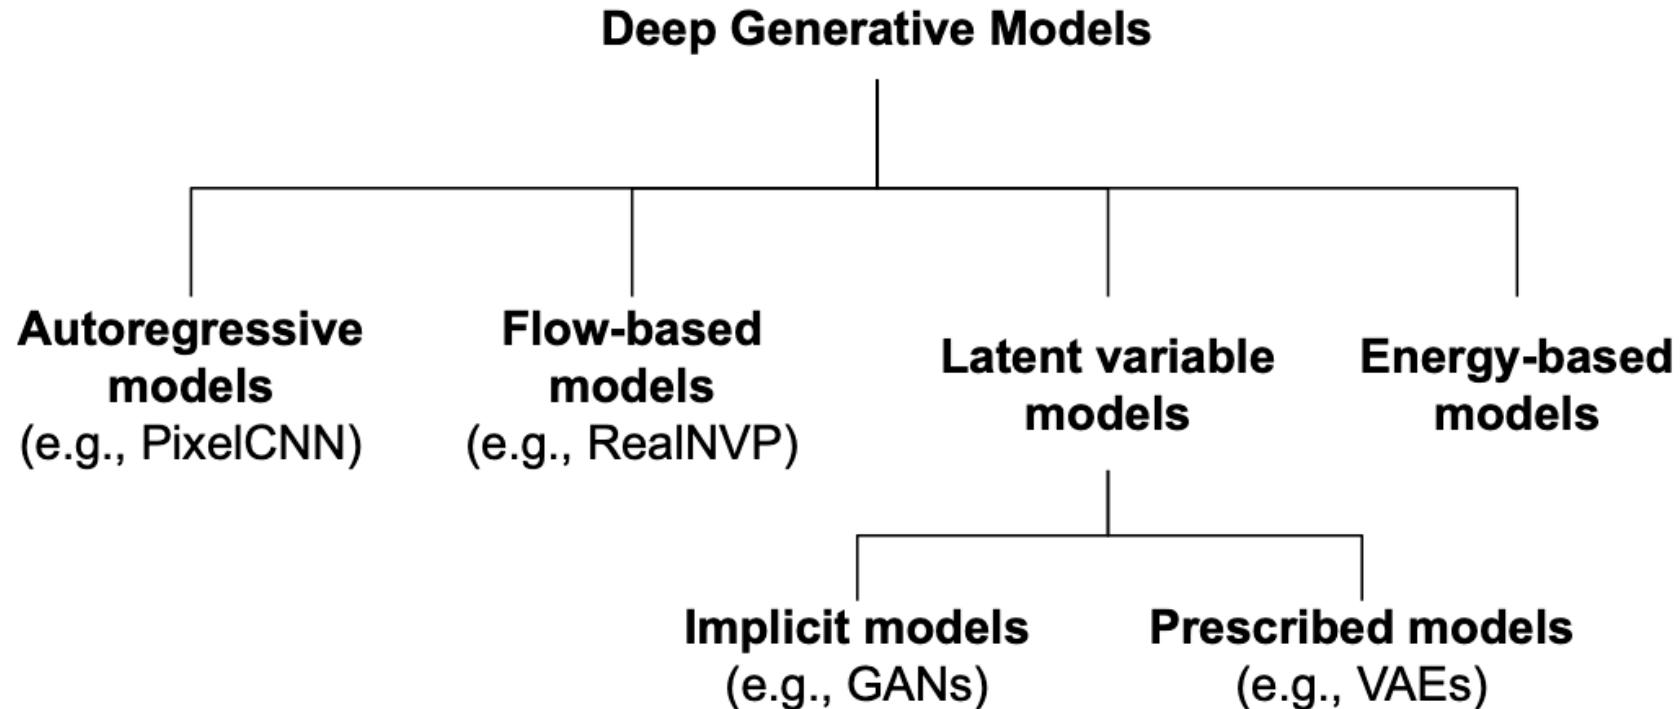


$p(\mathbf{x}, y) = p(y|\mathbf{x}) p(\mathbf{x})$

$p(\text{blue}|\mathbf{x})$ is high
and $p(\mathbf{x})$ is low
= uncertain decision!

Thus, learning the conditional is only a part of the story!
How can we learn $p(\mathbf{x})$?

DEEP GENERATIVE MODELING: HOW WE CAN FORMULATE IT?



DEEP GENERATIVE MODELING: HOW WE CAN FORMULATE IT?

Generative models	Training	Likelihood	Sampling	Compression	Representation
Autoregressive models	stable	exact	slow	lossless	no
Flow-based models	stable	exact	fast/slow	lossless	yes
Implicit models	unstable	no	fast	no	no
Prescribed models	stable	approximate	fast	lossy	yes
Energy-based models	stable	unnormalized	slow	rather not	yes

DEEP GENERATIVE MODELING: WHERE CAN WE USE IT?

“ i want to talk to you . ”

“i want to be with you . ”

“i do n’t want to be with you . ”

i do n’t want to be with you .

she did n’t want to be with him .

he was silent for a long moment .

he was silent for a moment .

it was quiet for a moment .

it was dark and cold .

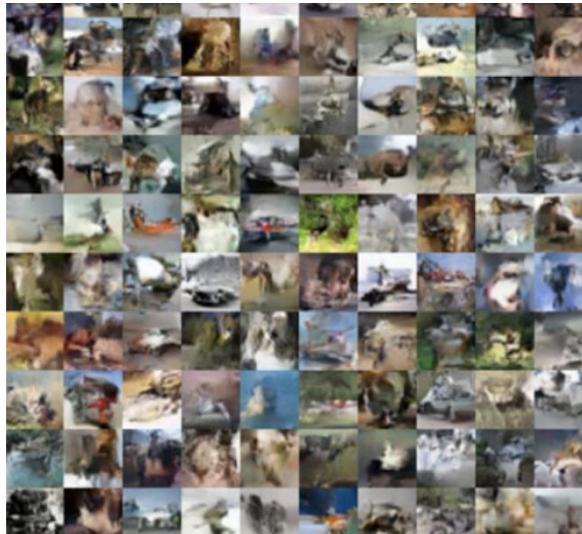
there was a pause .

it was my turn .

Text

Text synthesis (e.g., chatbots)

DEEP GENERATIVE MODELING: WHERE CAN WE USE IT?



Images

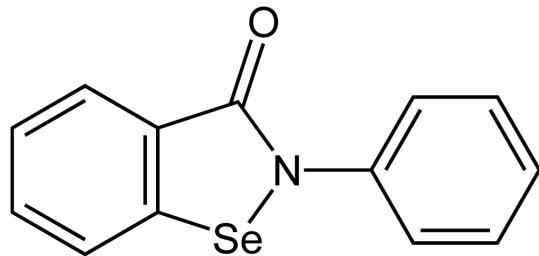


Audio

Speech synthesis (e.g., chatbots)

Deep fakes

DEEP GENERATIVE MODELING: WHERE CAN WE USE IT?



Graphs

Drug discovery

3D structure discovery

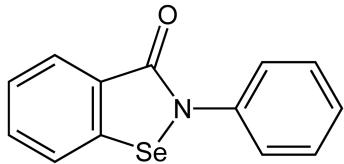
Molecular dynamics

DEEP GENERATIVE MODELING: WHERE CAN WE USE IT?

“ i want to talk to you . ”
“ i want to be with you . ”
“ i do n’t want to be with you . ”
“ i do n’t want to be with you . ”
she did n’t want to be with him .

he was silent for a long moment .
he was silent for a moment .
it was quiet for a moment .
it was dark and cold .
there was a pause .
it was my turn .

Text



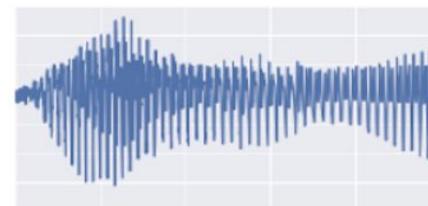
Graphs



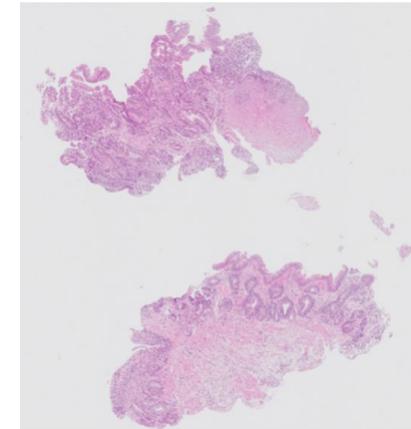
Reinforcement learning



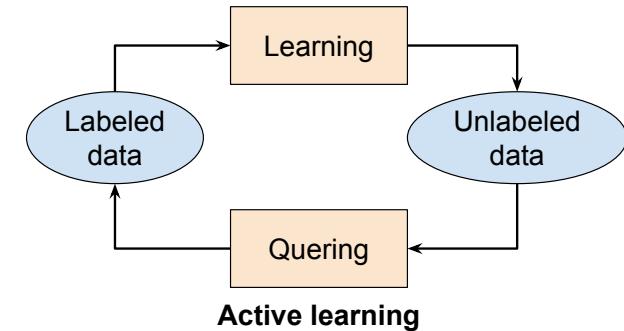
Images



Audio



Medical imaging



BASIC RULES FOR DEEP GENERATIVE MODELING

BASIC RULES FOR DEEP GENERATIVE MODELING

Two rules of probability theory:

- **Sum rule:**

$$p(x) = \sum_y p(x, y)$$

- **Product rule:**

$$p(x, y) = p(x | y)p(y)$$

or

$$p(x, y) = p(y | x)p(x)$$

BASIC RULES FOR DEEP GENERATIVE MODELING

The objective (typically): **the log-likelihood function**

Given **iid data**: $\mathcal{D} = \{x_1, x_2, \dots, x_N\}$

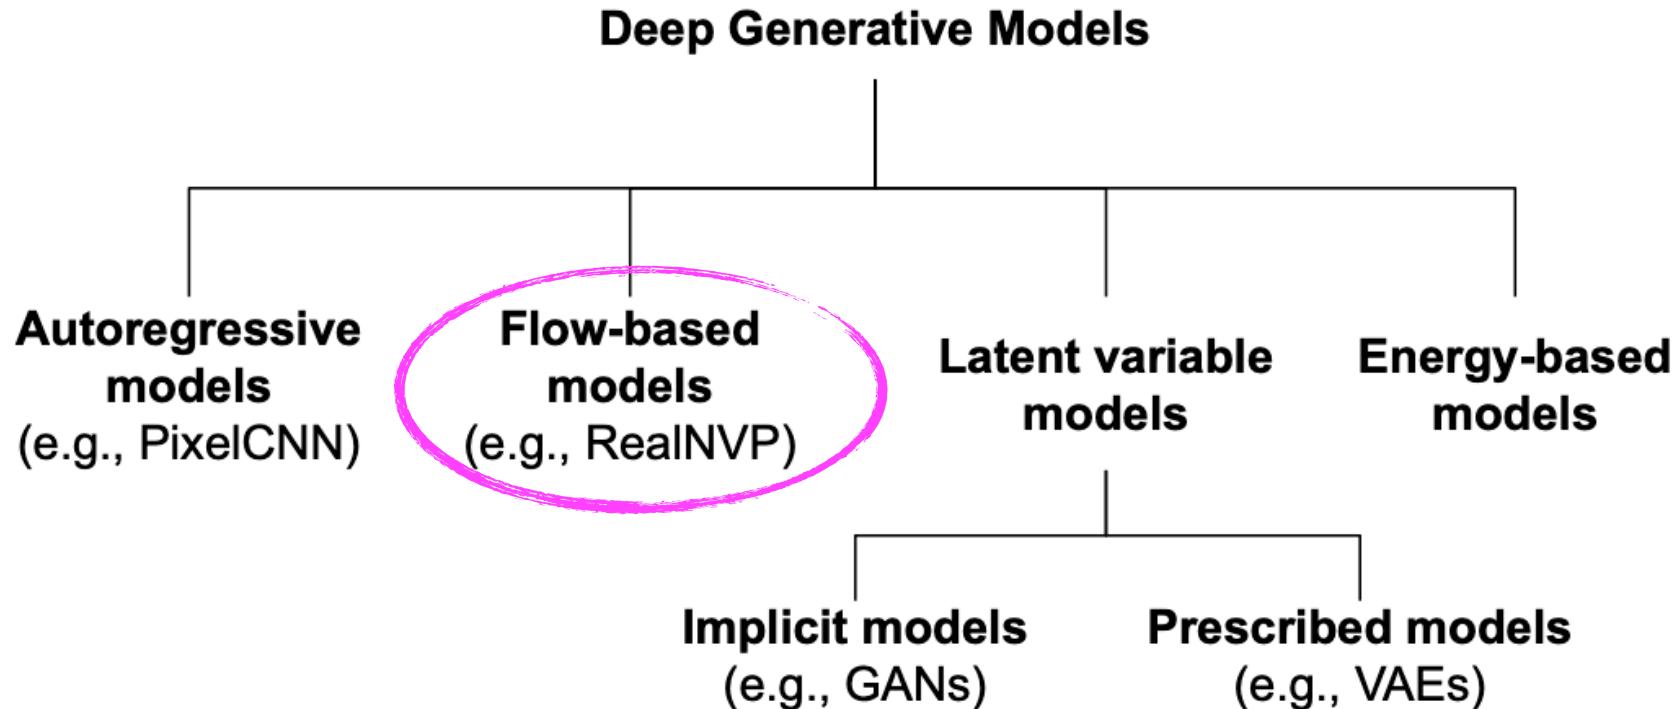
A **model** $p(x | \theta)$.

The log-likelihood function is:

$$\ln p(\mathcal{D} | \theta) = \ln \prod_{n=1}^N p(x_n | \theta)$$

$$= \sum_{n=1}^N \ln p(x_n | \theta)$$

DEEP GENERATIVE MODELING: HOW WE CAN FORMULATE IT?



FLows (Flow-based Models)

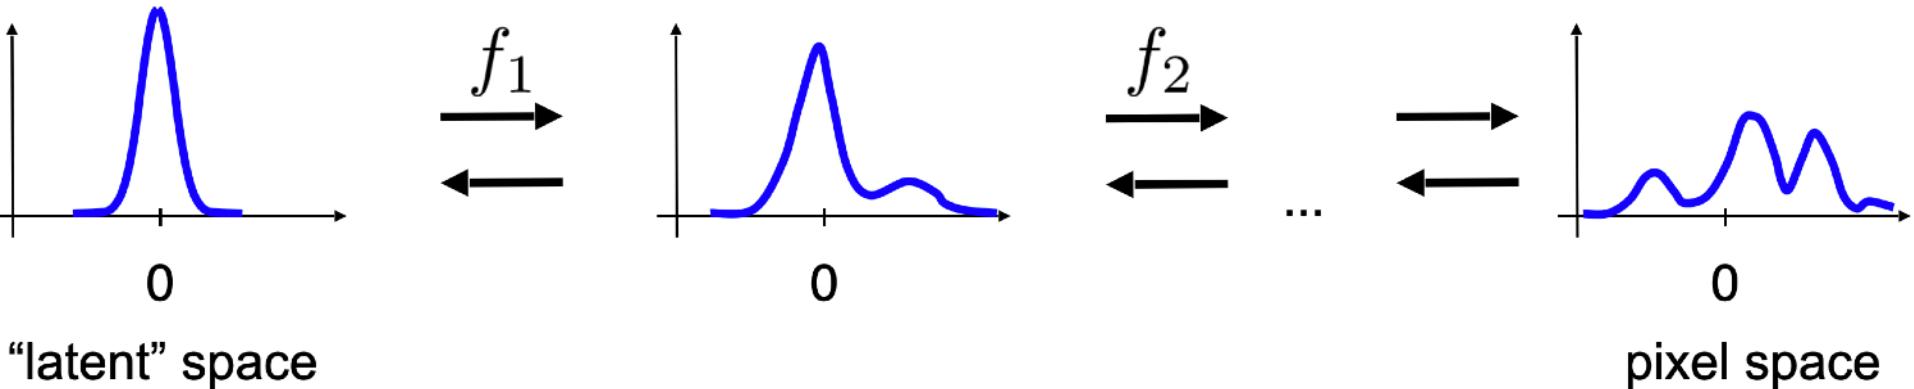
We change a random variable \mathbf{x} to another random variable \mathbf{z} using **invertible** transformations, $\mathbf{x}, \mathbf{z} \in \mathbb{R}^D$:

$$p(\mathbf{x}) = \pi(\mathbf{z}_0 = f^{-1}(\mathbf{x})) \prod_{i=1}^K \left| \mathbf{J}_{f_i}(z_{i-1}) \right|^{-1}$$

FLows (Flow-based Models)

We change a random variable \mathbf{x} to another random variable \mathbf{z} using **invertible** transformations, $\mathbf{x}, \mathbf{z} \in \mathbb{R}^D$:

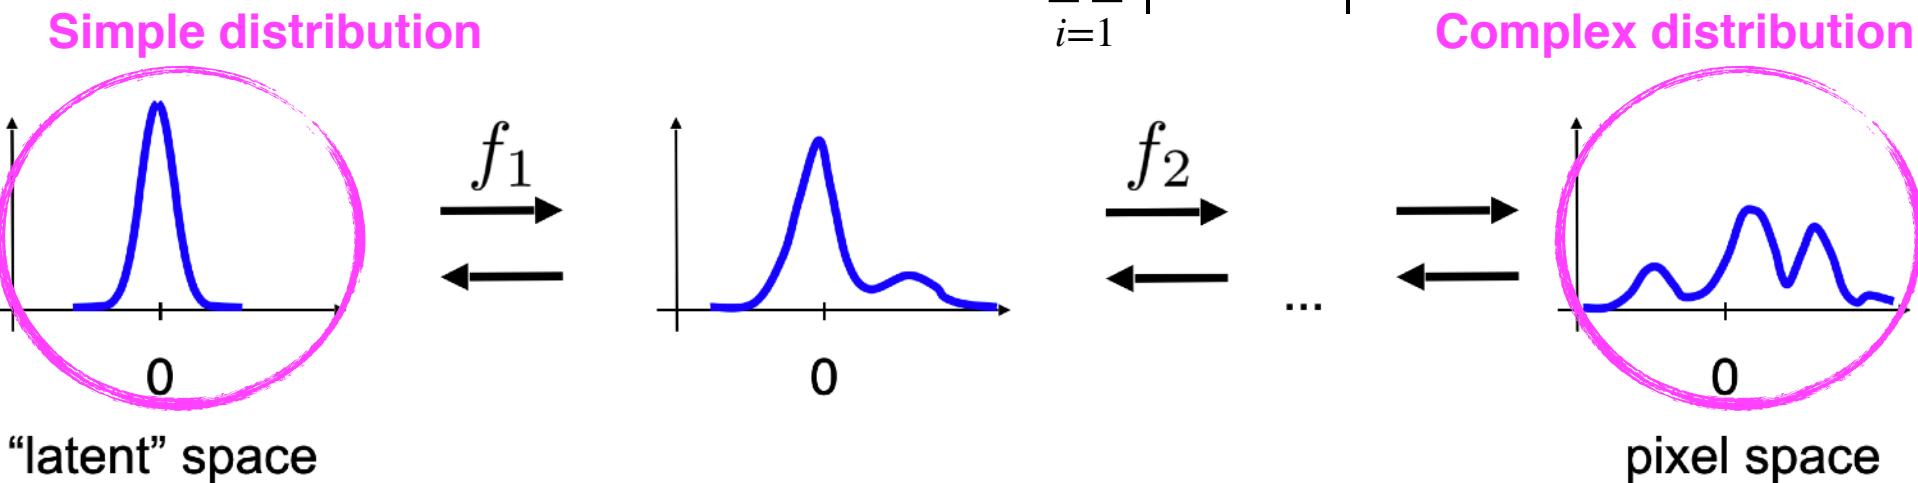
$$p(\mathbf{x}) = \pi(\mathbf{z}_0 = f^{-1}(\mathbf{x})) \prod_{i=1}^K \left| \mathbf{J}_{f_i}(z_{i-1}) \right|^{-1}$$



FLows (FLOW-BASED MODELS)

We change a random variable \mathbf{x} to another random variable \mathbf{z} using **invertible** transformations, $\mathbf{x}, \mathbf{z} \in \mathbb{R}^D$:

$$p(\mathbf{x}) = \pi(\mathbf{z}_0 = f^{-1}(\mathbf{x})) \prod_{i=1}^K \left| \mathbf{J}_{f_i}(z_{i-1}) \right|^{-1}$$

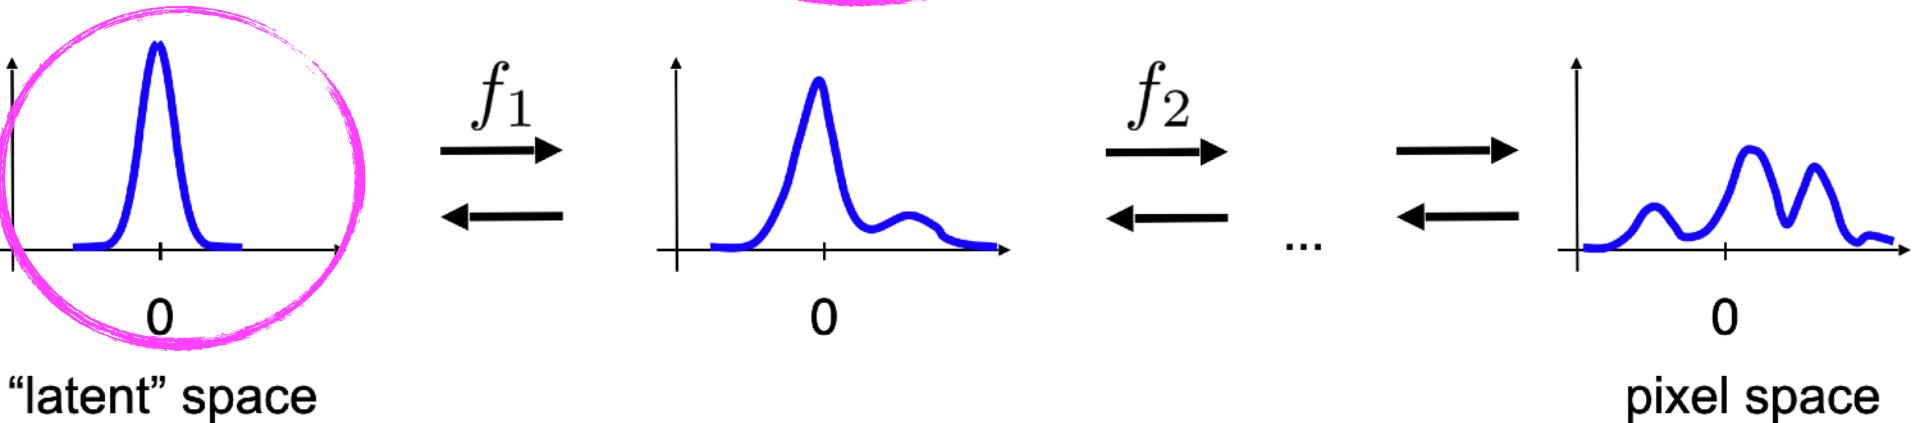


FLows (FLOW-BASED MODELS)

We change a random variable \mathbf{x} to another random variable \mathbf{z} using **invertible** transformations, $\mathbf{x}, \mathbf{z} \in \mathbb{R}^D$:

Known, e.g., Gaussian

$$p(\mathbf{x}) = \pi(\mathbf{z}_0 = f^{-1}(\mathbf{x})) \prod_{i=1}^K \left| \mathbf{J}_{f_i}(z_{i-1}) \right|^{-1}$$

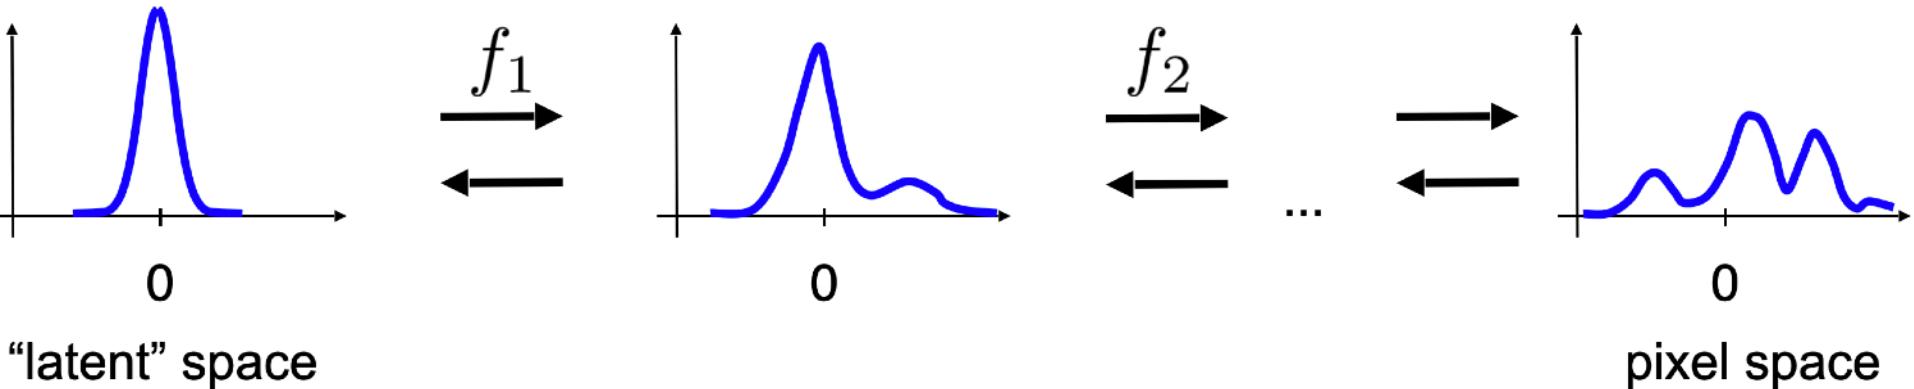


FLows (Flow-based Models)

We change a random variable \mathbf{x} to another random variable \mathbf{z} using **invertible** transformations, $\mathbf{x}, \mathbf{z} \in \mathbb{R}^D$:

Jacobian must be tractable

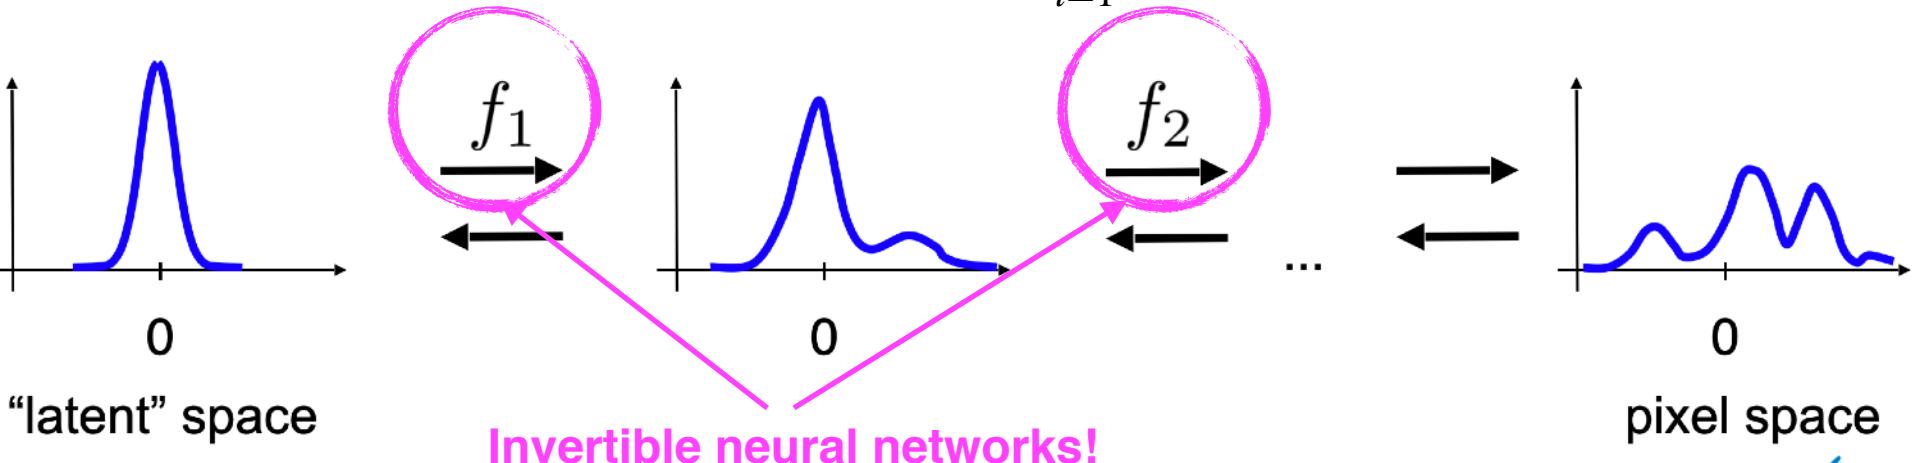
$$p(\mathbf{x}) = \pi(\mathbf{z}_0 = f^{-1}(\mathbf{x})) \prod_{i=1}^K \left| \mathbf{J}_{f_i}(z_{i-1}) \right|^{-1}$$



FLows (Flow-based Models)

We change a random variable \mathbf{x} to another random variable \mathbf{z} using **invertible** transformations, $\mathbf{x}, \mathbf{z} \in \mathbb{R}^D$:

$$p(\mathbf{x}) = \pi(\mathbf{z}_0 = f^{-1}(\mathbf{x})) \prod_{i=1}^K \left| \mathbf{J}_{f_i}(z_{i-1}) \right|^{-1}$$



FLows (Flow-based Models)

We change a random variable \mathbf{x} to another random variable \mathbf{z} using **invertible** transformations, $\mathbf{x}, \mathbf{z} \in \mathbb{R}^D$:

$$p(\mathbf{x}) = \pi(\mathbf{z}_0 = f^{-1}(\mathbf{x})) \prod_{i=1}^K \left| \mathbf{J}_{f_i}(z_{i-1}) \right|^{-1}$$

Training objective:

$$\ln p(\mathbf{x}) = \ln \pi(\mathbf{z}_0 = f^{-1}(\mathbf{x})) - \sum_{i=1}^K \ln \left| \mathbf{J}_{f_i}(z_{i-1}) \right|$$

FLows (Flow-based Models): INVERTIBLE LAYERS

Two main components

1) Coupling layer:

$$\begin{aligned}\mathbf{y}_a &= \mathbf{x}_a \\ \mathbf{y}_b &= \exp\left(s\left(\mathbf{x}_a\right)\right) \odot \mathbf{x}_b + t\left(\mathbf{x}_a\right)\end{aligned}$$

is invertible by design:

$$\mathbf{x}_b = (\mathbf{y}_b - t(\mathbf{y}_a)) \odot \exp(-s(\mathbf{y}_a))$$

$$\mathbf{x}_a = \mathbf{y}_a$$

2) Permutation layer

FLows (Flow-based Models): INVERTIBLE LAYERS

Two main components

1) Coupling layer:

$$\begin{aligned}\mathbf{y}_a &= \mathbf{x}_a \\ \mathbf{y}_b &= \exp\left(s(\mathbf{x}_a)\right) \odot \mathbf{x}_b + t(\mathbf{x}_a)\end{aligned}$$

is invertible by design:

$$\mathbf{x}_b = (\mathbf{y}_b - t(\mathbf{y}_a)) \odot \exp(-s(\mathbf{y}_a))$$

$$\mathbf{x}_a = \mathbf{y}_a$$

Jacobian is tractable!

$$\det(\mathbf{J}) = \prod_{j=1}^{D-d} \exp\left(s(\mathbf{x}_a)\right)_j = \exp\left(\sum_{j=1}^{D-d} s(\mathbf{x}_a)_j\right)$$

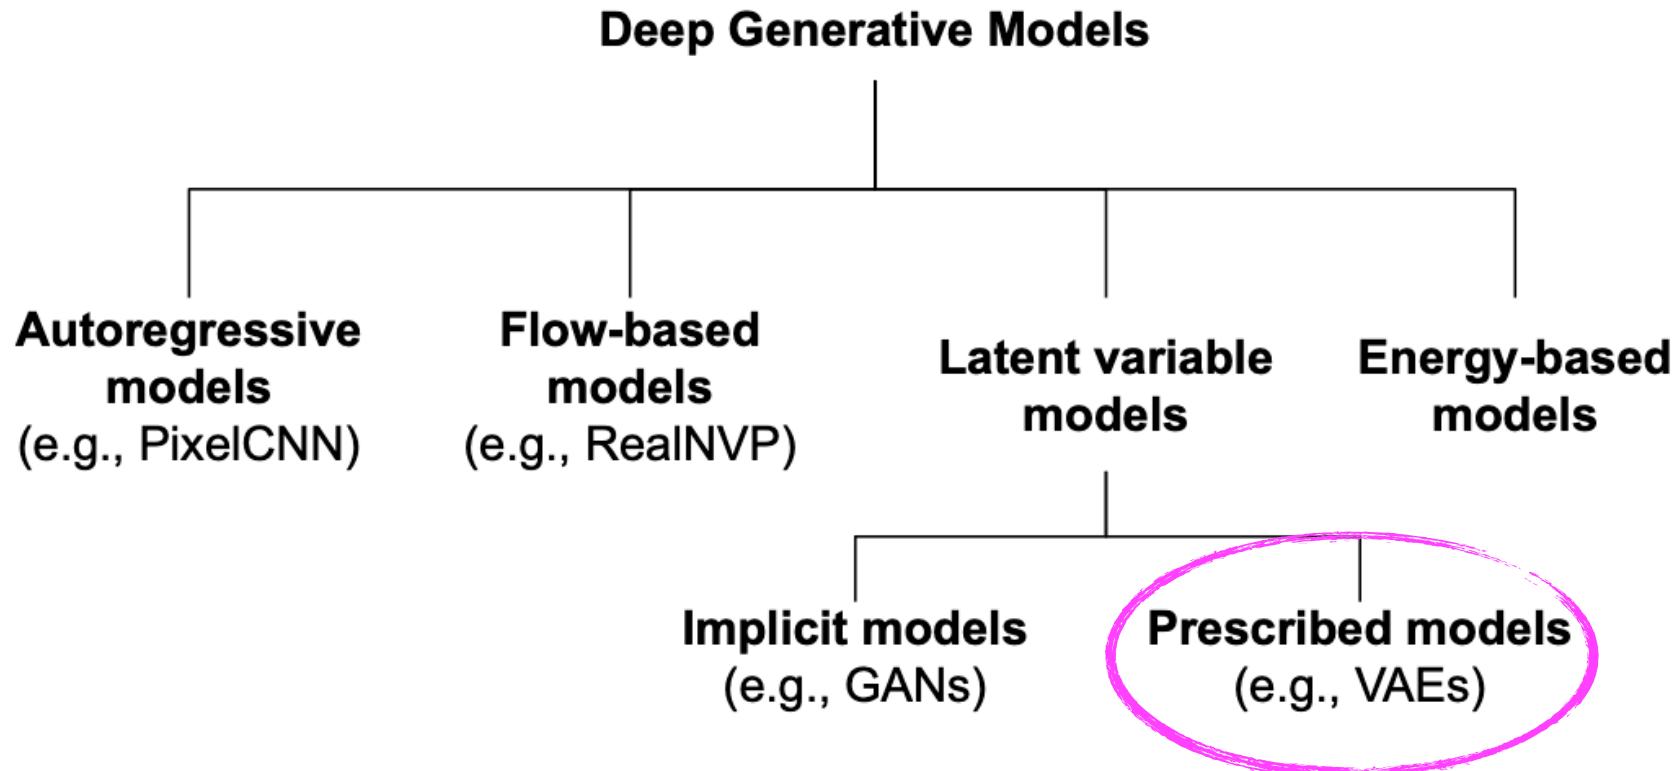
2) Permutation layer

$$\det(\mathbf{J}) = 1$$

FLows (FLOW-BASED MODELS)



DEEP GENERATIVE MODELING: HOW WE CAN FORMULATE IT?

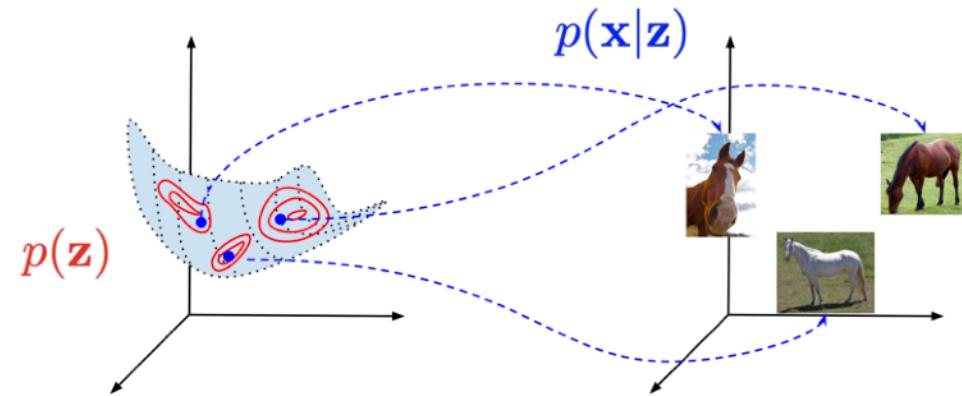


VARIATIONAL AUTO-ENCODERS

Let's consider a **latent variable model** where we distinguish:

- **latent variables** $\mathbf{z} \in \mathcal{Z}^M$
- **observable variables** $\mathbf{x} \in \mathcal{X}^D$

Latent variables lie on a **low-dimensional manifold**.



Generative process:

1. $\mathbf{z} \sim p(\mathbf{z})$
2. $\mathbf{x} \sim p(\mathbf{x} | \mathbf{z})$

VARIATIONAL AUTO-ENCODERS

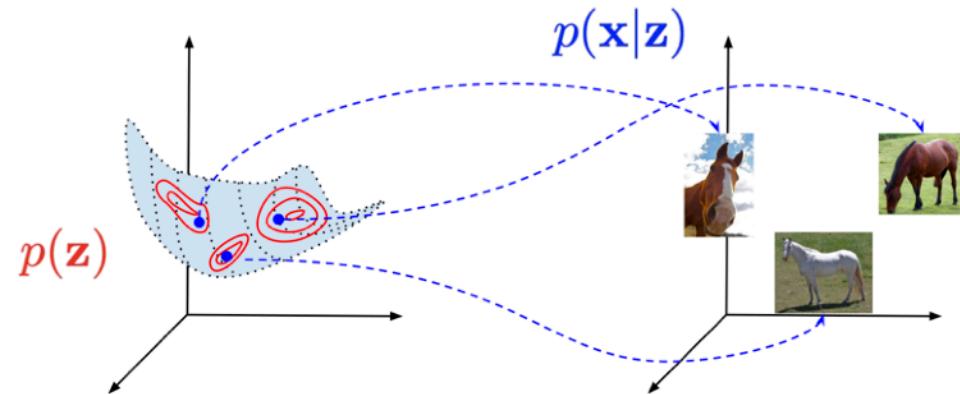
Let's consider a **latent variable model** where we distinguish:

- **latent variables** $\mathbf{z} \in \mathcal{Z}^M$
- **observable variables** $\mathbf{x} \in \mathcal{X}^D$

Latent variables lie on a **low-dimensional manifold**.

The objective function:

$$\ln p(\mathbf{x}) = \ln \int p(\mathbf{x} | \mathbf{z}) p(\mathbf{z}) d\mathbf{z}$$



Generative process:

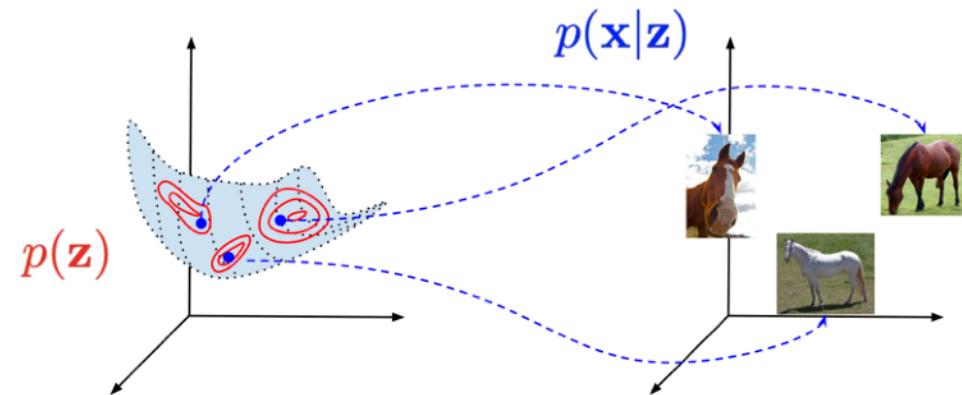
1. $\mathbf{z} \sim p(\mathbf{z})$
2. $\mathbf{x} \sim p(\mathbf{x} | \mathbf{z})$

VARIATIONAL AUTO-ENCODERS

Let's consider a **latent variable model** where we distinguish:

- **latent variables** $\mathbf{z} \in \mathcal{Z}^M$
- **observable variables** $\mathbf{x} \in \mathcal{X}^D$

Latent variables lie on a **low-dimensional manifold**.



Generative process:

1. $\mathbf{z} \sim p(\mathbf{z})$
2. $\mathbf{x} \sim p(\mathbf{x} | \mathbf{z})$

The objective function:

$$\ln p(\mathbf{x}) = \ln \int p(\mathbf{x} | \mathbf{z}) p(\mathbf{z}) d\mathbf{z}$$

The integral is intractable...

VARIATIONAL AUTO-ENCODERS

$$\ln p(\mathbf{x}) \geq \mathbb{E}_{\mathbf{z} \sim q_\phi(\mathbf{z}|\mathbf{x})} [\ln p(\mathbf{x}|\mathbf{z})] - \mathbb{E}_{\mathbf{z} \sim q_\phi(\mathbf{z}|\mathbf{x})} [\ln q_\phi(\mathbf{z}|\mathbf{x}) - \ln p(\mathbf{z})]$$

ELBO: Evidence Lower Bound

VARIATIONAL AUTO-ENCODERS

$$\ln p(\mathbf{x}) \geq \mathbb{E}_{\mathbf{z} \sim q_\phi(\mathbf{z} | \mathbf{x})} [\ln p(\mathbf{x} | \mathbf{z})] - \mathbb{E}_{\mathbf{z} \sim q_\phi(\mathbf{z} | \mathbf{x})} [\ln q_\phi(\mathbf{z} | \mathbf{x}) - \ln p(\mathbf{z})]$$

We consider **amortized inference**: $q_\phi(\mathbf{z} | \mathbf{x})$

In other words, a single parameterization for each new input \mathbf{x} .

VARIATIONAL AUTO-ENCODERS

$$\ln p(\mathbf{x}) \geq \mathbb{E}_{\mathbf{z} \sim q_\phi(\mathbf{z} | \mathbf{x})} [\ln p(\mathbf{x} | \mathbf{z})] - \mathbb{E}_{\mathbf{z} \sim q_\phi(\mathbf{z} | \mathbf{x})} [\ln q_\phi(\mathbf{z} | \mathbf{x}) - \ln p(\mathbf{z})]$$

We consider **amortized inference**: $q_\phi(\mathbf{z} | \mathbf{x})$

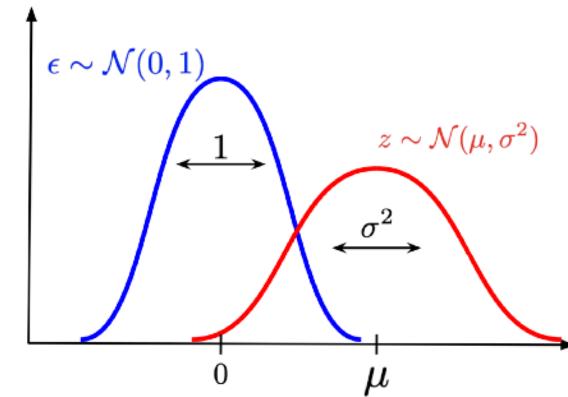
In other words, a single parameterization for each new input \mathbf{x} .

Moreover, we use **reparameterization trick**:

Every Gaussian variable could be defined as:

$$z = \mu + \sigma \cdot \varepsilon$$

where $\varepsilon \sim \mathcal{N}(0, 1)$



VARIATIONAL AUTO-ENCODERS

$$\ln p(\mathbf{x}) \geq \mathbb{E}_{\mathbf{z} \sim q_\phi(\mathbf{z} | \mathbf{x})} [\ln p(\mathbf{x} | \mathbf{z})] - \mathbb{E}_{\mathbf{z} \sim q_\phi(\mathbf{z} | \mathbf{x})} [\ln q_\phi(\mathbf{z} | \mathbf{x}) - \ln p(\mathbf{z})]$$

We consider **amortized inference**: $q_\phi(\mathbf{z} | \mathbf{x})$

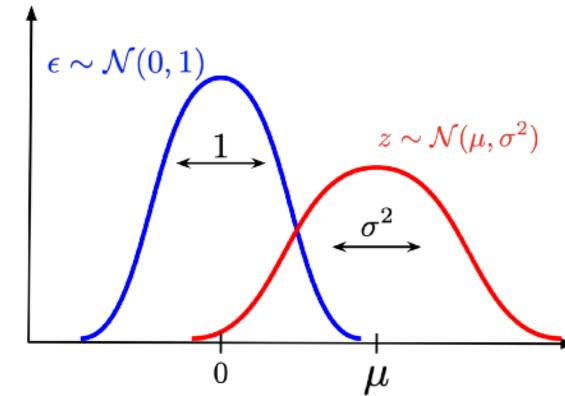
In other words, a single parameterization for each new input \mathbf{x} .

Moreover, we use **reparameterization trick**:

It reduces the variance of the gradients.

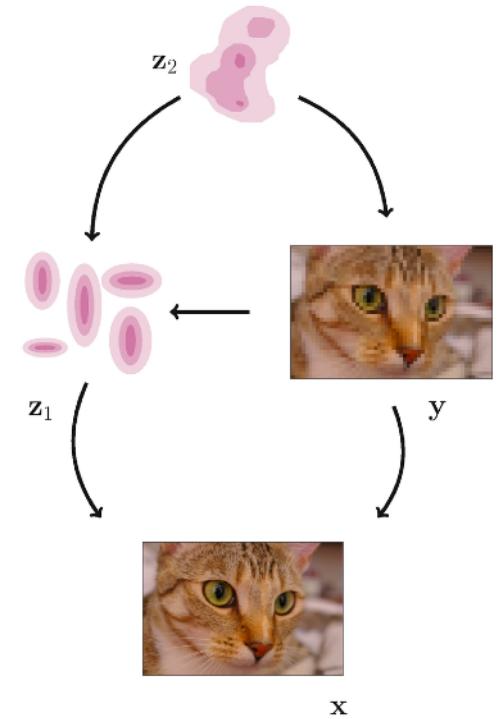
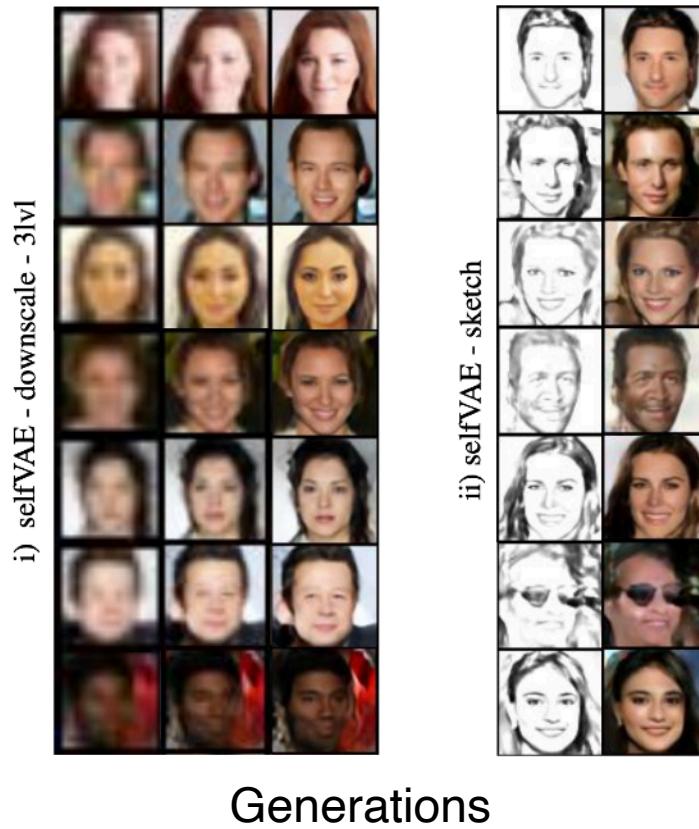
It allows to get randomness outside \mathbf{z} .

$$z = \mu + \sigma \cdot \epsilon$$



⁶⁷ Kingma, D.P., and Welling, M.. "Auto-encoding variational bayes." *ICLR 2014*

HIERARCHICAL VARIATIONAL AUTO-ENCODERS

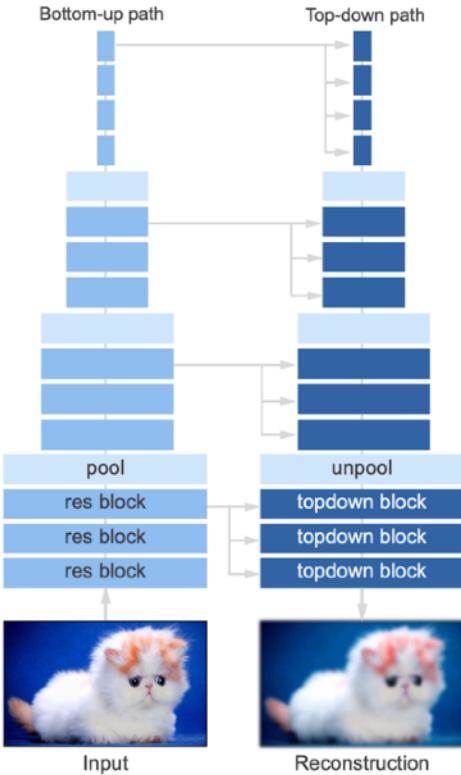


Hierarchical VAE

TOP-DOWN VARIATIONAL AUTO-ENCODERS



Generations



Very Deep VAE

⁶⁹ Child, R. "Very Deep VAEs Generalize Autoregressive Models and Can Outperform Them on Images." *ICLR 2021*

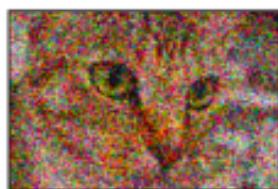
DIFFUSION-BASED DEEP GENERATIVE MODELS

$$q_{\phi}(\mathbf{z}_i \mid \mathbf{z}_{i-1}) = \mathcal{N}\left(\mathbf{z}_i \mid \sqrt{1 - \beta_i} \mathbf{z}_{i-1}, \beta_i \mathbf{I}\right)$$

Forward diffusion (FIXED!)



\mathbf{x}



\mathbf{z}_1



\mathbf{z}_2



\mathbf{z}_3



\mathbf{z}_4

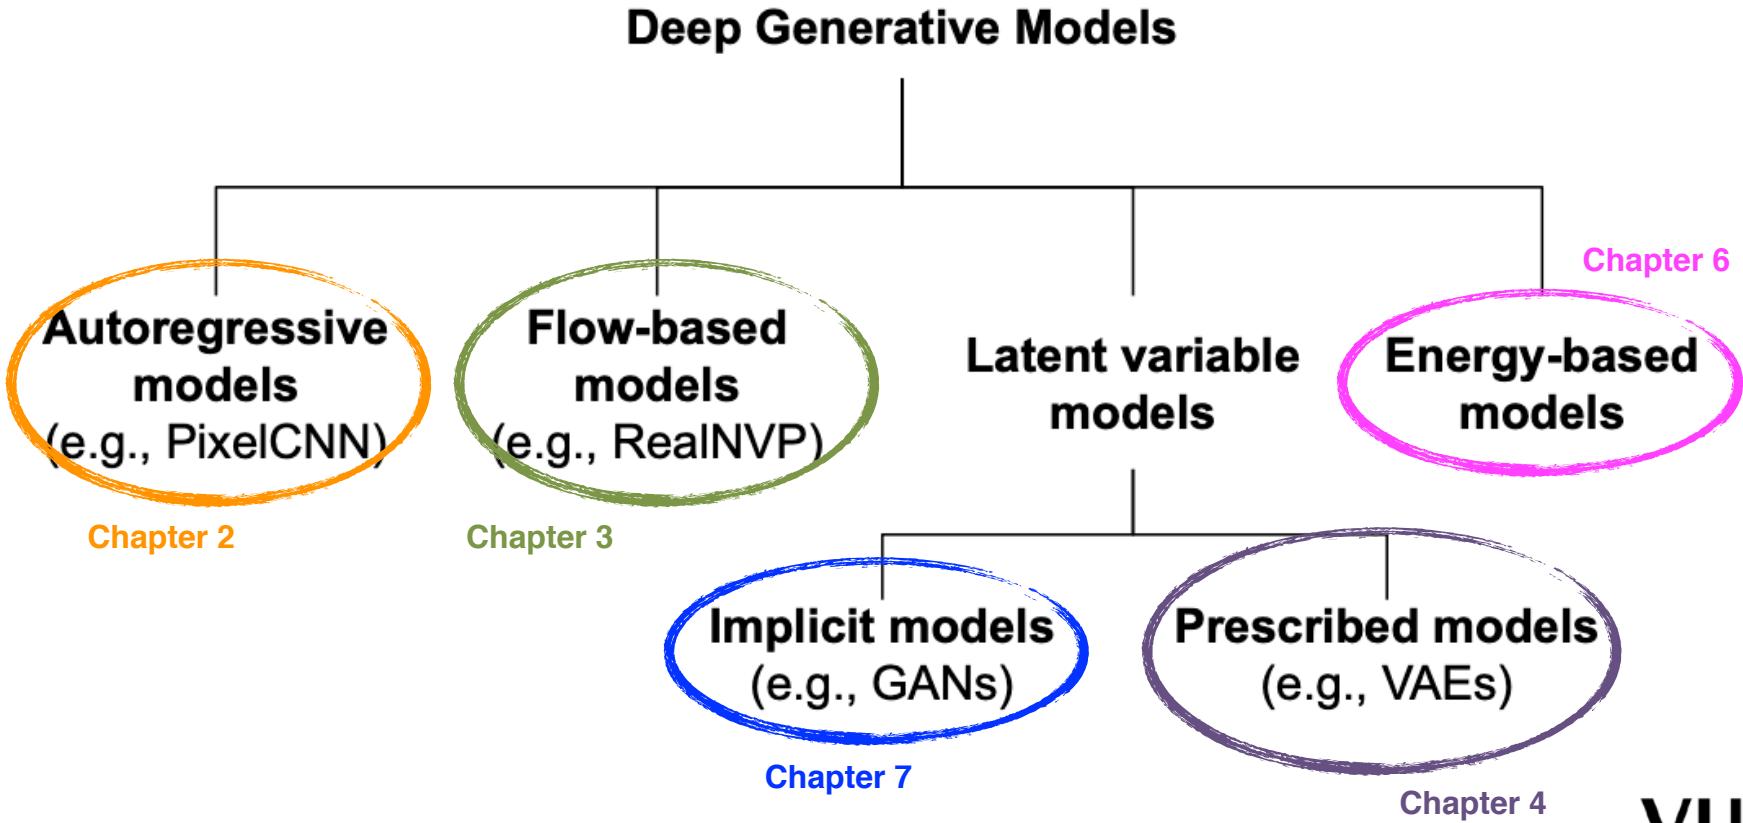


\mathbf{z}_5

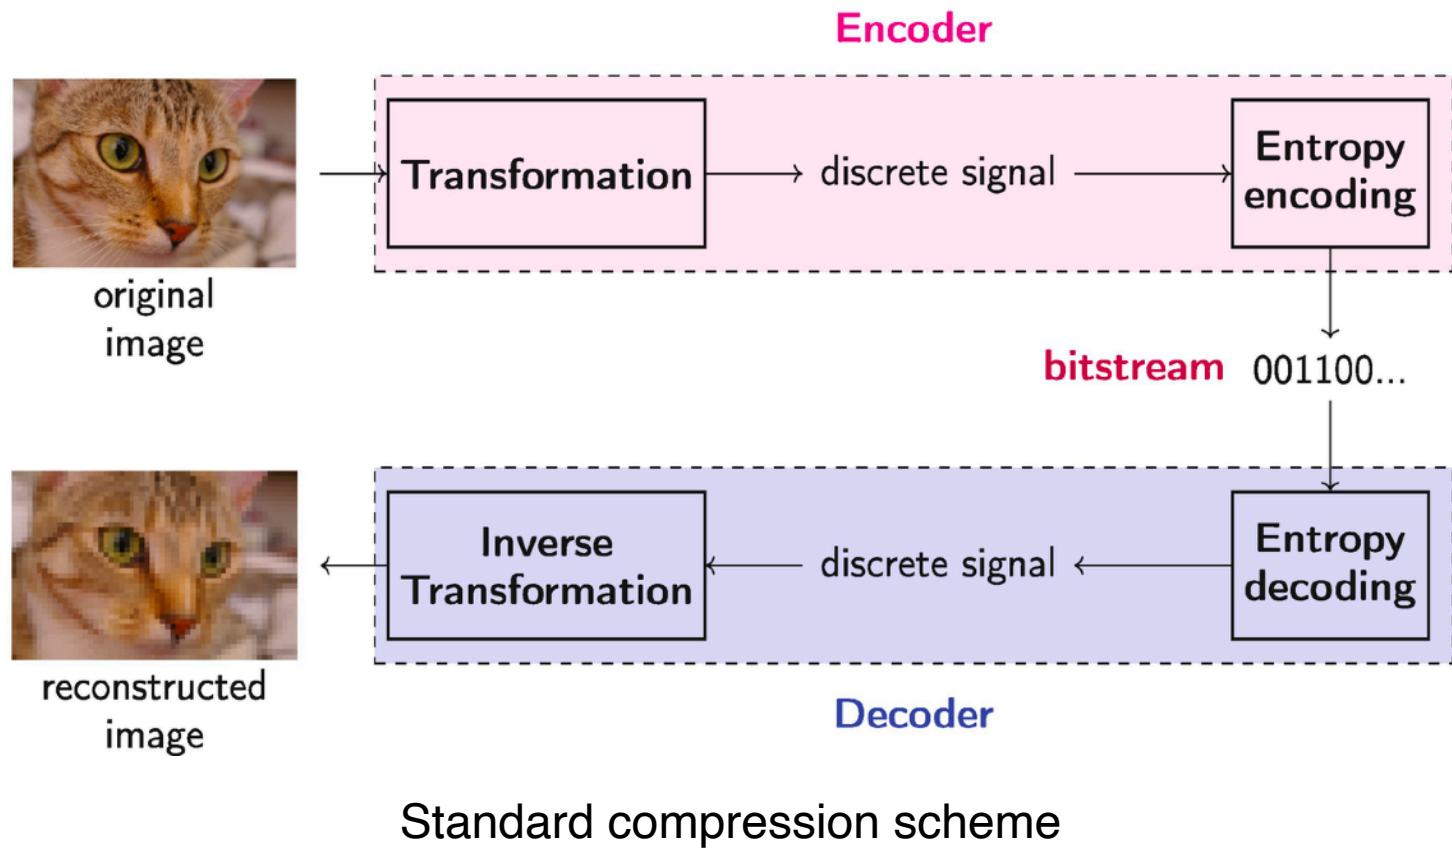


Backward diffusion (learnable)

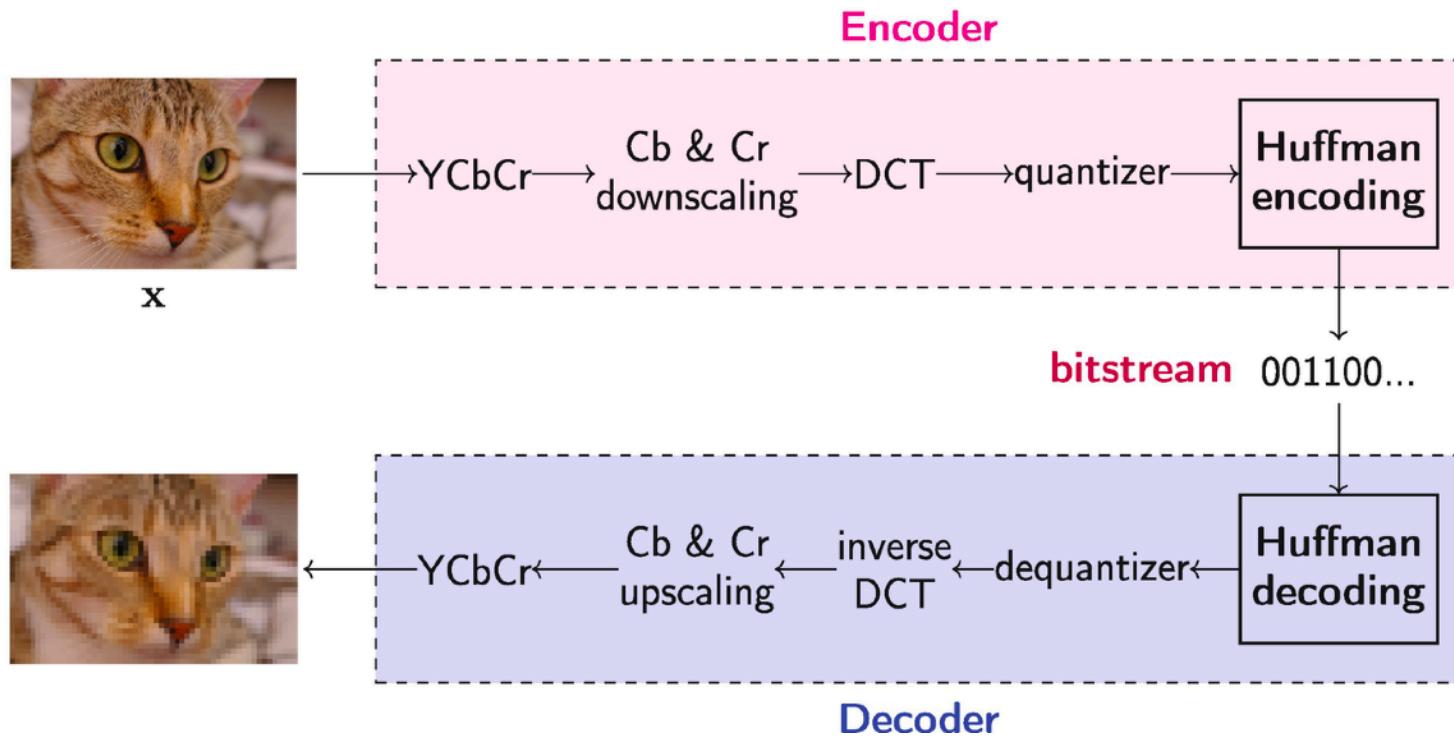
DEEP GENERATIVE MODELING: HOW WE CAN FORMULATE IT?



APPLICATIONS BEYOND DATA SYNTHESIS: NEURAL COMPRESSION

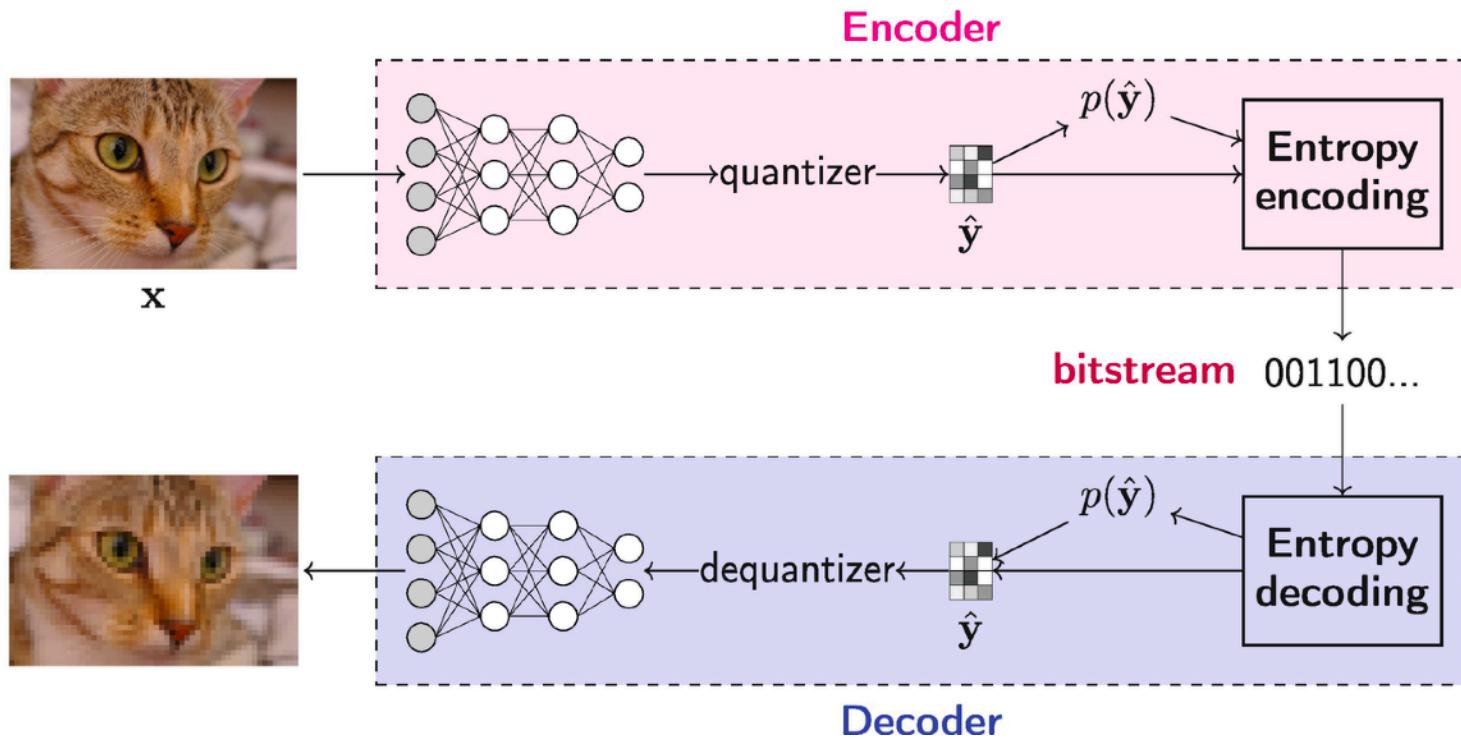


APPLICATIONS BEYOND DATA SYNTHESIS: NEURAL COMPRESSION



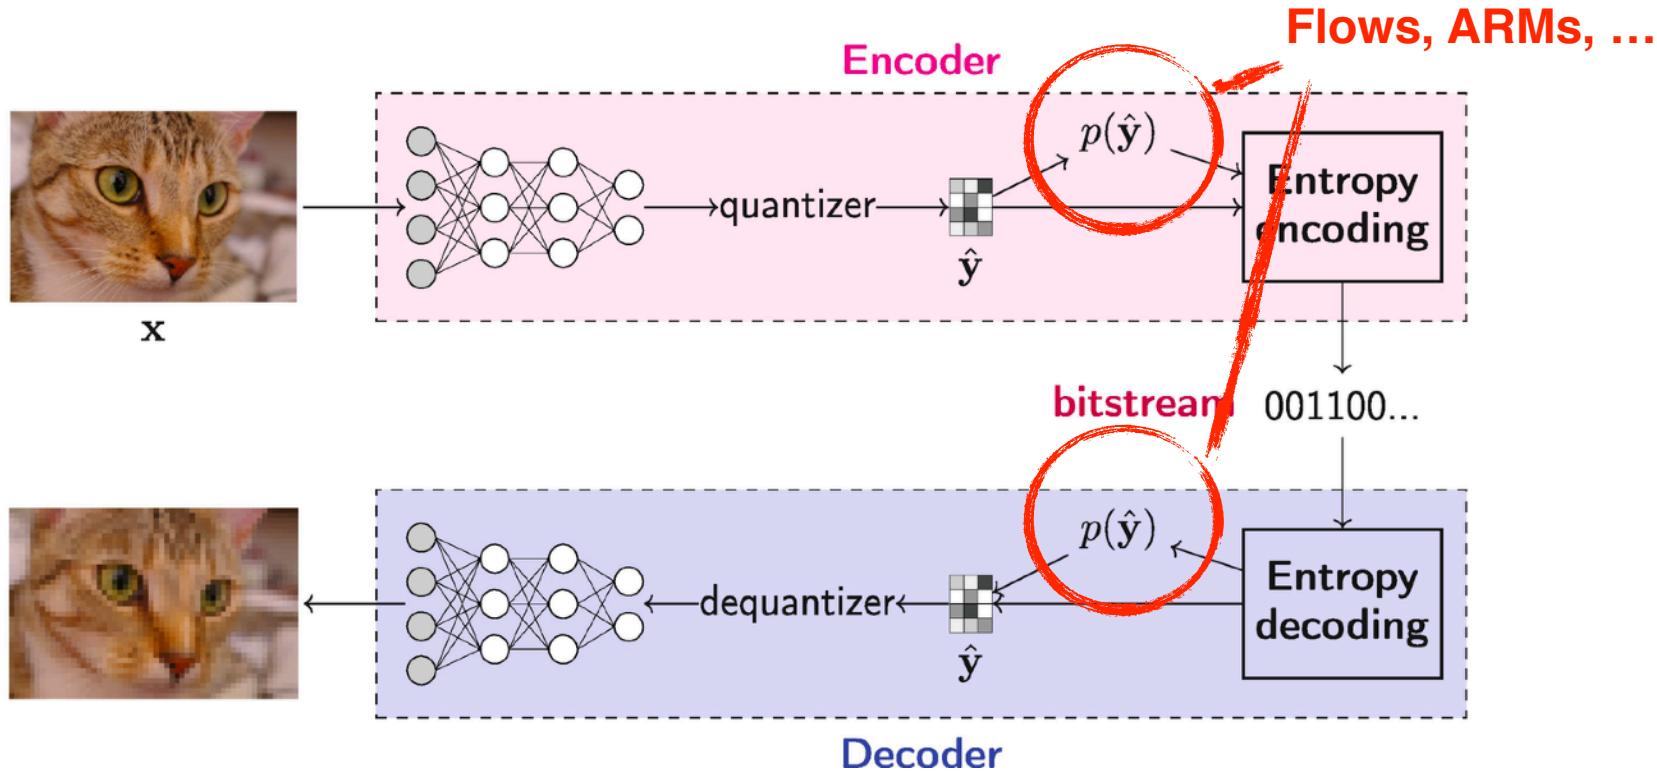
JPEG compression scheme

APPLICATIONS BEYOND DATA SYNTHESIS: NEURAL COMPRESSION



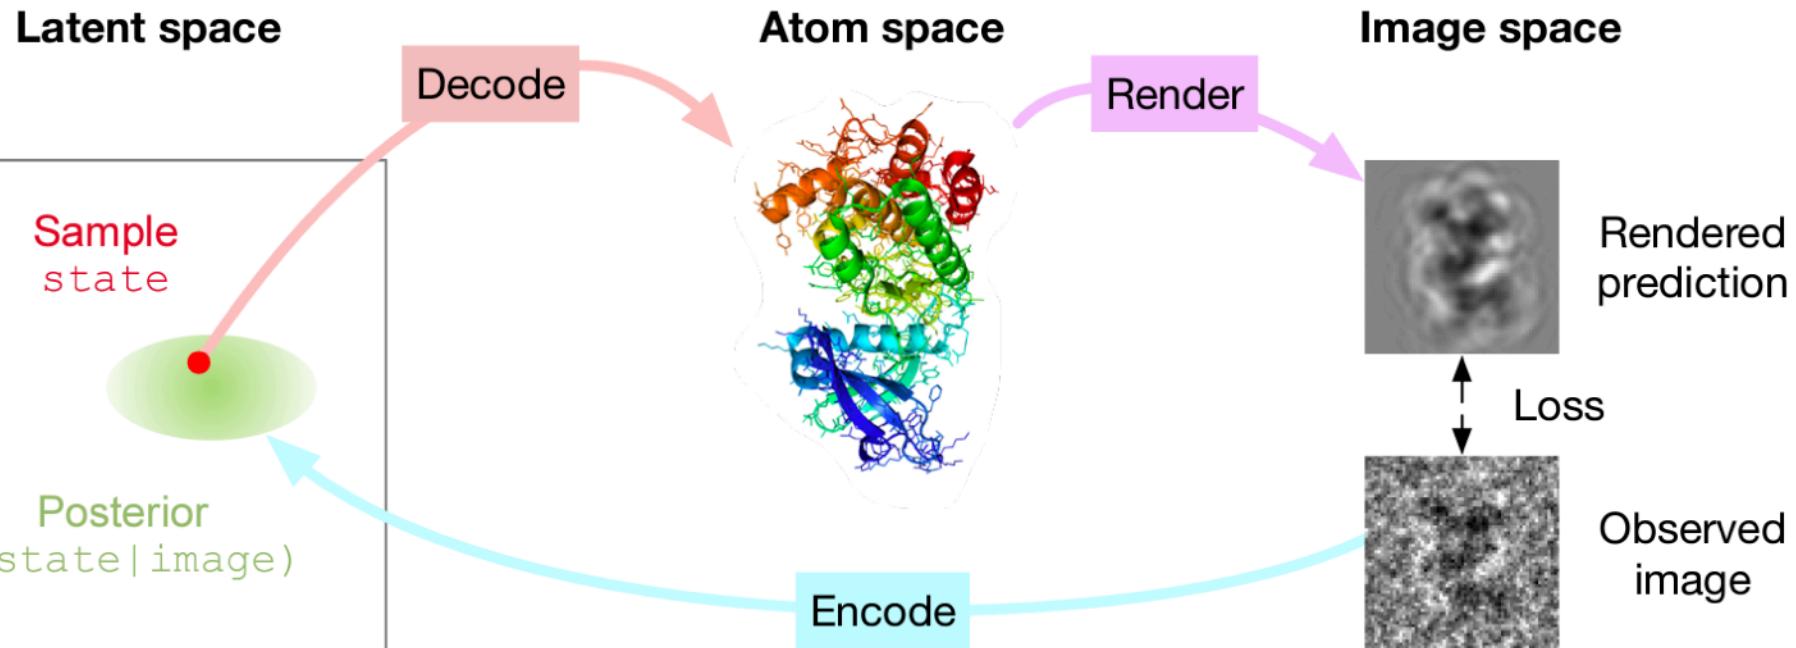
Neural compression scheme

APPLICATIONS BEYOND DATA SYNTHESIS: NEURAL COMPRESSION

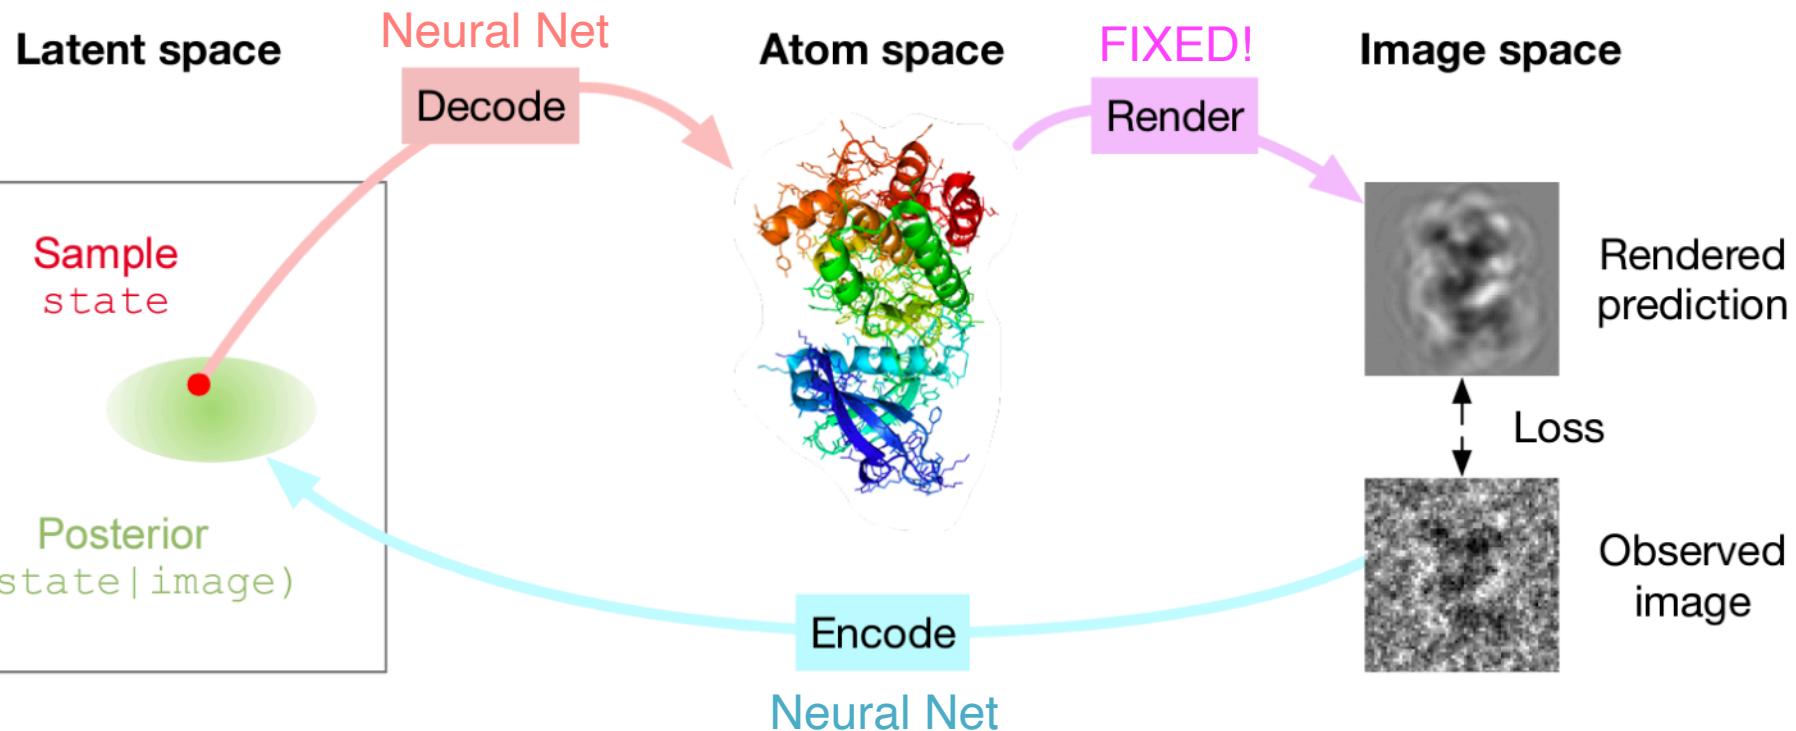


Neural compression scheme

APPLICATIONS BEYOND DATA SYNTHESIS: MOLECULAR MODELING

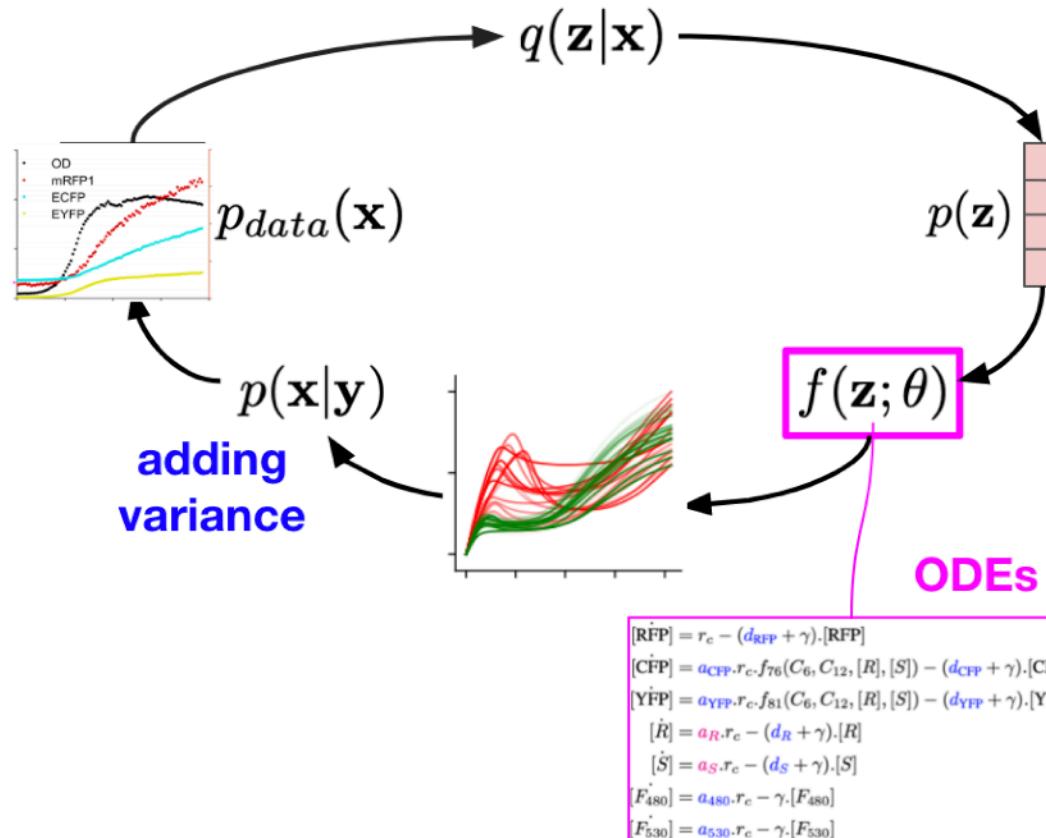


APPLICATIONS BEYOND DATA SYNTHESIS: MOLECULAR MODELING

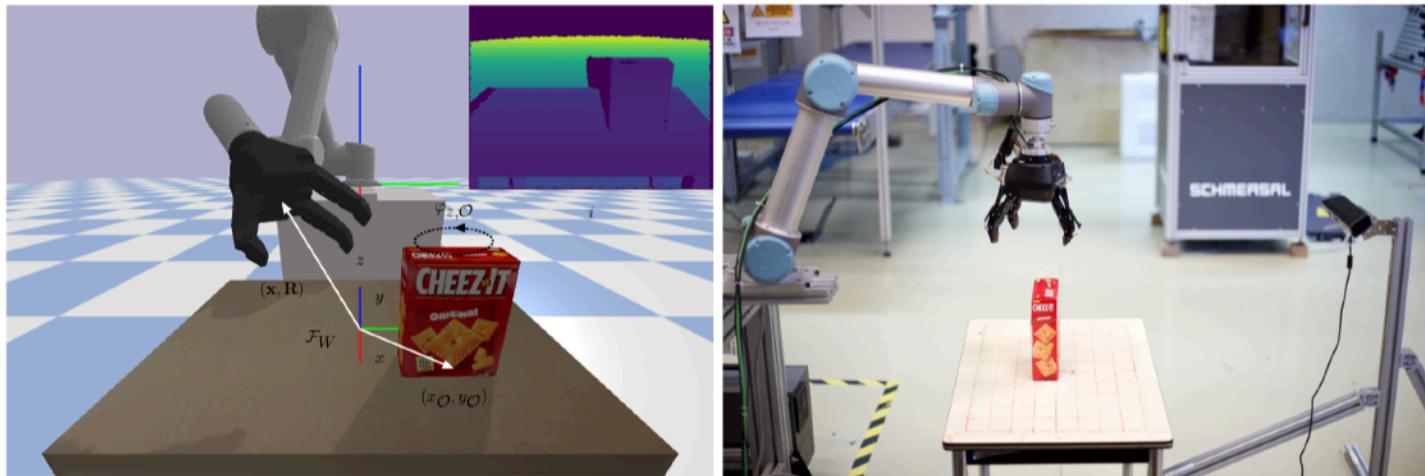


⁷⁷ Rosenbaum, Dan, et al. Inferring a Continuous Distribution of Atom Coordinates from Cryo-EM Images using VAEs. arXiv preprint arXiv:2106.14108, June 26, 2021

APPLICATIONS BEYOND DATA SYNTHESIS: GENETICS



APPLICATIONS BEYOND DATA SYNTHESIS: ROBOT LEARNING

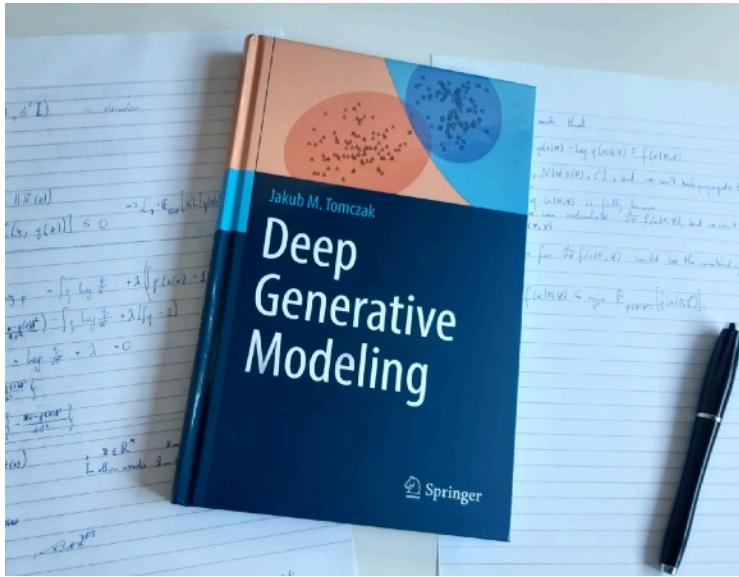


Generative model for generating trajectories.

MORE ABOUT DEEP GENERATIVE MODELING

Blog: <https://jmtomczak.github.io/blog.html>

Book: <https://link.springer.com/book/10.1007/978-3-030-93158-2>



CONCLUSION

- Why generative modeling?

$$p(\mathbf{x}, y) = p(y | \mathbf{x}) p(\mathbf{x})$$



CONCLUSION

- Why generative modeling?

$$p(\mathbf{x}, y) = p(y | \mathbf{x}) p(\mathbf{x})$$

- How to model the joint?

→ Hybrid modeling (Chapter 5)



CONCLUSION

- Why generative modeling?

$$p(\mathbf{x}, y) = p(y | \mathbf{x}) p(\mathbf{x})$$

- How to model the joint?

→ Hybrid modeling (Chapter 5)

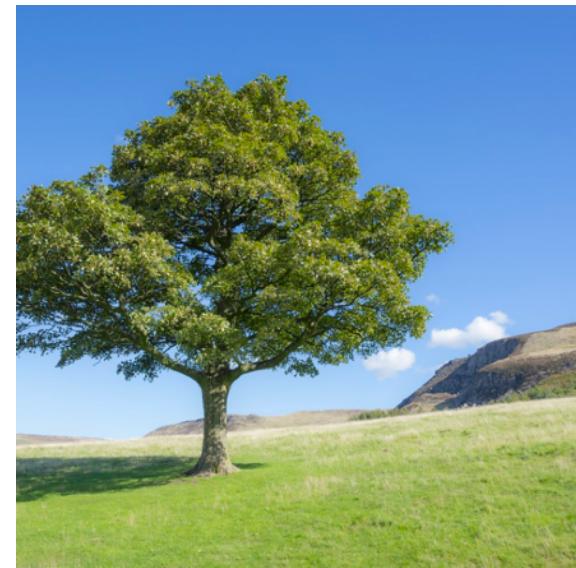
- Important directions:

→ Better uncertainty quantification

→ New parameterization (new neural networks)

→ Out-of-Distribution

→ Continual learning



THANK YOU FOR YOUR ATTENTION

Jakub M. Tomczak
Computational Intelligence group
Vrije Universiteit Amsterdam

Webpage: <https://jmtomczak.github.io/>

Github: <https://github.com/jmtomczak>

Twitter: <https://twitter.com/jmtomczak>