

WHY DO WE NEED DEEP GENERATIVE MODELING?

Jakub M. Tomczak
13 January 2020

Introduction

IS GENERATIVE MODELING IMPORTANT?

We learn a neural network to classify images:

IS GENERATIVE MODELING IMPORTANT?

We learn a neural network to classify images:



IS GENERATIVE MODELING IMPORTANT?

We learn a neural network to classify images:



$$p(\text{panda}|x) = 0.99$$

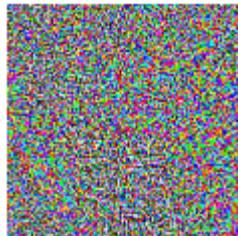
...

IS GENERATIVE MODELING IMPORTANT?

We learn a neural network to classify images:



+



=



$$p(\text{panda} | x) = 0.99$$

noise

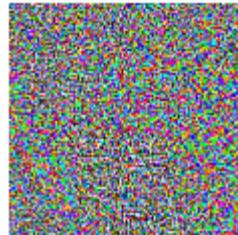
...

IS GENERATIVE MODELING IMPORTANT?

We learn a neural network to classify images:



+



=



$$p(\text{panda}|x) = 0.99$$

...

noise

$$p(\text{panda}|x) = 0.01$$

...

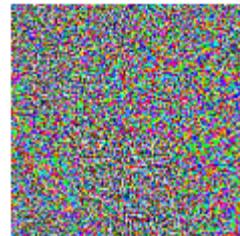
$$p(\text{dog}|x) = 0.9$$

IS GENERATIVE MODELING IMPORTANT?

We learn a neural network to classify images:



+



=



$$p(\text{panda}|x) = 0.99$$

...

noise

$$p(\text{panda}|x) = 0.01$$

...

$$p(\text{dog}|x) = 0.9$$

There is no semantic understanding of images.

IS GENERATIVE MODELING IMPORTANT?

This simple example shows that:

- A discriminative model is (probably) **not enough**.
- We need a notion of **uncertainty**.
- We need to **understand** the reality.

IS GENERATIVE MODELING IMPORTANT?

This simple example shows that:

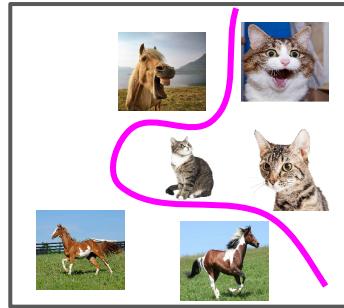
- A discriminative model is (probably) **not enough**.
- We need a notion of **uncertainty**.
- We need to **understand** the reality.

A possible solution is **generative modeling**.

IS GENERATIVE MODELING IMPORTANT?

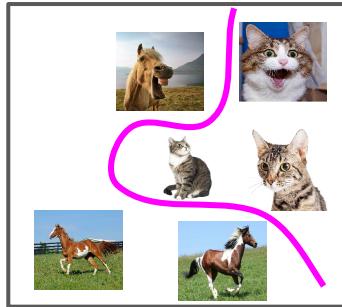


IS GENERATIVE MODELING IMPORTANT?

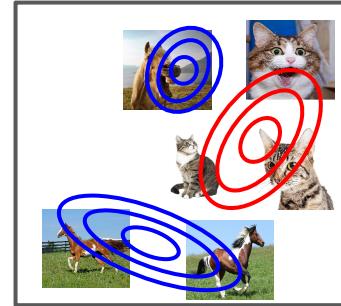


$$p_{\theta}(y|x)$$

IS GENERATIVE MODELING IMPORTANT?

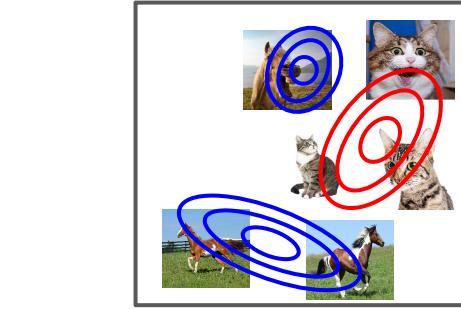
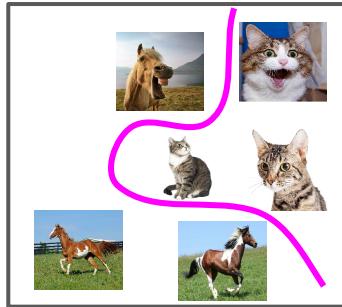


$$p_{\theta}(y|x)$$

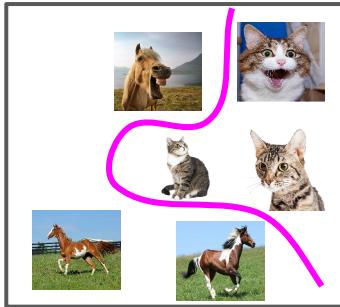


$$p_{\theta}(x,y) = p_{\theta}(y|x) \cdot p_{\theta}(x)$$

IS GENERATIVE MODELING IMPORTANT?

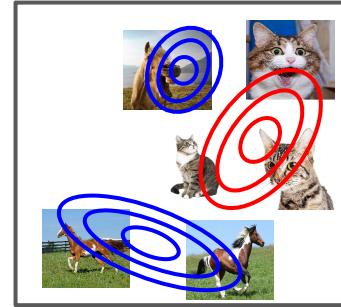


IS GENERATIVE MODELING IMPORTANT?



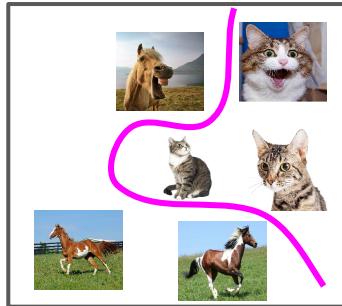
$$p_{\theta}(y|x)$$

High probability
of a **horse**.
=
Highly probable
decision!

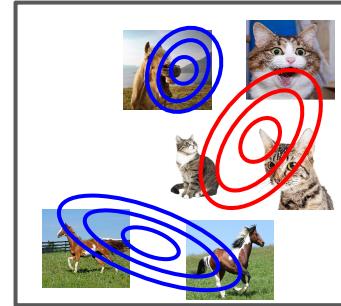


$$p_{\theta}(x, y) = p_{\theta}(y|x) \ p_{\theta}(x)$$

IS GENERATIVE MODELING IMPORTANT?

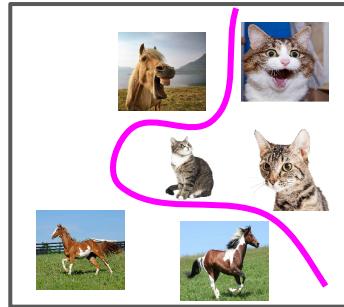


High probability
of a **horse**.
=
**Highly probable
decision!**



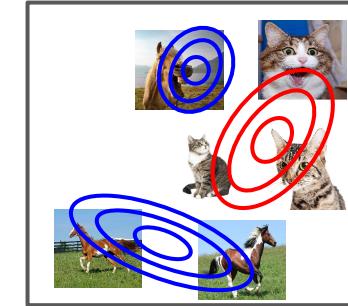
High probability of
a **horse**.
X
Low probability of
the **object**
=
**Uncertain
decision!**

IS GENERATIVE MODELING IMPORTANT?



$$p_{\theta}(y|x)$$

High probability
of a **horse**.
=
**Highly probable
decision!**



$$p_{\theta}(x, y) = p_{\theta}(y|x) \circ p_{\theta}(x)$$

High probability of
a **horse**.
x
Low probability of
the **object**
=
**Uncertain
decision!**

WHERE DO WE USE DEEP GENERATIVE MODELING?

" i want to talk to you . "

" i want to be with you . "

" i do n't want to be with you . "

i do n't want to be with you .

she did n't want to be with him .

he was silent for a long moment .

he was silent for a moment .

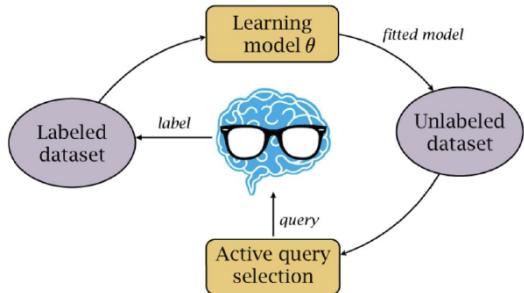
it was quiet for a moment .

it was dark and cold .

there was a pause .

it was my turn .

Text analysis



Active Learning

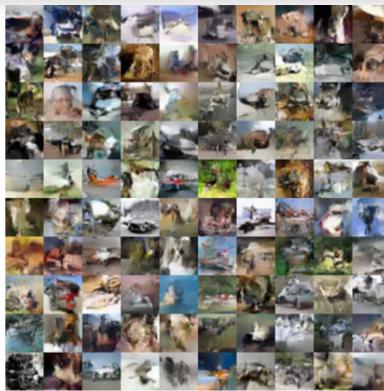
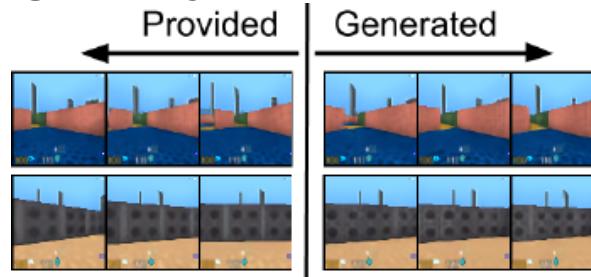
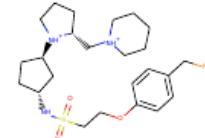


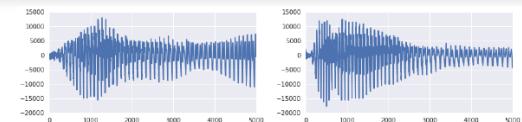
Image analysis



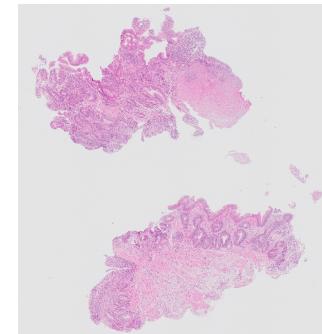
Reinforcement Learning



Graph analysis



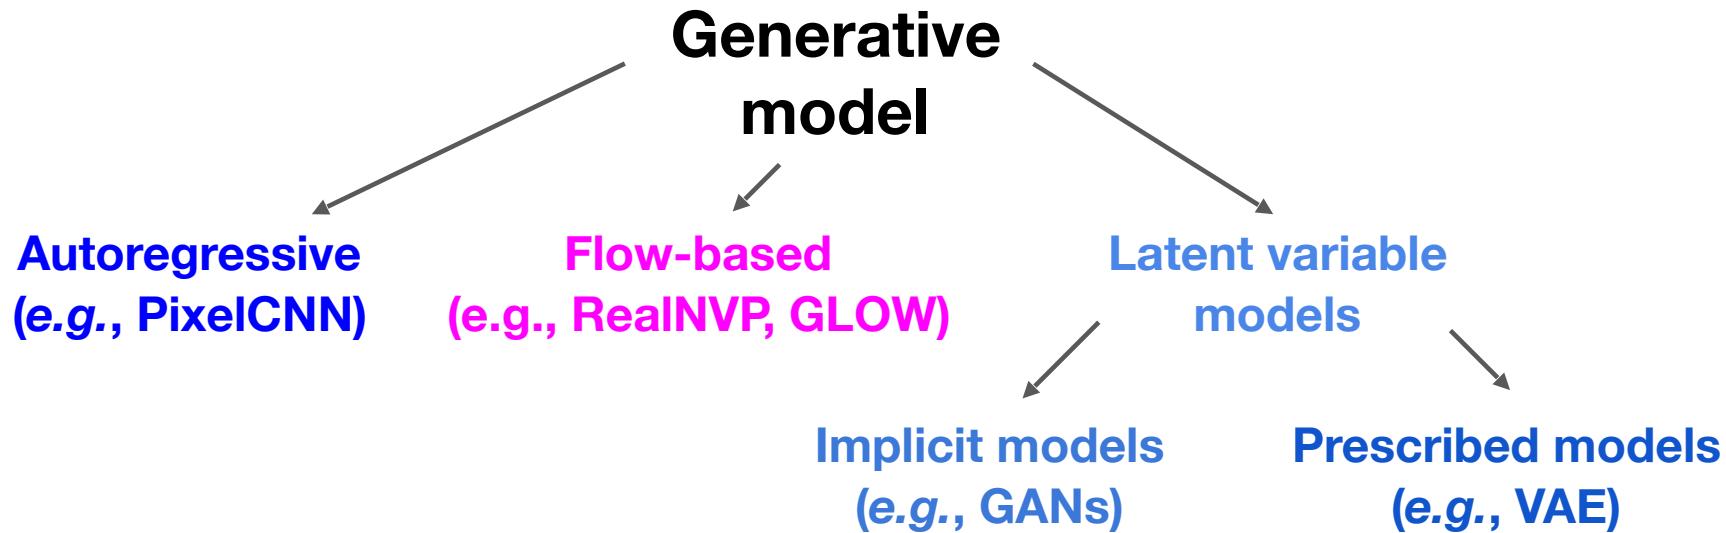
Audio analysis



Medical data

and more...

HOW TO FORMULATE GENERATIVE MODELS?



HOW TO FORMULATE GENERATIVE MODELS?

	Training	Likelihood	Sampling	Compression
Autoregressive models (e.g., PixelCNN)	Stable	Exact	Slow	No
Flow-based models (e.g., RealNVP)	Stable	Exact	Fast/Slow	No
Implicit models (e.g., GANs)	Unstable	No	Fast	No
Prescribed models (e.g., VAEs)	Stable	Approximate	Fast	Yes

HOW TO FORMULATE GENERATIVE MODELS?

	Training	Likelihood	Sampling	Compression
Autoregressive models (e.g., PixelCNN)	Stable	Exact	Slow	No
Flow-based models (e.g., RealNVP)	Stable	Exact	Fast/Slow	No
Implicit models (e.g., GANs)	Unstable	No	Fast	No
Prescribed models (e.g., VAEs)	Stable	Approximate	Fast	Yes

GENERATIVE MODELS AS (SPHERICAL) COWS



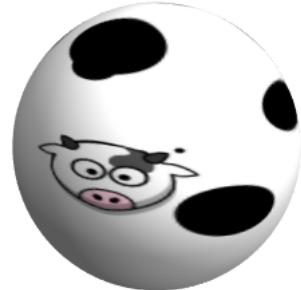
GENERATIVE MODELS AS (SPHERICAL) COWS

flow-based models



GENERATIVE MODELS AS (SPHERICAL) COWS

flow-based models



latent variable models



Deep latent variable models

GENERATIVE MODELING

Modeling in high-dimensional spaces is difficult.



GENERATIVE MODELING

Modeling in high-dimensional spaces is difficult.



GENERATIVE MODELING

Modeling in high-dimensional spaces is difficult.

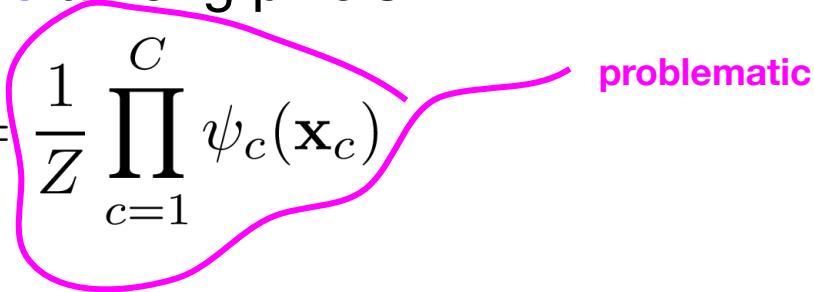
Modeling **all dependencies** among pixels:

$$p(\mathbf{x}) = \frac{1}{Z} \prod_{c=1}^C \psi_c(\mathbf{x}_c)$$

GENERATIVE MODELING

Modeling in high-dimensional spaces is difficult.

Modeling **all dependencies** among pixels:

$$p(\mathbf{x}) = \frac{1}{Z} \prod_{c=1}^C \psi_c(\mathbf{x}_c)$$


problematic

GENERATIVE MODELING

Modeling in high-dimensional spaces is difficult.

Modeling **all dependencies** among pixels:

$$p(\mathbf{x}) = \frac{1}{Z} \prod_{c=1}^C \psi_c(\mathbf{x}_c)$$

A pink oval highlights the term $\prod_{c=1}^C \psi_c(\mathbf{x}_c)$. A pink arrow points from this oval to the word "problematic" located to the right of the equation.

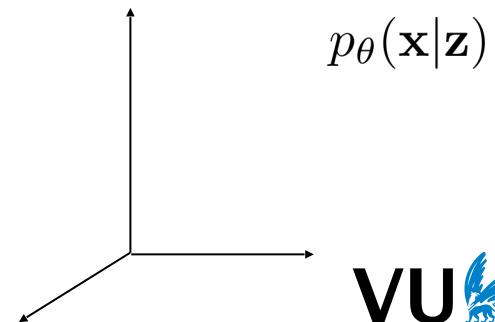
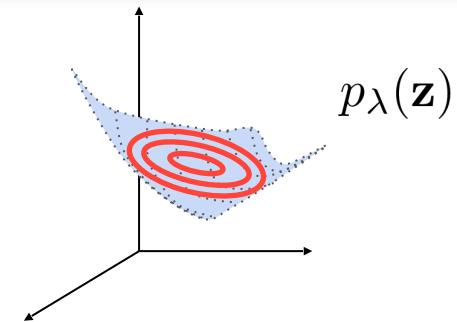
problematic

A possible **solution**: **Latent Variable Models!**

GENERATIVE MODELING WITH LATENT VARIABLES

Generative process:

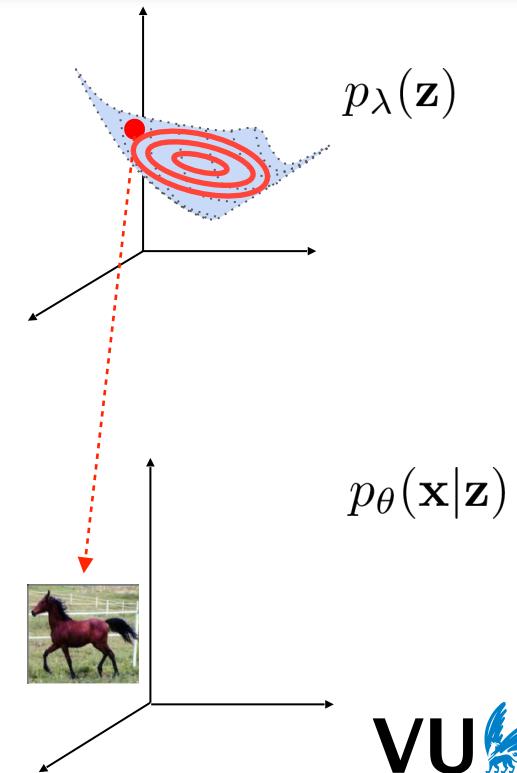
1. $\mathbf{z} \sim p_\lambda(\mathbf{z})$
2. $\mathbf{x} \sim p_\theta(\mathbf{x}|\mathbf{z})$



GENERATIVE MODELING WITH LATENT VARIABLES

Generative process:

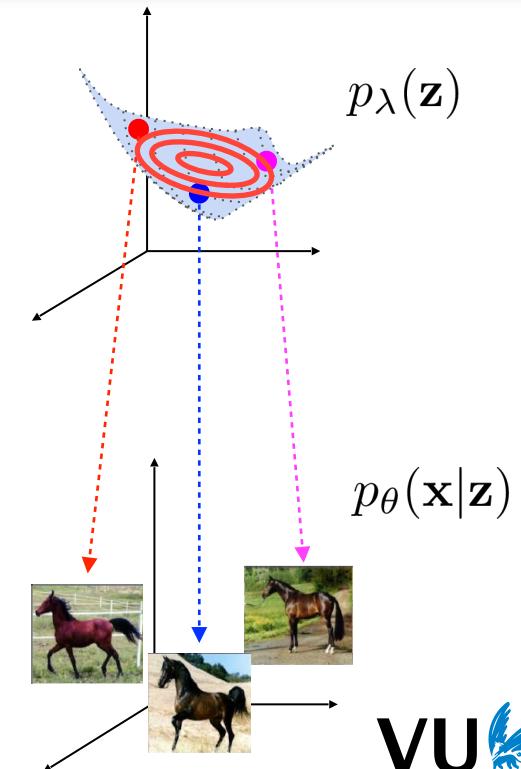
1. $\mathbf{z} \sim p_\lambda(\mathbf{z})$
2. $\mathbf{x} \sim p_\theta(\mathbf{x}|\mathbf{z})$



GENERATIVE MODELING WITH LATENT VARIABLES

Generative process:

1. $\mathbf{z} \sim p_\lambda(\mathbf{z})$
2. $\mathbf{x} \sim p_\theta(\mathbf{x}|\mathbf{z})$



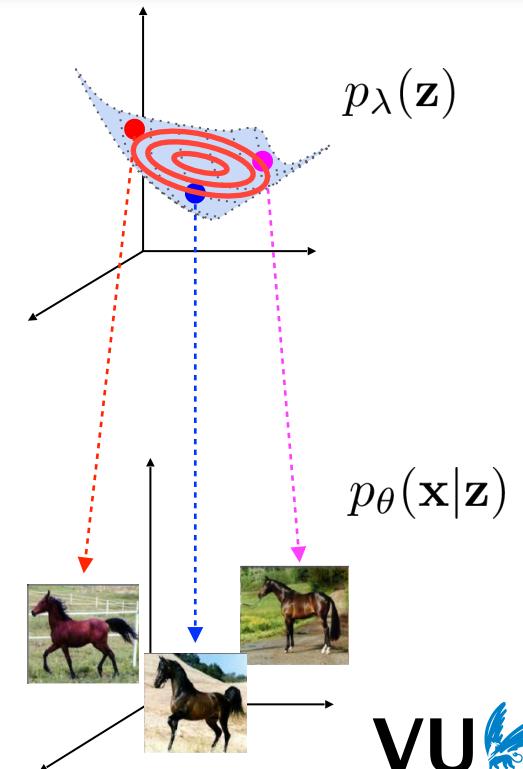
GENERATIVE MODELING WITH LATENT VARIABLES

Generative process:

1. $\mathbf{z} \sim p_\lambda(\mathbf{z})$
2. $\mathbf{x} \sim p_\theta(\mathbf{x}|\mathbf{z})$

Log of marginal distribution:

$$\log p_\vartheta(\mathbf{x}) = \log \int p_\theta(\mathbf{x}|\mathbf{z}) p_\lambda(\mathbf{z}) d\mathbf{z}$$



GENERATIVE MODELING WITH LATENT VARIABLES

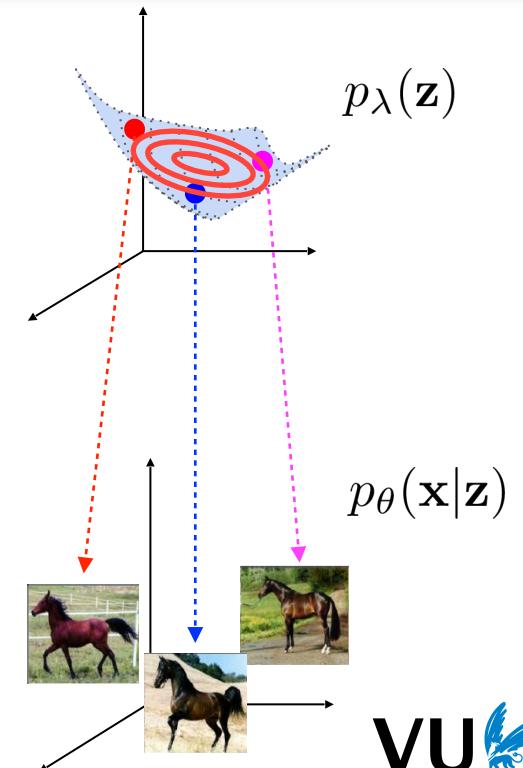
Generative process:

1. $\mathbf{z} \sim p_\lambda(\mathbf{z})$
2. $\mathbf{x} \sim p_\theta(\mathbf{x}|\mathbf{z})$

Log of marginal distribution:

$$\log p_\vartheta(\mathbf{x}) = \log \int p_\theta(\mathbf{x}|\mathbf{z}) p_\lambda(\mathbf{z}) d\mathbf{z}$$

How to train such model efficiently?



VARIATIONAL INFERENCE FOR LATENT VARIABLE MODELS

$$\begin{aligned}\log p_\vartheta(\mathbf{x}) &= \log \int p_\theta(\mathbf{x}|\mathbf{z}) \ p_\lambda(\mathbf{z}) d\mathbf{z} \\ &= \log \int \frac{q_\phi(\mathbf{z}|\mathbf{x})}{q_\phi(\mathbf{z}|\mathbf{x})} p_\theta(\mathbf{x}|\mathbf{z}) \ p_\lambda(\mathbf{z}) d\mathbf{z} \\ &\geq \int q_\phi(\mathbf{z}|\mathbf{x}) \log \frac{p_\theta(\mathbf{x}|\mathbf{z}) \ p_\lambda(\mathbf{z})}{q_\phi(\mathbf{z}|\mathbf{x})} d\mathbf{z} \\ &= \mathbb{E}_{\mathbf{z} \sim q_\phi(\mathbf{z}|\mathbf{x})} \left[\log p_\theta(\mathbf{x}|\mathbf{z}) \right] - \text{KL} \left(q_\phi(\mathbf{z}|\mathbf{x}) \middle\| p_\lambda(\mathbf{z}) \right)\end{aligned}$$

VARIATIONAL INFERENCE FOR LATENT VARIABLE MODELS

$$\begin{aligned}\log p_\theta(\mathbf{x}) &= \log \int p_\theta(\mathbf{x}|\mathbf{z}) p_\lambda(\mathbf{z}) d\mathbf{z} && \text{Variational posterior} \\ &= \log \int \frac{q_\phi(\mathbf{z}|\mathbf{x})}{q_\phi(\mathbf{z}|\mathbf{x})} p_\theta(\mathbf{x}|\mathbf{z}) p_\lambda(\mathbf{z}) d\mathbf{z} \\ &\geq \int q_\phi(\mathbf{z}|\mathbf{x}) \log \frac{p_\theta(\mathbf{x}|\mathbf{z}) p_\lambda(\mathbf{z})}{q_\phi(\mathbf{z}|\mathbf{x})} d\mathbf{z} \\ &= \mathbb{E}_{\mathbf{z} \sim q_\phi(\mathbf{z}|\mathbf{x})} \left[\log p_\theta(\mathbf{x}|\mathbf{z}) \right] - \text{KL}\left(q_\phi(\mathbf{z}|\mathbf{x}) || p_\lambda(\mathbf{z})\right)\end{aligned}$$

VARIATIONAL INFERENCE FOR LATENT VARIABLE MODELS

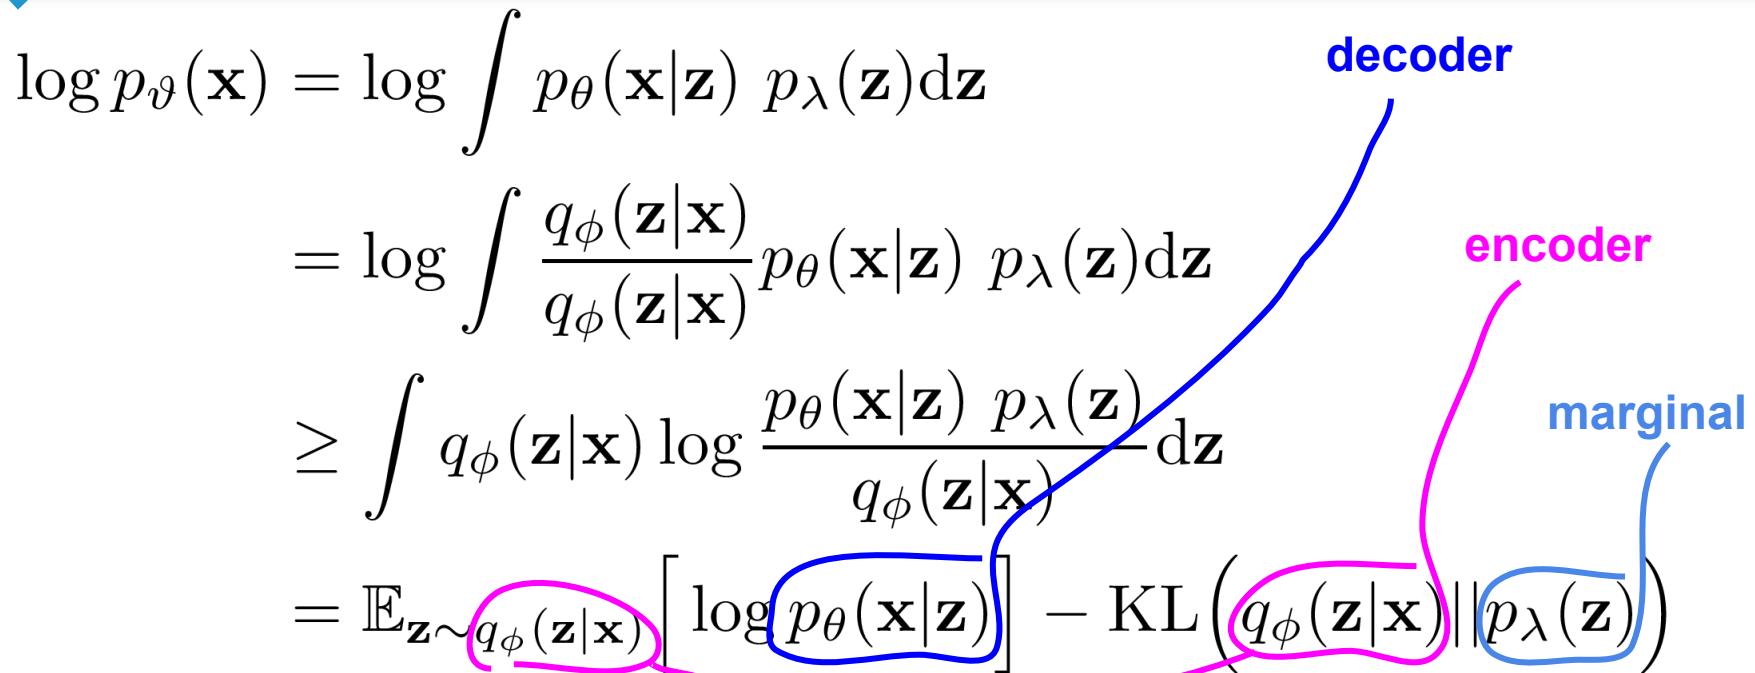
$$\begin{aligned}\log p_\theta(\mathbf{x}) &= \log \int p_\theta(\mathbf{x}|\mathbf{z}) p_\lambda(\mathbf{z}) d\mathbf{z} \\ &= \log \int \frac{q_\phi(\mathbf{z}|\mathbf{x})}{q_\phi(\mathbf{z}|\mathbf{x})} p_\theta(\mathbf{x}|\mathbf{z}) p_\lambda(\mathbf{z}) d\mathbf{z} \\ &\geq \int q_\phi(\mathbf{z}|\mathbf{x}) \log \frac{p_\theta(\mathbf{x}|\mathbf{z}) p_\lambda(\mathbf{z})}{q_\phi(\mathbf{z}|\mathbf{x})} d\mathbf{z} \\ &= \mathbb{E}_{\mathbf{z} \sim q_\phi(\mathbf{z}|\mathbf{x})} \left[\log p_\theta(\mathbf{x}|\mathbf{z}) \right] - \text{KL} \left(q_\phi(\mathbf{z}|\mathbf{x}) \middle\| p_\lambda(\mathbf{z}) \right)\end{aligned}$$

Jensen's inequality

VARIATIONAL INFERENCE FOR LATENT VARIABLE MODELS

$$\begin{aligned}\log p_\vartheta(\mathbf{x}) &= \log \int p_\theta(\mathbf{x}|\mathbf{z}) p_\lambda(\mathbf{z}) d\mathbf{z} \\ &= \log \int \frac{q_\phi(\mathbf{z}|\mathbf{x})}{q_\phi(\mathbf{z}|\mathbf{x})} p_\theta(\mathbf{x}|\mathbf{z}) p_\lambda(\mathbf{z}) d\mathbf{z} \\ &\geq \int q_\phi(\mathbf{z}|\mathbf{x}) \log \frac{p_\theta(\mathbf{x}|\mathbf{z}) p_\lambda(\mathbf{z})}{q_\phi(\mathbf{z}|\mathbf{x})} d\mathbf{z} \\ &= \underbrace{\mathbb{E}_{\mathbf{z} \sim q_\phi(\mathbf{z}|\mathbf{x})} \left[\log p_\theta(\mathbf{x}|\mathbf{z}) \right]}_{\text{Reconstruction error}} - \underbrace{\text{KL} \left(q_\phi(\mathbf{z}|\mathbf{x}) || p_\lambda(\mathbf{z}) \right)}_{\text{Regularization}}\end{aligned}$$

VARIATIONAL INFERENCE FOR LATENT VARIABLE MODELS

$$\begin{aligned}\log p_\vartheta(\mathbf{x}) &= \log \int p_\theta(\mathbf{x}|\mathbf{z}) p_\lambda(\mathbf{z}) d\mathbf{z} \\ &= \log \int \frac{q_\phi(\mathbf{z}|\mathbf{x})}{q_\phi(\mathbf{z}|\mathbf{x})} p_\theta(\mathbf{x}|\mathbf{z}) p_\lambda(\mathbf{z}) d\mathbf{z} \\ &\geq \int q_\phi(\mathbf{z}|\mathbf{x}) \log \frac{p_\theta(\mathbf{x}|\mathbf{z}) p_\lambda(\mathbf{z})}{q_\phi(\mathbf{z}|\mathbf{x})} d\mathbf{z} \\ &= \mathbb{E}_{\mathbf{z} \sim q_\phi(\mathbf{z}|\mathbf{x})} \left[\log p_\theta(\mathbf{x}|\mathbf{z}) \right] - \text{KL} \left(q_\phi(\mathbf{z}|\mathbf{x}) \parallel p_\lambda(\mathbf{z}) \right)\end{aligned}$$


VARIATIONAL INFERENCE FOR LATENT VARIABLE MODELS

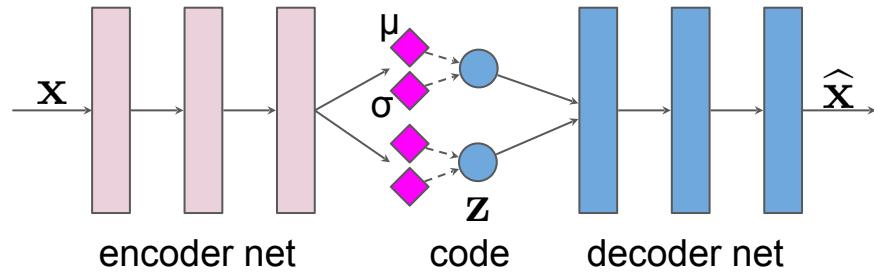
$$\begin{aligned}\log p_\vartheta(\mathbf{x}) &= \log \int p_\theta(\mathbf{x}|\mathbf{z}) p_\lambda(\mathbf{z}) d\mathbf{z} \\ &= \log \int \frac{q_\phi(\mathbf{z}|\mathbf{x})}{q_\phi(\mathbf{z}|\mathbf{x})} p_\theta(\mathbf{x}|\mathbf{z}) p_\lambda(\mathbf{z}) d\mathbf{z} \\ &\geq \int q_\phi(\mathbf{z}|\mathbf{x}) \log \frac{p_\theta(\mathbf{x}|\mathbf{z}) p_\lambda(\mathbf{z})}{q_\phi(\mathbf{z}|\mathbf{x})} d\mathbf{z} \\ &= \mathbb{E}_{\mathbf{z} \sim q_\phi(\mathbf{z}|\mathbf{x})} \left[\log p_\theta(\mathbf{x}|\mathbf{z}) \right] - \text{KL} \left(q_\phi(\mathbf{z}|\mathbf{x}) \parallel p_\lambda(\mathbf{z}) \right)\end{aligned}$$

The diagram illustrates the Variational Auto-Encoder (VAE) architecture. It features three main components: a blue curve labeled "decoder" that maps the latent variable \mathbf{z} to the observed variable \mathbf{x} ; a pink curve labeled "encoder" that maps \mathbf{x} back to \mathbf{z} ; and a blue line labeled "marginal" representing the true posterior distribution $p_\lambda(\mathbf{z})$. The equation shows the log probability of \mathbf{x} as the sum of the expected log probability under the decoder and the Kullback-Leibler divergence between the encoder's distribution and the marginal.

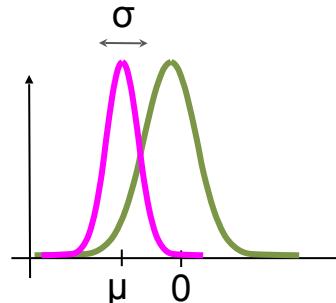
VARIATIONAL AUTO-ENCODERS

Variational posterior (**encoder**) and likelihood function (**decoder**) are parameterized by neural networks.

Reparameterization trick:
move the stochasticity to independent random variables



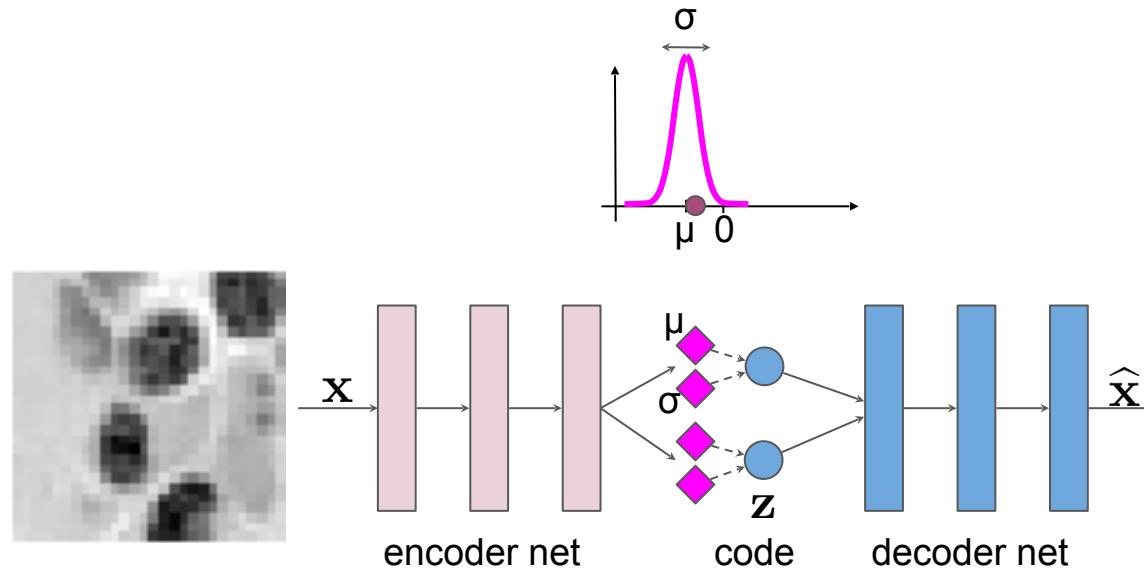
$$\mathbf{z} = f(\boldsymbol{\mu}, \boldsymbol{\sigma}; \boldsymbol{\varepsilon}), \text{ where } \boldsymbol{\varepsilon} \sim p(\boldsymbol{\varepsilon})$$



VARIATIONAL AUTO-ENCODERS

VAE copies input to output through a **bottleneck**.

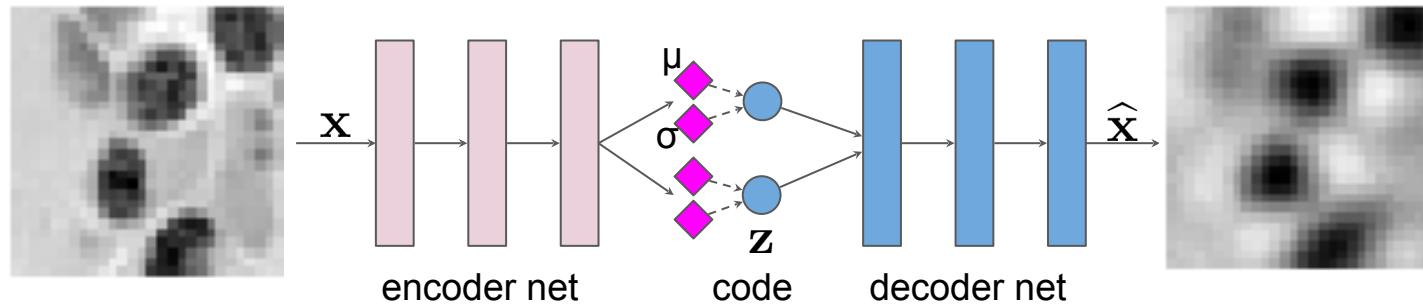
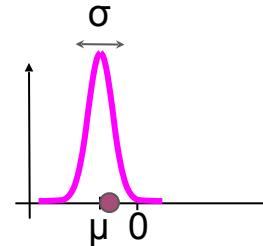
VAE learns a **code** of the data.



VARIATIONAL AUTO-ENCODERS

VAE copies input to output through a **bottleneck**.

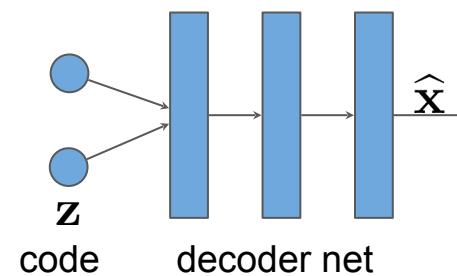
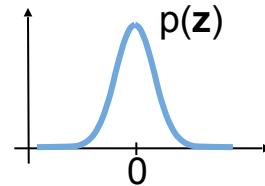
VAE learns a **code** of the data.



VARIATIONAL AUTO-ENCODERS

VAE has a **marginal** on the latent code.

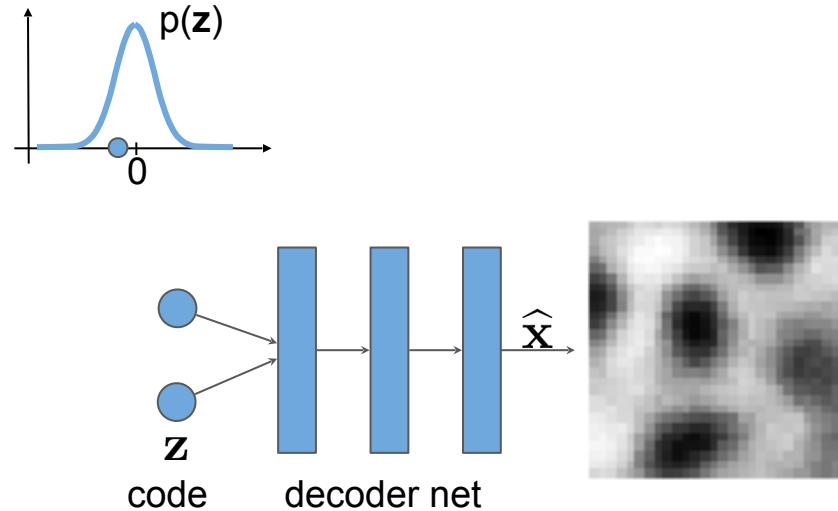
VAE can **generate** new data.



VARIATIONAL AUTO-ENCODERS

VAE has a **marginal** on the latent code.

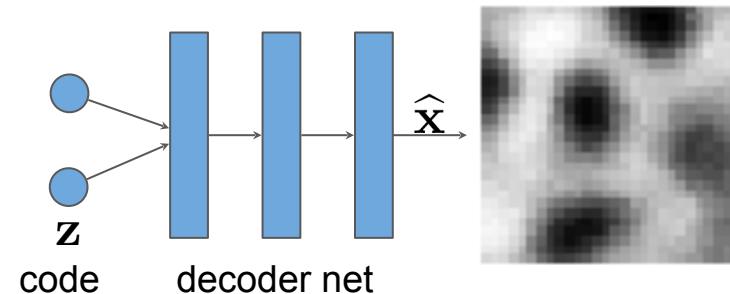
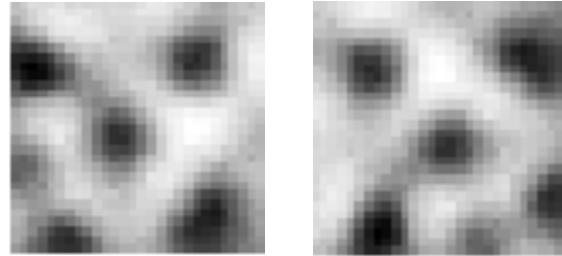
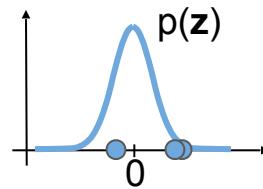
VAE can **generate** new data.



VARIATIONAL AUTO-ENCODERS

VAE has a **marginal** on the latent code.

VAE can **generate** new data.



COMMON ISSUES WITH VAEs

$$q_{\phi}(\mathbf{z}|\mathbf{x}) \propto p_{\theta}(\mathbf{x}|\mathbf{z}) p_{\lambda}(\mathbf{z})$$

Weak decoders → bad generations/reconstructions

Weak encoders → bad latent representation

Weak marginals → bad generations

Variational posteriors → what family of distributions?

Others...

COMPONENTS OF VAEs

$$q_{\phi}(\mathbf{z}|\mathbf{x}) \propto p_{\theta}(\mathbf{x}|\mathbf{z}) p_{\lambda}(\mathbf{z})$$



Resnets
DRAW
Autoregressive models
Normalizing flows

Autoregressive models
Normalizing flows
VampPrior
Implicit prior

COMPONENTS OF VAEs

$$q_{\phi}(\mathbf{z}|\mathbf{x}) \propto p_{\theta}(\mathbf{x}|\mathbf{z}) p_{\lambda}(\mathbf{z})$$

Normalizing flows

Discrete encoders

Hyperspherical dist.

Hyperbolic-normal dist.

Group theory

Resnets

DRAW

Autoregressive models

Normalizing flows

Autoregressive models

Normalizing flows

VampPrior

Implicit prior

COMPONENTS OF VAEs

$$q_{\phi}(\mathbf{z}|\mathbf{x}) \propto p_{\theta}(\mathbf{x}|\mathbf{z}) p_{\lambda}(\mathbf{z})$$

Normalizing flows

Discrete encoders

Hyperspherical dist.

Hyperbolic-normal dist.

Group theory

Resnets

DRAW

Autoregressive models

Normalizing flows

Autoregressive models

Normalizing flows

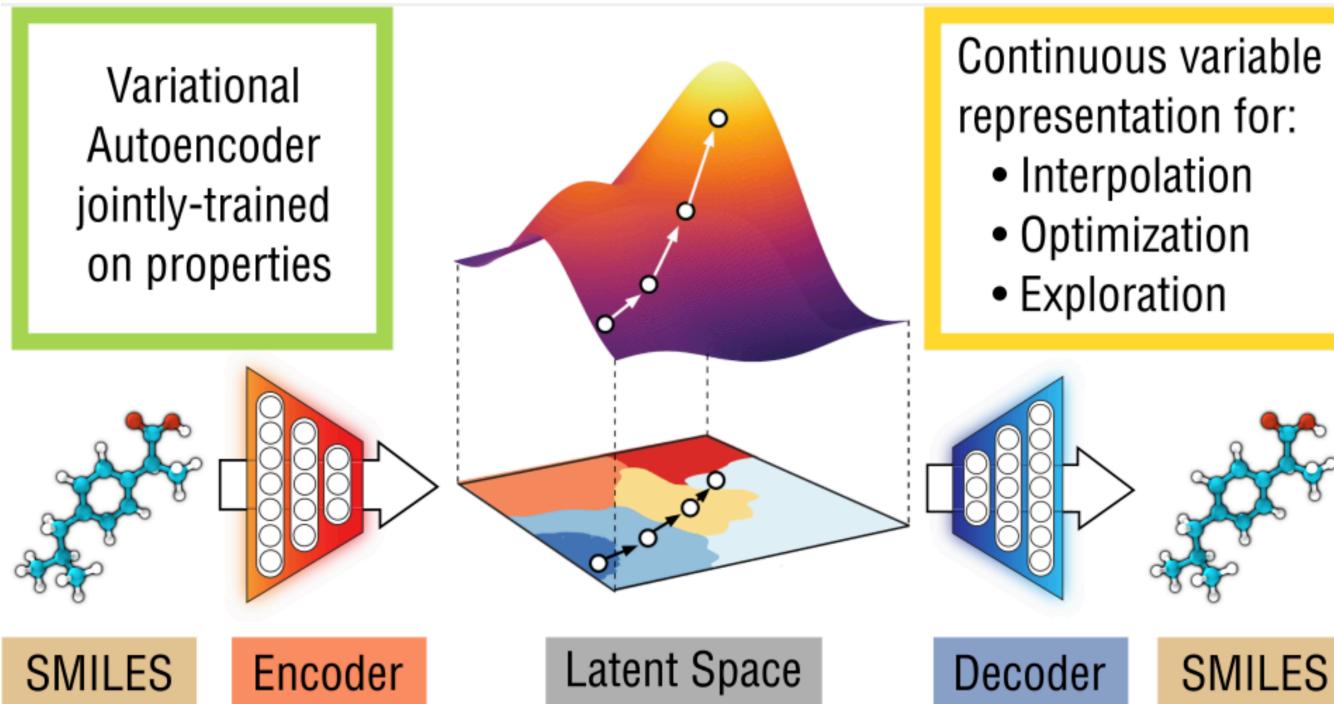
VampPrior

Implicit prior

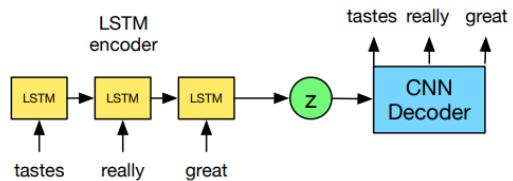
$$\text{ELBO}(\mathbf{x}; \theta, \phi, \lambda)$$

Adversarial learning
MMD
Wasserstein AE

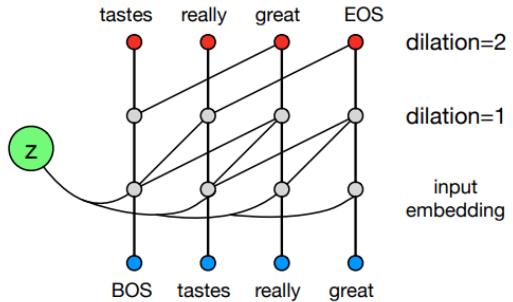
APPLICATIONS



APPLICATIONS

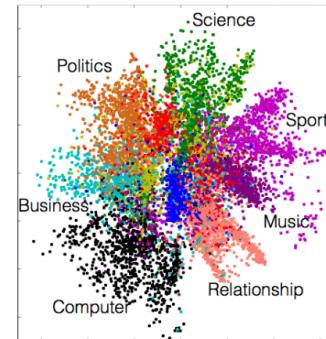


(a) VAE training graph using a dilated CNN decoder.

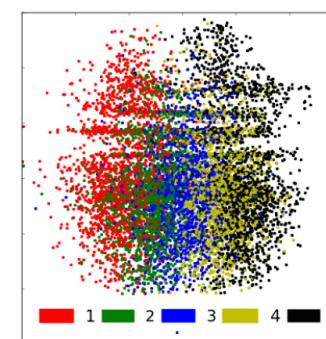


(b) Diagram of dilated CNN decoder.

- 1 star** the food was good but the service was horrible . took forever to get our food . we had to ask twice for our check after we got our food . will not return .
- 2 star** the food was good , but the service was terrible . took forever to get someone to take our drink order . had to ask 3 times to get the check . food was ok , nothing to write about .
- 3 star** came here for the first time last night . food was good . service was a little slow . food was just ok .
- 4 star** food was good , service was a little slow , but the food was pretty good . i had the grilled chicken sandwich and it was really good . will definitely be back !
- 5 star** food was very good , service was fast and friendly . food was very good as well . will be back !

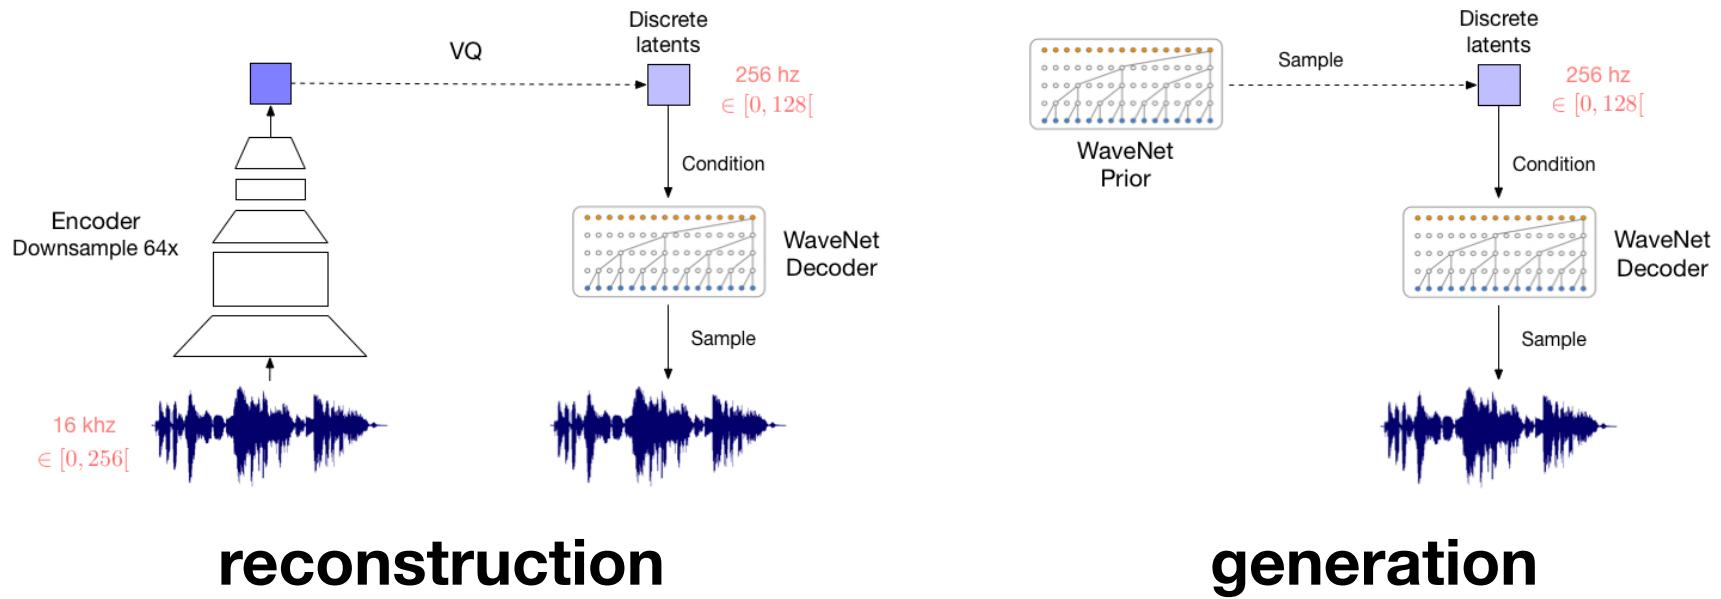


(a) Yahoo

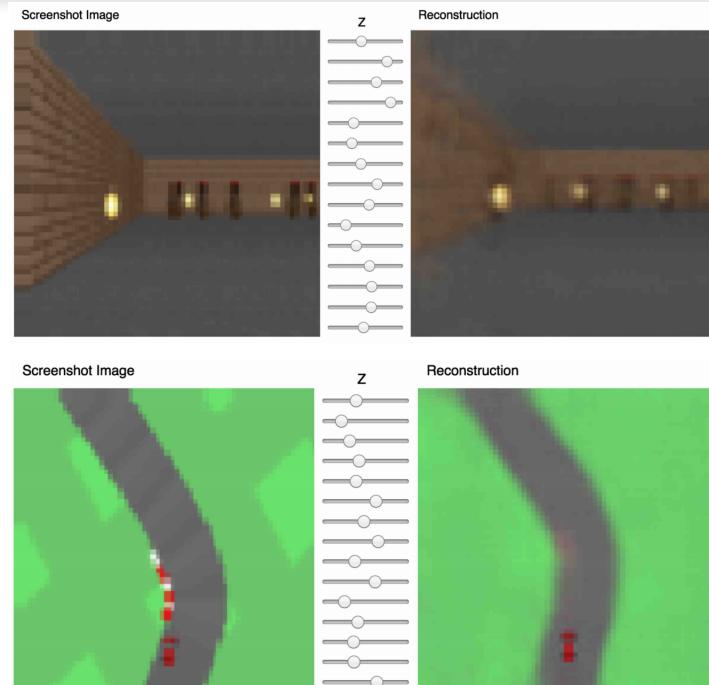
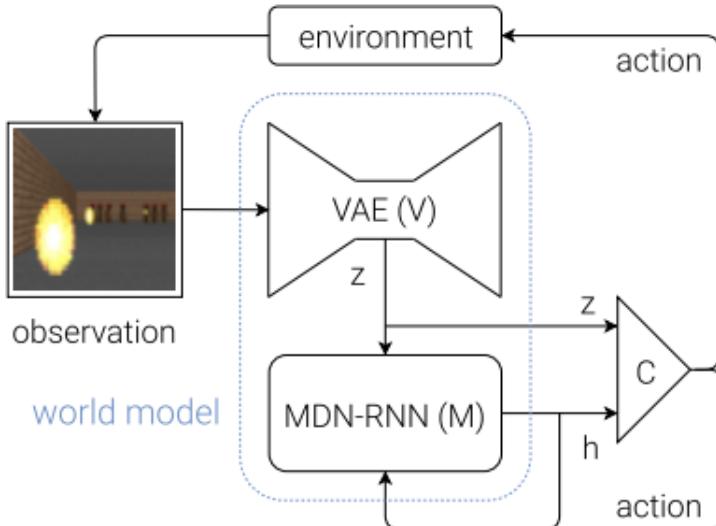


(b) Yelp

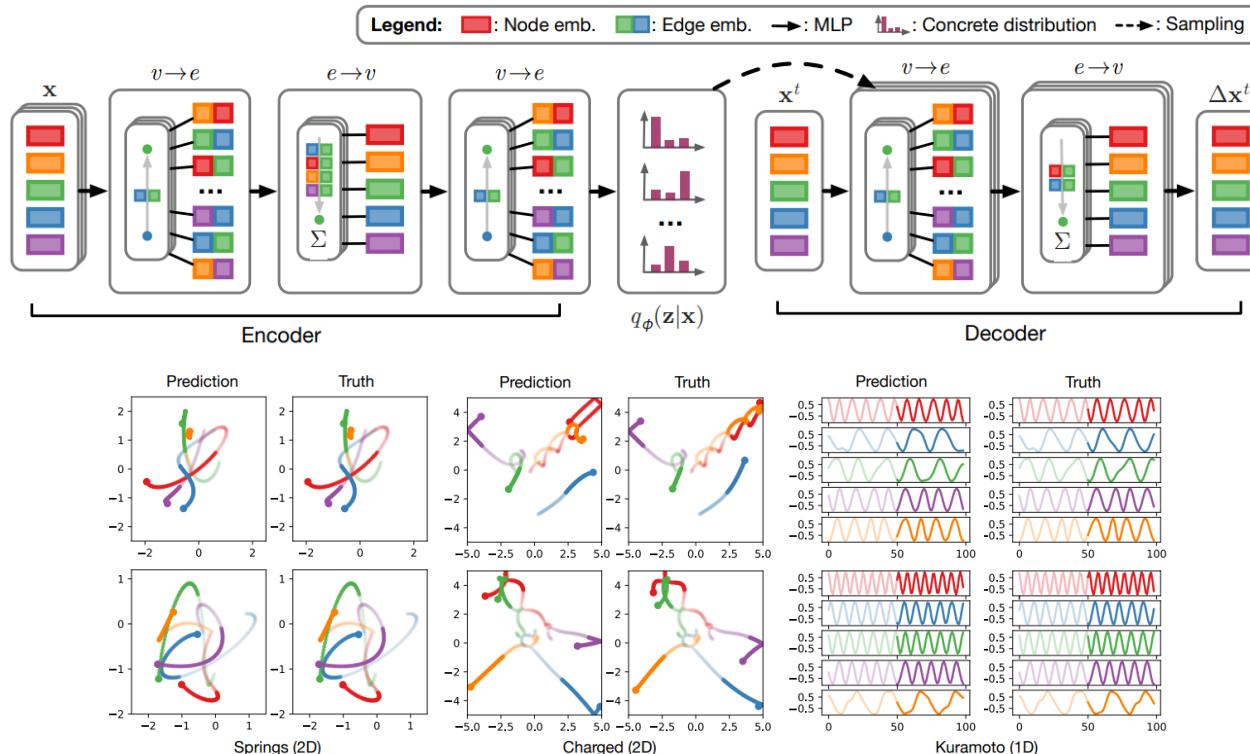
APPLICATIONS



APPLICATIONS



APPLICATIONS



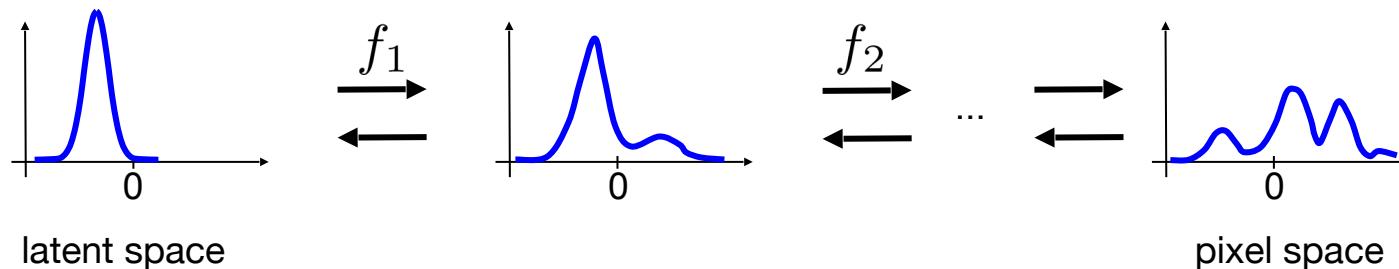
Flow-based models

THE CHANGE OF VARIABLES FORMULA

Let's recall the change of variables formula with invertible transformations:

$$p(\mathbf{x}) = \pi_0(\mathbf{z}_0) \prod_{i=1}^K \left| \det \frac{\partial f_i(\mathbf{z}_{i-1})}{\partial \mathbf{z}_{i-1}} \right|^{-1}$$

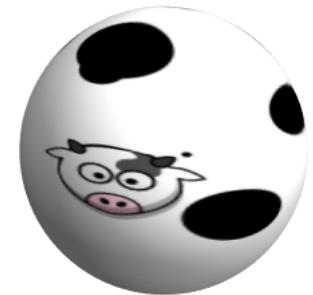
We can think of it as an invertible neural network:



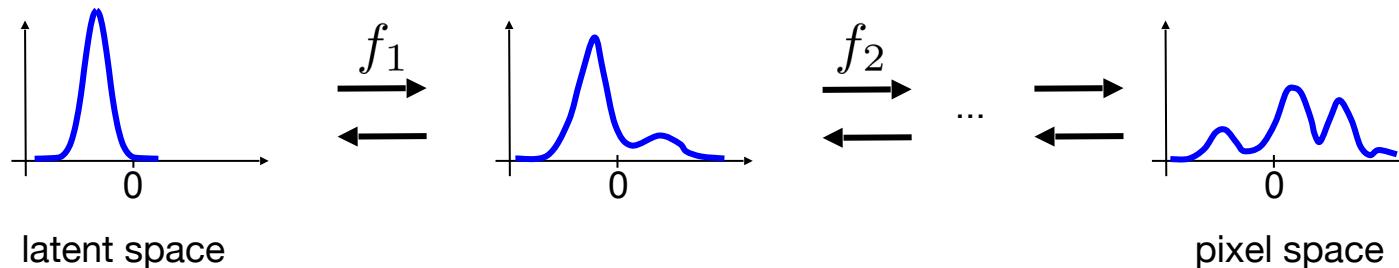
THE CHANGE OF VARIABLES FORMULA

Let's recall the change of variables formula with invertible transformations:

$$p(\mathbf{x}) = \pi_0(\mathbf{z}_0) \prod_{i=1}^K \left| \det \frac{\partial f_i(\mathbf{z}_{i-1})}{\partial \mathbf{z}_{i-1}} \right|^{-1}$$



We can think of it as an invertible neural network:



REALNVP

Design the invertible transformations as follows:

$$\mathbf{y}_{1:d} = \mathbf{x}_{1:d}$$

$$\mathbf{y}_{d+1:D} = \mathbf{x}_{d+1:D} \odot \exp(s(\mathbf{x}_{1:d})) + t(\mathbf{x}_{1:d})$$

REALNVP

Design the invertible transformations as follows:

$$\mathbf{y}_{1:d} = \mathbf{x}_{1:d}$$

$$\mathbf{y}_{d+1:D} = \mathbf{x}_{d+1:D} \odot \exp(s(\mathbf{x}_{1:d})) + t(\mathbf{x}_{1:d})$$

Invertible by design:

$$\left\{ \begin{array}{lcl} \mathbf{y}_{1:d} & = & \mathbf{x}_{1:d} \\ \mathbf{y}_{d+1:D} & = & \mathbf{x}_{d+1:D} \odot \exp(s(\mathbf{x}_{1:d})) + t(\mathbf{x}_{1:d}) \end{array} \right. \Leftrightarrow \left\{ \begin{array}{lcl} \mathbf{x}_{1:d} & = & \mathbf{y}_{1:d} \\ \mathbf{x}_{d+1:D} & = & (\mathbf{y}_{d+1:D} - t(\mathbf{y}_{1:d})) \odot \exp(-s(\mathbf{y}_{1:d})) \end{array} \right.$$

REALNVP

Design the invertible transformations as follows:

$$\mathbf{y}_{1:d} = \mathbf{x}_{1:d}$$

$$\mathbf{y}_{d+1:D} = \mathbf{x}_{d+1:D} \odot \exp(s(\mathbf{x}_{1:d})) + t(\mathbf{x}_{1:d})$$

Invertible by design:

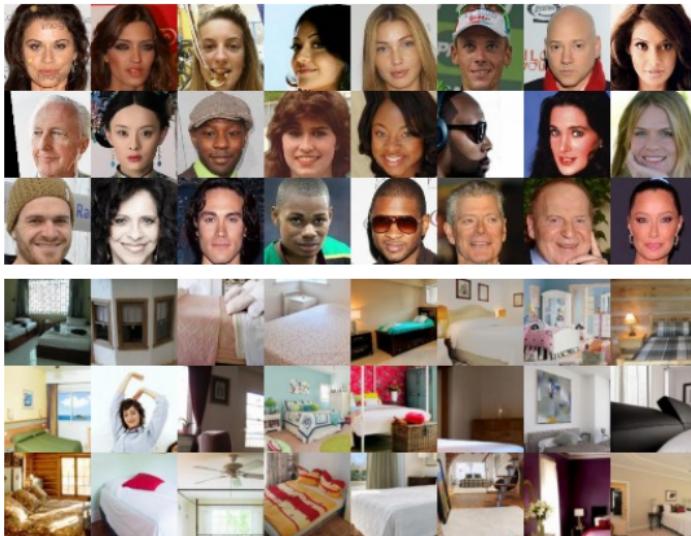
$$\begin{cases} \mathbf{y}_{1:d} = \mathbf{x}_{1:d} \\ \mathbf{y}_{d+1:D} = \mathbf{x}_{d+1:D} \odot \exp(s(\mathbf{x}_{1:d})) + t(\mathbf{x}_{1:d}) \end{cases} \Leftrightarrow \begin{cases} \mathbf{x}_{1:d} = \mathbf{y}_{1:d} \\ \mathbf{x}_{d+1:D} = (\mathbf{y}_{d+1:D} - t(\mathbf{y}_{1:d})) \odot \exp(-s(\mathbf{y}_{1:d})) \end{cases}$$

Easy Jacobian:

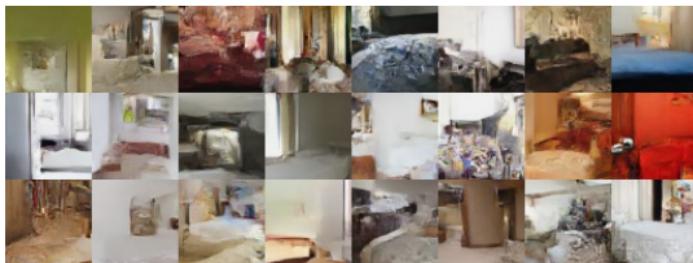
$$\mathbf{J} = \left[\begin{array}{cc} \mathbb{I}_d & \mathbf{0}_{d \times (D-d)} \\ \frac{\partial \mathbf{y}_{d+1:D}}{\partial \mathbf{x}_{1:d}} & \text{diag}(\exp(s(\mathbf{x}_{1:d}))) \end{array} \right]$$

$$\det(\mathbf{J}) = \prod_{j=1}^{D-d} \exp(s(\mathbf{x}_{1:d}))_j = \exp\left(\sum_{j=1}^{D-d} s(\mathbf{x}_{1:d})_j\right)$$

RESULTS



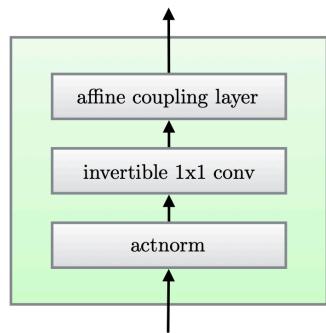
RESULTS



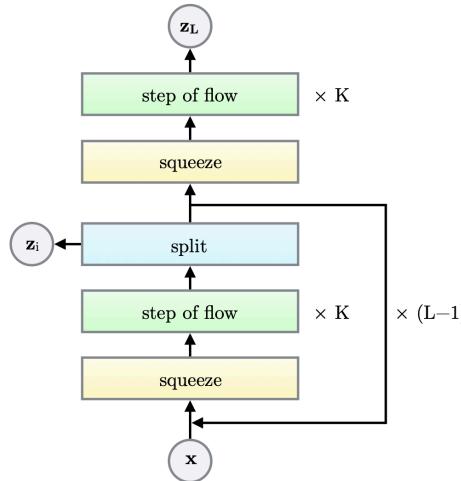
GLOW: REALNVP WITH 1X1 CONVOLUTIONS

A model contains ~1000 convolutions.

A new component: 1x1 convolution instead of a permutation matrix.

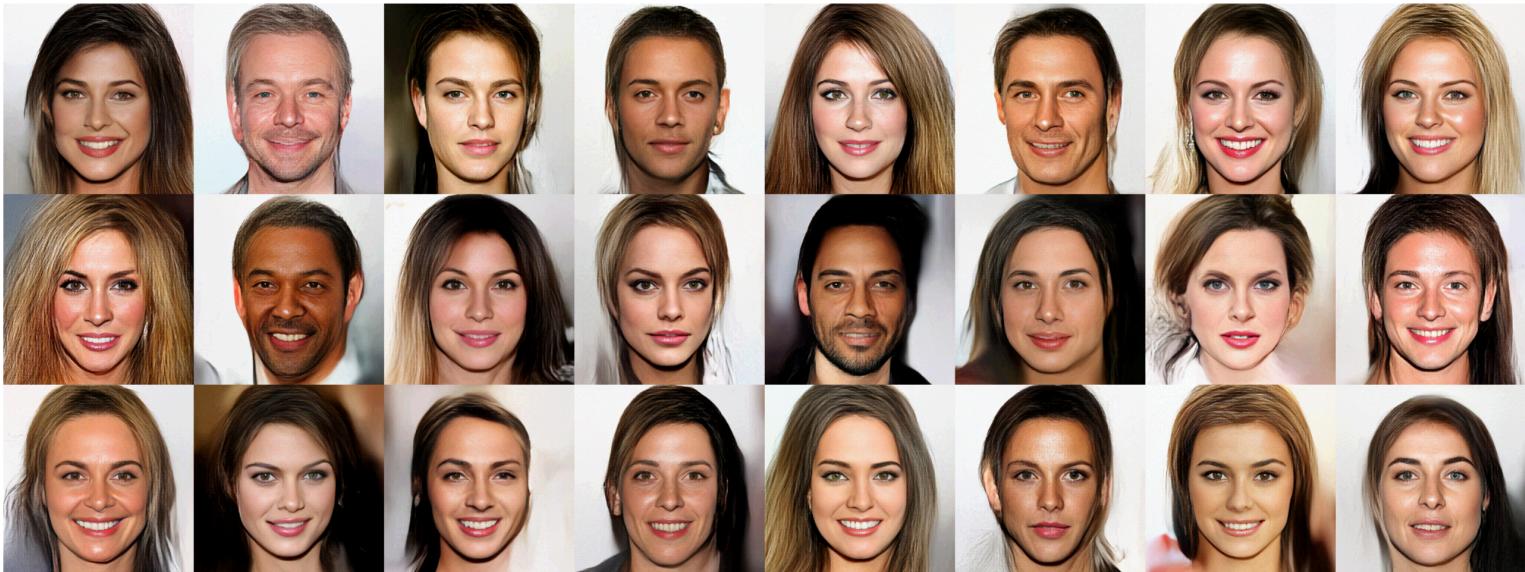


(a) One step of our flow.

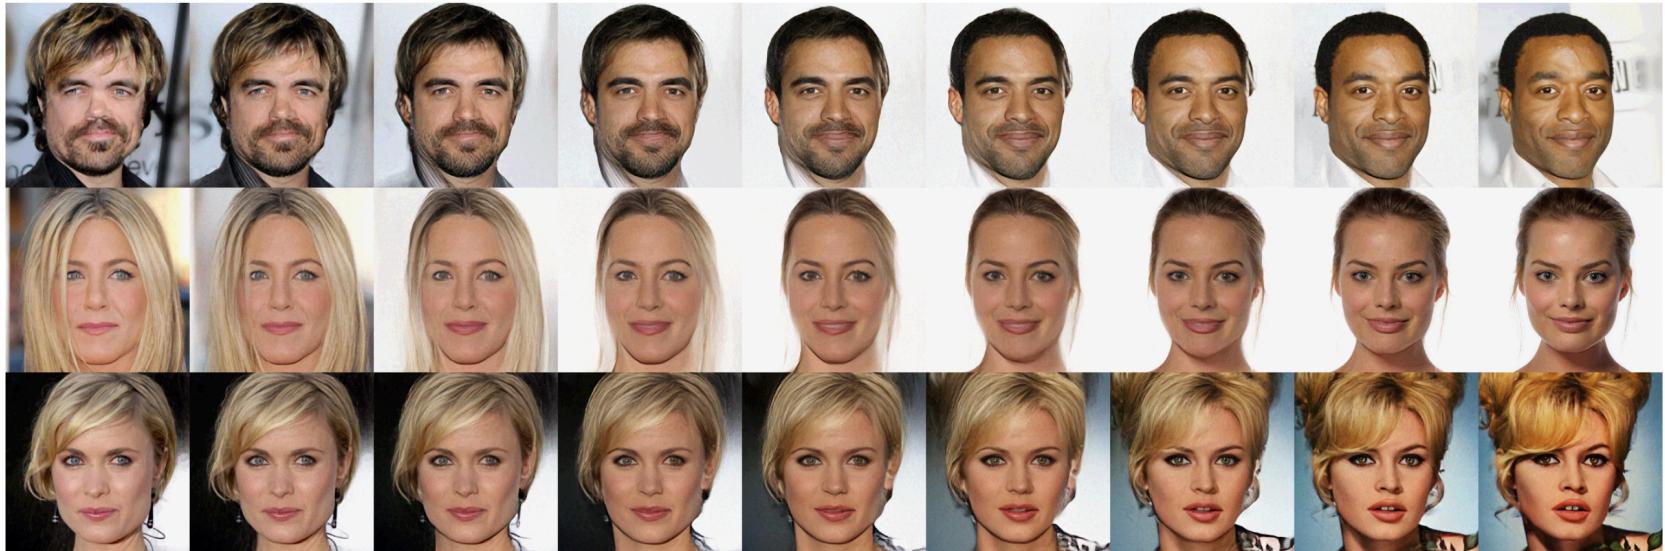


(b) Multi-scale architecture (Dinh et al., 2016).

GLOW: SAMPLES



GLOW: LATENT INTERPOLATION



INTEGER DISCRETE FLOW: NO NEED TO CALCULATE JACOBIAN!



~15%



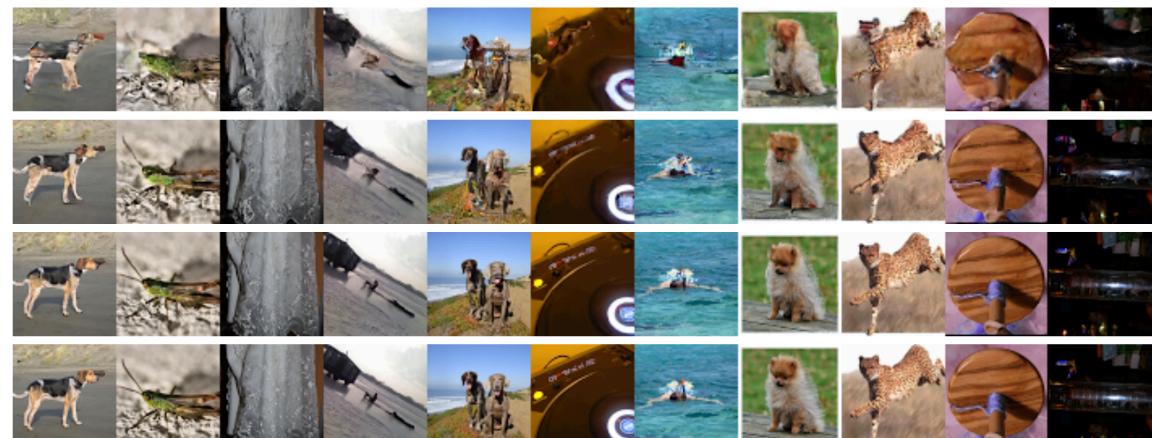
~30%



~60%



100%



Future directions

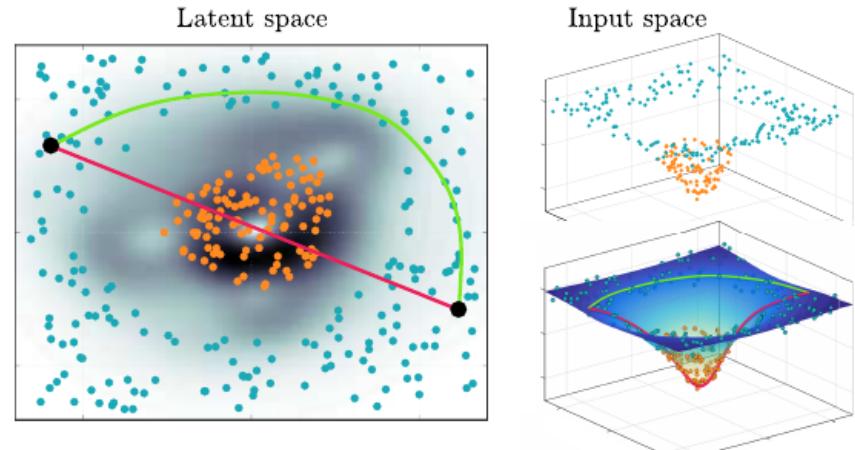
BLURRINESS AND SAMPLING IN VAEs

How to avoid sampling from **holes**?

Should we follow **geodesics in the latent space**?

How to use **geometry** of the latent space to build better **decoders**?

How to build **temporal decoders**?



COMPRESSION AND VAEs

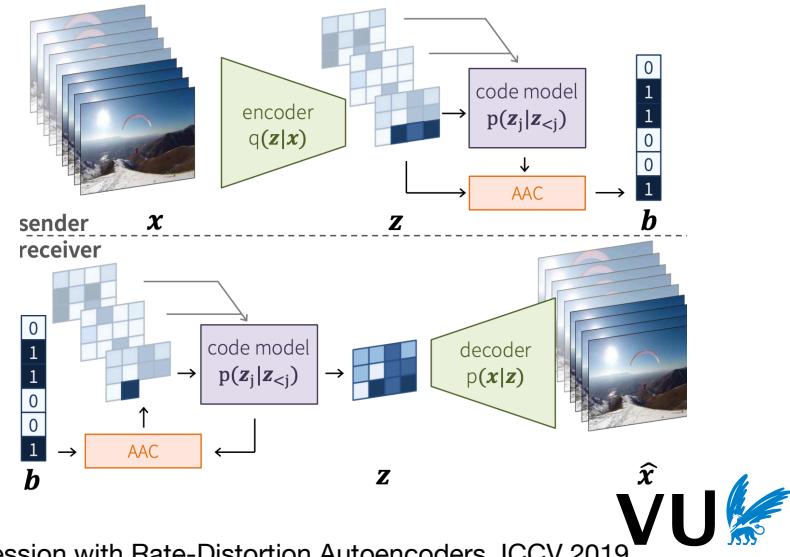
Taking a **deterministic & discrete encoder** allows to simplify the objective.

It is important to learn a **powerful prior**.
This is challenging!

Is it **easier** to learn a prior with **temporal dependencies**?

Can we alleviate some dependencies by using **hypernets**?

$$\begin{aligned} \text{RE}(x|z) - \text{H}[q(z|x)] - \text{CE}[q(z)||p(z)] \\ = \text{RE}(x|z) - \text{CE}[q(z)||p(z)] \end{aligned}$$

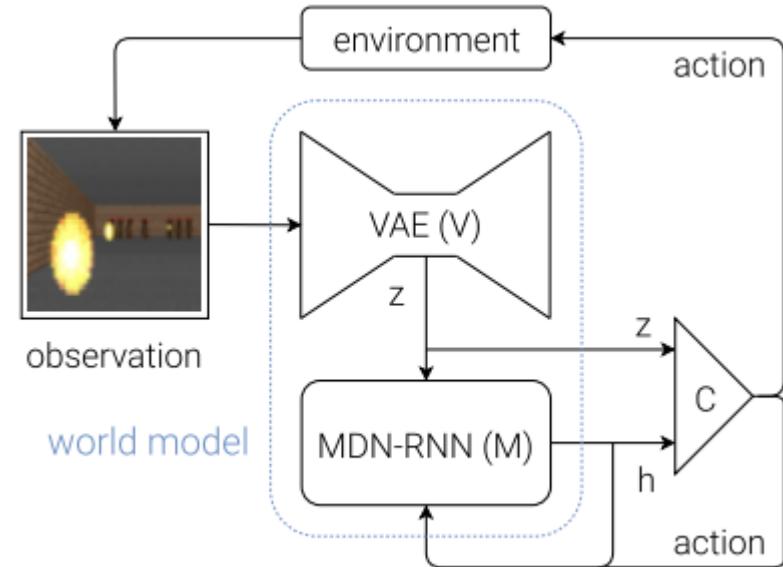


ACTIVE LEARNING/RL AND VAEs

Using **latent representation** to navigate and/or quantify uncertainty.

Formulating **policies** in the latent space entirely.

Do we need a better notion of **sequential dependencies**?



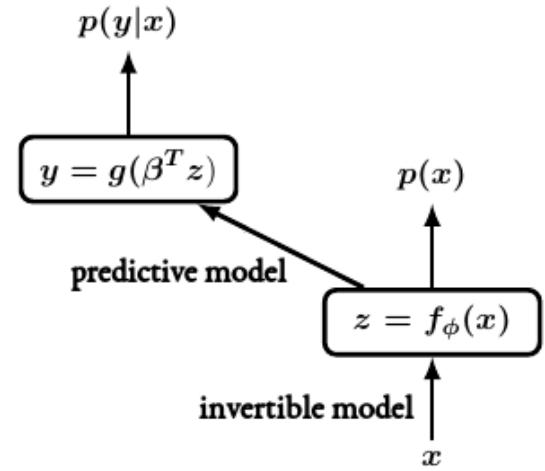
HYBRID AND FLOW-BASED MODELS

We need a **better understanding** of the latent space.

Joining an **invertible model** (flow-based model) with a **predictive model**.

Isn't this model an **overkill**?

How would it work in the **multi-modal learning** scenario?

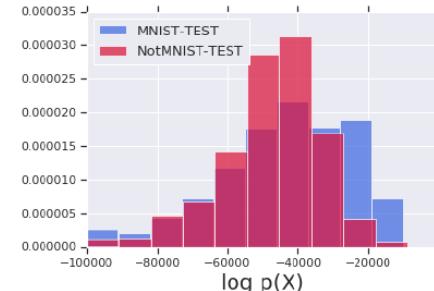


HYBRID MODELS AND OOD SAMPLE

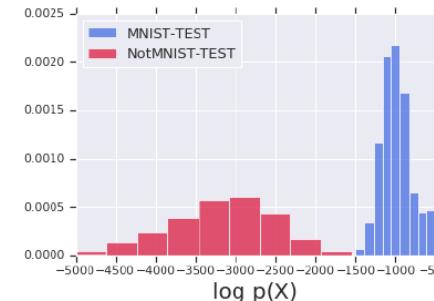
Going back to first slides, we need a good notion of $p(x)$.

Distinguishing **out-of-distribution (OOD)** samples is very important.

Crucial for **decision making, outlier detection, policy learning...**



(a) Discriminative Model ($\lambda = 0$)



(b) Hybrid Model

Thank you!

Webpage:

<https://jmtomczak.github.io/>

Code on github:

<https://github.com/jmtomczak/>

Contact:

jmk.tomczak@gmail.com